

FlowHub – CAS AISE Projektabgabe

CAS AI-Assisted Software Engineering (AISE) · W4B-C-AS001 · ZH-Sa-1 · FS26 **Student:** Andreas Imboden (`freaxnx01`) **Repository:** <https://github.com/freaxnx01/FlowHub-CAS-AISE> **Abgabedatum:** Mai 2026

Hinweis zum Aufbau dieses Dokuments

Diese Datei ist die zentrale Einreichungs-Seite für die CAS-AISE-Projektarbeit. Sie wird als PDF in Moodle hochgeladen und enthält bewusst **keine** vollständigen Inhalte — alle Artefakte (Architektur, ADRs, Block-Nachbereitungen, Reflexion, Bewertungskriterien-Selfcheck, ...) liegen im verlinkten GitHub-Repository und sind über das Inhaltsverzeichnis weiter unten direkt erreichbar. Sämtliche Links zeigen auf den `main`-Branch und sind aus dem PDF heraus klickbar.

1. Projektzusammenfassung

FlowHub ist ein KI-gestützter persönlicher Eingangskorb, der Informationsschnipsel aus dem Alltag (Filmtipps, Artikel, Belege, Bookmarks, Notizen) automatisch erkennt, klassifiziert und an die passenden Self-Hosted-Services im Homelab des Benutzers weiterleitet — ohne dass der Benutzer im Moment der Erfassung entscheiden muss, wohin die Information gehört.

Das adressierte Kernbedürfnis ist "**Capture without friction**": Statt heute fünf Schritte (Idee → App-Wahl → App öffnen → Kategorisieren → Ablegen) reduziert FlowHub die Erfassung auf einen einzigen Schritt — typischerweise eine Nachricht an einen Telegram-Bot. Die Klassifikation übernimmt ein **Skill-basiertes Routing-System** (Keywords, URL-Muster, lokales LLM als Fallback), die Ablage erfolgt in bestehende Homelab-Services wie Vikunja, paperless-ngx oder Wallabag.

Technisch ist FlowHub ein **Modular Monolith in .NET 10** mit Blazor-Frontend (MudBlazor, Interactive Server) und einer hexagonalen Schichtung innerhalb der Module. Die Implementierung erfolgt inkrementell über die fünf CAS-Blöcke (Einführung, Frontend, Service, Persistence, Deployment) — jeder Block schliesst mit einer dokumentierten Nachbereitung ab, die gegen die Moodle-Bewertungskriterien selbst geprüft wird. Ein expliziter Fokus liegt auf **KI-**

unterstützter Entwicklung: Skills, Agent-Instructions, Prompt-Hygiene und Reflexion über den AI-Workflow sind dokumentiert und Teil des Bewertungsumfangs.

Die ausführliche Projektbeschreibung mit Stakeholdern, Funktionsumfang, Architektur und Risiken liegt unter

[docs/projektbeschreibung/FlowHub_Projektbeschreibung_v4.md](#) .

2. Repository

- **GitHub:** <https://github.com/freaxnx01/FlowHub-CAS-AISE>
- **Branch (Abgabestand):** `main`
- **Lizenz / README:** [README.md](#)
- **Changelog:** [CHANGELOG.md](#)
- **Agent-Konventionen:** [CLAUDE.md](#)

3. Inhaltsverzeichnis der Abgabe-Artefakte

3.1 Projektbeschreibung & Architektur

- [Projektbeschreibung v4 \(Markdown\)](#) — aktueller Stand: Vision, Stakeholder, Funktionsumfang, Architektur
- [Projektbeschreibung v4 \(PDF\)](#)
- [Arc42-Architekturdokumentation v1.1 \(PDF\)](#) — vollständige Arc42-Sicht
- [Architektur-Übersicht v2 \(SVG\)](#)

3.2 Architecture Decision Records (ADRs)

- [ADR 0001 — Frontend Render Mode & Architecture](#)
- [ADR 0002 — Service Architecture & Async Communication](#)

3.3 Spezifikation & Design

- [System Context](#)
- [Use Cases](#)
- [Testing Strategy](#)

- UI-Design-Output (Wireframes & Flows pro Feature):
 - [Dashboard](#)
 - [Captures-List](#)
 - [Capture-Detail](#)
 - [New Capture](#)
 - [API](#)
- POC-Pläne:
 - [AI-Classification POC — Design](#)
 - [AI-Classification POC — Plan](#)

3.4 Block-Nachbereitungen (CAS-Module)

Jeder Block hat drei Phasen: **Vorbereitung** (vor dem Präsenztage), **PVA** (Präsenz-/Vor-Ort-Arbeit am Samstag) und **Nachbereitung** (Implementierung & Reflexion). Die **Nachbereitungen** sind die jeweils gegen die Bewertungskriterien geprüften Liefergegenstände.

Block 1 — Einführung

- [Vorbereitung%20Vorbereitung.md](#))
- [PVA%20PVA.md](#))
- **[Nachbereitung%20Nachbereitung.md](#)**)

Block 2 — Frontend

- [Vorbereitung%20Vorbereitung.md](#))
- [PVA%20PVA.md](#))
- **[Nachbereitung%20Nachbereitung.md](#)**)

Block 3 — Service

- [Vorbereitung%20Vorbereitung.md](#))
- [PVA%20PVA.md](#))
- **[Nachbereitung%20Nachbereitung.md](#)**)

Block 4 — Persistence

- [Vorbereitung%20Vorbereitung.md](#))
- [PVA%20PVA.md](#))
- **[Nachbereitung%20Nachbereitung.md](#)**)

Block 5 — Deployment

- [Vorbereitung%20Vorbereitung.md](#))
- [PVA%20PVA.md](#))
- **Nachbereitung**[%20Nachbereitung.md](#))

3.5 Projektarbeit & Reflexion

- [Idee FlowHub](#) — Konzept-Notizen
- [Dev](#) — Entwicklungsnotizen
- [Skills](#) — Skill-System
- [External Services](#) — Integrationen
- [Glossary](#)
- **Learnings CAS AISE** — persönliche Lessons Learned (AI-Instructions, Skill-Plugins, Context Hygiene, Code-Exploration)

3.6 Bewertungskriterien & Selbsteinschätzung

- [Bewertungskriterien \(Moodle-Rubric\)](#) — kanonische Quelle: 18 Items, 5 Buckets, max. 100 Punkte
- Selbst-Check pro Block: jeweils im unteren Abschnitt der Block-Nachbereitung (siehe 3.4) als Checkliste mit Punktgewichtung
- **N/A-Hinweis:** Das Quarkus-/Jakarta-EE-Programmierkriterium (max. 10 Pkt.) ist für FlowHub bewusst nicht anwendbar — der Stack ist .NET 10 / [ASP.NET Core](#); Begründung in den Block-Nachbereitungen vermerkt.

3.7 Knowledge Base (Hintergrund)

Vertiefende Notizen aus den Vorlesungen, nicht Teil der primären Bewertung, aber als Kontext referenziert:

- [Software Architecture](#)
 - [UML](#)
 - [Akronyme](#)
-

4. Hinweise für die Bewertung

- **Primärer Einstiegspunkt** ist diese Datei (`SUBMISSION.md`) — alle anderen Artefakte sind über die Links in Abschnitt 3 erreichbar.
- **Codebasis** liegt im selben Repository unter `source/` (Modulstruktur), Tests unter `tests/`.
- **Run-Anleitung** und Dev-Konventionen siehe `CLAUDE.md` (Abschnitt *Essential Commands*).
- **KI-Nutzung** ist explizit dokumentiert — sowohl auf Block-Ebene (Reflexionsabschnitt der jeweiligen Nachbereitung) als auch konsolidiert im Learnings-Dokument (Abschnitt 3.5).

Project description (v4)

Source: `docs/projektbeschreibung/FlowHub_Projektbeschreibung_v4.md`

FlowHub – Projektbeschreibung

CAS AI-Assisted Software Engineering (AISE) W4B-C-AS001 · ZH-Sa-1 · FS26

Student: Andreas Imboden **Datum:** Februar 2026

1. Vision

FlowHub ist ein KI-gestützter persönlicher Eingangskorb, der Informationsschnipsel aus dem Alltag automatisch erkennt, kategorisiert und an die richtigen Services weiterleitet.

Der digitale Alltag produziert ständig kleine Informationsfragmente: ein Film den man schauen möchte, ein Tech-Artikel zum Lesen, ein Foto eines Kassenbelegs. Heute landen diese Schnipsel verstreut in verschiedenen Apps, Chats oder werden schlicht vergessen.

FlowHub schafft einen einzigen Eingang für all diese Inputs und erledigt die Ablage automatisch – mit KI-Unterstützung und minimalem Aufwand für den Benutzer.

2. Stakeholder

Primär: Homelab-Betreiber (Persona: Andreas)

- Technisch versierter Anwender
- Betreibt Self-Hosted Services im Heimnetzwerk (Proxmox, Docker)
- Nutzt bereits: paperless-ngx, Passbolt, GitLab, Vikunja, Wallabag
- Problem: Informationsschnipsel landen überall, kein einheitlicher Eingang
- Ziel: Schnelle, reibungslose Erfassung von Alltagsinformationen via Telegram

Sekundär: Digital Hoarders

- Sammeln Artikel, Bookmarks, Notizen in vielen Tools
- Frustriert durch manuelles Kategorisieren und Sortieren
- Würden von automatisierter Ablage profitieren

Tertiär: CAS-Dozenten (FFHS)

- Erwarten: Verteilte Web-Applikation mit KI-Einsatz
- Erwarten: Saubere Software-Architektur-Dokumentation
- Erwarten: Reflexion über KI-unterstützte Entwicklung

3. Kundenbedürfnis

Das adressierte Kernbedürfnis ist "**Capture without friction**": Der Benutzer möchte eine Information festhalten, ohne in diesem Moment entscheiden zu müssen, wohin sie gehört und wie sie abgelegt wird.

Heute:

Idee → Welche App? → App öffnen → Kategorisieren → Ablegen = 5+ Schritte, Kontextwechsel, oft vergessen

Mit FlowHub:

Idee → Telegram Bot → Fertig = 1 Schritt, KI übernimmt den Rest

4. Externe Services

FlowHub integriert ausschliesslich Self-Hosted Services aus dem eigenen Homelab (kein Cloud-SaaS):

Service	URL	Zweck
Vikunja	https://vikunja.io/	Task-Management für Bücher, Filme/TV
paperless-ngx	https://docs.paperless-ngx.com/	Dokumenten-Management-System (DMS)
Wallabag	https://wallabag.org/	Read-Later Service für Artikel
GitLab (Self-Hosted)	https://gitlab.freaxnx01.ch/	Repository mit Obsidian Markdown Files (Knowledge Base)

5. Funktionsübersicht

5.1 Skill-System

FlowHub verwendet ein **Skill-basiertes Routing**: Jeder eingehende Input wird einem Skill zugewiesen, der die Ablage in den passenden Ziel-Service übernimmt. Die Erkennung erfolgt über Keywords, URL-Muster und – falls nötig – ein lokales LLM.

ArticleSkill - Artikel zum späteren Lesen (→ Wallabag)

Erkennung: URL von Nachrichtenportalen, Blogs, Fachzeitschriften

Beispiele:

- <https://www.heise.de/select/ct/2026/4/2533109542020998570>
- <https://www.heise.de/ratgeber/Fluechtige-SSH-Schluessel-mit-opkssh-und-OpenID-Connect-generieren-10639864.html>

HomelabSkill - Homelab-Services zum Ausprobieren (→ Vikunja)

Erkennung: URL von Software/Service-Websites; Keywords: homelab, ausprobieren, install, self-host, try

Beispiele:

- <https://adguard.com/de/welcome.html>
 - <https://jellyfin.org/>
-

BookSkill - Bücher & IT-Bücher (→ Vikunja)

Erkennung: URL von Buchshops (exlibris, galaxus, amazon, orellfuesli);

Keywords: buch, lesen, bestellen, isbn

Beispiele:

- <https://www.exlibris.ch/de/buecher-buch/english-books/linus-torvalds/just-for-fun/id/9780066620732/>
 - <https://www.galaxus.ch/de/s18/product/eine-kurze-geschichte-der-menschheit-deutsch-yuval-noah-harari-2015-sachbuecher-7003551>
-

MovieSkill - Filme & TV-Serien (→ Vikunja)

Erkennung: Keywords: schauen, watch, film, movie, serie, TV; Google-Share-URLs (share.google/...)

Beispiele:

- The Imitation Game – Ein streng geheimes Leben + <https://share.google/1vwEtMUiRriic4nNi>
 - Star Trek + <https://share.google/338RxCPCv9Ytm0wgA>
-

KnowledgeSkill - Knowledge Base Einträge (→ Obsidian via GitLab)

Erkennung: Artikel-URL mit explizitem Wissensbasis-Kontext; Keywords: notiz, wissen, merken, knowledge, speichern

Beispiele:

- <https://www.heise.de/select/ct/2026/5/2532311091092661684>
 - <https://www.heise.de/select/ct/2026/5/2534615175182620957>
-

DocumentSkill - Dokumente (→ paperless-ngx)

Erkennung: Dateianhang (PDF, Foto); Keywords: quittung, rechnung, beleg, dokument

Beispiele:

- Foto einer Quittung
- PDF einer Rechnung

QuoteSkill - Zitate mit KI-Anreicherung (→ Obsidian via GitLab)

Erkennung: Keywords: zitat, quote, gesagt von; Anführungszeichen im Text; Share-Links auf Zitate

Enrichment via KI: Autor und AuthorInfo (Kontext, Werk, Biografie-Snippet) werden automatisch ergänzt

Beispiel:

- Shannons Zitat zur Informationstheorie:

<https://g.co/gemini/share/917675e7a359>

GenericSkill - Fallback (→ FlowHub Inbox / PostgreSQL)

Greift wenn kein anderer Skill passt. Input landet in der Inbox zur späteren manuellen Klassifizierung.

Skill-Vorschlag bei unbekanntem Input

Wenn die Kategorisierung keinen bestehenden Skill zuweisen kann, schlägt das System dem Benutzer einen neuen Skill vor:

```
Bot: "Ich konnte diesen Input keinem Skill zuweisen.  
Möchtest du einen neuen Skill erstellen?  
[Ja, Skill erstellen] [In Inbox ablegen] [Verwerfen]"
```

Bei Bestätigung wird eine SKILL.md-Vorlage generiert und der Input zwischengespeichert. Der neue Skill kann anschliessend konfiguriert und aktiviert werden.

5.2 MVP - CAS Projektarbeit

Der MVP-Scope ist bewusst eng gefasst. Komplexität wird durch saubere Architektur beherrscht, nicht durch Featureumfang.

Input-Kanal

Feature	Beschreibung	Status
Telegram Bot (Text)	Textnachrichten empfangen und verarbeiten	☐ MVP
Telegram Bot (Datei/Foto)	Dateianhänge empfangen und verarbeiten	☐ MVP

Skill-System (automatische Erkennung und Routing)

Skill	Erkennung	Ziel-Service	Status
ArticleSkill	URL (News/Blog/Fachmagazin)	Wallabag	☐ MVP
HomelabSkill	URL (Software/Service) + Keywords	Vikunja	☐ MVP
BookSkill	URL (Buchshop) + Keywords	Vikunja	☐ MVP
MovieSkill	Keywords + Google-Share- URL	Vikunja	☐ MVP
DocumentSkill	Dateianhang (Foto/PDF)	paperless-ngx	☐ MVP
KnowledgeSkill	URL + Keywords	Obsidian (GitLab)	Future
QuoteSkill	Keywords + Anführungszeichen	Obsidian (GitLab)	Future
GenericSkill	Kein anderer Skill greift	FlowHub Inbox (PostgreSQL)	☐ MVP
Skill-Vorschlag	Unbekannter Input → neuen Skill vorschlagen	-	Future

Ziel-Services (Integrationen)

Service	Typ	Integration	MVP-Umfang
Wallabag	Self-Hosted	REST API	URL speichern
Vikunja	Self-Hosted	REST API	Task/Item erstellen
paperless-ngx	Self-Hosted	REST API	Dokument hochladen
Obsidian / GitLab	Self-Hosted	Git API	Markdown-Datei erstellen
FlowHub Inbox	Eigene DB	PostgreSQL	Fallback-Speicher

Benutzerinteraktion

Feature	Beschreibung	Status
Direkte Verarbeitung	Klarer Input → sofortige Aktion	☐ MVP
Rückfrage mit Auswahl	Unklarer Input → Bot fragt mit 2-3 Optionen zurück	☐ MVP
Bestätigungs-Feedback	Bot antwortet mit Ergebnis	☐ MVP
Skill-Vorschlag	Unbekannter Input → neuen Skill vorschlagen	Future

KI-Integration

Feature	Beschreibung	Status
Keyword-basierte Erkennung	Zuverlässiger Fallback ohne KI-Kosten	☐ MVP
Microsoft.Extensions.AI + Ollama	Lokales LLM für Skill-Erkennung	☐ MVP (wenn Zeit reicht)
Confidence-Score	Schwellwert bestimmt ob Rückfrage nötig	☐ MVP (wenn Zeit reicht)
Enrichment (QuoteSkill)	Autor, AuthorInfo via KI ergänzen	Future

5.3 Future - Spätere Versionen (nach CAS)

Diese Features sind architekturell vorbereitet, werden aber nicht für die Abgabe implementiert. Sie belegen die Zukunftsfähigkeit der Architektur.

Erweiterte Skill-Funktionalität

Skill	Erweiterung
MovieSkill	IMDB/OMDB Metadaten (Titel, Rating, Jahr, Genre)
ArticleSkill	Zusätzlich: Obsidian-Notiz mit KI-Summary erstellen
DocumentSkill	AI OCR-Analyse (Betrag, Datum, Händler, Tags)
QuoteSkill	Vollständige Enrichment-Pipeline: Autor, Werk, Kontext
KnowledgeSkill	RAG-Integration: Obsidian als Wissensbasis für den Assistenten
Skill-Vorschlag	Automatische SKILL.md-Generierung via KI

Erweiterte Integrationen

Service	Erweiterung
Obsidian / GitLab	RAG-Quelle für den Assistenten
paperless-ngx	AI-gestützte Tagging und Metadaten-Extraktion
Vikunja	Erweiterte Metadaten (Cover, Rating, Genre für Filme/Bücher)

Erweiterte KI-Features

Feature	Beschreibung
Lerneffekt	Nutzerverhalten verbessert Confidence über Zeit
RAG	Obsidian als Wissensbasis für den Assistenten
Enrichment	Automatische Anreicherung von Quotes, Büchern, Filmen

Weitere Input-Kanäle

Kanal	Beschreibung
Email	Emails als Input verarbeiten
Web Upload	Browser-Extension oder Web-UI für Uploads
API	Direkte REST-API für Drittanwendungen

Deployment-Evolution


Phase	Beschreibung
Phase 1 (MVP)	Docker Compose auf Proxmox Homelab
Phase 2 (Future)	Migration zu k3s (Kubernetes)

6. Systemarchitektur

6.1 Überblick

Hinweis — Konzept-/Zielarchitektur, nicht Abgabestand. Das folgende Übersichtsbild zeigt die *angestrebte* Gesamtarchitektur und enthält bewusst auch noch **nicht implementierte** Bausteine (z. B. Ollama-LLM, Telegram-Kanal, Redis). Die tatsächlich gebaute, als-built-Architektur ist der **Modular Monolith** aus sechs Projekten mit MassTransit-Pipeline — maßgeblich sind dafür die C4-/Hexagonal-Diagramme in `docs/design/perspectives.md`, das ER-Diagramm in `docs/design/db/er.md` und die ADRs (insb. ADR 0002/0003/0005, jeweils mit „As built“-Notiz). Siehe auch `docs/spec/system-context.md` → „Current state (Block 5)“.

Ist-Architektur (Block 5, as built) — der tatsächlich gebaute Modular Monolith:



Zielbild (Konzeptphase) — enthält noch nicht gebaute Bausteine:

FlowHub Systemarchitektur .NET 10 / C# / [ASP.NET](#) Core · Blazor SSR · [Microsoft.Extensions.AI](#)

□ INPUT-LAYER Telegram Bot · Text + Datei + Foto Einziger Eingang · "Capture without friction" ⚙ FLOWHUB BACKEND .NET 10 / C# / ASP.NET Core □ Telegram Bot Service Webhook Handler □ SKILL-SYSTEM SkillRegistry · lädt SKILL.md Files · Dependency Injection □ ArticleSkill → Wallabag □ HomelabSkill → Vikunja □ BookSkill → Vikunja □ MovieSkill → Vikunja □ DocumentSkill → paperless-ngx □ KnowledgeSkill → Obsidian/GitLab □ QuoteSkill → Obsidian/GitLab □ GenericSkill → Inbox (PostgreSQL) □ Skill-Vorschlag Unbekannter Input → neuen Skill vorschlagen · SKILL.md generieren SkillDispatcher · Keyword-Match → AI Confidence-Score (0.0–1.0) → Direkt / Rückfrage □ AI-LAYER M.E.AI Abstraction Ollama (lokal) Llama 3.x Anthropic API (Fallback) Confidence-Scoring Provider via Config wechselbar □ REST CLIENTS via Refit (typesicher) WallabagClient VikunjaClient PaperlessClient GitLabClient ISkillHandler → Port-Abstraktion □ Blazor SSR · Admin-Dashboard Server-Side Rendering · .NET-native · kein JavaScript-Framework Inbox-Verwaltung · Skill-Konfiguration · Verarbeitungs-Logs Skill-Vorschlag Review · Pending Inputs □

PERSISTENCE PostgreSQL 16 FlowHub Inbox · EF Core · Migrations Redis 7
Pending Inputs · Session-State □ HOMELAB · Proxmox Self-Hosted □ Wallabag
Read-Later wallabag.org □ Vikunja Task-Management vikunja.io □ paperless-ngx
DMS paperless-ngx.com □ GitLab + Obsidian Knowledge Base gitlab.freaxnx01.ch
Legende:

Input Layer Backend Skills AI-Layer REST Clients Blazor SSR Persistence Homelab
(Self-Hosted) FlowHub – CAS AI-Assisted Software Engineering (FFHS FS26) ·
Andreas Imboden · Februar 2026 Alle Services Self-Hosted auf Proxmox ·
Anthropic API als KI-Fallback · Docker Compose Deployment

6.2 Hybrid Skill-System (Kernarchitektur)

FlowHub verwendet einen **Hybrid-Ansatz** für das Skill-System, der Code und Konfiguration kombiniert:

SKILL.md (Deklarativ, versioniert):

```
---
name: article-skill
description: Saves articles for later reading to Wallabag
handler: FlowHub.Skills.Handlers.ArticleSkillHandler
metadata:
  flowhub:
    triggers:
      keywords: [artikel, lesen, read later, später lesen]
      urlPatterns: [heise.de, ct.de, arstechnica.com]
    config:
      wallabag_tag: flowhub
---
# Dokumentation des Skills (für Menschen und KI lesbar)
```

ArticleSkillHandler.cs (Business-Logic, typsicher):

```

public class ArticleSkillHandler : ISkillHandler
{
    private readonly IWallabagClient _wallabag;

    public ArticleSkillHandler(IWallabagClient wallabag)
    {
        _wallabag = wallabag;
    }

    public async Task<SkillResult> ExecuteAsync(InputItem input, SkillConfi
    {
        await _wallabag.SaveEntryAsync(input.Url, config.WallabagTag);
        return new SkillResult { Success = true };
    }
}

```

Vorteile dieses Ansatzes:

- Konfiguration ohne Neustart änderbar (SKILL.md)
- Typsichere Business-Logic (C#)
- Selbstdokumentierend (Markdown human readable)
- Zukunftsfähig: KI kann SKILL.md Files lesen und generieren

6.3 Technologie-Stack

Schicht	Technologie	Begründung
Backend	.NET 10 / C# / ASP.NET Core	LTS-Release Nov. 2025, vertrauter Stack, hohe Performance
Web-UI	Blazor SSR (Server-Side Rendering)	.NET-native, kein JS-Framework nötig, SEO-freundlich
KI-Integration	Microsoft.Extensions.AI + Ollama / Anthropic SDK	Abstraktion über LLM-Provider, Ollama: lokal & kostenlos
Datenbank	PostgreSQL 16	Bewährt, EF Core Integration
Cache/State	Redis 7	Session-State, Pending Inputs
API-Framework	ASP.NET Core Minimal APIs	Schlank, performant, .NET-nativ
REST Clients	Refit	Typsicher, deklarativ, Interface-basiert
ORM	Entity Framework Core 10	.NET-Standard, Code-First Migrations
Telegram	Telegram.Bot (NuGet)	Offiziell unterstützte .NET-Library
Deployment	Docker Compose	Einfach, Proxmox-kompatibel

6.4 Infrastruktur (Proxmox Homelab)

```
Proxmox VE
+-- VM: Docker Host
|   +-- FlowHub Backend (.NET 10)
|   +-- PostgreSQL
|   +-- Redis
|   +-- Ollama (Llama 3.x)
|
+-- VM/LXC: Wallabag          (Self-Hosted Read-Later)
+-- VM/LXC: paperless-ngx    (bereits vorhanden)
+-- VM/LXC: Vikunja          (Task-Management / Kanban)
+-- VPS (extern, DE):
    +-- GitLab                (Obsidian Markdown Repo)

Cloud:
+-- Anthropic API             (Claude, Fallback für KI)
```

7. Architekturentscheidungen (ADR)

Dieses Kapitel listet die **frühen Plattform- und Strategie-Entscheidungen** aus der Konzeptphase (vor Block 1). Sie definieren den Rahmen für die spätere Implementierung. Während der Umsetzung sind sechs **lightweight Implementation-ADRs** (Context → Decision → Consequences) im Repository unter [docs/adr/](#) entstanden, die dort einzeln einsehbar sind:

ADR	Titel	Status
0001	Frontend Render Mode and Architecture	Accepted
0002	Service Architecture and Async Communication	Accepted
0003	Async Pipeline (MassTransit)	Accepted
0004	AI Integration — Provider Abstraction (MEAI)	Accepted
0005	Persistence — EF Core + PostgreSQL	Accepted
0006	Vector Search — pgvector + Mistral Embeddings	Accepted

Die nachfolgenden Einträge **PE-1 bis PE-7** (Plattform-Entscheidungen) sind bewusst von den Implementation-ADRs entkoppelt — sie sind Strategie, nicht

Lösungsstruktur.

PE-1: C# / .NET 10 als primäre Sprache und Plattform

Entscheidung: C# mit .NET 10 als primäre Sprache.

Begründung:

- Vertrauter Stack aus dem beruflichen Hintergrund (geringere Lernkurve)
- .NET 10 als aktueller LTS-Release (erschienen November 2025) mit langfristigem Support
- Starkes Ökosystem: EF Core, [ASP.NET](#) Core, Refit, Blazor – alles aus einem Guss
- Null-Safety via Nullable Reference Types (C# 8+)
- In der PVA explizit bestätigt: freie Wahl des Technologie-Stacks

Konsequenz: Kein Wechsel zu Quarkus/Kotlin nötig; vorhandenes C#-Wissen direkt produktiv einsetzbar.

PE-2: Blazor SSR für das Web-Dashboard

Entscheidung: Blazor Server-Side Rendering (SSR) für die Admin-Ansicht.

Begründung:

- .NET-native, kein separates JavaScript-Framework (React, Angular, Vue) nötig
- SSR liefert fertig gerendertes HTML – schnell, SEO-freundlich, kein WASM-Download
- Komponenten-Modell vertraut aus dem .NET-Ökosystem
- Volle C#-Typsicherheit auch im Frontend

Konsequenz: Blazor SSR eignet sich optimal für das Admin-Dashboard; für komplexe interaktive Elemente kann Blazor Interactive Server Mode ergänzt werden.

PE-3: Hybrid Skill-System ([SKILL.md](#) + C# Handler)

Entscheidung: Kombination aus deklarativer Konfiguration (`SKILL.md`) und typsicherer Business-Logic (C#).

Begründung:

- Pure Code: Gut testbar, keine Runtime-Konfiguration
- Pure Config: Flexibel, aber keine Typsicherheit
- Hybrid: Best of Both Worlds

Konsequenz: Zwei Artefakte pro Skill (`SKILL.md` + Handler), aber maximale Flexibilität.

PE-4: Docker Compose jetzt, k3s als Future Option

Entscheidung: Docker Compose für MVP-Deployment.

Begründung:

- Schnellerer Start (1 Tag Setup vs. 1 Woche k8s)
- Einfacheres Debugging während Entwicklung
- Passt zu FFHS Block 5 Timeline (k8s im Juni)
- App ist von Beginn an k8s-ready (12-Factor, Health Checks, Stateless)

Konsequenz: Migration zu k3s als Block 5 Projekt möglich, ohne Code-Änderungen.

PE-5: Ollama (lokal) als primäres LLM

Entscheidung: Llama 3.x via Ollama lokal, Anthropic API als Fallback.

Begründung:

- Ollama: Kostenlos, privat, offline-fähig
- Anthropic API: ~\$6/Monat für Homelab-Nutzung, nur für hochwertige Aufgaben
- `Microsoft.Extensions.AI` : Provider-Wechsel nur via Config, kein Code-Änderung

Konsequenz: 4.7 GB Model-Download, 4 GB RAM Bedarf für Ollama.

PE-6: paperless-ngx für Dokumenten-MVP als reines Upload-Ziel

Entscheidung: Im MVP fungiert paperless-ngx als Dokumenten-Archiv ohne KI-Analyse.

Begründung:

- Fokus auf funktionierende Grundstruktur, nicht Feature-Tiefe
- OCR und Metadaten-Extraktion sind komplex (separates Lernziel)
- paperless-ngx hat eigene gute Basis-Kategorisierung via Regeln

Konsequenz: DocumentSkill macht nur: Erkennen → Upload. Keine Tags, keine Metadaten-Extraktion im MVP. **Future:** KI-gestützte Analyse des OCR-Texts (Betrag, Datum, Händler) via `Microsoft.Extensions.AI`.

PE-7: Kein Cloud-SaaS - ausschliesslich Self-Hosted Services

Entscheidung: Alle Ziel-Services laufen Self-Hosted im Homelab. Todoist (Cloud SaaS) wird nicht integriert.

Begründung:

- Datensouveränität: Alle persönlichen Daten bleiben im eigenen Homelab
- Kostenfreiheit: Keine laufenden SaaS-Gebühren
- Vikunja ersetzt Todoist vollständig für Task/Item-Management und übernimmt auch die Kanban-Sicht für Homelab-Services

Konsequenz: Alle 4 Ziel-Services (Wallabag, Vikunja, paperless-ngx, GitLab) sind Self-Hosted auf Proxmox.

8. CAS-Kursstruktur und Projektfortschritt

Mapping: FFHS Blöcke → FlowHub Entwicklung

Block	Datum	FFHS Thema	FlowHub Deliverable
1	Feb 21	Konzeption, Einführung	□ Projektbeschreibung, Architekturentwurf, ADRs
2	Mär 21	Frontend, Web-Präsentation	Telegram Bot UI, Blazor SSR Dashboard (Admin-Ansicht)
3	Apr 25	Services, REST, MCP	Skill-System, REST Clients via Refit (Wallabag, Vikunja, paperless)
4	Mai 23	Persistence, Datenbanken	PostgreSQL Schema via EF Core, Redis State, Repository-Pattern
5	Jun 20	Docker, Kubernetes	Docker Compose finalisiert, optional: k3s Migration
Abgabe	Jul 6	-	Vollständiges MVP + Dokumentation + Git-Repository

9. KI-Einsatz in der Entwicklung (Reflexion)

Detaillierte Auswertung pro Block / Slice: `docs/ai-usage.md` (531 Zeilen, gegliedert pro Block mit *generated-vs-handwritten*-Tabellen, Korrektur-Geschichten und Reflexion). Dieser Abschnitt fasst die Erkenntnisse über alle fünf Blöcke zusammen.

KI-Tools im Entwicklungsprozess

Werkzeug	Einsatzbereich
Claude Code (Opus 4.7, 1M-Kontext)	Brainstorming, Spec/Plan-Erstellung, ADR-Drafts, Controller für Subagent-Dispatches
Claude Sonnet 4.6 (Subagents)	Implementer + Spec-Reviewer + Code-Quality-Reviewer unter dem <code>superpowers:subagent-driven-development</code> -Workflow
Claude Haiku 4.5 (Subagents)	Mechanische Tasks (csproj/Markdown/YAML), ca. 60 % Token-Ersparnis gegenüber Sonnet
<code>/ultrareview</code> (Multi-Agent Branch-Review)	Architektonisches Review über ganzen Feature-Branch — fängt System-Invarianten ab, die Per-Task-Review nicht sieht
Microsoft.Extensions.AI + Anthropic/OpenRouter	KI im Produkt (Classifier, Quote-Enricher, Embeddings)
Eigenentwickelte CAS-Skills	<code>cas-aise-todo-list</code> , <code>cas-aise-grade-self-check</code> , <code>sync-ai-instructions</code> — Rubrik-Verankerung, Block/Phase-Steuerung, Cross-Project-Reuse

Workflow-Entwicklung (Block 1 → Block 5)

Mit jedem Block wurde ad-hoc Chat gegen strukturierteren, isolierteren Subagent-Dispatch getauscht. Der Mehraufwand für Struktur hat sich durch weniger Mid-Task-Eskalationen und seltenere "Agent ist abgedriftet"-Momente bezahlt gemacht.

Block	Workflow-Verschiebung
1 — Einführung	Ad-hoc Claude.ai-Chat für Architektur; Copilot inline für Code.
2 — Frontend	Erster sustained Claude-Code-CLI-Einsatz mit phasenbasierter Disziplin (<code>/ui-brainstorm</code> → <code>/ui-flow</code> → <code>/ui-build</code> → <code>/ui-review</code>).
3 — Service	Erster vollständiger <code>superpowers:brainstorming</code> → <code>writing-plans</code> → <code>subagent-driven-development</code> -Slice. Zweistufiges Review (Spec + Quality) pro Task.
4 — Persistence	Sonnet als Default, Haiku für Mechanik; Per-Task-Quality-Review fallen gelassen zugunsten einer Branch-Review am Slice-Ende — ca. 60 % Token-Burn reduziert ohne Qualitätsverlust.
5 — Deployment	<code>/ultrareview</code> (Multi-Agent Branch-Review) und rubrik-gegrundetes <code>cas-aise-grade-self-check</code> ergänzt; jede Block-Abschluss-Prüfung läuft durch die Skill.

Bis Block 5 verantworten Implementer-Subagents ganze Slices end-to-end. Die menschliche Arbeit konzentriert sich an zwei Stellen: **Spec** (wo Verträge festgenagelt werden) und **Review** (wo System-Invarianten geprüft werden). Der Engpass wandert vom Tippen zu diesen beiden Punkten.

Wiederkehrende Fehlerquellen

1. **Single-Pass-AI-Review übersieht System-Invarianten.** Block-5-Beispiel: AI-generierter Code, der Embedding-Generation in `EfCaptureService.SubmitAsync` integrierte, war lokal korrekt (eine Transaktion) — global aber falsch, weil ein langsamer Provider das NfA-09 p95 < 200 ms Submit-Budget sprengt. Per-Task-Quality-Review konnte das nicht fangen; `/ultrareview` wurde dafür eingeführt und hat es beim ersten Lauf gefunden.
2. **Deployment-Shape-Failures sind systematisch unter-getestet.** Compose- `${X:-}` -Interpolation substituiert leere Strings (nicht null), wodurch `??` -Defaults still no-op'en. `.editorconfig` nicht in den Docker-Build-Kontext kopiert → Analyzer-Suppressions werden zu Build-Errors. Casing-Mismatch `EMBEDDINGS__APIKEY` vs. `Embeddings__ApiKey` schaltete ein ganzes Feature aus. Fünf solcher Defekte in einem Nachmittag durch `just smoke-prod` gefangen.

3. **Training-Data-Lag bei aktuellen Libraries.** `Pgvector.EntityFrameworkCore` wurde als in-the-box-Npgsql-10-Feature angenommen (ist aber ein separates Paket). `MassTransit.Testing` wurde als separates Paket angenommen (ist Teil der Haupt-Assembly). API-Behauptungen der KI müssen mechanisch gegen das tatsächlich installierte Paket verifiziert werden.
4. **Offene Prompts inflationieren den Scope.** "Füge X hinzu" ohne Plan produziert eine Feature-Suite. Das Gegenmittel war `superpowers:writing-plans` — kleinteilige TDD-geordnete Tasks mit vollständigem Inline-Code, exakten Pfaden und Commit-Messages. Subagents gegen wohldefinierte Tasks meldeten DONE; Subagents gegen offene Prompts gingen explorieren.
5. **Pläne aus unvollständigem Repo-Wissen.** Block-3-Slice-B-Plan übersah `CaptureServiceStubTests.cs` und die existierenden `ChannelKind`-Werte; Implementer mussten mid-Task adaptieren. Brainstorming/Planning liest seitdem jede referenzierte Datei *bevor* der Plan eingefroren wird.

Wiederkehrende Stärken

1. **Der Plan ist der Vertrag.** Enthält der Plan exakte Pfade, exakten Code, exakte Commit-Messages und exakte Verifikationskommandos, operiert der Implementer mit engem Judgment-Spielraum und meldet DONE statt NEEDS_CONTEXT. Die 95 %+ "AI-drafted"-Anteile im Produktionscode sind nur deshalb erreichbar.
2. **Review-Kadenz muss zur Fehlerklasse passen.** Per-Task-Spec + Per-Task-Quality war für Block 3 Slice B richtig (erster SDD-Lauf). Ab Block 4 fing Per-Task-Quality nichts Neues mehr, während ein Branch-weiter `/ultrareview` am Slice-Ende genau die Architektur-Fehler fand, die Per-Task-Review nicht sehen konnte.
3. **Eigene Skills als Memory-Layer des Projekts.** `cas-aise-todo-list`, `cas-aise-grade-self-check`, `sync-ai-instructions` decken Concerns ab, die das Upstream-`superpowers`-Plugin nicht behandelt: Rubrik-Verankerung, kalendergetriebene Priorisierung, Cross-Project-Instruction-Reuse. Ohne `cas-aise-grade-self-check` driftete die Rubrik-Abdeckung in langen Sessions — die Skill fand fehlende Evidenz in genau dieser Reflexion vor Abgabe.
4. **KI für rigide-Struktur-Artefakte, Mensch für architektonisches Urteil.** GitHub-Actions-YAML, ADR-Scaffolds, OpenAPI-ProblemDetails, EF-Migrationen — rigide Strukturen, in denen der 90-%-korrekte AI-Erstwurf reale Zeit spart. Architektur-Entscheidungen (hexagonaler Split für AI, In-Process vs. RabbitMQ Transport, Dispatcher-Event-Carries-Description vs. Inline-Skill-Call) bleiben beim Menschen; die KI listet Alternativen, sie wählt nicht.

5. Korrektur-Geschichten sind die Rubrik-Evidenz, die zählt.

`docs/insights/block-5.md` "Defects Found by the Smoke Run" und die "Ultrareview-driven correction" / "Smoke-driven correction"-Sektionen in `docs/ai-usage.md` sind auditierbar: konkreter Defekt, fixender Commit, Lehre. Pauschale "KI hat geholfen"-Aussagen schliessen keinen Rubrik-Loop; "KI produzierte X, Smoke fing Y, Fix landete in Commit Z" schliesst ihn.

Was ich anders machen würde

- **Rubrik-Verankerung ab Block 1, nicht Block 3.** `cas-aise-grade-self-check` entstand mittendrin; Block 1-2 hatten keine formalen `use-cases.md` / `nfa.md` / `acceptance-criteria.md`. Retroaktive Doku-Arbeit in spätem Block 3 / frühem Block 4 kostete reale Zeit.
- **/ultrareview ab Block 3, nicht Block 5.** Branch-weites Multi-Agent-Review fängt die Klasse von Architektur-Fehlern, die Per-Task-Review nicht sieht. Günstiger nach Slice B als nach einem 21-Task-Beta-MVP.
- **Ein just smoke -Rezept pro Block, nicht nur Block 5.** Fünf latente Bugs in einem Nachmittag durch `just smoke-prod`. Ein einfacheres Smoke (App booten, Feature-Pfad curlen) hätte mindestens drei davon früher gefangen.
- **InternalsVisibleTo - und EF-Core-Test-Host-Pattern im Plan-Template kodifizieren.** Block 3 Slice D und Block 4 hatten beide Implementer, die bei Compile-Errors per Default Sichtbarkeit aufweiteten. Das richtige Pattern ist mechanisch — sobald im Plan ausgeschrieben verschwindet die Friktion.

Hat die Ausgangshypothese gestimmt?

Die Block-1-Annahme war: Claude ist "ein schneller Tipper mit guter Library-Kenntnis" — nützlich für Boilerplate, suspekt bei Architektur, im Wesentlichen ein Produktivitäts-Multiplikator und keine Workflow-Änderung. Bis Block 5 ist das in beide Richtungen weit daneben.

Boilerplate ist schneller als erwartet. Der 95 %+ AI-drafted-Anteil im Produktionscode ist nicht "KI hat 95 % der Arbeit gemacht" — die 5 % Menschen-Input konzentrieren sich auf High-Judgment-Punkte (Scope, Verträge, Trade-offs). Das eigentliche Tippen ist aber wirklich ein kleiner Bruchteil der verstrichenen Zeit. Der Engpass wandert von der Implementation zu Spec und Review — exakt die Umkehrung, auf die SDD optimiert.

Architektur ist gefährlicher als erwartet. Der Embedding-on-Submit-Fall ist ein Ein-Absatz-Rewrite, der im Nachhinein offensichtlich ist: KI produzierte lokal korrekten Code, der eine System-Invariante brach, und ein Single-Pass-Review

hätte ihn ausgeliefert. Die Deployment-Shape-Failures haben die gleiche Form: lokal korrekt, global falsch. Die menschliche Rolle an der Architektur-Grenze ist nicht geschrumpft, sie ist gewachsen — weil der Agent jetzt genug Code schnell genug produziert, dass nur noch der Filter an Design und Review Sinn ergibt.

KI hat keine Rolle im Workflow ersetzt. Sie hat verschoben, welche Rolle der Engpass ist. Für das nächste Projekt: zuerst die Spec-Dokumente, `/ultrareview` ab Tag eins, ein Smoke-Target vor dem ersten Feature. Implementations-Durchsatz ist nicht mehr die Beschränkung — Design- und Review-Durchsatz sind es.

10. Abgrenzung: Was FlowHub nicht ist

- ☐ Kein allgemeiner Chatbot – FlowHub hat eine klar definierte Aufgabe
 - ☐ Kein Ersatz für paperless-ngx / Vikunja / Wallabag – FlowHub ergänzt, ersetzt nicht
 - ☐ Kein IFTTT/n8n Clone – FlowHub ist code-basiert, nicht no-code
 - ☐ Keine Multi-User-Plattform – Single-User Homelab-Tool
 - ☐ Keine Cloud-Abhängigkeiten – ausschliesslich Self-Hosted
-

11. Risiken und Mitigationen

Risiko	Wahrscheinlichkeit	Mitigation
Scope Creep (zu viele Skills)	Hoch	Klare MVP-Liste, Future klar abgegrenzt
Blazor SSR Lernkurve	Niedrig	Bekanntes C#-Ökosystem; gute Microsoft-Dokumentation
Ollama zu langsam (kein GPU)	Mittel	Claude API als Fallback; Keyword-Detection als Basis
Telegram API Änderungen	Niedrig	Abstraktion via ISkillHandler Interface
Vikunja API-Änderungen	Niedrig	Refit-Abstraktion, versionierte API-Clients
Zeitdruck (Abgabe Juli)	Mittel	MVP bewusst schlank, Future Features dokumentiert

12. Glossar

Begriff	Bedeutung
Skill	Handler für einen bestimmten Input-Typ (z.B. ArticleSkill)
SKILL.md	Konfigurationsdatei eines Skills (YAML Frontmatter + Markdown)
ISkillHandler	C# Interface mit der Business-Logic eines Skills
SkillRegistry	Lädt alle SKILL.md Files und zugehörige Handler
Skill-Vorschlag	Feature: System schlägt neuen Skill vor wenn kein bestehender passt
Pending Input	Input der auf Benutzer-Bestätigung wartet (in Redis)
Confidence	Wie sicher die KI bei der Skill-Erkennung ist (0.0–1.0)
Enrichment	KI-gestützte Anreicherung (z.B. Autor/AuthorInfo bei QuoteSkill)
Homelab	Selbst betriebene Server-Infrastruktur zu Hause (Proxmox)
Blazor SSR	Server-Side Rendering mit Blazor (.NET); HTML wird serverseitig gerendert
EF Core	Entity Framework Core – .NET ORM für Datenbankzugriffe
Refit	Typsichere REST-Client-Library für .NET (deklarativ via Interfaces)
M.E.AI	Microsoft.Extensions.AI – Abstraktion für LLM-Provider in .NET
Vikunja	Self-Hosted Task-Management / Kanban (ersetzt Todoist)
Wallabag	Self-Hosted Read-Later Service für Artikel
Knowledge Base	Obsidian Markdown Files, versioniert in GitLab (Self-Hosted)

Erstellt mit Unterstützung von Claude (Anthropic) – gemäss FFHS Richtlinien für KI-Einsatz in Projektarbeiten.

Architecture Decision Records

Source: `docs/adr/README.md`

Architecture Decision Records

Index of all ADRs for the FlowHub project.

ADR	Title	Status	Date
0001	Frontend Render Mode and Architecture	Accepted	2026-04-09
0002	Service Architecture and Async Communication	Accepted	2026-04-17
0003	Async Pipeline (MassTransit)	Accepted	2026-05-02
0004	AI Integration — Provider Abstraction	Accepted	2026-05-03
0005	Persistence — EF Core + PostgreSQL	Accepted	2026-05-04
0006	Vector Search — pgvector + Mistral Embeddings	Accepted	2026-05-07
0007	LLM Hosting	Accepted	2026-05-24
0008	Logging Policy	Accepted	2026-05-24
0009	Telemetry & PII Policy	Accepted	2026-05-24

Dates match each ADR's own header. ADRs 0001-0006 are inlined in the submission bundle; 0007-0009 live in this directory in the repository.

Why ADRs

FlowHub is a single-developer CAS thesis project that nonetheless commits to architectural rigor: decisions that shape the codebase for more than one block are captured as ADRs so the reasoning survives the next session, the next slice, and the grader's review. Each ADR records the **context** that forced the decision, the **decision** itself with its alternatives, and the **consequences** (positive and negative) that follow.

ADRs are referenced from three places:

- **Code** — inline comments at the affected seam (e.g. `// see ADR 0005 §"Repository pattern"`).
- **Specification** — `docs/spec/use-cases.md`, `docs/spec/nfa.md`, and `docs/spec/acceptance-criteria.md` cross-link to ADRs for the "why" behind a requirement's realization.
- **Per-block reflection** — each block's *Nachbereitung* in `vault/Blöcke/` cites the ADR(s) accepted during that block.

When to write a new ADR

A new ADR is created whenever a decision:

- Changes a public contract (driving port, REST endpoint, persistence schema).
- Selects a third-party dependency the project couples to (EF Core, MassTransit, MEAI, pgvector).
- Defines a policy that more than one feature must obey (e.g. the AI fallback contract in ADR 0004 §D5).

Smaller, reversible choices live in commit messages or inline comments — not every preference needs an ADR.

Format

ADRs follow the lightweight Nygard format adapted for this project: **Context** → **Decision** → **Consequences**, with explicit sections for **Alternatives considered** (and why they were rejected) and **Open questions** carried forward.

Status values: `Proposed` · `Accepted` · `Deprecated` · `Superseded by ADR-XXXX`

Stack-mismatch annotations

The CAS AISE coursework occasionally prescribes a Quarkus / Jakarta-EE-flavored realization (Hibernate / Panache, Jakarta Data, MicroProfile REST Client).

FlowHub's .NET 10 stack uses functional equivalents: EF Core 10 + `DbSet<T>` + LINQ, `IHttpClientFactory` + typed clients, MassTransit. ADRs that depart from the prescribed stack call this out explicitly in a "Stack mismatch — N/A justification" section, with the .NET-native counterpart documented alongside (precedent set in ADRs 0002, 0003, 0004, 0005).

ADR 0001 — Frontend Render Mode & Architecture

Source: `docs/adr/0001-frontend-render-mode-and-architecture.md`

ADR 0001 — Frontend Render Mode and Architecture

- **Status:** Accepted
 - **Date:** 2026-04-09
 - **Block:** Block 2 (Frontend) — Nachbereitung
 - **Decider:** freax
 - **Affects:** `source/FlowHub.Web/` , `source/FlowHub.Core/` , eventual `source/FlowHub.Api/`
-

Context

The Block 2 Moodle Auftrag (*Projektarbeit: Frontend*) requires us to:

- Decide on a frontend technology and rendering strategy (CSR vs. SSR vs. mix).
- Specify the frontend with wireframes and page-flow diagrams.
- Implement the frontend against services with **static / generated test data** (real persistence comes in Block 4).
- Generate AI-assisted unit tests with bUnit.

The repo's `CLAUDE.md` already locks in **Blazor + MudBlazor** as the UI stack. This ADR fills in the rest: which Blazor render mode, how the UI relates to the (future) REST API, how authentication works, and what high-level Pages we expect.

FlowHub is **operator-facing**: a single user (or small team) running it as a self-hosted service in a homelab. There is no public-facing surface, no SEO requirement, and no anonymous access. Captures arrive primarily from a Telegram bot today; the Web UI must also act as a Channel that can submit new Captures (see *Decision 4*).

The relevant glossary terms (`Capture` , `Channel` , `Skill` , `Integration` , `Page` , `Component` , `Card` , `Widget` , `Render Mode`) are defined in `Projektarbeit/Glossary.md` in the CAS Obsidian vault.

Decision

1. Render mode: Interactive Server (Blazor 8+ Server interactivity over SignalR) as the default

Every routable Page in `source/FlowHub.Web/` declares `@rendermode InteractiveServer` unless there's a documented reason to use a different mode. The Blazor component runs **in the ASP.NET Core process**, UI events flow over a long-lived SignalR circuit, and DOM diffs are pushed back to the browser.

Per-Page exceptions allowed but rare:

Page type	Render mode	Why
Login / OIDC callback	Static SSR	No interactivity, no client state, must work before circuit established
Health / status (<code>/_status</code>)	Static SSR	Plain text/HTML, scrapeable, no interactivity needed
Everything else	Interactive Server	Default — see below

WebAssembly / Auto are explicitly out of scope for Block 2. They can be revisited later if a specific component justifies the WASM runtime cost (e.g. an offline-capable Capture composer, a heavy in-browser editor). Adding them later is cheap because Blazor 8+ allows per-component render-mode declarations.

2. The Web UI does not consume FlowHub's REST API

Because Interactive Server components run in-process, the Web UI `@inject`s application services (`CaptureService` , `SkillRegistry` , `IntegrationHealth` , ...) and calls them **directly**. No HTTP, no JSON serialization, no fetch, no controller round-trip for the UI's own data.

The REST API (built in Block 3) exists for **non-UI consumers**:

- `FlowHub.Telegram` (bot, separate process)
- External integrations / automation
- CLI tools, future mobile clients
- Webhook receivers from upstream services

This means: the Block 3 *Services* assignment is **not** "build the API the frontend needs". It's "build the API the *other clients* need." That distinction shapes how we model the API's resource design.

3. Authentication: OIDC against the existing homelab Authentik instance

FlowHub is a new client of the user's existing **Authentik** SSO (already operated in the homelab — see `homelab/_uncategorized/Authentik.md`). **ASP.NET** Core's `Microsoft.AspNetCore.Authentication.OpenIdConnect` provider is wired in `Program.cs`; tokens land in a cookie that the SignalR circuit can read.

Rejected: **ASP.NET Core Identity (local user store)** — would create a second user database to keep in sync, contradicts the homelab "one IdP for all services" pattern, and adds Block 5 deployment complexity (secrets, password reset flows, MFA bootstrapping) that Authentik already solves.

The OIDC client registration in Authentik is a **runtime config concern**, not a code concern. Client ID, client secret, and issuer URL all come from environment variables (`Auth__OIDC__*`) per the 12-factor rule already in `CLAUDE.md`.

4. The Web UI is itself a Channel

Per the Glossary, a `Channel` is an inbound source of Captures. The Web UI is one of them. Concretely:

- A `WebChannel` exists alongside `TelegramChannel`.
- The dashboard has a **persistent quick-add field** (always-visible, in the `MudAppBar` or just below it) — paste a URL, type a quote, drop an image — *Enter* sends it to `CaptureService.Submit(new RawCapture(source: WebChannel, content: ...))`.
- A dedicated `/captures/new` Page exists for the longer-form variant: image upload, multi-line text, manual category override before classification.
- Both routes flow through the **same** `CaptureService.Submit(...)` entry point. The channel is metadata on the Capture, not a code path.

This satisfies the Moodle Lernziel *"Quarkus für Web-Formulare nutzen"* (translated to our stack: [ASP.NET](#) Core / Blazor for web forms) without inventing a contrived form just for the assignment.

5. Initial Page inventory (directional, refined in `/ui-brainstorm`)

#	Path	Purpose	Render mode
1	<code>/</code>	Dashboard — recent Captures, Skill health, Integration health, quick-add field	Interactive Server
2	<code>/captures</code>	Captures list — filter, search, sort, lifecycle filter (raw / classified / orphan / unhandled)	Interactive Server
3	<code>/captures/{id}</code>	Capture detail — content preview, classification, override category, retry routing	Interactive Server
4	<code>/captures/new</code>	Long-form capture submission (Web Channel)	Interactive Server
5	<code>/skills</code>	Skill list with status, recent activity, success/failure counts	Interactive Server
6	<code>/integrations</code>	Integration list with health, last successful write per integration	Interactive Server
7	<code>/auth/login</code> , <code>/auth/callback</code>	OIDC redirect endpoints (Authentik)	Static SSR

Not in scope for Block 2 Nachbereitung (parked for later blocks):

- Settings / preferences page
- Skill suggestion review queue (for unhandled Captures)
- Audit log viewer
- Multi-user / RBAC

Each Page above gets its own `/ui-brainstorm` → `/ui-flow` → `/ui-build` → `/ui-review` cycle in the order listed. Pages 1, 4, and 2 are the **MVP path** that should

be implementable end-to-end against Faker test data within the 26 h Block 2 budget; pages 3, 5, 6 are stretch.

6. Test data: Bogus (the .NET equivalent of Faker)

Stub services in `source/FlowHub.Web/Stubs/` return Bogus-generated Captures, Skills, and Integrations. The same stubs are used by bUnit tests so component behaviour can be asserted against deterministic seeded data. In Block 3 the stubs are replaced with real `CaptureService` calls; in Block 4 the in-memory store is replaced with EF Core + PostgreSQL. The component code does not change — only the DI registration.

Alternatives Considered

A. Blazor WebAssembly (CSR) only

- ☐ Forces a REST API to exist *for the UI's sake*, which inverts the design intent (the API should serve external clients, not the UI).
- ☐ First-load WASM runtime download (~1–3 MB) on a single-user dashboard is a poor trade.
- ☐ Loses in-process service injection — every read becomes a fetch + serialize + deserialize round-trip for a UI that has zero scaling pressure.
- ☐ Live Capture stream from the server requires SSE/WebSockets anyway → effectively re-implements what SignalR gives you for free.

B. Static SSR + JavaScript sprinkles (htmx-style, or Blazor Static SSR with form posts)

- ☐ Captures/dashboard are inherently *live* — the operator wants new Captures to appear without refreshing. Static SSR forces full page reloads or polling.
- ☐ The Moodle Lernziele explicitly call out applying CSR *and* SSR concepts. Going pure-Static loses half the conceptual coverage.
- ☐ Loses the "no REST API for UI" benefit because every form post still goes through HTTP routing.

C. Interactive Auto (Server-then-WASM swap)

- ☐ Reasonable but premature. Adds complexity (component runs in two contexts, must be tested in both, no in-process service injection on the WASM side) for a single-user dashboard where the WASM upgrade has no measurable benefit.
- ☐ Can be adopted later per-component if a specific case justifies it (e.g. an offline draft composer). Per-component render modes are a Blazor 8+ feature.

D. **ASP.NET** Core Identity (local user store)

- ☐ Second user database to keep in sync with the homelab IdP.
- ☐ Contradicts the established homelab "one IdP, many services" pattern.
- ☐ Adds Block 5 deployment work (password reset flows, MFA, account lockout, secret rotation) that Authentik already solves.

E. No auth at all (network-level isolation only)

- ☐ Even in a homelab, FlowHub will eventually be reachable via the Cloudflare tunnel / reverse proxy. App-level auth is the only honest answer.
- ☐ Fails the Block 5 *Deployment* expectations and the *Validierung* grading dimension.

Consequences

Positive

- **No HTTP layer for UI data flow.** Less code, less serialization, less to test, better diagnostics (server logs see the full call stack including the Blazor event handler).
- **Live Capture stream is free** — when a new Capture lands (from any Channel), the server already knows about it and can push a UI patch over the open SignalR circuit. No polling, no SSE, no extra plumbing.
- **REST API stays focused** on its real consumers (Telegram, integrations, automation) — better resource design, clearer versioning story.
- **Auth is solved at the platform layer** — Authentik handles MFA, password reset, account lifecycle.

- **Block 3 (Services) is conceptually cleaner** — the API exists for external clients, not as scaffolding for the UI's own data needs.

Negative

- **Always-on connection requirement.** A dropped SignalR circuit means lost UI state. Mitigation: use `PersistentComponentState` for anything that must survive a reconnect; show a clear reconnect indicator.
- **Server-side state per active user.** Fine for single-operator FlowHub; would need rethinking at hundreds of concurrent users.
- **Blazor Server has a wider attack surface during the SignalR circuit** — the server holds component state that's tied to the user's session. Mitigation: standard [ASP.NET](#) Core authorization on every event handler.
- **WASM offline mode is not on the table** without revisiting this decision.
- **bUnit tests run components in isolation**, but won't catch SignalR-specific issues like state-after-reconnect. E2E tests in Block 5 will need to cover that.

Neutral

- The UI's render mode is explicit per Page — adding a WASM component later is a single attribute change plus the WASM project wiring.
- The decision is reversible at the per-Page level. If `/captures/new` later needs offline drafting, it can become Interactive Auto without affecting the rest of the app.

Implementation Notes (for Block 2 Nachbereitung)

The implementation work this ADR unblocks:

1. **Scaffold** `source/FlowHub.Web/` — empty today except `.gitkeep`. Add Blazor App template (Interactive Server), wire MudBlazor, wire OIDC against Authentik (env-var driven), add `MudLayout` with `MudAppBar` quick-add field placeholder.
2. **Define stub services** in `source/FlowHub.Web/Stubs/` returning Bogus-generated Captures/Skills/Integrations. These implement interfaces that will live in `FlowHub.Core` so Block 3 can swap real implementations in.

3. **Per-page work** via the mandatory UI workflow: `/ui-brainstorm` → `/ui-flow` → `/ui-build` → `/ui-review`, in the MVP order: Dashboard → New Capture → Captures list. Stretch: Capture detail → Skills → Integrations.
4. **Component tests** via bUnit against the Bogus stubs.
5. **No auth wiring against a real Authentik client** in Block 2 — a placeholder OIDC config + a dev-mode bypass is enough; the real client registration happens in Block 5 (Deployment) when the service is exposed.
 - **Dev bypass implementation:** a `DevAuthHandler` : `AuthenticationHandler<AuthenticationSchemeOptions>` returning a fixed `ClaimsPrincipal` ("Dev Operator", configurable roles via `Dev__Auth__Roles` env var). Registered **only** when `builder.Environment.IsDevelopment()` — production code path goes through the OIDC scheme. **Never** use `[AllowAnonymous]` sprinkles to bypass auth in dev — the real auth pipeline must run in dev so `[Authorize]`, `User.IsInRole(...)`, and `AuthorizationStateProvider` behave identically to production. bUnit tests use bUnit's own `AuthorizationContext` test helper and do not depend on `Program.cs` at all.

Each numbered item above will be its own short brainstorm → plan → implementation cycle. This ADR is the umbrella decision; it does not replace the per-feature planning.

References

- Moodle: *W4B-C-AS001.AISE.ZH-Sa-1.PVA.FS26: Projektarbeit: Frontend* — assignment text
- `CLAUDE.md` — repo conventions (Blazor + MudBlazor lock-in, hexagonal architecture, MudBlazor component preferences)
- `.ai/cas-instructions.md` — block schedule and grading dimensions
- Vault: `Projektarbeit/Idee FlowHub.md` — original concept and management summary
- Vault: `Projektarbeit/Glossary.md` — Capture, Channel, Skill, Integration, Render Mode definitions
- Vault: `Projektarbeit/Skills.md` — running list of Skill ideas and example Captures per category
- Vault: `homelab/_uncategorized/Authentik.md` — existing SSO infrastructure

- Microsoft Learn: [Blazor render modes \(.NET 8+\)](#)

ADR 0002 — Service Architecture & Async Communication

Source: `docs/adr/0002-service-architecture-and-async-communication.md`

ADR 0002 — Service Architecture and Async Communication

- **Status:** Accepted
- **Date:** 2026-04-17
- **Block:** Block 3 (Services) — Vorbereitung / early Nachbereitung
- **Decider:** freax
- **Affects:** `source/FlowHub.Core/`, `source/FlowHub.Api/` (to be scaffolded), `source/FlowHub.Skills/`, `source/FlowHub.Integrations/`, `source/FlowHub.AI/`, `source/FlowHub.Web/`, future `infra/docker-compose.*.yaml`

Context

The Block 3 Moodle Auftrag (*W4B-C-AS001.AISE.ZH-Sa-1.PVA.FS26: RESTful API and Leseauftrag: Microservices- und Service-based-Architekturen*) requires us to:

- Demonstrate understanding of microservice / service-based architectures.
- Build a service that exposes a synchronous REST API.
- Demonstrate at least one form of asynchronous, event-based communication between services.
- Reflect on the trade-offs of synchronous vs asynchronous styles.

ADR 0001 already established that:

- FlowHub is a single-operator, self-hosted application running in a homelab.
- The Web UI runs **in-process** (Blazor Interactive Server) and does **not** consume FlowHub's REST API — the API exists for *other* clients (Telegram, integrations, automation).
- The repository follows a flat `source/FlowHub.<Capability>/` layout (not `src/Modules/<Module>/`).

What ADR 0001 did **not** decide is how the capabilities (`Core` , `AI` , `Skills` , `Integrations` , `Telegram` , `Persistence`) talk to each other once they are filled in, and whether they live in one process or several. That is the question Block 3 forces us to answer.

The Block 3 Vorbereitung built a sandbox (`poc/restful-api-playground/`) with `OrderService` + `NotificationService` + RabbitMQ + gRPC to internalise the trade-offs. The reflection from that sandbox (`poc/restful-api-playground/REFLECTION.md`) concluded:

- **MQ** earns its place where producers and consumers must be decoupled, where the producer's response time should not depend on the consumer, where bursts must be absorbed, and where retry / DLQ are needed for flaky downstream calls.
- **gRPC** earns its place only when there is a real cross-service call between services we own and we want a typed, streaming-capable, low-latency contract.
- For a single-language, single-operator monolith calling external REST APIs, gRPC adds tooling without solving a problem we have.

This ADR codifies how those conclusions land in FlowHub's actual code layout.

Decision

1. Keep the Modular Monolith — do not split into physical microservices for Block 3

FlowHub stays a single deployable process. Each capability remains its own .NET project under `source/FlowHub.<Capability>/` :

- `FlowHub.Core` — domain types, driving ports, domain services
- `FlowHub.AI` — classification adapter
- `FlowHub.Skills` — skill registry + skill implementations
- `FlowHub.Integrations` — outbound adapters to Wallabag, Vikunja, ...
- `FlowHub.Persistence` — EF Core (Block 4)
- `FlowHub.Telegram` — inbound channel adapter
- `FlowHub.Api` — REST API surface for non-UI consumers (new, scaffolded in this block)

- `FlowHub.Web` — Blazor host (already exists)

Cross-capability communication is via **interfaces declared in `FlowHub.Core`**, implemented by the appropriate capability project, and resolved through the .NET DI container. No project references between sibling capabilities; the only allowed inbound dependency is on `FlowHub.Core`.

As built (Block 5 update). The capability list above is the Block-2 plan; the shipped solution (`FlowHub.slnx`) contains **six** projects — `FlowHub.Core`, `FlowHub.AI`, `FlowHub.Skills`, `FlowHub.Persistence`, `FlowHub.Api`, `FlowHub.Web`. Two refinements: (1) the Wallabag and Vikunja outbound adapters landed in `FlowHub.Skills` as `ISkillIntegration` implementations — the separately sketched `FlowHub.Integrations` project was not needed and is **not** part of the source tree. (2) `FlowHub.Telegram` is likewise **not yet scaffolded** — the implemented capture channels are the Web Quick-Capture field and the REST API. Both are planned, intentionally absent until built, and documented as such in `docs/spec/system-context.md` → "Current state (Block 5)".

This satisfies the Moodle learning objectives at the **logical** level: clear service boundaries, ports and adapters, dependency inversion, and a documented reason for each boundary. It does **not** introduce the operational complexity (separate deployments, network calls, distributed tracing across processes, eventual consistency between data stores) that a physical microservice split would require.

2. Introduce an async message bus for the Capture-Enrichment and Skill-Routing pipelines

The flow is:

1. A `Capture` arrives from any Channel (Web, Telegram, future API).
2. `CaptureService.Submit(...)` persists the Capture **synchronously** and returns. The caller (UI, Telegram, API client) sees a `Created` result with the Capture id.
3. `CaptureService` then **publishes** a `CaptureCreated` event onto an in-process bus.
4. Subscribers run independently:
 - `AI.ClassificationHandler` — calls the AI adapter, writes the classification back, publishes `CaptureClassified`.
 - `LinkPreviewHandler` — fetches an OpenGraph preview for URL captures.

- (future) `TagSuggestionHandler`, `DuplicateDetectionHandler`, ...
- 5. `Skills.RoutingHandler` subscribes to `CaptureClassified` and dispatches to the matching `Skill`. Each `Skill` then calls the relevant `Integration` (Wallabag, Vikunja). Failure handling (retry, DLQ) lives at this layer.

Implementation:

- Use **MassTransit** as the abstraction so the same producer / consumer code works against multiple transports.
- **Default transport in dev and test:** MassTransit's **in-memory transport**. No broker dependency for `just run` / `just watch` / `dotnet test`.
- **Production / staging transport: RabbitMQ**, configured via `Bus__Transport=RabbitMq` and `Bus__RabbitMq__Host=...`. Wired up in Block 5 (Deployment) alongside Docker Compose.
- The transport choice is a `Program.cs` registration concern. Producer and consumer code is transport-agnostic.

This gives Block 3 a clean answer to *"demonstrate asynchronous, event-based communication"* without forcing every developer to run RabbitMQ locally.

3. Outbound calls to external Skill integrations stay synchronous REST

Wallabag, Vikunja, etc. expose REST APIs. Skill implementations call them synchronously inside an MQ consumer. The consumer is the resilience boundary:

- HTTP failures from external integrations are caught at the consumer and trigger MassTransit's retry policy.
- After retries are exhausted the message moves to a dead-letter queue (RabbitMQ in prod, the in-memory transport simulates this via a fault consumer).
- Operator-visible failures surface on the Skill detail page (`/skills`) and the Capture detail page (`/captures/{id}`).

The fact that the *external* call is synchronous is fine — the *FlowHub-internal* event flow upstream of it is async, so the user-facing POST does not wait for Wallabag to respond.

4. Build a REST API in `source/FlowHub.Api/` for non-UI consumers

Per ADR 0001, the API exists for Telegram, integrations, automation, CLI, and webhook receivers — not for the Blazor UI. This block scaffolds it:

- Hosted in the same process as `FlowHub.Web` for now (separate [ASP.NET](#) Core endpoint groups, shared DI container).
- Minimal API endpoints, FluentValidation at the boundary, ProblemDetails (RFC 9457) for errors, Scalar UI at `/scalar`.
- Initial endpoints (refined during implementation):
 - `POST /api/captures` — submit a Capture from a non-UI channel.
 - `GET /api/captures` / `GET /api/captures/{id}` — read, with filters mirroring the UI list page.
 - `POST /api/captures/{id}/retry` — re-publish a `CaptureCreated` event for a failed Capture.
 - `GET /api/skills` / `GET /api/integrations` — read-only health and metadata.

The API is deliberately **separate** from any internal interfaces. Internal contracts (the methods on `CaptureService` etc.) can change freely; the public REST contract is versioned and breaks visibly.

5. Do not introduce gRPC in FlowHub for Block 3

Reasons (carried over from `poc/restful-api-playground/REFLECTION.md`):

- FlowHub is a single-language (.NET / C#) stack. The polyglot benefit of gRPC's code-generated multi-language clients does not apply.
- Without a physical service split (Decision 1), there are no FlowHub-internal RPC calls to make. In-process method calls already give us typed, low-latency communication for free.
- External integrations (Wallabag, Vikunja) speak REST. The integration boundary is fixed by them.
- Adding gRPC just to satisfy the Moodle reading list would be architecture-for-its-own-sake. The Block 3 reflection covers gRPC in writing instead, with the playground POC as the worked example.

6. Versioning and contract stability

- **Internal events** (`CaptureCreated` , `CaptureClassified` , `SkillRouted` , ...) live in `FlowHub.Core` and may evolve freely while everything is in one process. Once the bus crosses a process boundary (i.e. when / if we ever split), additive-only changes become mandatory and a `v2` event type is introduced for breaking changes.
 - **Public REST API contracts** are versioned via URL prefix (`/api/v1/...`) from day one. A `v2` prefix is added when a breaking change ships.
 - **Internal interfaces** (`ICaptureRepository` , `ISkill` , ...) carry no versioning — change them with the call sites, in one PR.
-

Alternatives Considered

A. Split FlowHub into physical microservices for Block 3

- ☐ Operational cost (multiple deployments, separate databases or schema partitioning, distributed tracing setup, inter-service auth) is wildly out of scale for a single-operator homelab app.
- ☐ ADR 0001 already established that the Web UI runs in-process and reads via direct service calls. Splitting would force the UI onto HTTP for its own data, contradicting that decision.
- ☐ The PVA grading dimension *"Realistische Projektarbeit"* would suffer — a contrived multi-service split for one operator is not realistic.
- ☐ The decision is reversible: a capability that genuinely outgrows the monolith can be lifted out into its own process later, with the bus already in place to absorb the change.

B. Use a custom in-process event aggregator (e.g. MediatR notifications) instead of MassTransit

- ⚠ Works for the in-memory case but locks us in. Swapping to a real broker later would mean rewriting every publisher and consumer.
- ⚠ No retry / DLQ semantics out of the box — we would re-implement what MassTransit gives us.
- ☐ Simpler dependency surface today.

- → Rejected because the Block 3 Lernziele explicitly cover async / event-based architectures, and MassTransit makes the broker upgrade path a config change rather than a code change. The cost of MassTransit in dev (one extra package, no broker required thanks to the in-memory transport) is acceptable.

C. Use gRPC for FlowHub-internal communication

- ☐ See Decision 5. No service split → no internal RPC calls to make.
- ☐ Single-language stack removes the polyglot benefit.

D. Expose the FlowHub REST API by also pointing the Blazor UI at it

- ☐ Directly contradicts ADR 0001 Decision 2.
- ☐ Adds HTTP serialization for every UI read with no benefit (single operator, no scaling pressure).

E. RabbitMQ as the only transport (no in-memory option)

- ☐ Forces every developer to run Docker for `just watch` and every CI run to spin up a broker container.
- ☐ Slows the test suite.
- → Rejected. MassTransit's in-memory transport is purpose-built for this case.

Consequences

Positive

- **One deployment unit** — `just run` still starts everything. CI stays simple. Docker Compose in Block 5 has one app container plus its dependencies (DB, broker), not seven.
- **Clear logical boundaries** with the option to split later — ports in `Core`, adapters in capability projects, no sibling project references.
- **Async pipeline answers the Moodle Lernziel** for event-based / asynchronous communication with a worked example (`CaptureCreated` → enrichment fan-out → routing).

- **Resilience to flaky external integrations** — MQ retry / DLQ wraps every outbound Skill call.
- **UI stays responsive** — `POST /api/captures` and the Web Channel quick-add return immediately; enrichment happens asynchronously.
- **Transport flexibility** — switch dev / prod transports via env var. No code changes in producers / consumers.

Negative

- **Eventual consistency in the enrichment chain.** The caller cannot tell from the `Created` response whether classification succeeded. Mitigation: surface enrichment status on the Capture detail page; expose a `lifecycle` field that progresses `Raw → Classified → Routed → Completed | Failed`.
- **In-memory transport in dev hides certain bugs** that only show up against a real broker (serialization, connection drops, queue backlog). Mitigation: a dedicated `make integration-test` target spins up RabbitMQ and runs a small set of end-to-end tests against it. Block 5 makes RabbitMQ part of the staging deployment.
- **MassTransit is a non-trivial dependency** (additional NuGet packages, learning curve, opinionated patterns). Mitigation: keep its surface area limited to publishers, consumers, and `Program.cs` registration; do not adopt MassTransit's higher-level features (sagas, state machines) until there is a concrete need.
- **The "microservices" framing in the Moodle Auftrag is answered with a logical split, not a physical one.** This is a defensible position but must be explained clearly in the PVA — write-up risk if it reads as ducking the assignment. Mitigation: this ADR is the explanation, plus a short reflection in the Block 3 Nachbereitung notes.

Neutral

- The decision to skip gRPC is documented and reversible. If a later block introduces a real cross-process call between FlowHub-owned services, gRPC is the obvious tool — the playground POC shows we know how to wire it up.
- The REST API layer is now a thin facade over the same services the UI uses. Any feature added to the UI is implicitly available to the API once an endpoint is added, and vice versa.

Implementation Notes (for Block 3 Nachbereitung)

The implementation work this ADR unblocks:

1. **Scaffold** `source/FlowHub.Api/` — empty placeholder today. Add Minimal API endpoint groups (`captures` , `skills` , `integrations`), `FluentValidation`, `ProblemDetails`, `Scalar UI`, and the `/api/v1/...` URL prefix.
2. **Scaffold** `source/FlowHub.Skills/` **and** `source/FlowHub.Integrations/` — both are placeholders. Define `ISkill` and `IIntegration` ports in `FlowHub.Core` , then move at least one implementation (Wallabag is the easiest external service to mock) into the Skills / Integrations projects.
3. **Wire MassTransit** in `source/FlowHub.Web/Program.cs` (the host project for Block 3).
 - Default `Bus__Transport=InMemory` in `appsettings.Development.json` .
 - `Program.cs` reads the env var and selects the transport at startup.
4. **Define internal events** in `source/FlowHub.Core/Events/` : `CaptureCreated` , `CaptureClassified` , `CaptureRouted` , `CaptureCompleted` , `CaptureFailed` .
5. **Replace the Bogus stubs incrementally** — `CaptureService` already exists as a stub; route its `Submit(...)` through the bus instead of returning a fake immediately. `AI.ClassificationHandler` becomes the first real consumer.
6. **Add Bruno collections** under `bruno/api/` for the new REST endpoints, mirroring the URL structure.
7. **Add an integration test** under `tests/FlowHub.Web.ComponentTests/` (or a new `tests/FlowHub.Api.IntegrationTests/`) that publishes `CaptureCreated` against the in-memory transport and asserts the consumer pipeline runs end-to-end.
8. **Document the boundary in the PVA write-up** — Decision 1 is the load-bearing claim and needs a short explicit paragraph in the Block 3 Nachbereitung notes.

Each numbered item above will be its own short brainstorm → plan → implementation cycle. This ADR is the umbrella decision; it does not replace per-feature planning.

References

- Moodle: *W4B-C-AS001.AISE.ZH-Sa-1.PVA.FS26: RESTful API* — assignment text
- Moodle: *Leseauftrag: Microservices- und Service-based-Architekturen* — reading list
- ADR 0001: `docs/adr/0001-frontend-render-mode-and-architecture.md`
- POC: `poc/restful-api-playground/` — RabbitMQ + gRPC playground
- POC reflection: `poc/restful-api-playground/REFLECTION.md`
- `CLAUDE.md` — repo conventions (modular monolith, no cross-module project references, MudBlazor, ProblemDetails, Scalar)
- `.ai/cas-instructions.md` — block schedule and grading dimensions
- Vault: `Blöcke/03 Service/03 Service - a) Vorbereitung.md` — Block 3 prep checklist and Leitfragen
- MassTransit docs: [In-Memory transport](#) and [RabbitMQ transport](#)

ADR 0003 — Async Pipeline

Source: `docs/adr/0003-async-pipeline.md`

ADR 0003 — Async Pipeline: MassTransit Topology, Retry, and Fault Handling

- **Status:** Accepted
- **Date:** 2026-05-02
- **Block:** Block 3 (Services) — Nachbereitung · Slice B
- **Decider:** freax
- **Affects:** `source/FlowHub.Core/Events/` , `source/FlowHub.Core/Classification/` , `source/FlowHub.Core/Skills/` , `source/FlowHub.Web/Pipeline/` , `source/FlowHub.Web/Program.cs` , `docker-compose.yml` , `tests/FlowHub.Web.ComponentTests/Pipeline/`

Context

ADR 0002 established the overall architecture: FlowHub stays a Modular Monolith, MassTransit is the message bus abstraction (in-memory in dev/test, RabbitMQ in production), and the `CaptureCreated → enrichment → routing` flow is the worked example for the Block 3 requirement to *"demonstrate asynchronous, event-based communication"*. What ADR 0002 deliberately deferred was the concrete topology: which specific flows go through the bus, which events exist, how retry and dead-letter handling behave, how faults surface back to the lifecycle state machine, and how the deployment story satisfies the rubric's *"Sub-Systeme als unabhängige Container deploybar"* dimension (5 pts).

This ADR closes those open questions. It covers the implementation decisions made during Block 3 Nachbereitung Slice B (the working MassTransit pipeline, first consumers, and tests). It is a narrower, more focused record than ADR 0002 — the brainstorming narrative lives in `docs/superpowers/specs/2026-04-30-async-pipeline-design.md` ; this ADR distils the durably decided items from decisions D1 through D13 of that spec.

The Block 3 Moodle Auftrag explicitly names *"asynchrone Kommunikation mittels einer Queue"* as the primary deliverable. The rubric's *KI / Sub-Systeme / Reflexion* bucket additionally awards 5 pts for *"Sub-Systeme als unabhängige Container deploybar"*. Both are addressed without splitting the codebase.

Decision

1. Async surface = capture-submit and routing (flows 1 + 2 only)

The pipeline covers exactly two flows: (1) capture submit triggers an enrichment consumer via `CaptureCreated`, and (2) enrichment success triggers a routing consumer via `CaptureClassified`. No other flows are made asynchronous in Slice B.

This is narrower than the four-flow table sketched in earlier drafts. Flow 3 (a dedicated routing event separate from enrichment output) was dropped because it adds an unnecessary extra hop without giving the consumer more information — routing happens inline at the tail of the enrichment consumer, which already has the classification result. Flow 4 (manual retry as an async re-publish) is handled by a synchronous REST endpoint that re-publishes `CaptureCreated` for a given Capture id; it does not need its own event type or its own flow designation. The result is a simple two-event, two-consumer pipeline that is easy to test exhaustively and easy to explain in the PVA write-up.

As built (Block 5 update). Slice B is the two-consumer core described here. Later slices added two more consumers around the same two events plus a fault observer, so `FlowHub.Web/Pipeline/` now registers **five** message handlers: `CaptureEnrichmentConsumer`, `SkillRoutingConsumer` (Slice B), plus `CaptureEmbeddingConsumer` (vector search, ADR 0006), `CaptureNotificationConsumer` (optional ntfy.sh), and `LifecycleFaultObserver` (maps `Fault<T>` to lifecycle state). The two-event vocabulary is unchanged.

2. Event vocabulary: `CaptureCreated` and `CaptureClassified` only

The final event set for Slice B is exactly two records:

```
// source/FlowHub.Core/Events/CaptureCreated.cs
public sealed record CaptureCreated(
    Guid CaptureId,
    string Content,
    ChannelKind Source,
    DateTimeOffset CreatedAt);

// source/FlowHub.Core/Events/CaptureClassified.cs
public sealed record CaptureClassified(
    Guid CaptureId,
    IReadOnlyList<string> Tags,
    string MatchedSkill,
    DateTimeOffset ClassifiedAt);
```

Three candidate events from earlier drafts are dropped: `SkillRoutingRequested` (routing happens inline, not via a separate event), `IntegrationCallFailed` (failures surface as `Fault<CaptureClassified>`, which is handled by the fault observer), and `SkillRouted` (the terminal success state is recorded by the routing consumer directly on the Capture entity — no downstream subscriber needs to observe it in Slice B). ADR 0002 §6 notes that events may evolve freely while all consumers are in-process; additive changes become mandatory only when the bus crosses a process boundary in a future block.

3. Classification port: `IClassifier` with `KeywordClassifier` stub in Slice B

Classification is abstracted behind a port declared in `source/FlowHub.Core/Classification/`:

```
public interface IClassifier
{
    Task<ClassificationResult> ClassifyAsync(string content, CancellationTo
}

public sealed record ClassificationResult(
    IReadOnlyList<string> Tags,
    string MatchedSkill);
```

The Slice B implementation is `KeywordClassifier`: URL-pattern content → `Tags=["link"]`, `MatchedSkill="Wallabag"`; content containing `"todo"` or `"task"` (case-

insensitive) → `Tags=["task"]`, `MatchedSkill="Vikunja"`; everything else → `Tags=["unsorted"]`, `MatchedSkill=""` (handled per Decision 6). Slice C swaps in `AiClassifier : IClassifier` without touching consumer code — the hexagonal port shape isolates the substitution entirely. This also earns the *"Code lesbar, nach Layer, Modulen und Sub-Systemen strukturiert"* rubric dimension (max 7 pts).

4. Skill integration port: `ISkillIntegration` with `LoggingSkillIntegration` stubs

Outbound integration calls are abstracted behind a second port in `source/FlowHub.Core/Skills/`:

```
public interface ISkillIntegration
{
    string Name { get; }
    Task WriteAsync(Capture capture, IReadOnlyList<string> tags, Cancellati
}
```

Slice B registers two `LoggingSkillIntegration` instances (one named `"Wallabag"`, one named `"Vikunja"`). Each stub logs `"would write to {Name} for capture {CaptureId}"` and returns successfully. The routing consumer resolves the correct integration via `IEnumerable<ISkillIntegration>` and selects by `Name` — if no match is found, it marks the Capture as `Unhandled` directly (Decision 6). Real adapters (Wallabag, Vikunja HTTP clients) replace these stubs in Block 4/5 without any consumer changes.

5. Per-consumer retry policy: enrichment Intervals(100, 500), routing Intervals(500, 2000, 5000)

Retry intervals are set per consumer rather than via a shared global policy:

```
x.AddConsumer<CaptureEnrichmentConsumer>(c =>
    c.UseMessageRetry(r => r.Intervals(100, 500)));

x.AddConsumer<SkillRoutingConsumer>(c =>
    c.UseMessageRetry(r => r.Intervals(500, 2000, 5000)));
```

The enrichment consumer calls `IClassifier`, which in Slice B is an in-process keyword match. Two retries at 100 ms and 500 ms is enough budget for a transient fault without introducing meaningful latency. The routing consumer calls `ISkillIntegration`, which in Block 4/5 will be an outbound HTTP call to a real external service. Three retries at 500 ms, 2 s, and 5 s reflects a realistic cross-process latency profile and gives a slow upstream service room to recover. Separating the policies avoids over-waiting for the cheap operation and under-waiting for the expensive one.

6. Empty classification → direct Orphan; fault observer maps `Fault<T>` to lifecycle state

Two entry points into `Orphan` and two entry points into `Unhandled` are made explicit:

Trigger	Path	Final stage
<code>KeywordClassifier</code> returns <code>MatchedSkill=""</code>	Enrichment consumer sets stage directly	<code>Orphan</code>
Enrichment consumer throws past retry budget	<code>Fault<CaptureCreated></code> → <code>LifecycleFaultObserver</code>	<code>Orphan</code>
No <code>ISkillIntegration</code> registered for matched skill	Routing consumer sets stage directly	<code>Unhandled</code>
Routing consumer throws past retry budget	<code>Fault<CaptureClassified></code> → <code>LifecycleFaultObserver</code>	<code>Unhandled</code>

`LifecycleFaultObserver` is a single class implementing both `IConsumer<Fault<CaptureCreated>>` and `IConsumer<Fault<CaptureClassified>>`. It maps each fault to the appropriate terminal stage and logs at `LogLevel.Error` with `CaptureId`, originating message id, and `ExceptionInfo` from the `Fault.Exceptions` collection. The observer is best-effort: if it throws, the message rots in the bus's secondary error queue and the Capture's `LifecycleStage` is unchanged — operator must reconcile manually. No recursive retry.

Without this observer, captures that exhaust their retry budget would stay visually at `Raw` or `Classified` on the Dashboard while silently rotting in the error queue. The observer makes the pipeline's resilience story visible in the operator UI.

Empty-classification lands in `Orphan` without going through the fault path because it is a successful outcome of the enrichment consumer (no exception is thrown; the classifier returned a valid but empty result). Mixing the "no skill match" case into the fault path would mean the enrichment consumer deliberately throws to signal a business condition, which is an anti-pattern.

7. `KebabCaseEndpointNameFormatter` for stable queue names

```
x.SetKebabCaseEndpointNameFormatter();
```

With this formatter, `CaptureEnrichmentConsumer` becomes the endpoint name `capture-enrichment`, and `SkillRoutingConsumer` becomes `skill-routing`. These names are stable whether the transport is in-memory (Slice B) or RabbitMQ (Block 5). Without an explicit formatter MassTransit derives the queue name from the full class name, which can differ between .NET versions and is harder to write operator alerts against. Kebab-case also matches the existing REST URL convention (`/api/captures` , `/api/skills`).

8. `IBus` factory for singleton consumers that need `IPublishEndpoint`

`CaptureServiceStub` (the in-memory Capture repository from Block 2) publishes `CaptureCreated` via `IPublishEndpoint`. Because `CaptureServiceStub` is registered as a singleton and `IPublishEndpoint` is Scoped in MassTransit's DI model, a direct injection would produce a captive dependency (singleton holding a scoped service). The workaround for Slice B is to inject `IBus` instead:

```
// IBus is Singleton in MassTransit's DI model – safe to hold in a singleton
public CaptureServiceStub(IBus bus, ILogger<CaptureServiceStub> logger)
```

This resolves the captive-dependency problem because `IBus` is itself singleton-lifetime in MassTransit. The trade-off is a slightly heavier abstraction (`IBus` vs `IPublishEndpoint`). Block 4 will revisit this when the EF Core implementation of `ICaptureService` becomes Scoped, at which point `IPublishEndpoint` can be injected directly and `IBus` is no longer needed here.

9. One process in code; multi-container topology documented in `docker-compose.yml`

FlowHub stays a single deployable process per ADR 0002 Decision 1. `just run` still starts everything with no broker dependency. However, a `docker-compose.yml` sketch is committed alongside the code that demonstrates the multi-container topology for Block 5:

```
services:
  flowhub.web:    image: flowhub-web:dev      (Bus__Transport: RabbitMq)
  flowhub.api:    image: flowhub-api:dev      (Bus__Transport: RabbitMq)
  rabbitmq:       image: rabbitmq:3-management-alpine
```

This compose file is **not** invoked by `just run` (which remains single-process via `dotnet run`). Its purpose is to satisfy the Bewertungskriterien dimension "*Sub-Systeme als unabhängige Container deploybar*" (max 5 pts) without introducing the operational complexity of a physical process split during development. Block 5 fleshes out the full deployment story with real image builds and an override file.

10. Tests in `tests/FlowHub.Web.ComponentTests/Pipeline/`

Pipeline tests live alongside the existing component tests rather than in a dedicated test project. The MassTransit Test Harness (`ITestHarness`) integrates cleanly with xUnit's `IAsyncLifetime` pattern; the existing project already references xUnit, FluentAssertions, and NSubstitute. Splitting into a separate `FlowHub.Pipeline.Tests` project would add overhead (new `.csproj`, CI step, solution entry) for no real isolation benefit at this stage. The split can happen later if the suite grows past ~20 tests or if CI build times make parallelism worthwhile.

Retry intervals in tests are tightened to `10ms` so the suite stays under one second per case:

```
x.AddConsumer<CaptureEnrichmentConsumer>(c =>
    c.UseMessageRetry(r => r.Intervals(10, 10)));

x.AddConsumer<SkillRoutingConsumer>(c =>
    c.UseMessageRetry(r => r.Intervals(10, 10, 10)));
```

Alternatives Considered

A. Make all four flows asynchronous

The earlier spec table sketched four flows: submit, enrichment, routing, and manual retry. Routing as a separate third event (after `CaptureClassified` rather than inline with it) would require a `SkillRoutingRequested` event and a third consumer. Manual retry as an event would require a `CaptureRetryRequested` type.

Rejected because neither third nor fourth flow gains anything from the event hop. Routing already has the classification result in `CaptureClassified` — introducing `SkillRoutingRequested` just to forward it is noise. Manual retry is a single synchronous endpoint that re-publishes `CaptureCreated`; wrapping it in an event doesn't improve decoupling. Two flows is the minimum that meaningfully demonstrates the async pipeline story.

B. Only the enrichment flow async (no routing consumer)

An even thinner scope would wire only `CaptureCreated` → `CaptureEnrichmentConsumer` and call the integration synchronously inside the enrichment consumer.

Rejected because it removes the two-hop resilience story (retry on classification is different from retry on integration) and collapses the state machine to three states instead of five. The *Bewertungskriterien* rubric rewards *"intelligente und flexible Services"* (max 6 pts) — a single-consumer pipeline with synchronous integration writes doesn't demonstrate that. The routing consumer and its per-consumer retry policy are the part of the implementation that earns those points.

C. Global retry policy shared across all consumers

MassTransit supports a global retry policy set at the bus configuration level:

```
x.AddBus(provider => Bus.Factory.CreateUsingInMemory(cfg =>
    cfg.UseMessageRetry(r => r.Interval(3, 500))));
```

Rejected because a global policy is a lie. The enrichment consumer calls in-process code and benefits from fast retries (100 ms, 500 ms). The routing consumer calls external services in Block 4/5 and needs longer back-off (500 ms, 2 s, 5 s). A shared policy would either over-wait for the cheap operation or under-wait for the expensive one. Per-consumer configuration forces the policy to match the actual cost profile of each consumer — it is also the honest answer to the reviewer question "why these intervals".

D. Fault observer omitted; rely on the `_error` queue alone

MassTransit's in-memory transport and RabbitMQ both move exhausted-retry messages to an error queue automatically. It would be possible to skip `LifecycleFaultObserver` entirely and let the operator inspect the error queue directly.

Rejected because silent failures are invisible to the operator UI. Without the observer, a capture that fails all retries shows `LifecycleStage=Raw` on the Dashboard indefinitely — the operator has no idea it is stuck. The observer costs one class and two `IConsumer<Fault<T>>` implementations; the benefit is that the Dashboard and API `?stage=Orphan` / `?stage=Unhandled` filters actually reflect reality. The `△` states (`Orphan` , `Unhandled`) exist precisely to surface these failure cases.

E. Outbox pattern now (alongside the in-memory implementation)

An outbox pattern guarantees that database writes and message publishes happen atomically — the canonical solution to "what if the process crashes between writing the Capture and publishing `CaptureCreated`". MassTransit ships an `EntityFrameworkOutbox` that hooks into EF Core's `SaveChangesAsync`.

□ Deferred to Block 4. There is no EF Core persistence in Slice B — all state is in-memory (Block 2 stub). An outbox without a durable store is meaningless; the prerequisite for correctness (persistence) is the prerequisite for the outbox. Revisiting this in Block 4, once `FlowHub.Persistence` lands, is the right sequence. Noted in Block 3 *Nachbereitung* as an explicit open item.

Consequences

Rubric coverage

This ADR and its implementation directly address four Bewertungskriterien dimensions:

- **Entwurf: Lösungsansatz und Architektur beschrieben** (max 7 pts) — the ASCII architecture diagram in the spec, plus the capture-lifecycle state machine and the pipeline sequence diagrams in `docs/design/perspectives.md` and `docs/design/sequences/`, cover both bildlich and textuell.
- **Programmierung: Code lesbar, nach Layer, Modulen und Sub-Systemen strukturiert** (max 7 pts) — `IClassifier` + `ISkillIntegration` ports in `FlowHub.Core`, consumers in `FlowHub.Web/Pipeline/`, transport config isolated to `Program.cs`.
- **KI / Sub-Systeme als unabhängige Container deploybar** (max 5 pts) — `docker-compose.yml` demonstrates the multi-container topology; Decision 9 explains why the code stays single-process during development.
- **KI / Intelligente und flexible Services** (max 6 pts) — the pipeline itself (async enrichment, classification port, fault observer, retry policy) is the worked example of an intelligent, failure-tolerant service.

The Quarkus / Jakarta EE criterion (max 10 pts) remains N/A — see ADR 0002 and the submission PDF note.

Complexity tax

MassTransit adds a meaningful learning curve. The three areas where it bites soonest:

1. **Captive dependency** (Decision 8). MassTransit's DI lifetime model has `IBus` as singleton and `IPublishEndpoint` as scoped. Any singleton service that needs to publish must either inject `IBus` or use a factory. This is not obvious from the docs and requires a comment in the code to avoid future confusion.
2. **EventId namespacing convention**. Serilog event ids follow project-local ranges, scoped by where the `LoggerMessage` source-gen partial method physically lives:
 - **1000–1999 — Pipeline** (consumers + fault observer in `source/FlowHub.Web/Pipeline/`). Slice B uses `1001`

```
( CaptureEnrichmentConsumer.LogOrphan ), 1002
( SkillRoutingConsumer.LogUnhandled ), 1003
( LifecycleFaultObserver.LogObserverFailed ).
```

- **2000–2999 — Skills** (adapters in `source/FlowHub.Skills/`). Slice B uses 2001 (`LoggingSkillIntegration.LogStubWrite`).
- Higher ranges are unallocated; new modules pick a free range when they land.

Not enforced by tooling — followed by discipline and reflected here + in commit messages.

3. **Test harness retry intervals.** Tests use `10ms` retry intervals to keep the suite fast. If a test is flaky, the first thing to check is whether the retry budget is being exhausted faster than the assertions can observe the state transition. The `ITestHarness` `await harness.Consumed<T>()` helper handles most of this, but it is worth knowing.

In-memory transport caveats

The in-memory transport is purpose-built for dev and test: no broker dependency, fast startup, simple setup. Its limitation is durability — if the process crashes between `ICaptureService.SubmitAsync` publishing `CaptureCreated` and the enrichment consumer consuming it, the message is lost. The Capture exists in the in-memory store (it was written first) but will never be enriched; its `LifecycleStage` stays at `Raw` permanently.

This is acceptable for Slice B because:

- The in-memory store itself is ephemeral. A process restart resets all state anyway, so lost messages are no worse than lost captures.
- The outbox pattern (Decision 5 / alternative E) is the correct fix, and it depends on Block-4 persistence.

Operators running the Block-5 RabbitMQ configuration do not have this problem — RabbitMQ persists messages across process restarts.

RabbitMQ ops note: secondary error queue requires an operator playbook

When Block 5 lands RabbitMQ, MassTransit routes exhausted-retry messages to `<consumer>_error`. If `LifecycleFaultObserver` itself throws on one of those

messages, MassTransit routes the fault-of-the-fault to `<consumer>_error_error` — the secondary error queue. There is no automatic reconciliation: messages that land there rot indefinitely, and the corresponding Captures show stale `LifecycleStage` values in the UI.

⚠ Block 5 must include a Prometheus alert on the depth of `capture-enrichment_error_error` and `skill-routing_error_error`. A queue depth above 0 for more than five minutes should page the operator. Without this alert, the failure mode is invisible until a user notices a capture stuck at `Raw`. The operator playbook for manual reconciliation (re-publish `CaptureCreated` via `POST /api/captures/{id}/retry`) must be documented alongside the Grafana dashboard.

Consequences for the next blocks

Block 4: Persistence + Outbox

- **Outbox pattern** — once `FlowHub.Persistence` lands with EF Core, wire `MassTransit.EntityFrameworkCoreOutbox` so that `SaveChanges` and `Publish(CaptureCreated)` are atomic. This eliminates the lost-message window described above.

As built (Block 5 — residual risk, still open). The `EntityFrameworkOutbox` was **not** wired, even though RabbitMQ is now the live transport. The lost-message window therefore still exists: if the process crashes between persisting a Capture and publishing `CaptureCreated`, that Capture stays at `Raw` and is never enriched. Current mitigation: the synchronous **manual retry** endpoint re-publishes `CaptureCreated` for a given Capture id, and the dashboard surfaces stuck captures. Wiring the outbox is the correct fix and is carried as a known, documented residual risk rather than an asserted feature.

- **Revisit `IBus` factory pattern** — with a Scoped `ICaptureService` implementation backed by EF Core, `IPublishEndpoint` can be injected directly (both are Scoped); the `IBus` workaround in Decision 8 can be removed.
- **Idempotency receiver** — RabbitMQ delivers at-least-once. Once persistence exists, add a `MessageDataRepository`-backed idempotency filter so that redelivered messages are no-ops rather than double-writes. The

consumers SHOULD (but currently don't) look up by `CaptureId` before mutating state — this is a correctness gap noted here for Block 4.

- **FailureReason sanitization** — `MarkOrphanAsync` and `MarkUnhandledAsync` accept a `reason` string that in the fault observer comes directly from `ExceptionInfo.Message`. Before persisting this string in a production DB, truncate it to a safe length (e.g. 500 chars) and strip any content that might contain sensitive operator-internal information (stack paths, connection strings). The in-memory stub has no persistence so the risk is low in Slice B.

Block 5: RabbitMQ deployment + observability

- **RabbitMQ deployment** — `docker-compose.yml` (committed in Slice B) becomes the runtime default. Add `docker-compose.override.yml` for local secret injection. The `Bus__Transport=RabbitMq` / `Bus__RabbitMq__Host=rabbitmq` env vars are already wired in `Program.cs` from Decision 9.
- **Docker Compose topology** — `flowhub.web` + `flowhub.api` + `rabbitmq`. The two app containers share the same bus because they share the same RabbitMQ instance. Whether `flowhub.api` also publishes or only consumes is a Block-5 decision.

As built (Block 5 update). The separate `flowhub.api` container above was **not** built. `FlowHub.Api` shipped as an in-process **class library** composed into the `flowhub.web` host (it registers the REST endpoints), so the running topology is `flowhub.web` + `flowhub.migrations` (init-job) + backing services (`postgres`, `rabbitmq`, `prometheus`, `grafana`). Only `flowhub-web` is built and pushed on release. This is consistent with the Modular-Monolith decision (ADR 0002): the API is a logical sub-system, not a separately deployed process.

- **OIDC** — the dev-only `DevAuthHandler` is replaced with a real OIDC provider. This is orthogonal to the pipeline but shares the same Block-5 milestone.
- **Observability** — Prometheus alert on `capture-enrichment_error_error` and `skill-routing_error_error` queue depths (see RabbitMQ ops note above). Grafana board tracking `LifecycleStage` distribution over time (how many captures reach `Routed` vs stay at `Orphan` / `Unhandled`). The OpenTelemetry integration in `FlowHub.Web/Program.cs` already exports traces and metrics — MassTransit's `UseOpenTelemetry()` extension makes pipeline spans appear automatically.

References

- Brainstorming spec: `docs/superpowers/specs/2026-04-30-async-pipeline-design.md` — full D1–D13 decision table and component sketches
- Implementation plan: `docs/superpowers/plans/2026-04-30-async-pipeline.md` — task-by-task execution record
- ADR 0001: `docs/adr/0001-frontend-render-mode-and-architecture.md` — render mode, Blazor Interactive Server
- ADR 0002: `docs/adr/0002-service-architecture-and-async-communication.md` — Modular Monolith, MassTransit selection, transport strategy
- API surface sketch: `docs/design/api/api-surface.md` — REST endpoint conventions referenced by event contract shape
- Block 3 Nachbereitung: `vault/Blöcke/03 Service/03 Service - c) Nachbereitung.md` — block checklist and Bewertungskriterien subset
- Bewertungskriterien: `vault/0rganisation/Bewertungskriterien.md` — canonical Moodle rubric (18 items, max 100 pts)
- MassTransit in-memory transport:
<https://masstransit.io/documentation/configuration/transports/in-memory>
- MassTransit RabbitMQ transport:
<https://masstransit.io/documentation/configuration/transports/rabbitmq>
- MassTransit Test Harness:
<https://masstransit.io/documentation/concepts/testing>

ADR 0004 — AI Integration in Services

Source: `docs/adr/0004-ai-integration-in-services.md`

ADR 0004 — AI Integration in Services: Provider, Abstraction, Prompt + Cost Strategy

- **Status:** Accepted
- **Date:** 2026-05-03
- **Block:** Block 3 (Services) — Nachbereitung · Slice C
- **Decider:** freax
- **Affects:** `source/FlowHub.AI/` ,
`source/FlowHub.Core/Classification/ClassificationResult.cs` ,
`source/FlowHub.Web/Program.cs` , `Directory.Packages.props` ,
`tests/FlowHub.AI.IntegrationTests/`

Context

ADR 0003 §3 pre-committed a hexagonal `IClassifier` port and shipped a deterministic `KeywordClassifier` for Slice B. Both ADR 0003 and the Block 3-c Nachbereitung checklist explicitly name "Slice C swaps in an AI-backed adapter" as the next step. The Block 3 Moodle Auftrag demands "intelligente und flexible Services" using AI; the Bewertungskriterien dimension "*Intelligente / flexible Services mit KI gebaut*" (max 6 pts) is the explicit deliverable, and the highest-weighted single rubric item ("*Wurden KI-Werkzeuge verwendet und deren Nutzung beschrieben*", max 12 pts) demands **production-runtime** AI use, not just AI as a coding assistant.

This ADR records the durable decisions for that adapter. The brainstorming narrative lives in `docs/superpowers/specs/2026-05-03-slice-c-ai-integration-design.md` ; this ADR distils D1-D10 from that spec.

The Quarkus / Jakarta-EE programming criterion remains **N/A** for FlowHub's .NET stack — same precedent as ADR 0002 / 0003. The course's nominal Spring-AI /

Koog vocabulary is the teacher's reference stack; FlowHub's .NET-native equivalents (`Microsoft.Extensions.AI` , optional Semantic Kernel) are presented in their own terms.

Decision

1. Slice scope = classifier + AI-generated title (one round-trip)

The AI does two jobs in a single structured-output call: classify (`Tags` , `MatchedSkill`) and produce a short title (`Title`). No summary, no skill-suggestion queue, no embeddings, no agent loop. Estimated effort 2 days.

2. Two interchangeable adapters: Anthropic native + OpenRouter aggregator

Both adapters implement `Microsoft.Extensions.AI.IChatClient` . The operator picks one at boot via `Ai__Provider=Anthropic|OpenRouter` . The two transports are deliberately different in shape:

- **Anthropic** uses the native Anthropic API via the `Anthropic.SDK` NuGet package's MEAI-compatible `IChatClient` implementation. Vendor-specific features available (prompt caching via `cache_control: ephemeral`).
- **OpenRouter** uses `Microsoft.Extensions.AI.OpenAI` against `https://openrouter.ai/api/v1` . OpenAI-compatible, so the same package reaches hundreds of upstream models (Anthropic, OpenAI, Google, Meta Llama, Mistral, Qwen, ...) behind one adapter shape.

This earns the "*flexible*" in "intelligente und flexible Services" — one vendor-native SDK, one aggregator gateway, one shared interface.

3. Abstraction = `Microsoft.Extensions.AI` (MEAI), not Semantic Kernel

The Block 3-c Nachbereitung demands "`Microsoft.Extensions.AI` oder Semantic Kernel als Abstraktion einbinden". MEAI is the right shape for FlowHub's single-shot classifier-with-typed-output use case: one interface (`IChatClient`), schema-driven structured output (`CompleteAsync<TResponse>` / `GetResponseAsync`),

decorator chaining (`UseOpenTelemetry` , custom delegates) via `ChatClientBuilder` .
Lightweight; no kernel/plugin/planner machinery.

Semantic Kernel is reserved for a hypothetical Block-5 agent loop (e.g. an "intelligent retry advisor" that decides between re-route / summarise-and-surface / escalate when an integration call fails repeatedly) and is explicitly NOT introduced in Slice C. SK consumes `IChatClient` natively, so adding it later is additive, not a replacement.

4. Port shape = extend `ClassificationResult` with `Title`?

`IClassifier` stays one method. The result record gains an optional third field:

```
public sealed record ClassificationResult(
    IReadOnlyList<string> Tags,
    string MatchedSkill,
    string? Title = null);
```

`KeywordClassifier` returns `Title=null` . `AiClassifier` returns both populated, or falls back to keyword on AI failure (in which case `Title=null` again). Single round-trip is meaningfully cheaper than two ports (`IClassifier` + `IEnricher`), and the model returns both fields naturally inside one schema. Block 4/5 can split the port if a real use case demands it.

5. Failure behaviour = graceful fallback to `KeywordClassifier`

`AiClassifier` wraps `KeywordClassifier` as a hard floor. The fallback triggers on either:

- **Any exception** thrown by `IChatClient.GetResponseAsync<...>` — network, timeout, parse, schema, anything else (catch `Exception` , since MEAI's exact taxonomy varies between provider adapters).
- **Defensive post-validation:** even when the call returns successfully, `AiClassifier` checks `MatchedSkill ∈ {"Wallabag", "Vikunja", ""}` . If not, treat as a failure (the schema should have prevented this — never trust the model).

In either case the adapter:

1. Logs `Warning ( EventId 3010 AiClassifierFellBackToKeyword )` with the failure reason (exception type or `"schema_violation"`) and elapsed milliseconds.
2. Calls `_keyword.ClassifyAsync(content, ct)` and returns its result with `Title=null` .
3. Does NOT rethrow.

Consequences: capture is always classified — AI outage degrades quality, never availability. The MassTransit retry budget from ADR 0003 §5 stays reserved for genuine pipeline / bus failures; the fault-observer path (ADR 0003 §6) is not entered by AI errors.

6. Default models: Anthropic Haiku 4.5 / OpenRouter Llama 3.1 70B Instruct

Provider	Model	Default env	Rationale
Anthropic	<code>claude-haiku-4-5-20251001</code>	<code>Ai__Anthropic__Model</code>	Cheapest current Claude; strong JSON-schema adherence; tool-use bridge for MEAI.
OpenRouter	<code>meta-llama/llama-3.1-70b-instruct</code>	<code>Ai__OpenRouter__Model</code>	Open-weights; 70B is reliable on schema (smaller variants flake). Sets up the "commercial vs open-weights" rubric narrative.

Per-call tokens: ~200 input / ~150 output → cost is negligible regardless of provider; narrative dominates over cost. Both env vars overridable for snapshot bumps.

7. Structured output = `CompleteAsync<T>` with JSON schema generated from the DTO

```
internal sealed record AiClassificationResponse(  
    [property: Description("1–5 short lowercase tags describing the snippet  
    [property: JsonPropertyName("tags")]  
    string[] Tags,  
  
    [property: Description("Wallabag, Vikunja, or empty string for none")]  
    [property: AllowedValues("Wallabag", "Vikunja", "")]  
    [property: JsonPropertyName("matched_skill")]  
    string MatchedSkill,  
  
    [property: Description("3–8 word title or null if content is too short"  
    [property: JsonPropertyName("title")]  
    string? Title);
```

MEAI generates the JSON schema from the record. `AllowedValues` translates to a JSON-schema `enum`; both Anthropic (tool-use under the hood) and OpenRouter (`response_format: json_schema`) honour it. `MatchedSkill` MUST match a registered `ISkillIntegration.Name` (`"Wallabag"`, `"Vikunja"`, or `""` → Orphan) or routing breaks downstream — defensive runtime check inside `AiClassifier` re-validates even though the schema enforces it.

8. No-key startup = silent fallback to `KeywordClassifier`

`AddFlowHubAi(IConfiguration)` inspects `Ai__Provider` and the matching `ApiKey` and registers `KeywordClassifier` as the active `IClassifier` whenever either is missing. Invalid `Ai__Provider` values throw `InvalidOperationException` at startup. `just run` works zero-config for a fresh `git clone` — same dev-friendly philosophy as `DevAuthHandler` for auth.

Boot logs: `EventId 3020 AiProviderRegistered` (Information) on the AI path, `EventId 3021 AiProviderNotConfigured` (Information) with reason on the keyword path. Emitted via a small `IHostedService` (`AiBootLogger`) so the extension stays a pure DI helper.

9. Testing layered: mocked unit tests + trait-gated live integration tests

- 10 mocked unit tests
(`tests/FlowHub.Web.ComponentTests/Ai/AiClassifierTests.cs`) via `NSubstitute<IChatClient>` and `NSubstitute<IClassifier>` — zero real API calls. Cover happy path, argument forwarding (CT, MaxOutputTokens, Temperature), the 5 fallback paths (`HttpRequestException`, `TaskCanceledException`, `JsonException`, schema violation, generic `Exception`), and the 3010 logging contract.
- 8 mocked unit tests for the D8 registration matrix
(`AiServiceCollectionExtensionsTests.cs`).
- 4 trait-gated live integration tests (`tests/FlowHub.AI.IntegrationTests/` , `[Trait("Category", "AI")]`) that talk to real providers — excluded from default `just test` , run via `just test-ai` when the operator has keys.

`justfile` filters: `just test` → `--filter "Category!=AI"` . CI runs `just test` only; `just test-ai` is operator-on-demand.

10. Active provider selection = explicit `Ai__Provider` env var

One env var swaps providers. No implicit precedence; no per-request runtime selection; no round-robin. Demoing "swap providers via config" is a one-line env change.

Consequences

Rubric coverage

This ADR + its implementation directly address six Bewertungskriterien dimensions:

- **Entwurf: Lösungsansatz und Architektur beschrieben** (max 7) — ASCII component diagram + data-flow diagrams in the spec.
- **Programmierung: Code lesbar, nach Layer / Modulen strukturiert** (max 7) — `FlowHub.AI` becomes a real project with single responsibility; `IClassifier` port stays in Core; adapters separated from prompt + DTO.

- **Programmierung: Erkenntnisse aus der Programmierung dokumentiert** (max 3) — this ADR + the Slice-C section in `docs/ai-usage.md`.
- **Validierung: Unit-Tests programmiert** (max 3) — 18 mocked unit tests + 4 trait-gated live integration tests.
- **KI: Wurden KI-Werkzeuge verwendet und deren Nutzung beschrieben** (max 12) □ — Slices A/B already covered KI-as-development-tool; this slice adds *production-runtime* AI use, which is what the rubric actually asks.
- **KI: Intelligente / flexible Services mit KI gebaut** (max 6) — explicit deliverable. Two adapters demonstrate flexibility; graceful keyword fallback demonstrates "intelligent" failure handling.

The Quarkus / Jakarta-EE criterion (max 10) remains N/A for the .NET stack.

EventId namespacing — 3000–3999 reserved for AI

Extends the ADR 0003 §EventId convention:

- 1000–1999 — Pipeline (Slice B)
- 2000–2999 — Skills (Slice B)
- 3000–3999 — AI (this slice)
 - 3001 `AiClassifierStarted` (Debug — Slice C optional, may be omitted)
 - 3002 `AiClassifierSucceeded` (Debug — same)
 - 3010 `AiClassifierFellBackToKeyword` (Warning, on each fallback)
 - 3020 `AiProviderRegistered` (Information, at startup)
 - 3021 `AiProviderNotConfigured` (Information, at startup)

`LoggerMessage` source-gen used throughout (CA1848 / CA1873 enforced by `Directory.Build.props`).

Cost guards

`ChatOptions.MaxOutputTokens=300` (~2× the schema's natural ~150 output tokens), `Temperature=0.2` (deterministic-ish), `HttpClient` 10s timeout.

Anthropic prompt cache: deferred to Slice D — `Anthropic.SDK 5.10.0`'s `IChatClient` bridge does not surface the `cache_control: ephemeral` marker (see `BuildChatClient` TODO in `AIServiceCollectionExtensions.cs`). The ~80% input-token discount on the system-prompt segment is forfeit until either (a) the bridge

surfaces a `ChatOptions.AdditionalProperties` key for caching, or (b) `AiClassifier` is rewritten to use the native `MessageParameters` API. Low priority given the system prompt is ~150 tokens; per-call delta is sub-cent at Haiku 4.5 pricing. OpenRouter prompt cache is not universally supported across upstream models — Slice C does not claim it. Asymmetry documented here as a real difference between adapters.

OpenTelemetry

`Microsoft.Extensions.AI.UseOpenTelemetry()` decorator emits `gen_ai.client.operation.duration` and token-count metrics on the existing OTEL pipeline. Will surface in the Block-5 Grafana dashboard alongside MassTransit traces (`AI calls / sec`, `p95 latency`, `fallback rate`).

Reflexion — `Microsoft.Extensions.AI` vs Semantic Kernel

MEAI = thin abstraction over chat / embedding / tool-use; sister of `Microsoft.Extensions.Logging` / `Microsoft.Extensions.Caching` . One interface (`IChatClient`), schema-driven structured output, decorator chaining. Right pick when the call site is "give the model a string, get a typed object back".

Semantic Kernel = full agent framework: `Kernel` , plugins, planners, memory, agent runtime. Right pick when an LLM **orchestrates** multi-step work.

Decision rule for FlowHub:

- Single-shot LLM call with typed output → MEAI. Slice C is exactly this.
- Multi-step LLM-driven workflow → add Semantic Kernel without removing MEAI.

A credible Block-5 fit for SK: an *intelligent retry advisor* that, on repeated integration failures, inspects the failure pattern + capture content + integration health and decides between re-route / summarise-and-surface / escalate. Slice C deliberately stops short.

Consequences for next blocks

Block 4 (Persistence):

- `IEmbeddingGenerator<string, Embedding<float>>` already exposed by MEAI → embedding work for Block 5 KI-search (pgvector via Npgsql) is incremental, not greenfield.
- AI audit fields on the persisted `Capture`: (`provider`, `model`, `duration_ms`, `was_fallback`) per classification. Earns the *"Test-Ergebnisse dokumentiert"* rubric item with real production data.

Block 5 (Deployment + KI-search):

- ADR 0006 (KI-Suche) builds on the `IEmbeddingGenerator` shape. Anthropic doesn't ship embeddings → the embedding provider becomes either OpenAI native, OpenRouter passing through, or a self-hosted sentence-transformer behind an OpenAI-compatible facade. Asymmetry to document there.
- Integration-test secret rotation: `Ai__Anthropic__ApiKey` / `Ai__OpenRouter__ApiKey` move from User Secrets into the Block-5 deployment secret store (Authentik or Docker secrets).
- OTEL: `gen_ai.*` metrics export already; Grafana panel for "AI calls / sec, p95 latency, fallback rate" comes nearly free.
- Optional Semantic-Kernel adoption — additive on top of MEAI.

References

- Brainstorming spec: `docs/superpowers/specs/2026-05-03-slice-c-ai-integration-design.md`
- Implementation plan: `docs/superpowers/plans/2026-05-03-slice-c-ai-integration.md`
- ADR 0001: `docs/adr/0001-frontend-render-mode-and-architecture.md`
- ADR 0002: `docs/adr/0002-service-architecture-and-async-communication.md` — module split
- ADR 0003: `docs/adr/0003-async-pipeline.md` — `IClassifier` port, `EventId` namespacing
- AI Usage living doc: `docs/ai-usage.md`
- Block 3 Nachbereitung: `vault/Blöcke/03 Service/03 Service - c) Nachbereitung.md`
- Bewertungskriterien: `vault/Organisation/Bewertungskriterien.md`
- `Microsoft.Extensions.AI`: <https://learn.microsoft.com/en-us/dotnet/ai/microsoft-extensions-ai>

- OpenRouter API reference: <https://openrouter.ai/docs>

ADR 0005 — Persistence

Source: `docs/adr/0005-persistence.md`

ADR 0005 — Persistence: ORM, Provider, Repository Pattern, Migrations Workflow

- **Status:** Accepted
- **Date:** 2026-05-04
- **Block:** Block 4 (Persistence) — landed early via Beta MVP slice
- **Decider:** freax
- **Affects:** `source/FlowHub.Persistence/` , `source/FlowHub.Web/Program.cs` , `tests/FlowHub.Persistence.Tests/` , `tests/FlowHub.Api.IntegrationTests/IntegrationTestFactory.cs` , `Directory.Packages.props` , `.config/dotnet-tools.json`


Context

Block 2 shipped `CaptureServiceStub` — a Bogus-backed in-memory `ICaptureService` implementation that survives only the lifetime of one `just run` . Block 4's Moodle Auftrag asks for "ein geeignetes Datenmodell entwerfen und über ORM-Abstraktion realisieren ... mit dynamischen Abfragen" — i.e., real persistence, with an ORM, backing the same driving port the UI already consumes.

The Beta MVP slice (`docs/superpowers/specs/2026-05-04-beta-mvp-design.md` , `docs/superpowers/plans/2026-05-04-beta-mvp.md`) landed a vertical end-to-end demo *before* Block 4 formally begins (2026-05-09). The Beta slice's primary purpose was to validate the architecture against the real homelab Wallabag/Vikunja, not to deliver the full Block-4 scope. That shaped several decisions in this ADR — they make sense for the Beta-scope demo and are explicit about which evolutions Block 4 will re-litigate.

The Quarkus / Hibernate / Panache / Jakarta-Data programming criterion remains **N/A** for FlowHub's .NET stack — same precedent as ADRs 0002, 0003, 0004. The course's nominal Hibernate / Panache / Jakarta Data vocabulary is the teacher's reference stack; FlowHub's .NET-native equivalents (EF Core 10, `DbSet<T>` + LINQ, `IQueryable<T>` + Expression Trees) are presented in their own terms and the mapping is recorded in `vault/Blöcke/04 Persistence/04 Persistence - c)` `Nachbereitung.md`'s intro.

Architecture at a glance



Legende

Notation	Meaning
Rectangle (UI , APP , PORT , REPO)	In-process .NET component / layer inside <code>FlowHub.*</code> projects.
Cylinder (CTX , PG)	Data-bearing element — <code>DbContext</code> (in-memory tracking graph) or the physical PostgreSQL database.
Hexagon (MIG)	Out-of-band lifecycle process — separate container, runs before the app, not part of request flow.
Solid arrow -->	Runtime call on the request path; control + data flow.
Dashed arrow -. ->	Compile-time / lifecycle relationship — "implemented by" or "applies before". Not on the per-request hot path.
Arrow label	Wire format or protocol crossing the boundary (domain types , LINQ → SQL , Npgsql).

The boundary that matters for testability is `APP → PORT` : tests substitute the EF adapter with an in-memory or Testcontainers-backed implementation without touching the application services.

Decision

1. ORM = `Microsoft.EntityFrameworkCore` 10.0.7

Single ORM, .NET-native, first-party. No Dapper (we don't need raw-SQL micro-ORM performance for FlowHub's ingest-and-classify load profile), no NHibernate (not .NET-native enough), no LINQ-to-DB (smaller community, fewer providers). EF Core 10 is the latest stable, ships with .NET 10 SDK, and matches the toolchain the rest of the codebase already pins (`global.json` 10.0.x).

The hexagonal seam stays unchanged: `ICaptureService` (driving port, in `FlowHub.Core`) is the only contract the Web/API layers know about; `EfCaptureService` (driven adapter, in `FlowHub.Persistence`) is the implementation. EF Core's `DbContext` is **not** exposed beyond the persistence assembly.

2. Provider = SQLite for the Beta, PostgreSQL for Block 4 production

The Beta MVP runs on SQLite (`Microsoft.EntityFrameworkCore.Sqlite` 10.0.7). Connection-string fallback is `Data Source=flowhub.db` (relative to working dir, which is `source/FlowHub.Web/` under `just run`). This is a deliberate scope decision:

- Zero infrastructure setup — no Docker container, no `just db-up` , no port conflicts. The user can `just run` and have a working app immediately.
- The EF Core API surface is provider-agnostic for the Beta's query shapes (no Postgres-specific operators, no `JSONB` columns, no full-text search yet). All 13 unit tests in `FlowHub.Persistence.Tests` run against `Microsoft.EntityFrameworkCore.InMemory` , which makes provider-specific behaviour invisible at the test layer — that's intentional for the Beta and will need to evolve in Block 4 (see §10).
- Block 4 swaps to PostgreSQL via `UseNpgsql` + `ConnectionStrings:Default` env var, with Testcontainers PostgreSQL replacing `InMemory` for repository tests (per the spec's "Testcontainers PostgreSQL bevorzugt vor SQLite-In-Memory wegen Provider-Quirks" guidance). Migration source code regenerates from `OnModelCreating` against the new provider — schema is identical except for column-type discriminators that EF handles transparently.

The asymmetry — SQLite-in-prod (Beta) vs PostgreSQL-in-prod (Block 4) — is documented here so it doesn't feel like a regression when Block 4 changes the default.

3. No Repository pattern layer; `EfCaptureService` is the `ICaptureService` adapter

As built (Block 4/5 — decision reversed). This Slice-1 stance was **not** carried through. The implemented persistence layer **does** use the repository pattern: six repository ports live in `FlowHub.Core` (`ICaptureRepository`, `IChannelRepository`, `ISkillRepository`, `IIntegrationRepository`, `ITagRepository`, `ISkillRunRepository`) with six `Ef*Repository` implementations in `FlowHub.Persistence/Repositories/`. `EfCaptureService` composes `ICaptureRepository` rather than touching `FlowHubDbContext` directly. The six-entity model that landed in Block 4 made per-aggregate repositories worth the ceremony (focused querying, cleaner testability, no `DbContext` leakage into services). Treat the reasoning below as the original argument, superseded by this note.

The Block 4 Moodle Auftrag mentions "Repositories" as one of the standard ORM patterns. FlowHub deliberately does NOT introduce a separate repository interface between `ICaptureService` and `DbContext`. Reasons:

- `DbSet<T>` already *is* a repository in EF Core's design (a `Queryable<T>` over one table with `Add`, `Remove`, `Update` semantics). Wrapping it in an `ICaptureRepository` would be a level of ceremony with no testability or abstraction value the existing `ICaptureService` doesn't already provide.
- The existing `ICaptureService` *is* the repository abstraction in a hexagonal sense: it's the driving port the application code consumes; whether the implementation talks to EF Core or to Bogus or to some future read-model is transparent to callers.
- The spec's "Repository-Pattern-Entscheid" rubric checkbox is fulfilled by this decision being recorded — repository-as-separate-layer is a common but not-mandatory pattern; choosing not to use it is a defensible architectural call.

If a future driven adapter needs the `DbContext` directly (e.g. a read-model projection or a complex multi-aggregate transaction that doesn't fit the `ICaptureService` shape), it lives in `FlowHub.Persistence` and gets its own driving port. We're not introducing a generic `IRepository<T>` — that's abstraction without addition.

4. Migrations workflow = dotnet-ef tool manifest + in-process MigrationRunner (Beta) → separate init container (Block 5)

The repo pins `dotnet-ef` via `.config/dotnet-tools.json` (created during the Beta slice). This means:

- `dotnet tool restore` after a fresh clone gives every developer the same EF version locally.
- `dotnet ef migrations add <Name> --project source/FlowHub.Persistence --startup-project source/FlowHub.Web --output-dir Migrations` is the canonical add-migration command.
- A design-time `IDesignTimeDbContextFactory<FlowHubDbContext>` (in `source/FlowHub.Persistence/FlowHubDbContextFactory.cs`, internal sealed) lets the EF tooling discover the context without booting the full host.

For migration *application*, the Beta uses a `MigrationRunner` `IHostedService` that calls `db.Database.MigrateAsync()` at app startup (EventIds 5010/5011). This is **convenient for just run** but is explicitly out-of-line with **12-Factor XII** ("Run admin/management tasks as one-off processes"). Block 5 will:

- Move migration application to a separate init container in the production Compose / K8s manifests.
- Generate migration SQL bundles (`dotnet ef migrations script`) for ops review before applying.
- Keep `MigrationRunner` in `Development` only (or remove entirely — TBD with the Block 5 deployment spec).

The Beta's in-process migration runner is acceptable scope debt: the demo path needs migrations to exist on first run, and the deployment-grade alternative requires Compose orchestration that isn't built yet.

5. Entity visibility = internal sealed + InternalsVisibleTo

`CaptureEntity` (and any future entity) lives at `internal sealed` visibility in `source/FlowHub.Persistence/Entities/`. `FlowHubDbContext.Captures` is `internal DbSet<CaptureEntity>`. Tests reach the entity via `[assembly: InternalsVisibleTo(...)]` in `source/FlowHub.Persistence/Properties/AssemblyInfo.cs` — currently exposing internals to:

- `FlowHub.Persistence.Tests` — for direct entity instantiation in unit-test seeding and `db.Captures.AddRange(...)` setup.
- `FlowHub.Api.IntegrationTests` — for `IntegrationTestFactory` deterministic seeding (10 rows: 2 Orphan, 2 Completed, 2 Raw, 2 Classified, 1 Routed, 1 Unhandled). Seeding through `ICaptureService.SubmitAsync` would publish `CaptureCreated` events into the test's MassTransit harness, causing non-deterministic state at test start.

The internal-visibility surface area is now two test assemblies. This is acceptable trade-off for keeping the entity off the public surface (preserves `ICaptureService` as the only API surface). Block 4's domain expansion (`Skill`, `SkillRun`, `Channel`, `Integration`, `IntegrationHealthSample`, `Tag`) follows the same pattern.

6. Cursor pagination = keyset on (`CreatedAt DESC`, `Id DESC`) with `limit+1` probe

`EfCaptureService.ListAsync` implements offset-free pagination:

```
WHERE c.CreatedAt < cursor.CreatedAt
      OR (c.CreatedAt == cursor.CreatedAt AND c.Id < cursor.Id)
ORDER BY c.CreatedAt DESC, c.Id DESC
LIMIT @limit + 1
```

The `+1` probe row determines whether a `next` cursor should be issued. The `(CreatedAt, Id)` composite handles ties on identical timestamps (which Bogus seeding can produce). Cursor format is hand-rolled URL-safe base64 of JSON `(CreatedAt, Id)` tuples — the same format `Slice A` introduced for the REST API. No EF-Core-specific feature is used; the same pattern works against PostgreSQL unchanged in Block 4.

`Limit` is clamped server-side to `[1, 200]` (`Math.Clamp`); the spec calls out 200 as the max.

7. Indexes = `IX_Captures_Stage` (filtered list, "Needs Attention") + `IX_Captures_CreatedAt_DESC` (Recent Captures + pagination)

Two non-unique indexes are declared in `FlowHubDbContext.OnModelCreating`:

- `IX_Captures_Stage` — drives the Dashboard "Needs Attention" count (`Stage IN ( 'Orphan', 'Unhandled' )`) and the lifecycle filter chips on `/captures` .
- `IX_Captures_CreatedAt_DESC` — drives Recent Captures (top-N descending) and cursor pagination on `/captures` . Declared as `IsDescending()` so SQLite / PostgreSQL pick the right scan direction.

No index on `MatchedSkill` yet — the matched-skill filter is not exposed in the UI. Block 4 will add `(Source, CreatedAt DESC)` once the channel filter chips land in the spec.

8. `EntityTypeConfiguration = inline OnModelCreating (Beta)` → per-entity classes (Block 4)

The Beta has one entity (`CaptureEntity`); inline configuration in `OnModelCreating` is the pragmatic choice. When Block 4 adds `Skill` / `SkillRun` / `Channel` / `Integration` / `IntegrationHealthSample` / `Tag` / `CaptureTag(join)`, each will get its own `IEntityTypeConfiguration<T>` class in `source/FlowHub.Persistence/Configurations/` and `OnModelCreating` will reduce to `modelBuilder.ApplyConfigurationsFromAssembly(typeof(FlowHubDbContext).Assembly)` . The refactor is mechanical — no public-API change.

9. AI audit fields on `Capture` = deferred to Block 4

ADR 0004 §"Consequences for next blocks" promised AI audit fields (`provider` , `model` , `duration_ms` , `was_fallback`) per classification. The Beta's `CaptureEntity` does NOT carry them yet — those columns are part of Block 4 scope when the Capture-CRUD use cases formalize what audit data is queryable. The schema migration in Block 4 will add the columns; existing rows in the homelab SQLite DB at the time of the Postgres switchover are migrated nullably.

10. Test strategy = `InMemory` for unit tests (Beta), `Testcontainers PostgreSQL` for repository tests (Block 4)

- **Beta:** All 13 `EfCaptureServiceTests` run against `Microsoft.EntityFrameworkCore.InMemory` . Each test gets a fresh DB via a unique GUID database name; no file I/O, no provider-specific behaviour. This is fast and reliable but **does not exercise SQLite-specific behaviour** — and more importantly, would not exercise PostgreSQL-specific behaviour either. Acceptable Beta-scope choice; the Beta MVP runs SQLite identically to its unit-test environment for the Capture aggregate's current query shapes.

- **Block 4:** Repository tests will move to **Testcontainers PostgreSQL** to exercise real provider behaviour (operator semantics, JSONB column shape, full-text search, etc.). The 13 existing tests will be parameterized over both `InMemory` (fast feedback) and Testcontainers PostgreSQL (provider fidelity), or split into two suites with `Category` traits — TBD with Block 4's testing-strategy update.
- **API integration tests** (`tests/FlowHub.Api.IntegrationTests/` , 17 tests) swap the production Sqlite registration for `InMemory` per the `IntegrationTestFactory` pattern documented in §5. The Block 4 evolution will swap this to Testcontainers PostgreSQL for parity with production, at the cost of slower test boot.

The provider-removal pattern in `IntegrationTestFactory.CreateHost` is non-trivial (EF Core 8+ registers `DbContextOptionsConfiguration<TContext>` as a second descriptor; removing only `DbContextOptions<TContext>` is insufficient or you get *"two providers registered"* at runtime). The pattern is documented inline in the test file and called out in `docs/ai-usage.md` 's Beta-MVP retrospective.

Alternatives considered (and rejected)

Dapper instead of EF Core

Dapper is the .NET community's idiomatic micro-ORM (raw SQL + manual mapping). Rejected because:

- The Block 4 Auftrag explicitly asks for "ORM-Abstraktion" — a micro-ORM doesn't satisfy that wording.
- FlowHub's query shapes are not Dapper's strength: we want LINQ + Expression Trees ("Criteria-API"-Äquivalent per the Auftrag) for dynamic filters; Dapper defers all that to hand-written SQL.
- Migration tooling is more mature in EF Core than in Dapper-based projects (no first-party Dapper migration story).

NHibernate instead of EF Core

NHibernate is the older .NET ORM, originally a Hibernate port. Rejected because:

- Smaller community than EF Core in 2026.

- Microsoft's first-party investment is in EF Core (and increasingly `Microsoft.Extensions.AI` -style abstractions); going off-piste here adds long-term maintenance cost.
- No measurable advantage for FlowHub's query shapes.

Repository-per-entity pattern (`ICaptureRepository`, `ISkillRepository`, ...)

Would add an explicit `IRepository<T>` layer between `ICaptureService` and `DbContext`. Rejected because:

- `ICaptureService` is already the driving port; adding `ICaptureRepository` duplicates abstraction without addition.
- `DbSet<T>` is already a repository in EF Core's design.
- Mocking-for-tests doesn't need it: `EfCaptureServiceTests` use the `InMemory` provider, which is more representative than a hand-rolled mock.
- Pattern adds ceremony per entity; YAGNI per the project guardrails in `CLAUDE.md`.

Auto-apply migrations in `Program.cs` `app.Run()` body (vs. `IHostedService`)

Both apply at startup. Chose `IHostedService` because:

- Cleaner cancellation semantics (`StartAsync(Cancellation.Token)`).
- Composes cleanly with the `RemoveAll` pattern in `IntegrationTestFactory` (the test host removes only the migration runner descriptor, not the entire hosted-service set).
- Matches the project's `LoggerMessage` source-gen pattern for `EventIds` 5010/5011 — same shape as `LifecycleFaultObserver` and the AI provider boot loggers.

Consequences

EventId namespace

Reserved 5000–5999 for persistence runtime / startup events. Currently used:

- **5010 LogApplyingMigrations** (`MigrationRunner.StartAsync` entry, Information)
- **5011 LogMigrationsApplied** (`MigrationRunner.StartAsync` exit, Information)

Future: 5020-range for connection-pool events (Block 4), 5030-range for repository-level slow-query traces (Block 5 with OTEL).

Beta MVP scope acknowledged

Items the Beta intentionally did NOT do, deferred to Block 4 or Block 5:

Item	Deferred to	Reason
PostgreSQL provider	Block 4	Beta scope = SQLite for zero-setup demo
<code>IEntityTypeConfiguration<T></code> per entity	Block 4	Refactor lands when domain expands beyond <code>Capture</code>
Repository-pattern interfaces in <code>FlowHub.Core</code>	(rejected — see above)	Architectural decision, not deferred
AI audit fields on <code>Capture</code> (<code>provider</code> , <code>model</code> , <code>duration_ms</code> , <code>was_fallback</code>)	Block 4	Tied to ADR 0004 §"Consequences"
Migrations as separate init container (12-Factor XII)	Block 5	Compose / K8s deployment scope
Soft-delete strategy	Block 4	Lifecycle stages already provide pseudo-deletion via <code>Orphan</code> / <code>Unhandled</code>
Full audit fields (<code>UpdatedAt</code> , <code>CreatedBy</code>)	Block 4	Single-user system; <code>CreatedAt</code> is sufficient until OIDC lands
Testcontainers PostgreSQL for repository tests	Block 4	Beta uses InMemory for speed; provider fidelity in Block 4
<code>just db-up</code> / <code>just db-migrate</code> Make targets	Block 4	Tied to PostgreSQL switch
<code>docs/insights/block-4.md</code>	Block 4	Captures Block-4 lessons learned
ER diagram of full domain	Block 4	Beta has only <code>Capture</code> ; full ER is Block 4 deliverable

Consequences for next blocks

Block 4 (Persistence):

- Provider switch to PostgreSQL: connection-string change, Testcontainers adoption, `IntegrationTestFactory` provider-removal pattern revalidated.
- Domain expansion: 6+ new entities, each with `IEntityTypeConfiguration<T>`, fluent-API config, indexes, FKs.
- Dynamic-query helper (`Expression<Func<T, bool>>` builder) for the `/captures` filter chips — currently the only dynamic predicate is the `CaptureFilter.Stages Contains(stage)` expression.
- Add AI audit fields per ADR 0004 §"Consequences for next blocks".
- Revisit the captive- `HttpClient` pattern in `AddHttpClient<T>` + `AddSingleton<ISkillIntegration>(sp => sp.GetRequiredService<T>())` flagged in the Beta-MVP final code review (orthogonal to persistence but in scope for cleanup).

Block 5 (Deployment):

- Move migration application out of the running app; init container in Compose / K8s.
- pgvector adoption for embeddings (ADR 0006 — KI-Suche). Postgres-specific.
- Backup / restore tooling discussion (out of scope per Block 5 Nachbereitung's geparkt list, but `pg_dump` workflow worth one paragraph).

References

- Brainstorming spec: `docs/superpowers/specs/2026-05-04-beta-mvp-design.md`
- Implementation plan: `docs/superpowers/plans/2026-05-04-beta-mvp.md`
- ADR 0001: `docs/adr/0001-frontend-render-mode-and-architecture.md`
- ADR 0002: `docs/adr/0002-service-architecture-and-async-communication.md`
- ADR 0003: `docs/adr/0003-async-pipeline.md` — EventId namespacing
- ADR 0004: `docs/adr/0004-ai-integration-in-services.md` — promised AI audit fields
- AI Usage living doc: `docs/ai-usage.md` (Block 4 prep — Beta MVP section)
- Block 4 Nachbereitung: `vault/Blöcke/04 Persistence/04 Persistence - c) Nachbereitung.md`
- Bewertungskriterien: `vault/Organisation/Bewertungskriterien.md`
- EF Core 10 docs: <https://learn.microsoft.com/en-us/ef/core/>
- 12-Factor XII (Admin processes): <https://12factor.net/admin-processes>

ADR 0006 — Vector Search

Source: `docs/adr/0006-vector-search.md`

ADR 0006 — Vector Search: pgvector + Mistral Embeddings

Date: 2026-05-07

Status: Accepted

Context

FlowHub accumulates Captures over time. Keyword search fails when the query doesn't share vocabulary with stored titles or content. Semantic (embedding-based) search finds by meaning. The CAS AISE rubric item "Intelligente Services mit KI" (6 pts) requires a KI-based search feature.

Decision

- **Vector store:** pgvector extension on the existing PostgreSQL instance. No additional service.
- **Column type:** `vector(1024)` — fixed by `mistral-embed` output dimensionality.
- **Index:** HNSW (`hnsw (Embedding vector_cosine_ops)`) — approximate nearest-neighbour, sub-millisecond at <1 M rows.
- **Embedding provider:** Mistral `mistral-embed` via MEAI `IEmbeddingGenerator<string, Embedding<float>>`.
- **Provider abstraction:** OpenAI-compatible base URL + model name in env vars. Switching to OpenAI = change 3 env vars + new migration for column if dimensions differ.
- **Generation pipeline:** Embedding is generated **off the request path** by `CaptureEmbeddingConsumer`, which subscribes to `CaptureCreated` (same shape as `CaptureEnrichmentConsumer`). `POST /api/v1/captures` returns as soon as the row is persisted and the message is published — embedding latency does not count against NF-09 (p95 < 200 ms). Retries: 3 intervals at 500 ms / 2 s / 5 s;

on exhaustion the Capture remains without an embedding and can be backfilled via `POST /api/v1/admin/embeddings/rebuild`.

- **Graceful degradation:** If `Embeddings__ApiKey` is absent, `IEmbeddingService` is a no-op and the consumer skips storage; Captures are persisted without embeddings. Search returns `503` when not configured.

Provider Configuration

```
Embeddings__BaseUrl=https://api.mistral.ai/v1
Embeddings__ApiKey=<key>
Embeddings__Model=mistral-embed
Embeddings__Dimensions=1024
Embeddings__TimeoutSeconds=10
```

`Embeddings__Dimensions` is forwarded to the provider via `EmbeddingGenerationOptions.Dimensions` — required when using OpenAI `text-embedding-3-*` to truncate to a non-native size. `Embeddings__TimeoutSeconds` is bounded on the underlying `OpenAIClient` (`NetworkTimeout`); on timeout the consumer retries per its policy.

Switching Providers

To switch to OpenAI (`text-embedding-3-small`, 1536 dims):

1. Change env vars: `BaseUrl=https://api.openai.com/v1`, `Model=text-embedding-3-small`, `Dimensions=1536`.
2. Add an EF Core migration: `ALTER TABLE "Captures" ALTER COLUMN "Embedding" TYPE vector(1536)`.
3. Run `POST /api/v1/admin/embeddings/rebuild` to backfill.

OpenRouter and Anthropic are **not supported** for embeddings — neither exposes an embeddings API.

Consequences

- + No new service to operate; pgvector is an extension, not a separate vector DB.
- + Provider-agnostic via env vars.


- – Dimension change requires migration + rebuild.
- – HNSW index is not exact; results may differ slightly from exhaustive scan.

Specification & design

Source: `docs/spec/system-context.md`

FlowHub — System Context (C4 Level 1)

Context Diagram



Key relationships

From	To	Protocol	Notes
Operator → FlowHub.Web	HTTP + SignalR (WebSocket)	Blazor Interactive Server; single long-lived circuit per session	
Operator → Telegram Bot API	HTTPS	Operator sends message to bot; Telegram forwards via webhook	
FlowHub.Web → FlowHub.Core	In-process DI	No HTTP — <code>@inject</code> services directly (per ADR 0001 §2)	
FlowHub.Telegram → FlowHub.Core	In-process DI	Same process, same pattern	
FlowHub.Core → FlowHub.AI	In-process	Classification service calls Ollama REST API for inference	
FlowHub.AI → Ollama	HTTP REST (local)	<code>http://ollama:11434</code> — runs on the same homelab, never leaves the network	
FlowHub.Skills → Integrations	HTTP REST / git	Each Skill writes to one or more downstream services via their APIs	
FlowHub.Web → Authentik	OIDC (HTTPS)	Auth code flow; tokens in cookie; SignalR circuit reads cookie	

Current state (Block 5 — submission)


The solution (`FlowHub.slnx`) contains six implemented projects; two folders remain intentional, not-yet-implemented placeholders.

- **Implemented (in the solution, with code):**

- `FlowHub.Web` — Blazor Web App (Interactive Server) + the REST API host

- `FlowHub.Core` — domain types and driving/driven ports
- `FlowHub.Api` — REST endpoint contracts for non-UI consumers
- `FlowHub.AI` — LLM-backed classifier behind the `IClassifier` port (provider abstraction + keyword fallback)
- `FlowHub.Persistence` — EF Core + PostgreSQL repositories and migrations
- `FlowHub.Skills` — Wallabag and Vikunja `ISkillIntegration` adapters
- **Planned, not yet scaffolded (no project in the source tree):** a Telegram inbound channel and a generic integrations layer. They are intentionally absent until built — the Wallabag/Vikunja adapters that an early draft placed under `FlowHub.Integrations` live in `FlowHub.Skills` (see ADR 0002 "As built").
- **Not yet wired:** Authentik OIDC (dev bypass only), Ollama-hosted inference (see ADR 0007), the Telegram channel.
- **REST API:** available since Block 3 (`/api/v1/captures` , exercised live on the public demo).

Deployment context (Block 5, future)



Persistence Layer

FlowHub uses PostgreSQL 17 as its primary database, accessed via EF Core 10 with the Npgsql provider. The schema follows a migrations-first workflow: all schema changes are expressed as EF Core migration files committed to Git and applied as an idempotent SQL script at deploy time.

The Repository pattern separates domain logic from database access. Repository interfaces are defined in `FlowHub.Core` (returning domain types), with EF Core implementations in `FlowHub.Persistence`. Application-layer services (`ICaptureService`, `ISkillRegistry`, `IIntegrationHealthService`) compose repositories; they never reference `FlowHubDbContext` directly.

Local development uses `docker compose up postgres` to start a PostgreSQL container. `just db-migrate` applies pending migrations. `just run` starts the application, which auto-migrates via `MigrationRunner` for convenience.

Use Cases

Source: `docs/spec/use-cases.md`

FlowHub — Use Cases & Requirements

Solution Vision

FlowHub is an intelligent integration hub that captures heterogeneous inputs (URLs, text, images) from multiple channels (Telegram, Web UI), classifies them via AI, and routes them to the correct downstream service (Wallabag, Vikunja, Paperless-ngx, Obsidian). It automates 80% of the manual workflow the operator currently performs by hand: receive input → decide where it goes → copy-paste into the right tool.

See `Projektarbeit/Idee FlowHub.md` in the CAS Obsidian vault for the original concept and management summary.

Use Cases

UC-01: Submit a Capture via Web UI (Quick)

Actor: Operator **Trigger:** Operator pastes a URL or types text in the AppBar quick-capture field. **Precondition:** Operator is authenticated (DevAuthHandler in dev, Authentik OIDC in prod). **Flow:**

1. Operator types or pastes content into the AppBar quick-capture field.
2. Operator presses Enter or clicks the submit icon.
3. System creates a new Capture with `source = Web`, `stage = Raw`.
4. System shows a success snackbar with a link to the Capture detail.
5. The quick-capture field clears, ready for the next input.

Postcondition: A new Capture exists in the system with `LifecycleStage.Raw`.

Error: Empty content → inline validation hint "Type something first". **Error:**

Service failure → snackbar "Capture failed: {reason}", field content preserved.

Akzeptanzkriterien:

- After typing content and pressing Enter, a `Captured ✓` snackbar appears within 2 s (Playwright happy-flow `HappyFlowTests.QuickCapture_TodoEntry_AppearsInCapturesListAndDetail`).
- The captured row appears on `/captures` and the detail page renders the same content (same Playwright test).
- Empty content → no row created; inline hint visible.

UC-02: Submit a Capture via Web UI (Long Form)

Actor: Operator **Trigger:** Operator navigates to `/captures/new` . **Precondition:** Authenticated. **Flow:**

1. System loads available Skills from `ISkillRegistry` .
2. Operator enters multi-line content in the text field.
3. Optionally, operator selects a Skill override from the dropdown (default: "Let AI decide").
4. Operator clicks Submit.
5. System validates (content required), then creates the Capture.
6. System shows success snackbar. Form clears; page stays for rapid multi-entry.

Postcondition: A new Capture exists. If Skill override was selected, the Capture may skip AI classification. **Error:** Skill registry fails to load → dropdown disabled, "Let AI decide" forced, Submit still works.

Akzeptanzkriterien:

- `/captures/new` renders a multi-line text area and a Skill dropdown.
- Submitting non-empty content yields a 201-equivalent service call and a success snackbar; the form clears for the next entry.
- If `ISkillRegistry.ListAsync` throws, the dropdown is disabled and a `"Let AI decide"` placeholder is forced.

UC-03: Submit a Capture via Telegram

Actor: Operator (via Telegram bot) **Trigger:** Operator sends a message to the FlowHub Telegram bot. **Precondition:** Telegram bot is configured and running

(Block 3+). **Flow:**

1. Telegram bot receives the message.
2. Bot creates a Capture with `source = Telegram`, `stage = Raw`.
3. Bot confirms receipt in the chat.

Postcondition: Capture exists, visible in the Web UI Dashboard and Captures list. **Note:** Not implemented in Block 2 — placeholder `source/FlowHub.Telegram/`. Once the Telegram module lands in Block 3+, the bot calls `POST /api/v1/captures` (UC-08) to submit the capture, inheriting the same validation and pipeline flow.

Akzeptanzkriterien:

- Bot reply confirms receipt within 5 s of the user message.
- Capture appears in `/captures` with `source = Telegram` and the bot user's identifier in metadata.
- *(Deferred — implementation lands post-Block-3; covered by UC-08 acceptance for the API path the bot uses.)*

UC-04: Monitor Capture health via Dashboard

Actor: Operator **Trigger:** Operator opens the Dashboard (/). **Flow:**

1. System loads failure counts, recent captures (10), skill health, and integration health in parallel.
2. Dashboard shows the "Needs Attention" widget with orphan and unhandled counts.
3. Dashboard shows the "Recent Captures" grid with lifecycle badges.
4. Dashboard shows Skill and Integration health cards.

Postcondition: Operator has an at-a-glance view of system health. **Empty state:** No captures → "No captures yet" with a call-to-action. **All-clear state:** Zero failures → calm "All captures routed successfully" message.

Akzeptanzkriterien:

- The Dashboard renders within NF-08 budget (`p95 < 1.5 s` first paint on a 100-capture DB).
- Empty DB shows the empty-state message (verified by `bUnit DashboardTests`).

- Mixed-state DB shows non-zero Orphan / Unhandled counts in the "Needs Attention" widget.

UC-05: Browse and filter Captures

Actor: Operator **Trigger:** Operator navigates to `/captures` (or clicks through from Dashboard). **Flow:**

1. System loads all Captures.
2. If `?lc=` query param is present, the lifecycle filter pre-selects.
3. Operator filters by lifecycle stage (chip bar), channel (chip bar), and/or content text (search field).
4. Grid shows filtered results with pagination.
5. Operator clicks a row to drill into the Capture detail.

Postcondition: Operator can find any Capture by stage, channel, or content.

Akzeptanzkriterien:

- Deep link `/captures?lc=Orphan` pre-selects the Orphan chip and shows only Orphan rows.
- Combining a lifecycle chip + a channel chip applies an AND filter.
- The grid pages at 10/25/50 rows (configurable, default 10).

UC-06: Inspect and act on a failed Capture

Actor: Operator **Trigger:** Operator opens a Capture detail page (`/captures/{id}`) for an Orphan or Unhandled Capture. **Flow (Orphan):**

1. System shows the failure reason alert.
2. Operator sees action buttons: Retry routing, Reassign skill, Ignore.
3. Operator clicks an action (Block 2: stubbed → snackbar "Coming in Block 3").

Flow (Unhandled):

1. System shows "No Skill matched" alert.
2. Operator sees action buttons: Assign skill, Ignore.
3. Operator clicks an action (Block 2: stubbed).

Postcondition (Block 3+): Capture is retried, reassigned, or marked as ignored. Dashboard counts update. **Block 2:** Actions show intent but do not mutate state.

Akzeptanzkriterien:

- Detail page shows the `FailureReason` of an Orphan and "No Skill matched" for an Unhandled.
- "Retry" button is enabled only for Orphan / Unhandled stages (UC-11 implements the action).
- Block 2 stubs trigger a "Coming in Block 3" snackbar instead of mutating state.

UC-07: View Skill and Integration health

Actor: Operator **Trigger:** Operator navigates to `/skills` or `/integrations`.

Flow:

1. System loads health data from the respective service.
2. Grid shows each Skill/Integration with status (healthy/degraded/down) and activity metrics.

Postcondition: Operator knows which Skills/Integrations are operational.

Akzeptanzkriterien:

- `/skills` renders one row per registered Skill with status badge (healthy/degraded/down).
- `/integrations` renders one row per `ISkillIntegration` adapter with `LastWriteAt` + `LastWriteDurationMs`.
- Each row exposes a "history" expansion showing recent samples (UC-16).

UC-08: Submit a Capture via REST API

Actor: Non-UI client (automation script, Telegram bot module, future mobile client)

Trigger: Client sends `POST /api/v1/captures` with a JSON body.

Precondition: Client holds a valid bearer token (DevAuthHandler in dev; Authentik OIDC token in prod). **Flow:**

1. Client sends `POST /api/v1/captures` with body `{ "content": "...", "source": "Telegram|Web|Api", "skillOverride": "<SkillId|null>" }`.
2. FluentValidation at the API boundary checks: `content` non-empty, `source` is a known enum value. On failure → `400 Bad Request` with RFC 9457 ProblemDetails body (`type` URI from `FlowHubProblemTypes`).

3. Application service creates a `Capture` with `stage = Raw`, persists it (Block 4+; in-memory stub in Block 3).
4. Service publishes `CaptureCreated` event on the internal bus.
5. API returns `201 Created` with `Location: /api/v1/captures/{id}` and the full `Capture` resource body.

Postcondition: A new `Capture` with `LifecycleStage.Raw` exists and a `CaptureCreated` event is on the bus, triggering the async enrichment pipeline (UC-09). **Error:** Missing/invalid `content` → `400 ProblemDetails` with `errors` map. **Error:** Unknown `source` value → `400 ProblemDetails`. **Error:** Auth missing or invalid → `401 Unauthorized`. **Note:** API surface documented in `docs/design/api/api-surface.md`; OpenAPI schema browsable at `/scalar` (ADR 0002).

Akzeptanzkriterien:

- `POST /api/v1/captures` with a valid body returns 201 within NF-09 (p95 < 200 ms server-side). Verified by `just smoke-prod` step [5/6] and `tests/FlowHub.Api.IntegrationTests/`.
- Response body matches the `Capture` schema and `Location: /api/v1/captures/{id}` header is present.
- Missing `content` / unknown `source` → `400 ValidationProblem ( type = validation.md )`.
- Bruno collection `bruno/captures/submit-capture.bru` round-trips against a live stack.

UC-09: AI-classify and route a Capture (async pipeline)

Actor: System (no human interaction) **Trigger:** `CaptureCreated` event published on the `MassTransit` in-process bus. **Precondition:** `Capture` exists in `Raw` stage; `MassTransit` bus is running. **Flow:**

1. Bus delivers `CaptureCreated` to `CaptureEnrichmentConsumer`.
2. `CaptureEnrichmentConsumer` calls `IClassifier.ClassifyAsync(capture)`.
3. If an AI provider key is configured, `AiClassifier` sends the capture content to the configured LLM (Anthropic Claude Haiku 4.5 by default) with `MaxOutputTokens=300`, `Temperature=0.2`. If not configured, `KeywordClassifier` is used directly.
4. `IClassifier` returns a `ClassificationResult` with `SkillId` (nullable) and confidence.

5. Consumer updates the Capture's `SkillId` and sets `stage = Classified`; publishes `CaptureClassified` event.
6. Bus delivers `CaptureClassified` to `SkillRoutingConsumer`.
7. `SkillRoutingConsumer` looks up the matched `ISkillIntegration` by `SkillId`.
 - If found: calls `ISkillIntegration.HandleAsync(capture)`. On success → `stage = Routed`.
 - If no skill matched (`SkillId == null`): `stage = Orphan`.
8. On integration error, MassTransit retries per the retry policy (NF-10). After retries exhausted: `LifecycleFaultObserver` sets `stage = Unhandled` with a `FailureReason`.

Postcondition: Capture reaches a terminal stage (`Routed`, `Orphan`, or `Unhandled`). Dashboard counts update on next load. **Reference:** ADR 0003 (async pipeline), ADR 0004 (AI classifier).

Akzeptanzkriterien:

- A URL-content Capture reaches `Completed` with `MatchedSkill = "Wallabag"` and a non-empty `ExternalRef` (`Skills.ContractTests + just test-beta`).
- A todo-content Capture reaches `Completed` with `MatchedSkill = "Vikunja"` and a non-empty `ExternalRef`.
- MassTransit harness tests (`tests/FlowHub.Web.ComponentTests/Pipeline/*`) prove the three consumer hops fire in order.

UC-10: Graceful AI-classifier fallback to keyword floor

Actor: System (no human interaction) **Trigger:** `AiClassifier.ClassifyAsync` encounters any failure: network exception, HTTP timeout, JSON parse error, schema-constraint violation, or generic unhandled exception. **Precondition:** AI provider is configured; `CaptureEnrichmentConsumer` is processing a `CaptureCreated` event. **Flow:**

1. `AiClassifier.ClassifyAsync` executes the LLM call inside a try/catch covering all exception types.
2. On any exception: `AiClassifier` logs at Warning level with EventId `3010` (`AiClassifierFallbackToKeyword`), including the exception message and Capture ID.
3. `AiClassifier` delegates immediately to `KeywordClassifier.ClassifyAsync(capture)` and returns its result. The returned

`ClassificationResult` has `Title = null` (keyword classifier does not produce a human-readable title).

4. `CaptureEnrichmentConsumer` continues normally — it receives a valid `ClassificationResult` regardless of which classifier produced it.

Postcondition: The Capture is always classified and the pipeline continues. An AI provider outage degrades classification quality but does not cause availability loss. The Warning log at EventId 3010 provides operational visibility. **Note:** Fallback rate target and monitoring defined in NF-11. Reference: ADR 0004, decision D5.

Akzeptanzkriterien:

- With `Ai__OpenRouter__ApiKey` deliberately invalid, a Capture submitted via UC-08 still reaches `Classified` stage — via `KeywordClassifier`. Verified by `tests/FlowHub.Web.ComponentTests/Ai/AiClassifierTests.cs`.
- EventId 3010 `AiClassifierFellBackToKeyword` is logged at Warning level with the exception type and Capture ID.
- `AiClassifier` never throws to its caller (`ClassifyAsync` always returns a `ClassificationResult`).

UC-11: Retry a failed Capture from the dashboard

Actor: Operator **Trigger:** Operator opens a Capture detail page (`/captures/{id}`) for a Capture in `Orphan` or `Unhandled` stage and clicks "Retry". **Precondition:** Capture is in `Orphan` or `Unhandled` stage. Operator is authenticated. **Flow:**

1. UI calls `POST /api/v1/captures/{id}/retry` (API surface from UC-08's REST layer).
2. Application service resets `stage` to `Raw`, clears `FailureReason`.
3. Service publishes a new `CaptureCreated` event for the existing Capture ID.
4. The full async pipeline (UC-09) runs again from the start.
5. API returns `202 Accepted`. UI shows snackbar "Capture queued for retry" and reloads the detail page.

Postcondition: Capture re-enters the enrichment pipeline. The Lifecycle stage is eventually updated to `Routed`, `Orphan`, or `Unhandled` again based on pipeline outcome. **Error:** Capture not found → `404 ProblemDetails`. **Error:** Capture is not in a retryable stage (`Raw`, `Classified`, `Routed`) → `409 Conflict ProblemDetails`. **Note:** UC-06 documents the UI entry point and the Block-2 stub state. UC-11

defines the Block-3 implementation that fulfils the "Retry routing" action stubbed in UC-06 Flow (Orphan) step 3 and Flow (Unhandled) step 2.

Akzeptanzkriterien:

- `POST /api/v1/captures/{id}/retry` on an Orphan Capture returns 202 Accepted and the body shows `stage = Raw`, `failureReason = null`. Covered by `tests/FlowHub.Api.IntegrationTests/CaptureRetryEndpointTests.cs`.
- Same call on a `Completed` Capture returns 409 (`type = capture-not-retryable.md`).
- Same call with an unknown id returns 404 (`type = capture-not-found.md`).

UC-12: Filter Captures by Lifecycle Stage

Actor: Operator

Trigger: Operator selects one or more stages in the Captures list filter.

Flow: System queries `ICaptureService.ListAsync` with the selected stages. List refreshes to show only matching captures.

Acceptance: Selecting "Orphan" returns only orphaned captures. Selecting multiple stages returns the union.

UC-13: Filter Captures by Tag

Actor: Operator

Trigger: Operator enters a tag value in the tag filter field.

Flow: System queries `ICaptureService.ListAsync` with `CaptureFilter.Tag` set. Returns only captures that have the tag attached.

Acceptance: A capture tagged "dotnet" appears when filter is "dotnet". Captures without the tag are excluded.

UC-14: Search Captures by Content or Title

Actor: Operator

Trigger: Operator types a search term in the search box.

Flow: System queries with `CaptureFilter.SearchTerm`. PostgreSQL ILIKE match against Content and Title (case-insensitive). Not full-text search — deferred to Block 5.

Acceptance: Searching "hexagonal" returns captures whose Content or Title contains that substring (any casing).

UC-15: View Skill-Run History for a Capture

Actor: Operator

Trigger: Operator opens a Capture detail and views the routing history tab.

Flow: System queries `ISkillRunRepository.GetByCaptureIdAsync`. Returns all routing attempts for that capture ordered by StartedAt DESC.

Acceptance: A capture that was routed twice shows two SkillRun entries.

UC-16: View Integration Health History

Actor: Operator

Trigger: Operator opens the Integrations page and expands an integration's history.

Flow: System queries `IIntegrationRepository.GetRecentSamplesAsync`. Returns the last N health samples.

Acceptance: Shows SampledAt, Status, and DurationMs for each sample.

Non-Functional Requirements — Quality-Attribute Catalogue (Block 2)

This is the original Block-2 quality-attribute table (each entry has a measurable target + verification, but **not** a full five-dimension SMART decomposition). The SMART-decomposed requirements are maintained in `nfa.md` (NfA-*), which supersedes this table and maps the `NF-*` ids to it. The `NF-*` numbers are kept here only as stable references from existing acceptance criteria and use cases.

ID	Requirement	Measurable target	Verification
NF-01	Response time — Dashboard loads within acceptable time on a single-operator workstation	< 500 ms from navigation to all 4 cards rendered (on localhost with Bogus stubs)	Manual observation, full timing
NF-02	Availability — FlowHub runs as a self-hosted service in the operator's homelab	99% uptime during operator's active hours (not 24/7 SLA)	Docker healthcheck in
NF-03	Concurrency — system supports a single concurrent operator	1 concurrent SignalR circuit without degradation	Architecture (Interactive horizontal scaling needed)
NF-04	Security — all pages require authentication	0 pages accessible without a valid auth session	[Authorize] on pages in dev, OIDC in prod
NF-05	Testability — all UI components testable in isolation	100% of page components renderable in bUnit without a running server	bUnit component suite grown to the Block-5 test docs/spec/testing-str
NF-06	Maintainability — code compiles with zero warnings	TreatWarningsAsErrors=true in Directory.Build.props	just build in CI

ID	Requirement	Measurable target	Verification
NF-07	Portability — runs on Linux (homelab Docker) and WSL2 (dev)	<code>just run</code> works on both environments	Manual verification
NF-08	Data privacy — no Capture content leaves the operator's infrastructure	0 external API calls for data processing (AI classification runs locally via Ollama in future blocks)	Architecture review
NF-09	API latency — REST endpoints respond within tight bounds for an interactive integration hub	<code>POST /api/v1/captures</code> p95 < 200 ms (server-side, excluding async bus-publish); <code>GET /api/v1/captures</code> p95 < 100 ms with cursor pagination	Block 5 load test (k6/n evidence: integration-end-to-end against in-
NF-10	Async pipeline retry budget — transient failures in the enrichment, embedding, or routing consumers are retried before a Capture is marked <code>Unhandled</code>	<code>CaptureEnrichmentConsumer</code> : 2 retries at 100 ms / 500 ms; <code>CaptureEmbeddingConsumer</code> : 3 retries at 500 ms / 2 000 ms / 5 000 ms (failure leaves the Capture without an embedding — backfill via admin rebuild); <code>SkillRoutingConsumer</code> : 3 retries at 500 ms / 2 000 ms / 5 000 ms; after exhaustion <code>LifecycleFaultObserver</code> marks <code>Unhandled</code>	MassTransit <code>TestHarness</code> tests/FlowHub.Web.Com (6 harness tests)

ID	Requirement	Measurable target	Verifi
NF-11	AI classifier fallback rate — an AI provider outage must not propagate to availability loss	< 5% of classifications fall back to keyword during normal provider availability (Anthropic Haiku 4.5 SLA ~99.5%); fallback always succeeds — <code>AiClassifier</code> never throws to the caller	EventId 3010 log volume <code>just test-ai</code> live integration demonstrate the success D5
NF-12	AI classification cost — per-capture cost is sub-cent to keep the homelab budget bounded	<code>MaxOutputTokens=300</code> , <code>Temperature=0.2</code> ; estimated ~200 tokens input + ~150 tokens output → < \$0.001 per classification on Haiku 4.5 pricing	Anthropic dashboard used by operator runs; cost guard <code>AddFlowHubAi(IConfiguration §"Cost guards"</code>
NF-13	OpenAPI versioning SLA — the REST API is URL-versioned from day one so clients are not broken by future changes	Breaking changes land only in a new major version (<code>/api/v2/...</code>); v1 is retained for at least one major-version overlap period	ADR 0002 D6; endpoint <code>docs/design/api/api-s</code> prefix verified in route

UC-17: Deploy via Docker Compose

Actor: Operator **Precondition:** Docker and Docker Compose v2 are installed.

Steps:

1. Operator creates `.env` file with required env vars (DB credentials, optional OIDC config, optional embedding key).
2. Operator runs `docker compose up -d`.
3. `flowhub.migrations` service applies EF Core migrations and exits successfully.
4. `flowhub.web` starts, connects to PostgreSQL and RabbitMQ, exposes port 5070.
5. Operator opens `http://localhost:5070` . **Postcondition:** FlowHub is running; all dependencies healthy. **Exception:** If migrations fail, `flowhub.web` does not

start (depends_on condition).

Akzeptanzkriterien:

- `docker compose up --build -d --wait` returns exit 0 with all `service_healthy` dependencies satisfied (verified by `just smoke-prod step [1/6]`).
- `flowhub.migrations` container reaches `service_completed_successfully` with exit code 0.
- `GET /health/live` from inside the compose network returns 200 within NF-D3 (30 s) of container start.
- `GET /metrics` returns a Prometheus exposition that contains at least one `^dotnet_` and one `^http_` series.

Note: this UC was previously numbered UC-10 (collision with the AI-fallback UC at line 150). Renumbered 2026-05-12.

UC-18: Semantic Search for Captures

Actor: Operator **Precondition:** `Embeddings__ApiKey` is configured; at least one Capture exists with a stored embedding. **Note on embedding generation:** Embeddings for new Captures are generated asynchronously by `CaptureEmbeddingConsumer` (subscribed to `CaptureCreated`), so a Capture submitted via UC-09 becomes searchable a moment after submission, not synchronously with the `POST /api/v1/captures` response. This keeps capture submission inside the NF-09 p95 < 200 ms budget regardless of embedding-provider latency. Backfill of captures stored before the provider was configured is via `POST /api/v1/admin/embeddings/rebuild`.

Steps:

1. Operator sends `GET /api/v1/captures/search?q=database+performance&limit=5`.
2. FlowHub embeds the query via the configured embedding provider.
3. FlowHub returns up to 5 Captures ranked by cosine similarity.

Postcondition: Response contains Captures semantically related to the query, even if query words don't appear in titles or content. **Exceptions:**

- Empty/whitespace `q` → `400 Bad Request` (`ValidationProblem`).
- Embedding service not configured → `503 Service Unavailable` with `ProblemDetails` body.
- `limit` is clamped to 1..200 (mirrors the list-endpoint contract).

Akzeptanzkriterien:

- `POST /api/v1/captures` followed by polling `Captures.Embedding` populates the column within 30 s when `Embeddings__ApiKey` is set (verified by `just smoke-prod` step [6/6] — actual measurement ~2 s).
- `GET /api/v1/captures/search?q=...` with a non-empty `q` returns 200 with an array body when at least one embedded Capture exists.
- Same endpoint with empty `q` returns 400 `ValidationProblem`; with `Embeddings__ApiKey` unset, returns 503 `ProblemDetails` (`type` ends with `/validation.md`).
- `POST /api/v1/admin/embeddings/rebuild` returns 200 `{ processed, skipped, failed }` when keys are present, 503 otherwise.

Note: this UC was previously numbered UC-11 (collision with the Retry-Capture UC at line 164). Renumbered 2026-05-12.

Non-functional requirements (SMART)

Source: `docs/spec/nfa.md`

Non-Functional Attributes (NfA) — SMART Criteria

Every requirement below is decomposed along all five SMART dimensions — **S**pecific, **M**easurable, **A**chievable, **R**elevant, **T**ime-bound — so each carries an unambiguous target and a stated way to verify it.

NfA	Category	Dimension covered
NfA-01	Performance	S · M · A · R · T
NfA-02	Performance	S · M · A · R · T
NfA-03	Operability	S · M · A · R · T
NfA-04	Scalability	S · M · A · R · T
NfA-05	Reliability	S · M · A · R · T
NfA-D1..D3	Deployment	S · M · A · R · T
NfA-O1	Observability	S · M · A · R · T
NfA-P1..P2	Privacy & Compliance	S · M · A · R · T

Relationship to the NF-01..NF-13 table. The NF-* list in `docs/spec/use-cases.md` is the original Block-2 quality-attribute catalogue (response time, availability, security, testability, ...). This NfA-* document is its SMART-decomposed successor and is **authoritative** where the two overlap: NfA-01 refines NF-01 / NF-09 (latency), NfA-01 refines NF-01 (observability), NfA-D1..D3 cover deployment (NF-02), NfA-P1 refines NF-08 (data privacy), and NfA-P2 adds the AI-Act transparency requirement. The NF-* numbers are retained only as stable references from existing acceptance criteria and use cases.

NfA-01: Query Latency

- Specific:** All Capture list queries (`ICaptureService.ListAsync`) with a limit ≤ 50 must complete within 100ms at p95 under normal load.
- Measurable:** Measured via OpenTelemetry span duration on the `ListAsync` span; threshold surfaced in Grafana.
- Achievable:** Index-backed queries on `Stage` , `CreatedAt` , and `MatchedSkill` columns; cursor pagination avoids full-table scans.
- Relevant:** Dashboard and Captures list are the two highest-traffic read paths.
- Time-bound:** Verified against Testcontainers PostgreSQL with 10k seeded rows in Slice 4.

NfA-02: Index Strategy

- Specific:** Every high-frequency filter column carries a dedicated B-tree index, so the hot read paths never fall back to a sequential scan:

Index	Column(s)
<code>IX_Captures_Stage</code>	<code>Stage</code>
<code>IX_Captures_CreatedAt_DESC</code>	<code>CreatedAt DESC</code>
<code>IX_Captures_MatchedSkill</code>	<code>MatchedSkill</code>
<code>IX_IntegrationHealthSamples_IntegrationName_SampledAt_DESC</code>	<code>(IntegrationName SampledAt DESC)</code>

- Measurable:** Each index is declared via `HasIndex` in the EF Core model and present in the committed migration; `\d captures` in psql lists them, and `EXPLAIN` on the four query patterns shows index scans, not `Seq Scan` .

Achievable: Indexes are created code-first through EF Core migrations — no manual DBA step.

Relevant: Directly underwrites the NfA-01 p95 latency target on the two highest-traffic read paths.

Time-bound: In place since the Block 4 persistence slice; verified against the Testcontainers PostgreSQL suite.

NfA-03: Migration Strategy

Specific: All schema changes ship as EF Core code-first migrations committed to Git; production applies an idempotent SQL script as a separate init step, never auto-migrating inside `app.Run()`.

Measurable: Migration files exist under `source/FlowHub.Persistence/Migrations/`; production deploy runs `dotnet ef migrations script --idempotent` (reviewed pre-deploy) via the migrations init-container; dev convenience is the `MigrationRunner` hosted service.

Achievable: Standard EF Core tooling plus a dedicated init-job in `docker-compose.yml`.

Relevant: Satisfies 12-Factor XII (admin processes) and keeps production deploys safe and repeatable.

Time-bound: Established in Block 4; exercised end-to-end in the Block 5 deployment.

NfA-04: Data Volume Assumptions

Specific: The design targets these row volumes — Captures $\leq 100,000$ (Block 4 scope; partitioning beyond $\sim 1\text{M}$ is out of scope), IntegrationHealthSamples $\leq 10,000$ per integration (90-day prune deferred to Block 5), SkillRuns $\leq 500,000$ (archival deferred to Block 5).

Measurable: Volume assumptions are documented per table; the NfA-01 load test seeds 10,000 Capture rows to validate behaviour under realistic data.

Achievable: The indexed schema (NfA-02) and cursor pagination handle the stated volumes without partitioning.

Relevant: Bounds the performance and scaling envelope the other NfAs are verified against.

Time-bound: Valid for the Block 4 scope; revisited if a table is projected to exceed its stated ceiling.

NfA-05: Connection Resilience

Specific: Database access uses the Npgsql connection pool (dev defaults min=0, max=100) with built-in transient-fault retries; production pool sizing is set via connection-string parameters.

Measurable: Pool and retry behaviour are configured through the connection string / Npgsql data-source options and asserted by the Testcontainers integration tests that run against real PostgreSQL.

Achievable: Relies on Npgsql's built-in pooling and retry policy — no custom resilience code.

Relevant: Keeps the single-instance dev deployment and the production instance resilient to transient connection faults.

Time-bound: In place from Block 4; confirmed under the Block 5 deployed stack.

NfA-D1: Container Build Time

Category: Deployment

Specific: The multi-stage Docker image for `flowhub.web` MUST build from scratch in under 5 minutes on a GitHub-hosted `ubuntu-latest` runner (2 vCPUs, 7 GB RAM).

Measurable: GitHub Actions `release.yml` workflow — `docker/build-push-action` step duration; target ≤ 300 seconds.

Achievable: Multi-stage Alpine build with layer caching keeps the cold build inside the budget.

Relevant: Fast CI builds keep the release pipeline (and the inner-dev loop) responsive.

Time-bound: Enforced from Block 5 onward on every `release.yml` run.

NfA-D2: Image Size

Category: Deployment

Specific: The published `flowhub-web` Docker image MUST be under 200 MB compressed.

Measurable: `docker image inspect ghcr.io/freaxnx01/flowhub-web:<version> --format='{{.Size}}'` (uncompressed) and GHCR layer sizes (compressed); target ≤ 200 MB compressed (≤ 400 MB uncompressed).

Achievable: `aspnet:10.0-alpine` runtime base plus a self-contained trimmed publish keep the image small.

Relevant: Small images cut pull time and registry/storage cost on the VPS.

Time-bound: Checked on each published image from Block 5 onward.

NfA-D3: Startup Time

Category: Deployment

Specific: After migrations complete, `flowhub.web` MUST reach a healthy `/health/live` state within 30 seconds on a cold container start.

Measurable: Docker Compose healthcheck (`interval: 10s`, `retries: 3` = max 30 s) — `/health/live` returns HTTP 200 within 30 seconds of container start.

Achievable: Stateless startup with migrations run as a separate init step (NfA-03) keeps cold start fast.

Relevant: Bounds rollout time and makes orchestrated restarts predictable.

Time-bound: Verified against the deployed compose stack in Block 5.

NfA-O1: Observability — Metrics Endpoint

Category: Observability

Specific: The running `flowhub.web` process MUST expose Prometheus-format metrics at `/metrics`.

Measurable: `curl http://localhost:5070/metrics` returns HTTP 200 with `Content-Type: text/plain; version=0.0.4` and at least `dotnet_*` and `http_*` metric families.

Achievable: OpenTelemetry's Prometheus exporter is wired at the composition root and scraped by the bundled Prometheus.

Relevant: Metrics are the backbone of the NfA-01 latency check and the runtime health dashboards.

Time-bound: Available from the Block 5 observability slice and confirmed on the live demo.

NfA-P1: Personendaten-Residenz

Category: Privacy & Compliance

Status (Block 5): ZIEL — noch nicht umgesetzt. Diese NfA beschreibt den angestrebten Zustand, nicht den Abgabestand. Aktuell läuft die Klassifikation über OpenRouter und die Embeddings über Mistral (Cloud) — Capture-Inhalte verlassen für die KI-Inferenz heute also das Homelab, entgegen dem Ziel unten. Es gibt **keinen** lokalen Ollama-Adapter und

keinen `OutboundCallAuditTests` (siehe ADR 0007 „As built“). Die Kriterien unten sind die Abnahmedefinition für die spätere Umsetzung.

Statement (Ziel): Capture-Inhalte (Body, URL, Metadaten) sollen auf vom Betreiber selbst betriebener Infrastruktur verarbeitet und persistiert werden. Outbound-Calls an Drittsysteme nur für explizite Skill-Targets (Vikunja, Wallabag); LLM-Inferenz im Zielzustand lokal via Ollama (ADR 0007). **Measurable (Abnahmedefinition, offen):**

1. Statisch: `dotnet list package` zeigt keinen Cloud-LLM-Client als transitive Default-Dependency. (*heute nicht erfüllt — Cloud-Provider sind aktiver Default*)
2. Dynamisch: `integration test OutboundCallAuditTests` verifiziert, dass im Default-Profil kein HTTP-Request an Cloud-LLM-Endpunkte erfolgt. (*Test noch nicht implementiert*)
3. Doku: `docs/design/data-flow.md` zeigt Capture-Pfad → Klassifikation → Skill-Routing mit Homelab-Boundary. (*vorhanden*) **Achievable:** Erfordert einen lokalen `Local /Ollama-Provider-Adapter` in `FlowHub.AI` plus `Embeddings:Provider=Local` als Default — beides noch zu bauen. **Relevant:** GDPR Art. 2(2)(c) Haushaltsausnahme + CH revDSG bleiben nur tragfähig, solange Capture-Inhalte das Homelab nicht verlassen. **Time-bound: Offen — nach der CAS-Abgabe.** Bis dahin als bewusste, dokumentierte Lücke geführt (Cloud-Inferenz im Demo/Default).

NfA-P2: KI-Transparenz (AI Act Art. 50)

Category: Privacy & Compliance

Status (Block 5): ZIEL — noch nicht umgesetzt. Es gibt aktuell **keine** `ClassificationSource` / `ClassifiedAt` / `ConfidenceScore` -Spalten, **keine** EF-Migration, **kein** AI-Badge in der UI und **keinen** Test `LifecycleBadgeTests.AiClassified_ShowsAiBadge`. Die Kriterien unten sind die Abnahmedefinition für die spätere Umsetzung, nicht der Abgabestand.

Statement (Ziel): Jeder Capture, dessen `MatchedSkill` durch eine LLM-gestützte Klassifikation gesetzt wurde, soll in der UI sichtbar als "AI-classified" gekennzeichnet sein und die Klassifikations-Provenienz tragen (`ClassificationSource`, `ClassifiedAt`, optional `ConfidenceScore`). **Measurable (Abnahmedefinition, offen):**

1. UI: bUnit-Test `LifecycleBadgeTests.AiClassified_ShowsAiBadge` rendert das KI-Badge. *(Test noch nicht implementiert)*
2. Datenmodell: `Capture` hat Spalten `ClassificationSource` (enum: `None` | `Heuristic` | `AI` | `Manual`), `ClassifiedAt`, `ConfidenceScore`; EF-Migration committed. *(noch nicht vorhanden)*
3. API: `GET /api/v1/captures/{id}` gibt diese Felder zurück. *(noch nicht vorhanden)* **Achievable:** Der vorhandene `LifecycleBadge`-Component wird um einen AI-Zweig erweitert; die `IClassifier`-Implementierung in `FlowHub.AI` setzt `ClassificationSource` — beides noch zu bauen. **Relevant:** EU AI Act Art. 50 verpflichtet Betreiber von KI-Systemen, KI-Output erkennbar zu machen. **Time-bound: Offen — nach der CAS-Abgabe.** Bis dahin als bewusste, dokumentierte Lücke geführt.

Acceptance criteria

Source: `docs/spec/acceptance-criteria.md`

FlowHub — Acceptance Criteria

Consolidated acceptance criteria for FlowHub's user-facing and system-facing behavior. Each criterion is testable and maps to a specific use case in `docs/spec/use-cases.md` and to the test, runbook, or live evidence that verifies it.

This document is the single source of truth for "what counts as done" per feature. The individual `Akzeptanzkriterien:` blocks inside `use-cases.md` remain authoritative per UC; this file aggregates them for review, grading, and submission.

Conventions

- **AC-XX-N** — acceptance criterion N for UC-XX.
 - **Verified by** — concrete artifact: a test class, `make` target, runbook step, or Bruno collection. If a criterion is currently deferred (Block 2 stub state, future block), it is marked **(deferred)** with the block where it will be verified.
 - Non-functional requirements live in `docs/spec/nfa.md` as the SMART-decomposed `NfA-*` catalogue (the successor to the `NF-01 ... NF-13` quality table in `docs/spec/use-cases.md` ; see the mapping note there). They are not duplicated here.
-

1. Capture submission (UI + API)

UC-01 — Quick capture via Web UI AppBar

ID	Criterion	Verified by
AC-01-1	After typing content and pressing Enter, a <code>Captured</code> ✓ snackbar appears within 2 s.	Playwright <code>HappyFlowTests.QuickCapture_TodoEntry_AppearsInCapturesList</code>
AC-01-2	The captured row appears on <code>/captures</code> and the detail page renders the same content.	Same Playwright test
AC-01-3	Empty content → no row created; inline hint visible.	bUnit <code>AppBarQuickCaptureTests</code>

UC-02 — Long-form capture via `/captures/new`

ID	Criterion	Verified by
AC-02-1	<code>/captures/new</code> renders a multi-line text area and a Skill dropdown.	bUnit <code>NewCapturePageTests</code>
AC-02-2	Submitting non-empty content yields a 201-equivalent service call and a success snackbar; the form clears for the next entry.	bUnit <code>NewCapturePageTests</code>
AC-02-3	If <code>ISkillRegistry.ListAsync</code> throws, the dropdown is disabled and a "Let AI decide" placeholder is forced.	bUnit <code>NewCapturePageTests</code>

UC-03 — Capture via Telegram

ID	Criterion	Verified by
AC-03-1	Bot reply confirms receipt within 5 s of the user message.	(deferred — Block 3+)
AC-03-2	Capture appears in <code>/captures</code> with <code>source = Telegram</code> and the bot user's identifier in metadata.	(deferred — Block 3+)
AC-03-3	Telegram-originated submissions inherit the same validation and pipeline flow as the API path (UC-08).	Covered by UC-08 ACs once the Telegram module lands.

UC-08 — Submit a Capture via REST API

ID	Criterion	Verified by
AC-08-1	POST /api/v1/captures with a valid body returns 201 within NF-09 (p95 < 200 ms server-side).	just smoke-prod step [5/6]; tests/FlowHub.Api.IntegrationTests/
AC-08-2	Response body matches the Capture schema and Location: /api/v1/captures/{id} header is present.	tests/FlowHub.Api.IntegrationTests/CapturesEnc
AC-08-3	Missing content / unknown source → 400 ValidationProblem (type = validation.md).	tests/FlowHub.Api.IntegrationTests/CapturesEnc
AC-08-4	Bruno collection bruno/captures/submit-capture.bru round-trips against a live stack.	just smoke-prod ; manual Bruno run

2. Async pipeline (classification + routing)

UC-09 — AI-classify and route a Capture

ID	Criterion	Verified by
AC-09-1	A URL-content Capture reaches <code>Completed</code> with <code>MatchedSkill = "Wallabag"</code> and a non-empty <code>ExternalRef</code> .	<code>Skills.ContractTests ; just test-beta</code>
AC-09-2	A todo-content Capture reaches <code>Completed</code> with <code>MatchedSkill = "Vikunja"</code> and a non-empty <code>ExternalRef</code> .	<code>Skills.ContractTests ; just test-beta</code>
AC-09-3	MassTransit harness tests prove the three consumer hops fire in order (Created → Classified → Routed).	<code>tests/FlowHub.Web.ComponentTests/Pipeline/*</code> (6 harness tests)

UC-10 — Graceful AI-classifier fallback

ID	Criterion	Verified by
AC-10-1	With <code>Ai__OpenRouter__ApiKey</code> deliberately invalid, a Capture submitted via UC-08 still reaches <code>Classified</code> stage via <code>KeywordClassifier</code> .	<code>tests/FlowHub.Web.ComponentTests/Ai/AiC</code>
AC-10-2	EventId <code>3010</code> <code>AiClassifierFellBackToKeyword</code> is logged at Warning level with the exception type and Capture ID.	Same test class
AC-10-3	<code>AiClassifier</code> never throws to its caller — <code>ClassifyAsync</code> always returns a <code>ClassificationResult</code> .	Same test class

UC-11 — Retry a failed Capture

ID	Criterion	Verified by
AC-11-1	<code>POST</code> <code>/api/v1/captures/{id}/retry</code> on an Orphan Capture returns 202 Accepted with <code>stage = Raw</code> , <code>failureReason = null</code> .	<code>tests/FlowHub.Api.IntegrationTests/Captur</code>
AC-11-2	Same call on a <code>Completed</code> Capture returns 409 (<code>type = capture-not-retryable.md</code>).	Same test class
AC-11-3	Same call with an unknown id returns 404 (<code>type = capture-not-found.md</code>).	Same test class

3. Browse, filter, search

UC-04 — Dashboard

ID	Criterion	Verified by
AC-04-1	The Dashboard renders within NF-08 budget (<code>p95 < 1.5 s</code> first paint on a 100-capture DB).	Manual perf check; future Playwright timing
AC-04-2	Empty DB shows the empty-state message.	bUnit <code>DashboardTests</code>
AC-04-3	Mixed-state DB shows non-zero Orphan / Unhandled counts in the "Needs Attention" widget.	bUnit <code>DashboardTests</code>

UC-05 — Browse and filter Captures

ID	Criterion	Verified by
AC-05-1	Deep link <code>/captures?lc=Orphan</code> pre-selects the Orphan chip and shows only Orphan rows.	bUnit <code>CapturesListPageTests</code>
AC-05-2	Combining a lifecycle chip + a channel chip applies an AND filter.	bUnit <code>CapturesListPageTests</code>
AC-05-3	The grid pages at 10/25/50 rows (configurable, default 10).	bUnit <code>CapturesListPageTests</code>

UC-06 — Inspect a failed Capture

ID	Criterion	Verified by
AC-06-1	Detail page shows the <code>FailureReason</code> of an Orphan and "No Skill matched" for an Unhandled.	bUnit <code>CaptureDetailPageTests</code>
AC-06-2	"Retry" button is enabled only for Orphan / Unhandled stages (UC-11 implements the action).	bUnit <code>CaptureDetailPageTests</code>
AC-06-3	Block 2 stubs trigger a "Coming in Block 3" snackbar instead of mutating state.	bUnit <code>CaptureDetailPageTests</code> (historical — superseded by UC-11 in Block 3)

UC-12 — Filter by Lifecycle Stage

ID	Criterion	Verified by
AC-12-1	Selecting "Orphan" returns only orphaned captures.	<code>tests/FlowHub.Persistence.Tests/CaptureRepositoryTests.cs</code>
AC-12-2	Selecting multiple stages returns the union.	Same test class

UC-13 — Filter by Tag

ID	Criterion	Verified by
AC-13-1	A capture tagged "dotnet" appears when filter is "dotnet".	CaptureRepositoryTests
AC-13-2	Captures without the tag are excluded.	Same

UC-14 — Search by Content or Title

ID	Criterion	Verified by
AC-14-1	Searching "hexagonal" returns captures whose Content or Title contains that substring (case-insensitive ILIKE).	CaptureRepositoryTests
AC-14-2	Full-text search is explicitly deferred — substring match is the documented current behavior.	ADR-0005 §"Search scope"

UC-18 — Semantic Search

ID	Criterion	Verified
AC-18-1	POST /api/v1/captures followed by polling Captures.Embedding populates the column within 30 s when Embeddings__ApiKey is set (actual ~2 s).	just smoke-prod step [6/6]
AC-18-2	GET /api/v1/captures/search?q=... with a non-empty q returns 200 with an array body when at least one embedded Capture exists.	tests/FlowHub.Api.IntegrationTests/9
AC-18-3	Empty q returns 400 ValidationProblem; with Embeddings__ApiKey unset, returns 503 ProblemDetails.	Same test class
AC-18-4	POST /api/v1/admin/embeddings/rebuild returns 200 { processed, skipped, failed } when keys are present, 503 otherwise.	Same test class

4. Skill & Integration observability

UC-07 — Health pages

ID	Criterion	Verified by
AC-07-1	<code>/skills</code> renders one row per registered Skill with status badge (healthy/degraded/down).	bUnit <code>SkillsPageTests</code>
AC-07-2	<code>/integrations</code> renders one row per <code>ISkillIntegration</code> adapter with <code>LastWriteAt</code> + <code>LastWriteDurationMs</code> .	bUnit <code>IntegrationsPageTests</code>
AC-07-3	Each row exposes a "history" expansion showing recent samples (UC-16).	Same

UC-15 — Skill-Run history per Capture

ID	Criterion	Verified by
AC-15-1	A capture that was routed twice shows two <code>SkillRun</code> entries ordered by <code>StartedAt</code> DESC.	<code>tests/FlowHub.Persistence.Tests/SkillRunRepositoryTests.cs</code>

UC-16 — Integration health history

ID	Criterion	Verified by
AC-16-1	Shows <code>SampledAt</code> , <code>Status</code> , and <code>DurationMs</code> for each sample.	<code>IntegrationRepositoryTests</code>

5. Deployment

UC-17 — Deploy via Docker Compose

ID	Criterion	Verified by
AC-17-1	<code>docker compose up --build -d --wait</code> returns exit 0 with all <code>service_healthy</code> dependencies satisfied.	<code>just smoke-prod</code> step [1/6]
AC-17-2	<code>flowhub.migrations</code> container reaches <code>service_completed_successfully</code> with exit code 0.	<code>just smoke-prod</code> step [1/6]
AC-17-3	<code>GET /health/live</code> from inside the compose network returns 200 within NF-D3 (30 s) of container start.	<code>just smoke-prod</code> step [2/6]
AC-17-4	<code>GET /metrics</code> returns a Prometheus exposition containing at least one <code>^dotnet_</code> and one <code>^http_</code> series.	<code>just smoke-prod</code> step [3/6]

Coverage summary

Use case	ACs defined	Implemented & verified	Deferred / historical
UC-01 quick capture	3	3	0
UC-02 long-form capture	3	3	0
UC-03 Telegram	3	0	3 (post-Block-3)
UC-04 dashboard	3	3	0
UC-05 browse/filter	3	3	0
UC-06 failed-capture detail	3	2	1 (Block-2 stub)
UC-07 health pages	3	3	0
UC-08 REST API submit	4	4	0
UC-09 async pipeline	3	3	0
UC-10 AI fallback	3	3	0
UC-11 retry	3	3	0
UC-12 filter by stage	2	2	0
UC-13 filter by tag	2	2	0
UC-14 substring search	2	2	0
UC-15 skill-run history	1	1	0
UC-16 integration history	1	1	0

Use case	ACs defined	Implemented & verified	Deferred / historical
UC-17 docker compose	4	4	0
UC-18 semantic search	4	4	0
Total	50	46	4

Deferred ACs are all Block-3+ scope (Telegram channel implementation) or superseded historical stubs — none represent regressions or unverified production behavior.

Persistence-Layer Coverage (Block 4)

The persistence layer underlies multiple user-facing UCs. The table below maps each Block-4-relevant UC to the Persistence test class that owns its data-access ACs, so the persistence acceptance criteria can be traced as a coherent set rather than scattered across the per-UC sections above.

UC	What persistence has to deliver	Verified by (tests/FlowHub.Persisten
UC-05 — Browse/filter Captures	Cursor pagination (CreatedAt DESC, Id DESC); .Include(Tags) round-trip; no N+1	EfCaptureRepositoryTests
UC-09 — Async pipeline (Classified stage)	MarkClassifiedAsync persists MatchedSkill , Title , VikunjaProject atomically	EfCaptureServiceTests
UC-11 — Retry a failed Capture	Stage reset (Failed → Raw); idempotent re-publish; no duplicate SkillRun rows	EfCaptureServiceTests , EfSkillRunRepositoryTests
UC-12 — Filter by Lifecycle Stage	Single-stage and multi-stage union filters on IX_Captures_Stage	EfCaptureRepositoryTests (A 12-2)
UC-13 — Filter by Tag	M-N join via CaptureTag ; tag-name filter	EfCaptureRepositoryTests (A 13-2)
UC-14 — Substring search	PostgreSQL ILIKE over Content and Title	EfCaptureRepositoryTests (A 14-2)
UC-15 — Skill-Run history per Capture	SkillRun rows ordered by StartedAt DESC ; one row per skill invocation	EfSkillRunRepositoryTests (
UC-16 — Integration health history	Latest-per-integration sample query on (IntegrationName, SampledAt DESC) index	EfIntegrationHealthServiceT
UC-18 — Semantic search	pgvector vector(1024) column + HNSW index; FromSqlRaw with float-array literal	EfCaptureRepositoryTests.Se

UC	What persistence has to deliver	Verified by (tests/FlowHub.Persisten
All migrations	0001_Initial through 0008_AddVikunjaProjectToCapture apply cleanly on an empty DB	MigrationSmokeTest

All Persistence tests run against real PostgreSQL 17 via Testcontainers — no in-memory provider drift. The 29 tests in the project are listed by class in `docs/insights/block-4.md` § "Persistence-Layer Coverage".

Cross-references

- Use cases (authoritative per-UC narrative): `docs/spec/use-cases.md`
- Non-functional requirements: `docs/spec/nfa.md`
- Testing strategy and tooling: `docs/spec/testing-strategy.md`
- ADRs: `docs/adr/0001..0006` (ADR 0005 = Persistence design)
- Smoke-test runbook: `just smoke-prod` (defined in `justfile`)
- Beta-MVP operator runbook: `docs/runbooks/beta-mvp-acceptance.md`
- Block-4 insights (test counts + per-class coverage): `docs/insights/block-4.md`

Testing strategy


Source: `docs/spec/testing-strategy.md`

FlowHub — Testing Strategy

Overview

FlowHub uses a **layered testing approach** where each layer targets a specific feedback speed and confidence level. The strategy is designed so that most issues are caught early (in fast, isolated tests) and only integration-level concerns require the full stack.

Test Layers



Layer	Framework	What it tests	When to run	Block
Unit	xUnit + FluentAssertions + NSubstitute	Domain types, service logic, validators, pure functions	Every build (just test)	Block 1 (all)
Component (bUnit)	bUnit + MudBlazor.Services	Razor components render correctly, props/events wire up, states (loading/empty/error) display properly	Every build (just test)	Block 2 (all)
Integration	WebApplicationFactory + Testcontainers (PostgreSQL)	Full HTTP pipeline, DI composition, auth middleware, EF Core queries, API endpoints	Before merge / CI	Block 3 (all, API, 4 v, DB, ba)
E2E	Playwright	Browser-level: page navigation, SignalR circuit, visual layout, cross-page flows	Before release / CI	Block 4 (w/ de)

Current State (Block 3 — through Slice C)

Reading the test counts across this submission. Different documents quote different totals because each is a **per-block, per-filter snapshot**, not a contradiction. The suite grew block by block, and the number depends on which `dotnet test` filter is applied:

Where	Block	Filter	Cour
Block 1/2 Nachbereitung	1-2	bUnit components only	31
This section / Block 3	3	Category!=AI (default suite)	99
Block 4 Nachbereitung (docs/insights/block-4.md)	4	FullyQualifiedName!~E2ETests (incl. AI + integration)	223 (6 skippe
Block 5 Nachbereitung (docs/insights/block-5.md)	5	Category!=AI&!=BetaSmoke&!=E2E (offline default suite)	171 (C skippe

Canonical figures at submission (Block 5): 171 offline tests green (the default CI suite) and **223** green when AI + live-service integration tests are included (E2E excluded). The lower historical numbers are earlier snapshots, retained in the per-block Nachbereitungen as a progress record.

What's implemented

- **99 default-suite tests** passing on `just test` (excluding `Category=AI`):
 - **82 component tests** in `tests/FlowHub.Web.ComponentTests/` covering Razor components, service stubs, classification, and the MassTransit pipeline
 - **17 API integration tests** in `tests/FlowHub.Api.IntegrationTests/` covering REST endpoints via `WebApplicationFactory<Program>`
- **Slice A (REST API):** 17 integration tests via `Microsoft.AspNetCore.Mvc.Testing` covering `/api/v1/captures` endpoints:
 - `SubmitCaptureTests` — POST with validation (valid, empty content, invalid skill)
 - `ListCapturesTests` — GET with cursor pagination, stage/source filters
 - `GetCaptureByIdTests` — GET single capture, 404 not found
 - `RetryCaptureTests` — POST retry with stage validation
 - `ProblemDetailsFormatTests` — RFC 9457 wire format validation
 - `SmokeTests` — end-to-end captures workflow

- **Slice B (Async Pipeline):** 6 MassTransit Test Harness tests covering event consumers:
 - `CaptureEnrichmentConsumerTests` — 2 tests (enrich keywords, handle invalid skill)
 - `SkillRoutingConsumerTests` — 2 tests (route by matched skill, handle orphan)
 - `LifecycleFaultObserverTests` — 2 tests (catch pipeline exceptions, log with EventId)
- **Slice B (Keywords):** 5 `KeywordClassifier` tests covering deterministic classification logic
- **Block-2 Regression:** 16 `CaptureServiceStub` tests verifying backward compatibility:
 - Submit flow (new capture, appends to list, rejects empty)
 - GetRecent / GetFailureCounts (Block-2 smoke test data)
 - Lifecycle transitions (Classified, Routed, Orphan, Unhandled)
 - List with filters and cursor pagination
 - Reset for retry
- **Slice C (AI Integration):** 25 mocked unit tests + 4 trait-gated live tests shipped in `tests/FlowHub.Web.ComponentTests/Ai/` and `tests/FlowHub.AI.IntegrationTests/`:
 - **3 AiPrompts tests** — `BuildMessages` structure (system-prompt first, user-message second) + system-prompt language guard (no German routing tokens)
 - **3 AiClassificationResponse tests** — `Description` / `AllowedValues` / nullable `Title` attributes for schema generation
 - **10 AiClassifier tests** — success path + 3 argument-forwarding tests (cancellation token, `MaxOutputTokens=300`, `Temperature=0.2`) + 5 fallback paths (`HttpRequestException`, `TaskCanceledException`, `JsonException`, schema-violation guard, `InvalidOperationException`) + `EventId 3010` logging contract
 - **8 AddFlowHubAi registration tests** — D8 matrix: no-provider, Anthropic without key, OpenRouter without key, Anthropic with key, OpenRouter with key, invalid value, case-insensitive parse, model override

- **1 KeywordClassifier test** — `ClassifyAsync_AnyContent_KeywordClassifierReturnsNullTitle` locks the `Title=null` contract
- **4 trait-gated live integration tests** in `tests/FlowHub.AI.IntegrationTests/` — Anthropic Haiku 4.5 (URL→Wallabag, todo→Vikunja) + OpenRouter Llama 3.1 (URL→Wallabag, todo→Vikunja)

What's deferred

Test type	Why deferred	Lands in
E2E tests (Playwright)	No deployed environment; SignalR testing needs a running browser	Block 5
Integration tests with real DB (EF Core + Testcontainers)	Database persistence module not yet wired; <code>CaptureServiceStub</code> remains in-memory	Block 4
Load/performance tests	Single-operator system, no scaling pressure	Not planned

Slice-C-specific patterns

Mocking `IChatClient` with `NSubstitute`

`Microsoft.Extensions.AI.IChatClient` defines both instance methods and extension methods. Extension methods (e.g. `GetResponseAsync<T>`) cannot be intercepted directly by `NSubstitute`. Instead, tests mock the underlying instance method that the extension calls internally:

```
var chat = Substitute.For<IChatClient>();
chat.GetResponseAsync(
    Arg.Any<IEnumerable<ChatMessage>>(),
    Arg.Any<ChatOptions?>(),
    Arg.Any<CancellationToken>())
    .Returns(new ChatResponse(new ChatMessage(ChatRole.Assistant,
        JsonSerializer.Serialize(new { tags = new[] { "x" }, matched_skill
```

`AiClassifier` calls the typed extension

`GetResponseAsync<AiClassificationResponse>(messages, options, ct)`, which internally invokes the instance method `GetResponseAsync(messages, options, ct)`. `NSubstitute` mocks the instance method; the typed extension's parser then sees the mock-supplied JSON and deserializes it.

Trait-gated live AI tests with `[SkippableFact]`

Live integration tests that require real API keys use the `[SkippableFact]` attribute (from `Xunit.SkippableFact` 1.4.13). Tests skip conditionally when environment variables are unset:

```
[Trait("Category", "AI")]
public sealed class AnthropicHaikuLiveTests
{
    private static IClassifier BuildClassifier()
    {
        var apiKey = Environment.GetEnvironmentVariable("Ai__Anthropic__Api
Skip.If(string.IsNullOrEmpty(apiKey), "Ai__Anthropic__ApiKey n
// ... build IClassifier from real configuration
    }

    [SkippableFact]
    public async Task ClassifyAsync_UrlContent_LiveAnthropicReturnsWallabag
    {
        var sut = BuildClassifier();
        var result = await sut.ClassifyAsync("https://en.wikipedia.org/wiki
result.MatchedSkill.Should().Be("Wallabag");
        result.Title.Should().NotNullOrWhiteSpace();
    }
}
```

- `just test` excludes the `Category=AI` trait and all live tests remain **Skipped** (not silently passed)
- `just test-ai` runs only that trait — intended for operator-on-demand validation when keys are loaded
- CI runs `just test` only; live tests are operator-local

Trait filtering strategy

Tests use `[Trait("Category", "AI")]` to partition:

- **Default suite** (`Category!=AI`): 99 tests, runs every build, no external dependencies
- **Live AI suite** (`Category=AI`): 4 tests, runs on-demand, requires real API credentials

Test Naming Convention

```
MethodName_StateUnderTest_ExpectedBehavior
```

Examples:

- `Render_Loading_ShowsSkeletons_WhenCountsIsNull`
- `SubmitAsync_RejectsEmptyContent`
- `Dashboard_NeedsAttention_ShowsOrphanAndUnhandledCounts`

The `CA1707` analyzer rule (no underscores in method names) is suppressed for test projects via `<NoWarn>CA1707</NoWarn>` in the test `.csproj`, since the [CLAUDE.md](#) convention requires underscores.

Test Tools

Tool	Version	Purpose
xUnit	2.9.3	Test framework + runner
bUnit	1.40.0	Blazor component test host — renders components in isolation without a browser
FluentAssertions	6.12.2	Readable assertion syntax (<code>cut.Markup.Should().Contain(...)</code>)
NSubstitute	5.3.0	Interface mocking for service dependencies
Bogus	35.6.1	Test data generation (used in stub services, not directly in tests)
MudBlazor.Services	(via MudBlazor 8.5.1)	Required in test DI for MudBlazor components to render
Xunit.SkippableFact	1.4.13	[<code>SkippableFact</code>] + <code>Skip.If(...)</code> — conditionally skip live AI tests when API keys are absent

bUnit Test Patterns

Basic component rendering

```
var cut = RenderComponent<NeedsAttentionCard>(p =>
    p.Add(c => c.Counts, new FailureCounts(3, 1)));

cut.Markup.Should().Contain("3 orphan captures");
```


Mocked service injection

```
var captureService = Substitute.For<ICaptureService>();
Services.AddSingleton(captureService);
var cut = RenderComponent<Dashboard>();
```

Real Bogus stub injection (smoke tests)

```
Services.AddSingleton<ICaptureService>(new CaptureServiceStub());
var cut = RenderComponent<Dashboard>();
// Verifies against realistic seeded data
```

MudBlazor popover requirement

Components using `MudSelect` or `MudMenu` require `MudPopoverProvider` pre-rendered:

```
RenderComponent<MudPopoverProvider>(); // in constructor
```

EventCallback verification

```
var clicked = false;
var cut = RenderComponent<NeedsAttentionCard>(p =>
{
    p.Add(c => c.Counts, new FailureCounts(2, 0));
    p.Add(c => c.OnOrphanClick, EventCallback.Factory.Create(this, () => cl
});
cut.FindAll("button").First(b => b.TextContent.Contains("orphan")).Click();
clicked.Should().BeTrue();
```

Acceptance Criteria Format

Each page's wireframe defines states that serve as implicit acceptance criteria:

State	What to verify
Default	Data renders, correct values from service
Loading	Skeletons or progress indicator shown
Empty	Friendly message + call-to-action
Error	Alert with error message + Retry button
Not found	Alert for invalid ID / missing resource

The smoke tests (`SmokeTests.cs`) automate verification of these states against the Bogus stubs for all 6 pages.

Running Tests

```
just test           # run all tests (99 default-suite, excl. AI)
just test-ai       # run AI trait-gated tests only (requires Ai__*__ApiKey
just test-watch    # watch mode for component tests
just build         # build also catches analyzer warnings (TreatWarningsAsE
```

Persistence Layer Testing

Why Testcontainers Instead of InMemory

The `FlowHub.Persistence.Tests` project previously used `Microsoft.EntityFrameworkCore.InMemory`. Starting in Block 4, all persistence integration tests use Testcontainers PostgreSQL for these reasons:

1. **Provider parity:** InMemory does not enforce FK constraints, does not support `EF.Functions.ILike`, and differs in edge cases for cursor pagination. A test that passes against InMemory can fail against PostgreSQL.
2. **Realistic query plans:** Index-backed queries behave differently under InMemory (no query optimizer). Tests against real PostgreSQL catch missing indexes early.
3. **Migration coverage:** InMemory doesn't apply migrations.
`MigrationSmokeTest` verifies all migrations apply cleanly against a real engine.

Test Infrastructure

`PostgresFixture` (`tests/FlowHub.Persistence.Tests/Fixtures/PostgresFixture.cs`) starts one PostgreSQL container per test session (shared via `[Collection("Postgres")]`). Each test gets its own database (`testdb_<guid>`) via `CreateFreshDbAsync()` , which creates the database and applies all migrations before returning the context.

Test Pyramid for Persistence

Level	Location	Provider	Co (Bl 4
Unit (service orchestration)	<code>FlowHub.Web.ComponentTests</code> + <code>FlowHub.Persistence.Tests</code>	NSubstitute mocks	~20
Integration (repository)	<code>FlowHub.Persistence.Tests</code>	Testcontainers PostgreSQL	15
Integration (API endpoints)	<code>FlowHub.Api.IntegrationTests</code>	WebApplicationFactory	17
E2E (browser)	Playwright (planned)	Full stack	plan Bloc

Acceptance Criteria (UC-09 to UC-13)

Use Case	Test	Location
UC-09 Filter by Stage	ListAsync_FiltersByStage	EfCaptureRepositoryTests
UC-10 Filter by Tag	ListAsync_FiltersByTag (add in Block 5)	—
UC-11 Search by Content	ListAsync_FiltersBySearchTerm (add in Block 5)	—
UC-12 SkillRun history	GetByCaptureIdAsync	EfSkillRunRepository (add test in Block 5)
UC-13 Integration health history	GetRecentSamplesAsync	EfIntegrationRepository (add test in Block 5)

E2E Testing Plan (Playwright)

End-to-end tests cover the happy path for each major UI flow. Not yet automated in CI (deferred to post-submission); manual test plan is documented here.

Test Cases

Scenario	Steps	Expected
Submit Capture via UI	Navigate to <code>/</code> , type content in Quick Capture field, click Submit	Capture appears in recent captures list on Dashboard
View Captures List	Navigate to <code>/captures</code>	Table shows paginated Capture list; stage filter works
Retry Orphan Capture	Navigate to <code>/captures/{id}</code> , click Retry	Capture moves to Raw stage; toast confirms
Semantic Search (API)	<code>GET /api/v1/captures/search?q=EF+Core+query</code>	Returns Captures with semantically related content

CI Test Results Reference

Test results from the `ci.yml` GitHub Actions workflow are uploaded as `.trx` artifacts on every run. Screenshot the latest passing run for inclusion in the submission PDF:

1. Go to `github.com/freaxnx01/FlowHub-CAS-AISE/actions` .
2. Open the latest `CI` run on `main` .
3. Expand `Test` step to see test count and pass/fail breakdown.
4. Download `test-results` artifact for `.trx` file with full details.

Database model — entities

Source: `docs/design/db/entities.md`

FlowHub — DB Model Sketch (Block 4 preview)


- **Status:** Mixed — `Capture` shipped via Beta MVP (2026-05-04); `Skill`, `Integration`, `CaptureTag` remain sketches for Block 4
- **Date:** 2026-05-03 (initial sketch), 2026-05-04 (Beta MVP alignment)
- **Block:** Block 3 (Services) — Nachbereitung · DB-model rubric deliverable; carried into Block 4 (Persistence)
- **Lands as:** EF Core entities + migrations in `source/FlowHub.Persistence/`. `Capture` entity is live; remaining entities scheduled for Block 4 (PVA 2026-05-23).

This document was originally a forward-looking sketch authored during Block 3 Nachbereitung. The Beta MVP slice (`docs/superpowers/specs/2026-05-04-beta-mvp-design.md`, ADR 0005) landed the `Capture` aggregate ahead of the formal Block-4 phase, so the Capture-related rows below now reflect what actually shipped. `Skill` / `Integration` / `CaptureTag` rows remain design sketches and will be re-litigated in Block 4 — column types, lengths, and constraints may evolve before implementation.

Implementation status (2026-05-04):

- ☐ `Capture` — shipped as `CaptureEntity` in `source/FlowHub.Persistence/Entities/CaptureEntity.cs`, with EF Core configuration in `FlowHubDbContext.OnModelCreating`. Initial migration: `Migrations/20260504120638_Initial.cs`.
- ☐ `Skill`, `Integration`, `CaptureTag` — sketch only; Block 4 implementation pending.
- ☐ Outbox / idempotency receiver — explicitly deferred (ADR 0003 §"Open items"); not in Block 4 scope.

ER overview



Entity catalogue

Capture (aggregate root) — shipped via Beta MVP

Column	Type	Constraints	Notes
Id	uuid	PK, generated by app (<code>Guid.NewGuid()</code>)	maps to <code>Guid Id</code> in <code>FlowHub</code>
Content	text	NOT NULL	freeform; can be a URL, pl
Source	varchar(32)	NOT NULL	<code>ChannelKind</code> persisted as <code>Source</code> does not generate CHECK values are constrained by <code>EfCaptureService.SubmitAs</code> and <code>Enum.Parse<ChannelKin</code> add a CHECK on the Postg
Stage	varchar(32)	NOT NULL	<code>LifecycleStage</code> persisted Raw/Classified/Routed/Con Same enum-string pattern
CreatedAt	timestampz	NOT NULL	UTC; maps to <code>DateTimeOffs</code> <code>EfCaptureService.SubmitAs</code> (<code>DateTimeOffset.UtcNow</code>).
MatchedSkill	varchar(64)	NULL	set by classifier; null until to <code>Skill.Name</code> in the Bet exist yet. App-level integri
FailureReason	text	NULL	populated on Orphan/Unha chars per ADR 0003 §"Ope <code>LifecycleFaultObserver</code> sc level.
Title	varchar(512)	NULL	populated by <code>AiClassifier</code> <code>KeywordClassifier</code> returns initial 255-char sketch to 5 AI-generated titles without
ExternalRef	varchar(256)	NULL	set on <code>MarkCompletedAsync</code> downstream system's ider

Column	Type	Constraints	Notes
			4711 , Vikunja task id 777 page can later link back. B UI's Capture Detail Metadata

Skill (read/status snapshot) — ☐ sketch only, Block 4 implementation

Column	Type	Constraints	Notes
Name	varchar(64)	PK	skill identifier, e.g. "Wallabag", "Vikunja"
Status	varchar(16)	NOT NULL	HealthStatus (Unknown/Healthy/Degraded/Dc stored as string
RoutedToday	int	NOT NULL, DEFAULT 0	counter reset daily; reset strategy (cron vs event) is a Block 4 decision
LastReset	timestampz	NULL	tracks when RoutedToday was last zeroed — preliminary, Block 4 decision

Integration (read/status snapshot) — ☐ sketch only, Block 4 implementation

Column	Type	Constraints	Notes
Name	varchar(64)	PK	integration identifier, e.g. "Wallabag"
Status	varchar(16)	NOT NULL	HealthStatus stored as string
LastWriteAt	timestamptz	NULL	maps to DateTimeOffset? LastWriteAt
LastWriteDurationMs	bigint	NULL	TimeSpan? LastWriteDuration stored as milliseconds; reconstructed in app

CaptureTag (join table — Block 4 decision) — ☐ sketch only

Column	Type	Constraints	Notes
CaptureId	uuid	PK (composite), FK -> Capture.Id CASCADE DELETE	
Tag	varchar(64)	PK (composite)	normalized lowercase; source: ClassificationResult.Tags

This join table is the straightforward option. The alternative — a `text[]` PostgreSQL array column on `Capture` — is simpler to query within a single row but less indexable. **Choice is deferred to Block 4.**

Indexes

Shipped (Beta MVP)

Index	Column(s)	Rationale
IX_Captures_Stage	Capture.Stage	Dashboard "needs attention" count + lifecycle filter chips on /captures
IX_Captures_CreatedAt_DESC	Capture.CreatedAt DESC	Recent-captures list + cursor pagination on /captures

Planned (Block 4)

Index	Column(s)	Rationale
ix_capture_matched_skill	Capture.MatchedSkill	RoutedToday aggregation per skill; join to Skill table once Skill entity ships
ix_capture_source	Capture.Source	GET /api/v1/captures? source= filter (D4 in api-surface.md)
ix_capture_tag_tag	CaptureTag.Tag	Tag-based filtering if surfaced in Block 4/5; depends on CaptureTag join-table decision

Lifecycle states (recap)

Stage	Meaning	Transitions to
Raw	just submitted	Classified (via <code>CaptureEnrichmentConsumer</code>) or Orphan
Classified	classifier returned skill match	Routed (via <code>SkillRoutingConsumer</code>) or Orphan
Routed	sent to skill integration (in-flight)	Completed (success) or Unhandled (failure)
Completed	terminal — integration write succeeded	—
Orphan	terminal — no skill match or classification failure	reassigned manually
Unhandled	terminal — integration exhausted retries	retried manually via <code>POST /api/v1/captures/{id}/retry</code>

What is deliberately deferred

- **Audit fields** (`UpdatedAt` , `Version` / `RowVersion` for optimistic concurrency) — Block 4
- **AI audit fields** (`AiProvider` , `AiModel` , `AiDurationMs` , `WasFallback`) per AI integration spec §"Block 4 (Persistence)" — Block 4; earns rubric points for documented test results
- **Tag storage strategy** — join table (`CaptureTag`) vs PostgreSQL `text[]` array column — Block 4 decision
- **`RoutedToday` reset mechanism** — scheduled cron job vs. event-driven vs. computed column — Block 4
- **Outbox pattern table** — `MassTransit.EntityFrameworkOutbox` integration deferred per ADR 0003 §"Outbox / Idempotency"
- **Idempotency receiver** — at-least-once redelivery guard (lookup by `CaptureId` before mutating) — ADR 0003 §"Open items"

- **Soft-delete / retention** — FlowHub is append-only; no delete in v1 ([api-surface.md](#) D6)
 - **Full-text search** — `tsvector` column on `Content` if KI-Suche lands in Block 5
-

References


- `source/FlowHub.Core/Captures/Capture.cs` — aggregate root (actual field list)
- `source/FlowHub.Persistence/Entities/CaptureEntity.cs` — shipped EF Core entity (Beta MVP, internal sealed)
- `source/FlowHub.Persistence/FlowHubDbContext.cs` — shipped DbContext + indexes
- `source/FlowHub.Persistence/Migrations/20260504120638_Initial.cs` — shipped initial migration
- ADR 0005 (`docs/adr/0005-persistence.md`) — Persistence ADR
- `source/FlowHub.Core/Health/SkillHealth.cs` , `IntegrationHealth.cs` — health snapshot records
- `source/FlowHub.Core/Classification/ClassificationResult.cs` — Tags + MatchedSkill + Title
- `docs/design/api/api-surface.md` §"Data model (reminder, from FlowHub.Core)"
- ADR 0002 (`docs/adr/0002-service-architecture-and-async-communication.md`) — modular monolith, single DB
- ADR 0003 (`docs/adr/0003-async-pipeline.md`) — async pipeline, outbox deferral, FailureReason sanitization
- AI integration design spec (`docs/superpowers/specs/2026-05-03-slice-c-ai-integration-design.md`) §"Block 4 (Persistence)" — AI audit fields

Database model — ER diagram

Source: `docs/design/db/er.md`

FlowHub Entity-Relationship Diagram

Schema lineage: Block 4 introduced the base relational model (Captures, Channels, Skills, SkillRuns, Integrations, IntegrationHealthSamples, Tags). Block 5 added the pgvector `Embedding` column on Captures and the HNSW index for semantic search (migration `0004_AddEmbedding`).



Legende

The diagram is a Mermaid `erDiagram` (Chen-style entity-relationship). Read it together with the FK Strategy table below — soft FKs are deliberately **not drawn** because they have no DB-level constraint.

Symbol / column marker	Meaning
Entity { ... }	Database table. Box title is the table name (always plural in this schema).
type Field PK	Column belonging to the primary key. Composite PKs list every member with PK .
type Field FK	Column carrying a hard foreign-key constraint (enforced by PostgreSQL).
type Field (no marker)	Regular column. For soft references (Capture.Source , Capture.MatchedSkill), this is intentional — see FK Strategy below.
"..." after a column	Inline note (e.g., the Block-5 Embedding column carries its provider note here).
`A	
◻ `	one-and-only-one cardinality on A 's side.
o{	zero-or-more cardinality on B 's side.
(no line drawn)	Soft reference — application-level only, no DB FK. Listed only in the FK Strategy table.

varchar(N) , text , uuid , timestamptz , bool , bigint , vector(N) are PostgreSQL types; int reads as int4 .

FK Strategy

Relationship	Type	Reason
Capture.Source → Channel.Name	Soft (no DB FK)	Channels can be deregistered without orphan failures
Capture.MatchedSkill → Skill.Name	Soft (no DB FK)	Consistent with Beta MVP pattern
SkillRun.SkillName → Skill.Name	Hard (RESTRICT)	SkillRun is audit trail; Skill must exist
SkillRun.CaptureId → Capture.Id	Hard (CASCADE)	Run is meaningless without its Capture
IntegrationHealthSample.IntegrationName → Integration.Name	Hard (CASCADE)	Sample is meaningless without its Integration
Tag.CaptureId → Capture.Id	Hard (CASCADE)	Tag is owned by Capture

Delete Strategy

FlowHub uses **hard delete** for owned entities (via the FK CASCADE rules in the table above). There is no `IsDeleted` column or query filter.

Soft-delete semantics are carried by `LifecycleStage` instead:

Concern	Mechanism
"Failed Capture I might retry"	<code>LifecycleStage.Orphan</code> + <code>FailureReason</code> — Capture stays in the table, surfaces in the Dashboard "Needs Attention" widget, retryable via <code>POST /api/v1/captures/{id}/retry</code> .
"Capture didn't match any skill"	<code>LifecycleStage.Unhandled</code> — same persistence story, different operator action (Assign Skill).
"Capture done"	<code>LifecycleStage.Completed</code> — terminal, stays for history.
"I actually want this row gone"	<code>DELETE</code> via API or direct SQL — CASCADE removes Tags + SkillRuns; SkillRuns retain audit value but only for the lifetime of the Capture they describe.

This is a deliberate Block-4 decision (ADR 0005 §6 area): the operational concerns soft-delete usually solves (retry, hide, undo) are already covered by the lifecycle state machine, so an `IsDeleted` flag would be ceremony without benefit. If a future requirement needs "undelete" of a hard-deleted Capture, soft-delete can be reintroduced as a column without breaking the existing API.

Vector Search (Block 5)

Column	Type	Index	
<code>Captures.Embedding</code>	<code>vector(1024)</code> (pgvector)	<code>captures_embedding_hnsw_idx</code> (HNSW, <code>vector_cosine_ops</code>)	Populate asynchronously via Mistral embed. No captures embed keyword

Design — Perspektiven (Struktur, Verhalten, Interaktion)

Source: `docs/design/perspectives.md`

FlowHub Architecture — Three Perspectives

This document maps FlowHub's architecture to the three classical perspectives required by the CAS AISE rubric (*Struktur · Verhalten · Interaktion*) and points to the canonical artefact for each.

Perspective	Question it answers	Canonical artefacts
Struktur (Structure)	<i>Welche Bausteine gibt es und wie sind sie geschichtet?</i>	Modular monolith layout, layered modules, deployment topology
Verhalten (Behaviour)	<i>Wie verändert sich der Zustand der zentralen Domänenobjekte über die Zeit?</i>	Capture lifecycle state machine, MassTransit retry semantics
Interaktion (Interaction)	<i>Wie spielen die Bausteine zur Laufzeit zusammen?</i>	End-to-end sequence flows, API contracts, async pipeline traces


1. Struktur (Structure)

One sentence: FlowHub is a flat-layout **Modular Monolith** in .NET 10 (per ADR 0001) deployed as six independent containers under docker-compose (per ADR 0002).

Module map

```
source/
├── FlowHub.Core/           ← domain types + driving ports (Capture, Skill,
├── FlowHub.Web/           ← Blazor Web App (Interactive Server, MudBlazor)
├── FlowHub.Api/           ← Minimal API endpoint definitions (CaptureEndpo
├── FlowHub.AI/            ← AiClassifier + AiEmbeddingService over Microso
├── FlowHub.Persistence/   ← EF Core + PostgreSQL + pgvector adapter for IC
├── FlowHub.Skills/        ← Wallabag + Vikunja ISkillIntegration adapters
└── FlowHub.Integrations/  ← (placeholder)
```

Deployment topology (`docker-compose.yml`)




References

- ADR 0001 — Frontend Render Mode & Architecture (Q1 = flat layout)
- `docs/projektbeschreibung/FlowHub_Architecture-v2.svg` — bildliche Architektur
- `docs/projektbeschreibung/FlowHub_Projektbeschreibung_v4.md` §6 — Systemarchitektur (Überblick + Hybrid Skill-System)

2. Verhalten (Behaviour)

One sentence: Every Capture moves through a six-state lifecycle driven by the async pipeline (per ADR 0003), with explicit terminal states for happy-path success, retryable failure, and "no matching skill" outcomes.

Capture lifecycle state machine



Stage	Bedeutung	Eingang	Ausgang
Raw	Just arrived, no classification yet	Submit / Retry	→ Classified, → Unhandled
Classified	AI/Keyword has assigned a target skill	CaptureEnrichmentConsumer	→ Routed
Routed	Skill processing in flight	SkillRoutingConsumer dispatched	→ Completed, → Orphan
Completed	Happy terminal state — write succeeded, ExternalRef set	Integration write returned 2xx	terminal
Orphan	Skill / integration failed (retryable)	All MassTransit retries exhausted; LifecycleFaultObserver sets stage	→ Raw via Retry
Unhandled	No matching skill found (retryable)	Classifier returned "" for MatchedSkill	→ Raw via Retry

Retry / error semantics

- **Per-consumer retry** (ADR 0003): CaptureEnrichmentConsumer retries with intervals [100ms, 500ms]. CaptureEmbeddingConsumer + SkillRoutingConsumer retry with [500ms, 2s, 5s].
- **Fault observer:** LifecycleFaultObserver consumes Fault<TMessage> and flips the Capture to Orphan with a FailureReason. No retry on the fault itself (would loop).
- **Embedding pipeline is best-effort:** AiEmbeddingService catches provider errors and stores the Capture without an embedding (search degrades to non-vector path).


References

- ADR 0003 — Async Pipeline (MassTransit)
- `source/FlowHub.Core/Captures/LifecycleStage.cs` — enum definition
- `source/FlowHub.Web/Pipeline/` — consumer implementations
- `docs/spec/use-cases.md` UC-09, UC-10, UC-11 — lifecycle transitions described as user flows

3. Interaktion (Interaction)

One sentence: External actors interact with FlowHub through three channels (Web UI, Telegram bot, REST API); internally, every Capture traverses an async event-driven pipeline whose events are observable via OpenTelemetry traces and Prometheus metrics.

Hot-path sequence — Submit Capture via REST → Skill write



API surface

Channel	Endpoint / mechanism	Documented in
REST	<code>POST /api/v1/captures</code> + <code>GET /api/v1/captures{,/}{id},/search}</code> + <code>retry / admin</code>	OpenAPI doc at <code>/openapi/v1.</code> Scalar UI at <code>/scalar</code> , Bruno c in <code>bruno/</code>
Web UI	Blazor Interactive Server pages (<code>/</code> , <code>/captures</code> , <code>/captures/{id}</code> , <code>/skills</code> , <code>/integrations</code>)	<code>source/FlowHub.Web/Component</code> design wireframes in <code>docs/design/<feature>/</code>
Telegram	Bot polling via long-poll loop (placeholder Channel in current build)	<code>vault/Projektarbeit/External</code> <code>Services.md</code>
Operations	<code>/health/live</code> , <code>/metrics</code> (Prometheus), Grafana dashboards	<code>docs/spec/nfa.md</code> NfA-D3, NfA

Observability

- **Traces:** OpenTelemetry [ASP.NET](#) Core + runtime instrumentation; MEAI chat/embedding spans (`UseOpenTelemetry()` on the `IChatClient`).
- **Metrics:** Prometheus scrape at `/metrics` — `dotnet_*` + `http_*` series verified by `just smoke-prod` step [4/6].
- **Logs:** Structured Serilog → stdout (12-Factor XI); event IDs in `source/FlowHub.*/` `LoggerMessage` attributes.

References

- ADR 0002 — Service Architecture & Async Communication
- ADR 0003 — Async Pipeline (MassTransit)
- ADR 0004 — AI Integration in Services
- `docs/design/capture-detail/` — sequence diagrams for the Capture Detail flow
- `bruno/` — runnable REST request collection (one `.bru` per endpoint)

Design — Datenfluss

Source: `docs/design/data-flow.md`


Data Flow — Personendaten-Residenz & Trust Boundaries

This document is the canonical evidence artifact for **NfA-P1** (Personendaten-Residenz) in `docs/spec/nfa.md`. It shows where Capture content lives, which component is allowed to call outbound, and where the Homelab trust boundary runs. The AI-Act-Transparenz-Pflicht aus NfA-P2 wird durch das `ClassificationSource`-Feld im Sequenzdiagramm sichtbar gemacht.

Stand: Block 5 (2026-05-24). Architektur entspricht dem Beta-MVP (`docs/superpowers/specs/2026-05-04-beta-mvp-design.md`) plus Klassifikator (Block 3 Slice C, ADR 0004).

A. Trust-Boundary-Übersicht (Strukturbild)

Statisches Bild: welche Datenkategorie überquert welche Grenze. Ein Pfeil bedeutet "Daten fließen in diese Richtung über die Boundary"; gestrichelte Pfeile sind Opt-in-Pfade, die nur bei expliziter Konfiguration aktiviert werden.



Legende:


Symbol	Bedeutung
Durchgezogener Pfeil	Default-Pfad, immer aktiv
Gestrichelter Pfeil	Opt-in-Pfad, nur bei <code>Embeddings__Provider != "Local"</code>
Rahmen "Homelab Trust Boundary"	Alles innerhalb läuft auf vom Betreiber kontrollierter Hardware (Proxmox-VM, Docker-Host)
Rahmen "Cloud (Opt-in)"	Drittanbieter ausserhalb der Boundary; betreten der Boundary erzwingt Auftragsverarbeitungs-Pfad
<code>[(...)]</code>	Persistente Datenhaltung
<code>[...]</code>	Prozess / Service

Invarianten:

- Es gibt genau drei Outbound-Konnektoren aus der Homelab-Boundary heraus: Vikunja, Wallabag und (optional) Cloud-LLM. Alle drei sind im `FlowHub.Skills` - bzw. `FlowHub.AI` -Modul kapselt; der Audit-Test `OutboundCallAuditTests` verifiziert, dass kein anderer Modul-Code HTTP-Aufrufe nach extern initiiert.
- Cloud-LLM-Pfad ist **per Default deaktiviert**. Aktivierung erfordert `Embeddings__Provider=OpenAI` + `Embeddings__ApiKey` als Environment-Variablen — bewusste Konfigurationsentscheidung mit DPA-Konsequenz.
- Capture-Bodies werden **nicht** an Vikunja oder Wallabag weitergereicht; Skill-Adapter übergeben nur den minimalen Nutzlast-Slice (Tag + URL). Dies hält die Outbound-Boundary auch dann sauber, wenn ein Skill-Target später durch einen Cloud-Dienst ersetzt würde.

B. Capture-Lebenszyklus (Sequenzbild)

Dynamisches Bild eines einzelnen Captures vom Eingangskanal bis zur Skill-Routing-Terminierung. Markiert, wann `ClassificationSource = "AI"` gesetzt wird (NfA-P2-Hook) und wo die Boundary überquert wird.



Invarianten:

- **ClassificationSource** ist beim Persistieren in der Raw-Stage **None** und wird ausschliesslich vom Classifier auf **AI** (oder ggf. **Heuristic**) gesetzt. Manuelle Zuweisungen erhalten **Manual**. Das Feld ist Pflicht- spalte und nullable=false (siehe EF-Migration **AddClassificationSource**, Block 5).
- Der Capture-Body verlässt die Homelab-Boundary nur, wenn der Cloud-LLM-Pfad opt-in aktiviert ist. Skill-Outbound-Calls senden nur den minimalen Slice (Tag + URL), nicht den Body.
- Das Sequenzbild zeigt den Default-Pfad mit lokalem Ollama. Der Cloud-LLM-Variant ist eine 1:1-Substitution des **LLM**-Participants und wird in **docs/adr/0006-llm-hosting.md** separat dokumentiert.

C. Verweise

- **NfA-P1 / NfA-P2:** **docs/spec/nfa.md**
- **Vault-Stub:** **vault/Knowledge/Datenschutz-und-AI-Act.md** (Abschnitt 3.2 + 3.4)
- **ADRs:** **docs/adr/0004-classifier.md**, **docs/adr/0006-llm-hosting.md** (Draft)
- **Audit-Test:** **tests/FlowHub.Web.IntegrationTests/OutboundCallAuditTests.cs** (Block 5)
- **Skill-Routing-Detail:** **docs/design/sequences/skill-routing.md**

Design — Sequenzdiagramme (Übersicht)

Source: `docs/design/sequences/README.md`

Sequence Diagrams — Capture Pipeline

This folder contains Mermaid sequence diagrams for the two async pipeline flows in FlowHub. They document the **Verhalten** dimension of the Block-3 Nachbereitung rubric item *"Struktur / Verhalten / Interaktion — Sequence-/Activity-Diagramme"*.

Diagrams

File	Flows documented
<code>capture-enrichment.md</code>	A. AI classifier success path · B. AI fallback to <code>KeywordClassifier</code>
<code>skill-routing.md</code>	A. Integration write success · B. Retry exhaustion → <code>LifecycleFaultObserver</code>

Reading guide


Read `capture-enrichment.md` first — it covers the first pipeline hop (`CaptureCreated` → `CaptureEnrichmentConsumer` → `CaptureClassified`) and explains the `IClassifier` port and its Slice-C AI adapter. Then read `skill-routing.md` for the second hop (`CaptureClassified` → `SkillRoutingConsumer` → `Routed` | `Unhandled`) including the fault-observer recovery path.

Both files include invariant callouts that complement the structural view in [ADR 0003](#) (retry policies, fault routing, lifecycle stages) and [ADR 0004](#) (AI fallback semantics, provider abstraction). The Slice-C AI adapter design is specified in [docs/superpowers/specs/2026-05-03-slice-c-ai-integration-design.md](#).

Sequenz — Capture-Intake

Source: `docs/design/sequences/capture-intake.md`

Capture Intake — Sequence Diagram (Verhalten)



Sequenz — Capture-Enrichment

Source: `docs/design/sequences/capture-enrichment.md`

Sequence Diagrams — Capture Enrichment

These diagrams document the behaviour of `CaptureEnrichmentConsumer` and the `IClassifier` port. The consumer sits at the head of the MassTransit pipeline: it receives every `CaptureCreated` event and turns it into a classified capture (happy path) or a terminal `Orphan` state (failure paths).

Implemented (Slice C, 2026-05-03): `IClassifier` is resolved from DI at startup based on configuration. When `Ai:Provider` is not set, `AddFlowHubAi` registers `KeywordClassifier` directly — no `AiClassifier` wrapper is in the chain. When `Ai:Provider` and the matching API key are set, `AiClassifier` wraps `KeywordClassifier` as its fallback floor. The port shape (`ClassifyAsync(content, ct) → ClassificationResult`) is unchanged regardless of which implementation is active; the consumer code never changed between Slice B and Slice C.


A. Happy path — AI classifier configured, call succeeds

This path applies when `Ai:Provider` and the matching API key are set. `AiClassifier` makes one structured-output call to the provider via `GetResponseAsync<AiClassificationResponse>(...)`, validates the schema, and returns a `ClassificationResult(Tags, MatchedSkill, Title)`. The consumer writes the result back to `ICaptureService` and publishes `CaptureClassified` for the routing consumer.

Invariants:

- `MatchedSkill` is validated against `{"Wallabag", "Vikunja", ""}` even after a successful AI call — the schema constrains it, but the code also checks defensively.
- `Title` is populated only by `AiClassifier`; `KeywordClassifier` always returns `Title = null`.

- An empty `MatchedSkill` (no skill matched) is a *valid* outcome that skips `CaptureClassified` and goes straight to `MarkOrphanAsync` — no exception is raised.




B. AI fallback path — provider call fails, KeywordClassifier floor activates

Any exception from `IChatClient` (network, timeout, JSON parse, or schema-violation guard) is caught inside `AiClassifier`. The adapter logs the failure at `Warning` level and delegates to `KeywordClassifier` as a hard floor. The consumer sees a valid `ClassificationResult` in both cases — it never receives an exception from the classifier port.

Invariants:

- The fallback never rethrows — AI outage degrades classification quality, not availability.
- `Title` is always `null` in the fallback result; the capture reaches `Classified` stage without a generated title.
- The MassTransit retry budget (`Intervals(100, 500)`) is reserved for bus/infrastructure failures, not AI call failures — AI errors are handled inside `AiClassifier` before returning.
- The fault-observer path (`Fault<CaptureCreated>`) is not triggered by AI failures.



Sequenz — Skill-Routing

Source: `docs/design/sequences/skill-routing.md`

Sequence Diagrams — Skill Routing

These diagrams document the behaviour of `SkillRoutingConsumer` and its retry / fault-observer path. The consumer is the second hop in the MassTransit pipeline: it receives `CaptureClassified` events and dispatches each capture to the matching `ISkillIntegration` adapter.

Current implementation (Slice B): `ISkillIntegration` is fulfilled by `LoggingSkillIntegration` stubs (one registered as `"Wallabag"`, one as `"Vikunja"`). Each stub logs the write intent and returns successfully. Real HTTP adapters replace the stubs in Block 4/5 without any consumer changes.


Retry policy: `Intervals(500, 2000, 5000)` — three retries at 500 ms, 2 s, and 5 s. The longer intervals reflect the anticipated latency of outbound HTTP calls to real integrations in later blocks.

A. Happy path — integration registered, write succeeds

The consumer resolves `ISkillIntegration` by name from the DI-registered collection, calls `WriteAsync`, and marks the capture as `Routed`. There is no separate `Completed` transition in the current consumer code — `Routed` is the happy terminal state written by the routing consumer.

Invariants:

- Resolution is by exact, ordinal string equality on `ISkillIntegration.Name`.
- If no integration matches the `MatchedSkill` (e.g. a skill registered in the classifier but not yet wired as an adapter), the capture is marked `Unhandled` directly — this is a *synchronous* terminal path, not a fault-observer path.
- `GetByIdAsync` throws `InvalidOperationException` if the capture is missing from the store; this is treated as a pipeline bug and will exhaust the retry budget before faulting.




B. Retry exhaustion path — integration call keeps failing

When `WriteAsync` throws on every attempt, MassTransit exhausts the per-consumer retry budget (`Intervals(500, 2000, 5000)`) and publishes a `Fault<CaptureClassified>` event. `LifecycleFaultObserver` (which implements `IConsumer<Fault<CaptureClassified>>`) receives this event and marks the capture as `Unhandled` with a formatted reason string extracted from the first exception in `Fault.Exceptions` .

Invariants:


- `LifecycleFaultObserver` is best-effort: if `MarkUnhandledAsync` itself throws, the error is swallowed (logged at `Error` , `EventId 1003`) so the observer does not recurse into `Fault<Fault<T>>` . Cancellation (`OperationCanceledException`) still propagates.
- The observer is registered with *no retry policy* — recursive retry on a fault observer would loop forever.
- `Stage=Unhandled` surfaces in the Dashboard's "needs attention" tile and is queryable via `GET /api/v1/captures?stage=Unhandled` .
- A stuck `Unhandled` capture can be re-queued from that state via `POST /api/v1/captures/{id}/retry` , which re-publishes `CaptureCreated` and resets the lifecycle to `Raw` .



Sequenz — Skill-Routing (Hot Path)

Source: `docs/design/sequences/skill-routing-hot-path.md`

Skill Routing Hot Path — Sequence Diagram (Verhalten + Interaktion)



Design — User Journeys

Source: `docs/design/journeys.md`

FlowHub — User Journeys

Single source of truth for end-to-end user journeys across the Web UI. Derived from the approved Phase-2 flow diagrams in `docs/design/<feature>/flow.md`.

Each journey has the JSON acceptance-criteria shape used by Playwright E2E tests:

```
{
  "id": "<short-slug>",
  "category": "functional | error | edge",
  "description": "<one-line user-facing intent>",
  "entryUrl": "<route the test starts on>",
  "preconditions": ["..."],
  "steps": ["..."],
  "expected": ["..."],
  "passes": false
}
```

`passes` flips to `true` once the corresponding Playwright spec is green.
`preconditions` describe seeded data the test fixture must provide. `expected` are the assertions; `steps` are user actions only.

Routes (from `@page` directives):

Route	Page
/	Dashboard
/captures	Captures list
/captures/new	New Capture form
/captures/{id:guid}	Capture detail
/skills	Skills page
/integrations	Integrations page

The AppBar `QuickCaptureField` is shared across every page (see Dashboard flow Diagram 3).

Catalog

J01 — Quick capture from the AppBar (happy path)

```
{
  "id": "J01",
  "category": "functional",
  "description": "User submits a capture from the AppBar quick-capture field",
  "entryUrl": "/",
  "preconditions": ["Dev auth signed in", "ICaptureService backed by Bogus"],
  "steps": [
    "Type 'https://example.com/article' into the AppBar quick-capture field",
    "Press Enter"
  ],
  "expected": [
    "Snackbar shows 'Captured ✓' with an Open action",
    "Quick-capture input is cleared",
    "ICaptureService.SubmitAsync was called once with ChannelKind.Web"
  ],
  "passes": false
}
```

J02 — Quick capture failure surfaces an error toast

```
{
  "id": "J02",
  "category": "error",
  "description": "When the capture service throws, the user sees an error t",
  "entryUrl": "/",
  "preconditions": ["ICaptureService.SubmitAsync throws InvalidOperationException"],
  "steps": [
    "Type 'foo' into the AppBar quick-capture field",
    "Press Enter"
  ],
  "expected": [
    "Snackbar shows 'Capture failed: ...'",
    "Input still contains 'foo'",
    "isSubmitting returns to false"
  ],
  "passes": false
}
```

J03 — Empty quick-capture submit shows inline hint

```
{
  "id": "J03",
  "category": "edge",
  "description": "Submitting an empty quick-capture input does not call the",
  "entryUrl": "/",
  "preconditions": [],
  "steps": ["Press Enter in the AppBar quick-capture field without typing a"],
  "expected": [
    "Snackbar shows 'Type something first'",
    "ICaptureService.SubmitAsync is NOT called"
  ],
  "passes": false
}
```

J04 — Dashboard renders all four cards with seeded data

```
{
  "id": "J04",
  "category": "functional",
  "description": "Dashboard loads and shows Needs Attention, Recent Capture",
  "entryUrl": "/",
  "preconditions": ["Bogus stub seeds 12 captures, 6 skills, 6 integrations"],
  "steps": ["Navigate to /"],
  "expected": [
    "All four cards render (no skeletons remaining)",
    "Recent Captures shows at least one row",
    "Skill Health lists Books, Movies, Quotes",
    "Integration Health lists Wallabag, Vikunja, Obsidian"
  ],
  "passes": false
}
```

J05 — Dashboard 'orphan count' click filters Captures list

```
{
  "id": "J05",
  "category": "functional",
  "description": "Clicking the orphan count on the Dashboard navigates to /",
  "entryUrl": "/",
  "preconditions": ["Bogus stub seeds ≥1 orphan capture"],
  "steps": ["Click the 'N orphan captures' button on the Needs Attention ca"],
  "expected": [
    "URL changes to /captures?lc=Orphan",
    "Captures list pre-selects the Orphan lifecycle chip",
    "Only orphan rows are visible"
  ],
  "passes": false
}
```

J06 — Dashboard 'unhandled count' click filters Captures list

```
{
  "id": "J06",
  "category": "functional",
  "description": "Clicking the unhandled count navigates to /captures filte",
  "entryUrl": "/",
  "preconditions": ["Bogus stub seeds ≥1 unhandled capture"],
  "steps": ["Click the 'N unhandled captures' button on the Needs Attention"],
  "expected": [
    "URL changes to /captures?lc=Unhandled",
    "Captures list pre-selects the Unhandled lifecycle chip"
  ],
  "passes": false
}
```

J07 — Dashboard recent-capture row click opens detail

```
{
  "id": "J07",
  "category": "functional",
  "description": "Clicking a row in the Recent Captures card opens that cap",
  "entryUrl": "/",
  "preconditions": ["Bogus stub seeds ≥1 capture"],
  "steps": ["Click the first row in the Recent Captures grid"],
  "expected": [
    "URL changes to /captures/{id}",
    "Capture detail page renders with the matching content"
  ],
  "passes": false
}
```


J08 — Dashboard 'View all' navigates to Captures list

```
{
  "id": "J08",
  "category": "functional",
  "description": "View-all button on the Recent Captures card navigates to",
  "entryUrl": "/",
  "preconditions": [],
  "steps": ["Click the 'View all' button on the Recent Captures card"],
  "expected": ["URL changes to /captures"],
  "passes": false
}
```

J09 — Dashboard 'Manage integrations' navigates to /integrations

```
{
  "id": "J09",
  "category": "functional",
  "description": "Manage-integrations button on the Integration Health card",
  "entryUrl": "/",
  "preconditions": [],
  "steps": ["Click the 'Manage integrations' button on the Integration Heal"],
  "expected": ["URL changes to /integrations"],
  "passes": false
}
```

J10 — Captures list filters by lifecycle

```
{
  "id": "J10",
  "category": "functional",
  "description": "Selecting a lifecycle chip narrows the captures grid to t",
  "entryUrl": "/captures",
  "preconditions": ["Bogus stub seeds 12 captures across all stages"],
  "steps": ["Click the 'orphan' chip in the Lifecycle chip set"],
  "expected": [
    "Grid shows only orphan-stage rows",
    "Results counter updates to match the orphan count"
  ],
  "passes": false
}
```

J11 — Captures list filters by channel

```
{
  "id": "J11",
  "category": "functional",
  "description": "Selecting a channel chip narrows the captures grid to tha",
  "entryUrl": "/captures",
  "preconditions": ["Bogus stub seeds captures from both Web and Telegram"],
  "steps": ["Click the 'Telegram' chip in the Channel chip set"],
  "expected": [
    "Grid shows only Telegram-source rows",
    "Results counter updates"
  ],
  "passes": false
}
```

J12 — Captures list search filters by content substring

```
{
  "id": "J12",
  "category": "functional",
  "description": "Typing in the search field filters the grid by case-insen",
  "entryUrl": "/captures",
  "preconditions": ["Bogus stub seeds a capture containing 'Inception'"],
  "steps": ["Type 'incep' into the search field"],
  "expected": ["Grid shows the Inception capture", "Other captures are filt",
  "passes": false
}
```

J13 — Captures list 'no match' empty state

```
{
  "id": "J13",
  "category": "edge",
  "description": "When filters produce zero results the user sees a no-matc",
  "entryUrl": "/captures",
  "preconditions": [],
  "steps": [
    "Type 'zzzzzzzz-no-such-string' into the search field"
  ],
  "expected": [
    "Empty state shows 'No captures match'",
    "Clicking 'Clear filters' restores the full grid"
  ],
  "passes": false
}
```

J14 — Captures list row click opens detail

```
{
  "id": "J14",
  "category": "functional",
  "description": "Clicking a row in the Captures grid opens that capture's",
  "entryUrl": "/captures",
  "preconditions": ["Bogus stub seeds ≥1 capture"],
  "steps": ["Click the first row in the captures grid"],
  "expected": [
    "URL changes to /captures/{id}",
    "Capture detail page renders"
  ],
  "passes": false
}
```

J15 — Capture detail (Completed) view-only

```
{
  "id": "J15",
  "category": "functional",
  "description": "A Completed capture renders content + metadata with no fa",
  "entryUrl": "/captures/{completedId}",
  "preconditions": ["Bogus stub has a Completed capture with known id"],
  "steps": ["Navigate to /captures/{completedId}"],
  "expected": [
    "Content and Metadata sections render",
    "No 'Routing failed' alert",
    "No Retry/Reassign/Ignore buttons"
  ],
  "passes": false
}
```

J16 — Capture detail (Orphan) shows retry/reassign/ignore actions

```
{
  "id": "J16",
  "category": "functional",
  "description": "An Orphan capture shows the failure reason and offers Ret
  "entryUrl": "/captures/{orphanId}",
  "preconditions": ["Bogus stub has an Orphan capture with a known FailureR
  "steps": ["Navigate to /captures/{orphanId}"],
  "expected": [
    "Warning alert 'Routing failed: <reason>' is visible",
    "Buttons 'Retry routing', 'Reassign skill', 'Ignore' are present"
  ],
  "passes": false
}
```

J17 — Capture detail (Unhandled) shows assign/ignore actions only

```
{
  "id": "J17",
  "category": "functional",
  "description": "An Unhandled capture shows the integration failure and of
  "entryUrl": "/captures/{unhandledId}",
  "preconditions": ["Bogus stub has an Unhandled capture"],
  "steps": ["Navigate to /captures/{unhandledId}"],
  "expected": [
    "Error alert 'Skill integration failed: <reason>' is visible",
    "Button 'Assign skill' is present",
    "Button 'Retry routing' is NOT present"
  ],
  "passes": false
}
```

J18 — Capture detail action buttons stub-toast in Block 2

```
{
  "id": "J18",
  "category": "functional",
  "description": "Clicking any action on Orphan/Unhandled detail surfaces t",
  "entryUrl": "/captures/{orphanId}",
  "preconditions": ["Bogus stub has an Orphan capture"],
  "steps": ["Click 'Retry routing'"],
  "expected": [
    "Snackbar contains 'will work once backend Skills are wired in Block 3'
  ],
  "passes": false
}
```

J19 — Capture detail back link returns to list

```
{
  "id": "J19",
  "category": "functional",
  "description": "The 'Back to Captures' link returns the user to /captures",
  "entryUrl": "/captures/{anyId}",
  "preconditions": [],
  "steps": ["Click '← Back to Captures'"],
  "expected": ["URL changes to /captures"],
  "passes": false
}
```

J20 — Capture detail unknown id shows not-found

```
{
  "id": "J20",
  "category": "edge",
  "description": "Visiting /captures/{unknown-guid} shows a not-found warni
  "entryUrl": "/captures/00000000-0000-0000-0000-000000000000",
  "preconditions": [],
  "steps": ["Navigate to /captures/{unknownId}"],
  "expected": ["Warning alert 'Capture not found' is visible"],
  "passes": false
}
```

J21 — New Capture submit (happy path)

```
{
  "id": "J21",
  "category": "functional",
  "description": "Submitting the New Capture form creates a capture and sho
  "entryUrl": "/captures/new",
  "preconditions": ["ISkillRegistry returns a non-empty list"],
  "steps": [
    "Type 'Read this article tomorrow' into the Content field",
    "Click Submit"
  ],
  "expected": [
    "Snackbar shows 'Captured ✓ – open' with an Open action",
    "Form is cleared (content empty, dropdown back to 'Let AI decide')"
  ],
  "passes": false
}
```

J22 — New Capture validation blocks empty submit

```
{
  "id": "J22",
  "category": "edge",
  "description": "Submitting with empty Content shows a validation error an",
  "entryUrl": "/captures/new",
  "preconditions": [],
  "steps": ["Click Submit without typing anything"],
  "expected": [
    "Required-field error 'Content is required' is shown",
    "ICaptureService.SubmitAsync is NOT called"
  ],
  "passes": false
}
```

J23 — New Capture: skills failed to load still allows submission

```
{
  "id": "J23",
  "category": "error",
  "description": "When ISkillRegistry throws, the form shows a warning but",
  "entryUrl": "/captures/new",
  "preconditions": ["ISkillRegistry.GetHealthAsync throws"],
  "steps": [
    "Navigate to /captures/new",
    "Type 'still works' into Content",
    "Click Submit"
  ],
  "expected": [
    "Inline warning 'Could not load skills' is visible",
    "Submit succeeds (snackbar 'Captured ✓')"
  ],
  "passes": false
}
```


J24 — New Capture cancel returns to dashboard

```
{
  "id": "J24",
  "category": "functional",
  "description": "The Cancel button on the New Capture form navigates back",
  "entryUrl": "/captures/new",
  "preconditions": [],
  "steps": ["Click Cancel"],
  "expected": ["URL changes to /"],
  "passes": false
}
```

J25 — Skills page lists every registered skill

```
{
  "id": "J25",
  "category": "functional",
  "description": "/skills shows a row per skill with name, status, and rout",
  "entryUrl": "/skills",
  "preconditions": ["ISkillRegistry returns the 6 seeded skills"],
  "steps": ["Navigate to /skills"],
  "expected": [
    "Grid contains Books, Movies, Articles, Quotes, Knowledge, Belege",
    "Each row shows a HealthDot and a status label"
  ],
  "passes": false
}
```

J26 — Skills page load-failure shows retryable error

```
{
  "id": "J26",
  "category": "error",
  "description": "When the skill registry throws, /skills shows an error al",
  "entryUrl": "/skills",
  "preconditions": ["ISkillRegistry.GetHealthAsync throws"],
  "steps": ["Navigate to /skills"],
  "expected": [
    "Error alert 'Could not load skills: <reason>' is visible",
    "Retry button is present"
  ],
  "passes": false
}
```

J27 — Integrations page lists every wired integration

```
{
  "id": "J27",
  "category": "functional",
  "description": "/integrations shows a row per integration with name, stat",
  "entryUrl": "/integrations",
  "preconditions": ["IIntegrationHealthService returns the 4 seeded integra"],
  "steps": ["Navigate to /integrations"],
  "expected": [
    "Grid contains Wallabag, Vikunja, Paperless, Authentik",
    "Each row shows a HealthDot and a status label"
  ],
  "passes": false
}
```

J28 — Integrations page load-failure shows retryable error

```
{
  "id": "J28",
  "category": "error",
  "description": "When the health service throws, /integrations shows an er",
  "entryUrl": "/integrations",
  "preconditions": ["IIntegrationHealthService.GetHealthAsync throws"],
  "steps": ["Navigate to /integrations"],
  "expected": [
    "Error alert 'Could not load integrations: <reason>' is visible",
    "Retry button is present"
  ],
  "passes": false
}
```

Summary

Surface	Journeys
AppBar / QuickCapture	J01, J02, J03
Dashboard	J04, J05, J06, J07, J08, J09
Captures list	J10, J11, J12, J13, J14
Capture detail	J15, J16, J17, J18, J19, J20
New Capture	J21, J22, J23, J24
Skills	J25, J26
Integrations	J27, J28

E2E status (snapshot)

All 28 journeys have a Playwright spec + JSON sidecar under `tests/FlowHub.Web.E2ETests/Journeys/`. Live result against `just watch` + the docker-compose backing services:

Status	Count	Journeys
☐ Green	28	J01-J28 — full suite green

Closed in this iteration

- ☐ Added `Immediate="true"` to `QuickCaptureField`'s `MudTextField` — fixes J01 (and `HappyFlowTests`). `bUnit` test updated to use `input.Input(...)` instead of `input.Change(...)`.
- ☐ Added a `#blazor-circuit-ready` sentinel to `MainLayout` (rendered only after `OnAfterRender(firstRender)`), and `JourneyTestBase.GotoAsync` waits for it via `WaitForSelectorState.Attached` — deterministic replacement for the previous fixed 2 s sleep.
- ☐ `JourneyTestBase.GotoAsync` also waits (best-effort, 3 s cap) for all `.mud-skeleton` placeholders to clear — proves `OnInitializedAsync` data load completed and event handlers on data-derived elements are bound.
- ☐ Row-click selectors changed from `.mud-table-row` (also matches `MudDataGrid` header) to `tbody .mud-table-row` — closes J14, J19.
- ☐ J13 assertion uses a `WaitForFunction` accepting either the no-match panel text OR a body-row count of zero.
- ☐ **Same parallel-DbContext bug as the earlier CaptureDetail fix surfaced in `Dashboard.LoadAsync`** — four services hitting the shared scoped EF context via `Task.WhenAll` silently threw "second operation on context instance", so all four cards rendered as null forever. Switched to sequential awaits → unblocked J05, J06, J07.
- ☐ J05/J06 URL regex made case-insensitive — Dashboard navigates to lowercase `?lc=orphan` / `?lc=unhandled`.
- ☐ **Investigation: not a wiring bug, just an empty DB.** EF queries on `Skills` / `Integrations` were correctly returning 0 rows. Added migration `0005_SeedSkillsAndIntegrations` with `HasData(...)` for the 6 baseline skills (Books, Movies, Articles, Quotes, Knowledge, Belege) and 4 baseline integrations (Wallabag, Vikunja, Paperless, Authentik). Legitimate catalog values, not test fixtures — they ship with every deployment. Closes J25, J27.
- ☐ **J15 via test-side fixture, not seed data.** Added `E2EDbHelpers.UpsertCompletedCaptureAsync` (idempotent INSERT ... ON

CONFLICT UPDATE) that the J15 spec calls before navigating. Stable Guid 1111...1111 . Direct Npgsql write — no dev-only product API endpoint, no fake Completed capture polluting the migration. The helper uses the same connection string as FlowHubDbContextFactory and is overridable via FLOWHUB_E2E_DB .

10. **J26 + J28 via opt-in fault injection.** FlowHub.Web/Testing/ adds a tiny IFaultInjector singleton plus decorator wrappers around ISkillRegistry and IIntegrationHealthService . Both decorators and the POST /test/faults/{name}/{arm|disarm} endpoints register only when FLOWHUB_E2E_FAULTS_ENABLED=true is set in the environment; in any other deployment the namespace is dead code. The env var is set in docker-compose.override.yml (dev/demo overlay) so the local stack ships ready-to-use; production compose files don't set it. Specs arm in setup, disarm in finally so a failing assertion can't poison subsequent tests.
11. **Legacy HappyFlowTests removed.** Investigation pinned the intermittent full-suite flake on the bootstrap fixture (not J07 as the earlier note suggested). HappyFlowTests used WaitUntilState.NetworkIdle and a bare tr selector — exactly the patterns we'd already replaced for the journey suite. Since J01 (QuickCapture submit) + J14 (captures-list row click) + J19 (back navigation) cover the same end-to-end flow with the deterministic sentinel + tbody .mud-table-row selector, the legacy fixture was redundant. Five consecutive full-suite runs pass at 28/28.

Still open

(none — full journey suite is stable.)

bUnit cross-coverage (tests/FlowHub.Web.ComponentTests/): still 126/126 green. The bUnit suite owns the component-level negative paths (forced exceptions, error alerts, validation messages) for J02, J23, J26, J28 — the E2E equivalents can't trigger those states against a healthy live system without a fault-injection hook.

AI usage

Source: `docs/ai-usage.md`

AI Tool Usage — FlowHub

Living document for the rubric criterion "Wurden KI-unterstützende Werkzeuge verwendet und deren Nutzung beschrieben (12)" (highest-weighted single item, see `vault/0rganisation/Bewertungskriterien.md`). Updated as work happens — not as a one-shot at submission.

Tools in use

Tool	Purpose	Where it shows up
Claude Code (Opus 4.7, 1M context)	Brainstorming, plan writing, ADR drafting, controller for subagent dispatches	<code>docs/superpowers/specs/</code> , <code>docs/superpowers/plans/</code> , <code>docs/adr/</code> , this file
Claude Sonnet 4.6 (subagents)	Implementer + spec reviewer + code-quality reviewer subagents under the superpowers SDD workflow	All <code>source/</code> and <code>tests/</code> changes from Block 3 Slice B
GitHub Copilot	Inline suggestions during editing	Sparingly — Claude Code drives sessions end-to-end
ChatGPT	Ad-hoc concept clarification, side checks	Only when Claude is mid-task on something else

Workflow used in Block 3

The Block 3 Slice-B work (async pipeline) ran end-to-end through the **superpowers** plugin's structured workflow:

1. `superpowers:brainstorming` — locked 13 design decisions via A/B/C trade-off questions; output is `docs/superpowers/specs/2026-04-30-async-pipeline-design.md`.
2. `superpowers:writing-plans` — produced the 16-task TDD-driven implementation plan at `docs/superpowers/plans/2026-04-30-async-pipeline.md`.
3. `superpowers:subagent-driven-development` — for each task: dispatch a fresh implementer subagent → spec-compliance review subagent → code-quality review subagent → mark complete in the controller's TaskList.
4. `superpowers:using-git-worktrees` — isolated all implementation work in `.worktrees/block3-async-pipeline/` on branch `feat/block3-async-pipeline`, leaving `main` untouched until the branch is reviewed.

Block 3 Slice B — async pipeline

Brainstorming + spec writing

Conversational design via Claude Code's brainstorming skill. ~13 decisions surfaced as A/B/C questions; reviewer (the human) corrected scope choice from "Slice A first" (REST API) to "Slice B (ADR 0003 first — pin scope before coding)" after a technical discussion of MassTransit's tradeoffs and the genuine value vs. overhead of a queue in a single-user system.

The user explicitly delegated the remaining design decisions during a 1.5-hour break ("pick recommendations and give a short summary on return"); Claude completed the spec, committed it locally, and presented a decision summary for review.

Final spec self-reviewed by Claude before commit (placeholder scan, internal consistency, ambiguity check, scope check).

Implementation plan

Authored by Claude Code via the `writing-plans` skill from the approved spec. 16 tasks total, bite-sized TDD ordering: packages → events → ports → adapters → service mark-methods → consumers (enrichment, routing, fault) → Program.cs wiring → ADR 0003 → [ai-usage.md](#) → docker-compose sketch → vault checklist.

Implementation execution

Subagent-driven. Pattern per task:

- Dispatch implementer subagent (general-purpose, Sonnet 4.6 for TDD/judgment, Haiku for mechanical) with the full task text + context — never make the subagent read the plan file.
- Implementer reports DONE / DONE_WITH_CONCERNS / BLOCKED / NEEDS_CONTEXT.
- Dispatch spec-compliance reviewer subagent — reads actual code, compares to plan, returns `[]` / `[]` with file:line citations.
- If `[]` spec, dispatch `superpowers:code-reviewer` agent — returns Strengths / Issues (Critical / Important / Minor) / Assessment.
- Apply minor fix-ups inline (controller-side) for trivial issues; surface critical/important ones back to the implementer.

Notable adaptations the implementers caught (real value, not hallucinations)

- `MassTransit.Testing` **is not a separate NuGet package** in MassTransit 8.5.9 — `AddMassTransitTestHarness()` lives in the main `MassTransit` package. The plan was wrong; the implementer corrected the package list and the test-project reference.
- **Captive-dependency trap on `IPublishEndpoint`** — MassTransit registers `IPublishEndpoint` as Scoped; `CaptureServiceStub` is Singleton. Plain `AddSingleton<ICaptureService, CaptureServiceStub>()` would fail at startup. Implementer surfaced the issue, applied the canonical fix (factory using `IBus`), and propagated the same pattern to `Program.cs`.
- **`harness.Consumed.Any<T>(predicate, TimeSpan)` overload** doesn't exist in MassTransit 8.5.9 — implementer used `using var cts = new CancellationTokenSource(TimeSpan.FromSeconds(5))` and passed `cts.Token` instead.
- **`appsettings.Development.json` is in `.gitignore`** (12-factor compliance per [CLAUDE.md](#)). The first T11 commit force-added it; the implementer flagged the issue in their report, and a follow-up commit removed it. The `Program.cs` `else` branch already defaults to `UsingInMemory`, making the file redundant.
- **Pre-existing `CaptureServiceStubTests.cs`** had 4 Block-2 regression tests not mentioned in the plan. Implementer merged the 6 new tests with the existing 4 (rather than overwriting) and updated `CapturesTests.cs` + `SmokeTests.cs` to pass an `IPublishEndpoint` substitute to the new constructor.

Generated vs. handwritten share (estimate, Slice B only)

Artifact	AI-drafted	Human-edited
Brainstorming spec	~95%	~5% (decisions, scope, factual corrections)
Implementation plan	~95%	~5% (verifying packages, paths)
ADR 0003	~95%	TBC after review
Production code (consumers, classifier, registration)	~95%	~5% (constructor scope fixes during review)
Tests	~95%	0% (TDD-first; tests were author-once)

The 5% human input is high-value: scope correction, package-name correction, factory-pattern surfacing, 12-factor compliance enforcement.

Reflexion — what worked, what didn't

□ What worked

- The brainstorming-skill A/B/C question format forced explicit decision-making instead of hand-waving. 13 decisions written down, each with rationale.
- Subagent-driven development per task gave clean isolation: each implementer started with a fresh context, the controller curated exactly what they needed, and the two-stage review (spec then quality) caught real issues twice.
- LoggerMessage source-gen caught CA1848/CA1873 analyzer rules early; the team established an EventId namespacing convention (Pipeline 1000–1999, Skills 2000–2999) before the third consumer landed.
- The smoke-test gate (`just run` + `curl /` returning 200) revealed the captive-dependency issue immediately when it would have been silent in a unit-test-only run.

⚠ What needed correction

- The plan was authored from incomplete repository knowledge (missed reading `CaptureServiceStubTests.cs`, the existing `ChannelKind` values, the

`MassTransit.Testing` package layout). Implementers had to adapt mid-task. Lesson: brainstorming + plan authoring should explicitly grep/read every file the plan references.

- The first T11 commit force-added a gitignored file. The plan didn't catch this; the implementer's self-review did. Lesson: any time the agent runs `git add -f`, that's a red flag worth surfacing in the dispatch prompt.
- The first T8 implementer used a verbose 50-line `NoopPublishEndpoint` hand-rolled mock when `Substitute.For<IPublishEndpoint>()` was already available. The plan's fallback note was missed. Lesson: spell out the simpler alternative inline, not as a footnote.

□ What didn't work / where humans had to intervene

- The .NET SDK pin in `global.json` (10.0.201) didn't match the installed SDK (10.0.104). The agent installed 10.0.201 to `~/dotnet` and symlinked it into `~/local/bin/dotnet`. Worked, but added a non-trivial environment change the user has to know about; documented in `~/bashrc` with a removal note.
- A mid-stream `git history` rewrite on `main` (PII scrub) invalidated all in-flight commit SHAs. The user had to coordinate the stop/restart manually. Workflow worked: stop after current todo, document HEAD + remaining tasks, resume on the rewritten branch using commit messages as identity.

Block 3 Slice A — REST API

Brainstorming + spec writing

Conversational design via Claude Code's brainstorming skill. The session built on the `docs/design/api/api-surface.md` sketch produced during Block 3 Vorbereitung (D1–D6 already locked). Slice A added 5 new decisions (D7–D11): scope sizing (captures-only vs. full surface), versioning strategy (`/api/v1/` prefix), cursor pagination format, ProblemDetails conformance level, and test project topology.

The user picked option β (Captures-only scope) over the full 7-endpoint surface (α) and the no-skills-endpoint variant (γ), explicitly to free implementation time for Slices C and D. Skills and integrations endpoints deferred to Slice D.

Final spec self-reviewed by Claude before commit: placeholder scan, internal consistency check, ambiguity check, scope check.

Implementation plan

Authored by Claude Code via the `writing-plans` skill from the approved spec. 15 tasks total, TDD-ordered: new project scaffold → endpoint stubs → validators → cursor pagination → list async → GET by id → retry handler → JSON enum config → test project → integration tests per endpoint → OpenAPI/Scalar wiring → problem details → MassTransit IBus injection → vault/ai-usage docs → final test pass.

Implementation execution

Subagent-driven. 15 tasks dispatched sequentially; Haiku for mechanical tasks (T1 scaffolding, T4 project registration, T13 vault checklist, T14 problem docs); Sonnet 4.6 for all TDD, judgment-heavy, and integration tasks.

New project `source/FlowHub.Api/` co-hosted via project reference in `FlowHub.Web`; new test project `tests/FlowHub.Api.IntegrationTests/` bootstrapped via `WebApplicationFactory<Program>`. Cursor format: hand-rolled URL-safe base64 of JSON `(CreatedAt, Id)` — no [ASP.NET](#) dependency introduced into `FlowHub.Core`.

Notable adaptations the implementers caught

- **JsonStringEnumConverter configuration:** Minimal API uses integer enum serialization by default; the T8 implementer added `ConfigureHttpJsonOptions` with `JsonStringEnumConverter` server-side, plus matching `JsonSerializerOptions` in test deserialization. This wasn't pre-specified in the plan.
- **Captures namespace shadow:** creating `tests/FlowHub.Web.ComponentTests/Captures/` for `CaptureCursorTests` shadowed the `FlowHub.Web.Components.Pages.Captures` Blazor component referenced in `CapturesTests.cs`. T5 implementer fixed via a `using CapturesPage =` alias at the top of the test file.
- **Consumer race in retry handler:** synthesizing the response capture from the post-reset state (rather than re-querying after `ResetForRetryAsync`) avoids a flaky test where the Slice-B in-memory consumer reclassifies the capture before the response is built (T11).
- **MassTransit package reference:** T11 implementer added `<PackageReference Include="MassTransit" />` to `FlowHub.Api.csproj` for `IBus` injection — the package version was already centrally pinned in `Directory.Packages.props` from Slice B, so no version conflict arose.

- **Missing** `using Microsoft.AspNetCore.Http;` : the T2 endpoint snippet in the plan omitted this import; the T2 implementer caught and added it before the build could fail.

Generated vs. handwritten share (estimate, Slice A only)

Artifact	AI-drafted	Human-edited
Spec	~95%	~5% (scope choice $\alpha/\beta/\gamma$ + cursor format choice)
Plan	~95%	~5% (verifying paths)
Production code (endpoints, validators, cursor, ListAsync)	~95%	~5% (constructor / namespace fixes during review)
Tests (integration via WebApplicationFactory)	~95%	0%

Reflexion — what worked, what didn't

What worked

- The api-surface sketch from Block 3 Vorbereitung removed ~80% of the design work for Slice A — the brainstorming session reduced to scope sizing and gap-filling 5 new decisions rather than designing from scratch.
- Implementer subagents caught real .NET defaults issues (`JsonStringEnumConverter`, namespace shadow) that the plan didn't anticipate — the two-stage review (spec-compliance then code-quality) continued to justify its overhead.
- The `WebApplicationFactory`-based test project gave end-to-end confidence that the Minimal API routes, validators, and JSON serialization all compose correctly in the real DI container — no mocking of infrastructure.

What needed correction

- Plan defects: the T2 endpoint snippet was missing `using Microsoft.AspNetCore.Http;`, and the JSON enum string-converter plumbing wasn't pre-specified. Both were caught by implementers mid-task. Lesson:

plan authoring should compile-check inline snippets mentally against the existing `_Imports` / `using` surface.

- Cursor format was underspecified in the initial plan draft (no mention of URL-safe base64 vs. plain base64). The T5 implementer picked URL-safe; the plan was updated retroactively.

What didn't work / blockers

N/A — no blockers; subagent dispatches went straight through without escalation.

Block 3 Slice C — AI integration (production-runtime)

Slice C is the first time AI moves from *development tool* to *production-runtime component*. The classifier port `IClassifier` (introduced in Slice B as a hexagonal seam with a deterministic `KeywordClassifier` adapter) now has a second adapter, `AiClassifier`, that calls a real LLM in the request path of the enrichment consumer.

Brainstorming + spec writing

Conversational design via Claude Code's brainstorming skill. 10 decisions surfaced as A/B/C questions covering: scope (classifier-only vs. classifier+title vs. multi- feature enrichment), provider mix (which two adapters), abstraction layer (`Microsoft.Extensions.AI` vs. Semantic Kernel), port shape (extend `ClassificationResult` vs. sibling `IEnricher`), failure semantics (graceful fallback vs. hard fail), default models, structured-output strategy, no-key startup behaviour, test layering (mocked unit + trait-gated live), and provider selection.

User picked the .NET-native MEAI abstraction over Semantic Kernel (Q3) and the graceful fallback over hard-fail (Q5) — both decisions explicitly motivated by the rubric phrasing *"intelligente und flexible Services"* (graceful = flexible).

Final spec self-reviewed by Claude before commit (placeholder scan, internal consistency, ambiguity check, scope check).

Implementation plan

Authored by Claude Code via the `writing-plans` skill from the approved spec. 15 tasks total, TDD-ordered: packages → `ClassificationResult` extension → [FlowHub.AI](#)

csproj → AiPrompts → AiClassificationResponse DTO → AiClassifier (10 unit tests) → AddFlowHubAi (8 unit tests covering D8 matrix) → Program.cs wiring → integration test project → 4 trait-gated live tests → justfile filter → [ai-usage.md](#) → ADR 0004 → [CLAUDE.md](#) placeholder note → vault checklist + CHANGELOG + final pass.

Implementation execution

Subagent-driven. Sonnet 4.6 for TDD/judgment-heavy tasks (T2, T4–T10); Haiku for mechanical tasks (T1 packages, T3 csproj scaffolding, T9 integration test scaffold, T11 justfile, T14 [CLAUDE.md](#) tweak, T15 vault checklist).

Production-runtime AI use

This slice is what the rubric *"KI: Wurden KI-Werkzeuge verwendet"* (max 12 pts) actually asks for: AI inside the running application, not just AI as a coding assistant. The `AiClassifier` calls a real LLM (Anthropic Haiku 4.5 by default, swappable to OpenRouter Llama 3.1 70B Instruct via one env var) on every capture that flows through the enrichment pipeline.

Cost guards: `MaxOutputTokens=300` , `Temperature=0.2` , 10s HTTP timeout. Anthropic prompt-cache marker deferred to Slice D — `Anthropic.SDK 5.10.0` exposes it only via its native API, not via the MEAI bridge.

Failure handling: any provider exception or schema violation logs `EventId 3010 AiClassifierFellBackToKeyword` at Warning and routes to the deterministic `KeywordClassifier` floor — capture is always classified.

Notable adaptations the implementers caught (real value, not hallucinations)

These are version-sensitivity catches and analyzer-rule compliance fixes, not hallucinations — they represent real value delivered by the implementer subagents when plan templates met the actual installed packages.

- **Anthropic SDK shape mismatch:** the plan's `BuildChatClient` pattern (`new AnthropicClient(apiKey).Messages.AsBuilder().Build().AsIChatClient(model)`) didn't compile against `Anthropic.SDK 5.10.0` — `MessagesEndpoint` already implements `IChatClient` directly, so `.AsIChatClient()` does not exist on it. Implementer adapted to: `new AnthropicClient(apiKey).Messages.AsBuilder().ConfigureOptions(o => o.ModelId = model).Build()` .

- **file -scoped class in test signature (CS9051):** the plan's `AiClassifierTests.cs` template placed `FakeLogger<T>` as a `file` -scoped class, but `file` -local types cannot appear in member signatures of non-file-local types. Implementer nested it as `internal sealed` private types inside `AiClassifierTests` — functionally identical, zero semantic change.
- **CA2263 under warnings-as-errors:** `AiClassificationResponseTests` used `prop.PropertyType.Should().Be(typeof(string))`, which triggers CA2263 in .NET 10's default analyzer set. Implementer swapped to `.Should().Be<string>()`.
- **CA2201 in test 9:** test 9 (`ClassifyAsync_GenericException_FallsBackToKeyword`) used `new Exception(...)`, which triggers CA2201 ("do not raise reserved exception types"). Implementer swapped to `new InvalidOperationException("anything else")` — semantically equivalent for the catch-all fallback path.
- **Explicit package references in integration-test project:** the live integration test project `FlowHub.AI.IntegrationTests` needed explicit `Microsoft.Extensions.Configuration`, `Microsoft.Extensions.DependencyInjection`, and `Microsoft.Extensions.Logging` package references. Unlike `FlowHub.Web.ComponentTests`, this project only has a project reference to `FlowHub.AI` and doesn't inherit the [ASP.NET](#) hosting surface that brings those packages in transitively.
- **GetValue<int?> without Configuration.Binder:** pulling `MaxOutputTokens` as a nullable int via `configuration.GetValue<int?>()` is only available with the `Microsoft.Extensions.Configuration.Binder` package, which wasn't in scope. Implementer used `int.TryParse(configuration["Ai:MaxOutputTokens"], out var parsed) ? parsed : 300` instead of adding another transitive dependency.
- **MEAI version resolved to 10.5.1:** the plan's fallback version guidance cited `9.4.0-preview.1.25164.4`; the package resolver pulled `10.5.1` from the feed. The `IChatClient` surface API (`GetResponseAsync`) matched — the plan called this out explicitly as the stable call path — so no code change was needed, but the implementer confirmed the method shape before proceeding.

Generated vs. handwritten share (estimate, Slice C only)

Artifact	AI-drafted	Human-edited
Brainstorming spec (10 Q&A decisions)	~95%	~5% (Q3 MEAI vs. SK choice, Q5 graceful-fallback choice)
Implementation plan (15 tasks + templates)	~95%	~5% (verifying package names, path sanity)
Production code (<code>FlowHub.AI/</code> , <code>Program.cs</code> wiring)	~85%	~15% (Anthropic SDK shape fix, analyzer-rule compliance)
Tests (25 new: unit + integration scaffolding)	~95%	~5% (CA2263/CA2201 swaps, FakeLogger nesting)
Docs (ADR 0004, ai-usage, CHANGELOG)	~90%	~10% (framing, date corrections, this Reflexion section)

The human 15% share in production code is higher than Slices A and B because the Anthropic SDK shape genuinely changed between plan authoring and execution — the implementer had to inspect the installed package and adapt the adapter construction. The remaining delta is analyzer-rule fixes that could be absorbed into plan templates if the linting baseline were codified earlier in the slice lifecycle.

Reflexion — what worked, what didn't

What worked

Subagent-driven development with the user's Sonnet/Haiku dispatch defaults moved through all 14 implementation tasks without re-dispatch loops or escalations — the plan templates were specific enough that implementers operated with narrow judgment space and reported DONE rather than NEEDS_CONTEXT. The TDD discipline was particularly well-suited to the failure-path tests: the plan enumerated all five fallback scenarios (network error, timeout, JSON parse failure, schema violation, generic exception) in advance, so the implementer just typed each case and watched it go red then green. The zero-config startup behaviour was validated immediately by the `just run` smoke test: EventId 3021 `AiProviderNotConfigured` appeared on the first boot with no

`Ai__Provider` env var set, exactly matching the dev-friendly philosophy from the spec.

What needed correction

The plan's Anthropic SDK adapter shape was approximately six months stale — the installed `Anthropic.SDK 5.10.0` had evolved its `IChatClient` integration beyond what the plan template assumed. The implementer had to inspect the installed package's public surface before adapting the construction chain. This is a predictable cost of including version-pinned SDK snippets in plan templates: they degrade as packages release. A mitigation is to keep adapter-construction snippets in a "verify before paste" comment rather than copy-paste-ready code. Additionally, four separate analyzer-rule fixes (CS9051 file-scoped class scope, CA2263 `typeof` vs. generic overload, CA2201 reserved exception type, and the `LoggerMessage` source-gen pattern) consumed implementer attention across multiple tasks. Most of these are mechanical and could be absorbed into a slice-level "linting baseline" check so plan templates compile clean on first attempt.

Honest note on live integration tests

The four trait-gated live integration tests (`just test-ai`) were not exercised in this implementation run — running them requires real Anthropic or OpenRouter API keys in the environment, making them an operator-only step. The mocked unit tests (18 of them) cover all failure paths and the full D8 registration matrix; the live tests exist as a fast sanity gate for whoever runs `just test-ai` with keys configured.

Prompts of note

(Captured here when surprising or high-leverage. Empty for now — most prompts followed standard skill conventions.)

Block 4 prep — Beta MVP (Web → AI → Wallabag/Vikunja)

Workflow

End-to-end vertical slice that landed Block-4 (persistence) and Block-5 (real Skills + Integrations) work early to validate the architecture on a real homelab demo

before continuing rubric-aligned work. Same superpowers SDD loop as Slice C.

- `superpowers:brainstorming` produced `docs/superpowers/specs/2026-05-04-beta-mvp-design.md` (D1-D8).
- `superpowers:writing-plans` produced this plan (`docs/superpowers/plans/2026-05-04-beta-mvp.md`).
- `superpowers:subagent-driven-development` drove the 21 tasks. 22 commits on `feat/beta-mvp` (21 task commits + 1 final-review fixup commit + 1 plan-write commit + 1 vault-update commit).

Branch state at end of automated execution: 138/138 tests pass under `just test` (excludes `Category=AI` and `Category=BetaSmoke`); Task 21 (operator-driven demo + push) is the only remaining item.

Notable adaptations the implementers caught (real value, not hallucinations)

These are catches the subagent dispatches made during the 21-task run. Each represents a place where the plan met real-world package behaviour or compiler/analyzer rules and the implementer adapted; the controller then either accepted, fixed up, or asked for a re-dispatch.

- **EF Core 8+ dual-provider trap (Task 10):**

`services.RemoveAll<DbContextOptions<FlowHubDbContext>>()` alone is insufficient when overriding the `DbContext` registration in `IntegrationTestFactory` . `AddDbContext` also registers `IDbContextOptionsConfiguration<TContext>` (an *internal* generic type) carrying the SQLite provider configuration. Without removing this second descriptor, the subsequent `AddDbContext(...UseInMemoryDatabase...)` leaves *both* provider configurations resolved, and EF's internal service provider throws *"two providers registered"* at first `DbContext` resolution. The implementer discovered this via test failure, then resolved `IDbContextOptionsConfiguration<FlowHubDbContext>` by name at runtime and called `RemoveAll` on it. Worth a comment in the file because the problem is silent until you swap providers.

- **Surgical `MigrationRunner` removal in tests (Task 10):** the plan flagged wholesale `RemoveAll<IHostedService>()` as a hazard (it would also drop MassTransit's bus-control hosted service, breaking consumer-pipeline tests). Implementer used surgical match `ImplementationType?.Name == "MigrationRunner"` . Final code review later upgraded this to

`ImplementationType == typeof(MigrationRunner)` (refactor-safe; `MigrationRunner` is internal but `InternalsVisibleTo` was already in place).

- **CaptureEntity + DbSet<CaptureEntity> initially public instead of internal (Task 8):** the plan deliberately specified `internal` sealed class `CaptureEntity` to keep persistence shape off the public surface — only `EfCaptureService` (same assembly) should talk to it. The Task 8 implementer ran into a wall: `EfCaptureServiceTests` directly instantiates `CaptureEntity` for seeding (`db.Captures.AddRange(new() { Id = ... })`), which doesn't compile when both the entity and the `DbSet` property are internal. The implementer's first instinct was to widen visibility to `public`. The controller caught this in spec-compliance review and dispatched a follow-up commit (`c5a97fe`) reverting to `internal` and adding `[assembly: InternalsVisibleTo("FlowHub.Persistence.Tests")]` via `Properties/AssemblyInfo.cs`. Same `InternalsVisibleTo` was extended to `FlowHub.Api.IntegrationTests` in Task 10 to support direct-entity seeding in `IntegrationTestFactory`. Lesson: when encapsulation is non-negotiable, the plan should name the visibility hatch (`InternalsVisibleTo`) explicitly so implementers don't default to widening visibility.
- **FlowHubDbContextFactory design-time discovery (Task 9):** EF tooling couldn't discover `FlowHubDbContext` for `dotnet ef migrations add Initial` because `AddFlowHubPersistence` wasn't yet wired in `Program.cs` (that's Task 10). Implementer added an `IDesignTimeDbContextFactory<FlowHubDbContext>` (internal sealed, never DI-registered) — the canonical EF design-time pattern. Not in the plan; correctly added.
- **Microsoft.EntityFrameworkCore.Design in FlowHub.Web.csproj (Task 9):** required because EF CLI resolves design-time tooling via the *startup* project, not just the project that owns the `DbContext`. Added with `PrivateAssets=all` so it doesn't leak transitively to consumers of `FlowHub.Web`.
- **CA1825 in generated migration (Task 9):** EF's migration generator emitted `new bool[0]` on one line of `20260504120638_Initial.cs`, which trips CA1825 under warnings-as-errors. Implementer swapped to `Array.Empty<bool>()` post-generation. This is recurring whenever EF generates migrations against a schema with empty-array constraints; worth absorbing into a post-generation lint pass.
- **.config/dotnet-tools.json placement (Task 9):** `dotnet new tool-manifest` writes the file to the repo root by default, not to `.config/`. Implementer manually moved it to the conventional location. Worth tracking — the canonical location is what `dotnet tool restore` looks for first.

- **Deterministic seeding via `InternalsVisibleTo` + `CaptureEntity` (Task 10):** removing the `CaptureServiceStub` broke pre-existing `FlowHub.Api.IntegrationTests` that depended on Bogus-seeded data. Implementer added direct entity-seeding (10 deterministic rows: 2 Orphan, 2 Completed, 2 Raw, 2 Classified, 1 Routed, 1 Unhandled) in `IntegrationTestFactory.CreateHost` after `base.CreateHost` returns, and added `InternalsVisibleTo("FlowHub.Api.IntegrationTests")`. The alternative — seeding through `ICaptureService.SubmitAsync` — would publish `CaptureCreated` events into the test's MassTransit harness, causing non-deterministic state at test start. Direct entity seeding is the correct call for the Beta; the trade-off (`InternalsVisibleTo` for two test projects) is accepted and documented.
- **`SkillRoutingConsumerTests` parallel-run timing flake (Task 5):** the bare `harness.Consumed.Any<CaptureClassified>()` is flaky under parallel-run load — its 3-second default timeout was too tight when all three new tests ran concurrently. Implementer added an explicit 10-second `CancellationTokenSource`, matching the existing `LifecycleFaultObserverTests` idiom. Pattern that already exists in the repo; the plan didn't surface it.
- **`LifecycleFaultObserverTests.RoutingExhaustedRetries_MarksUnhandled` API drift (Task 5):** pre-existing test still used the old `ISkillIntegration.WriteAsync` API. The Task 5 plan didn't enumerate every test that referenced the old shape. Implementer caught it during the build and updated the test alongside the new ones.
- **`SkillsServiceCollectionExtensionsTests.Build` return type (Task 13):** the helper was authored to return `IServiceProvider`, but CA1859 (under warnings-as-errors in tests) flagged that the only callers use `GetServices<T>` extension methods which work on the concrete `ServiceProvider` just as well. Implementer changed the return type to `ServiceProvider`. Trivial but worth noting because it shows up across the codebase wherever a helper *could* return concrete.
- **CA1859 implications for csproj packages (Task 6):** Haiku implementer added `Microsoft.Extensions.Configuration.Abstractions` and `Microsoft.Extensions.DependencyInjection.Abstractions` to `Directory.Packages.props` because the new `FlowHub.Persistence.csproj` declared them as `<PackageReference Include="..." />` without versions, and central package management failed without the central pin. Not in the plan; correctly added.
- **MassTransit reference in `FlowHub.Persistence` (Task 8):** `EfCaptureService` depends on `IPublishEndpoint` (MassTransit); plan didn't list `MassTransit` in

the Persistence csproj's `<PackageReference>` block. Implementer added it.

- `flowhub.db` `.gitignore` **entry (Task 10)**: sensible add not in plan; SQLite file lands in `source/FlowHub.Web/` working dir on `just run` and shouldn't be committed.
- **Rate-limit interruption mid-Task-12 (Wallabag implementer dispatch)**: the Anthropic plan rate limit hit during the implementer subagent's run. The subagent had created all three files (`WallabagOptions.cs` , `WallabagSkillIntegration.cs` , `WallabagSkillIntegrationTests.cs`) but hadn't reached the `git commit` step before the rate cap. Files survived on disk; controller verified each one matched the plan content, ran the tests (7/7 pass, 138/138 full suite), committed manually. Recovery pattern: subagent work products are durable to rate-limit cutoffs because the filesystem is the boundary; the missing step is just the commit, which the controller can complete from outside. Worth knowing for any long-running plan execution.
- **Final-review-only catch — `CaptureDetail` Unhandled message**: the per-task spec-compliance review wouldn't have caught this; it's a UX issue visible only end-to-end. The original page rendered `"No Skill matched this Capture. The AI classifier could not determine a category."` for every Unhandled capture — but `Unhandled` actually means a skill *did* match, the integration was *attempted*, and the integration *failed*. (The "no skill matched" case lives in `Orphan` , written by `CaptureEnrichmentConsumer` .) Reviewer caught the semantic mismatch and proposed a one-line razor change to surface `FailureReason` correctly. Same fix updated the corresponding bUnit smoke test, which was asserting on the now-incorrect string. The lesson is that branch-wide review caught what 21 per-task reviews wouldn't have.

Generated vs. handwritten share (estimate, Beta MVP only)

Artifact	AI-drafted	Human-edited
Brainstorming spec (D1-D8)	~95%	~5% (D5 default-Anthropic-Haiku choice; D8 single-consumer-no-event-split confirmation)
Implementation plan (21 tasks + templates)	~95%	~5% (verifying paths against actual repo state)
Production code (Persistence + Wallabag + Vikunja + UI)	~85%	~15% (dual-provider trap fix, <code>InternalsVisibleTo</code> re-architecting, <code>MigrationRunner</code> typeof match, design-time <code>DbContextFactory</code>)
Tests (39 new: 13 Persistence + 13 Skills + 7 extension + 3 <code>SkillRoutingConsumer</code> + 2 <code>SkippableFact</code> + 1 enrichment-consumer)	~95%	~5% (CA1859 return-type swap, parallel-run cancellation token, deterministic seed)
Docs (CHANGELOG, CLAUDE.md , this file, vault Block-3 + Block-4 cross-references)	~90%	~10% (Beta-MVP follow-up admonitions, partial-tick scope notes)
Final review fixups (<code>CaptureDetail</code> Unhandled alert + <code>IntegrationTestFactory</code> typeof)	~50%	~50% (one-liners; reviewer named the change, controller applied)

The 15% human share in production code is concentrated in the test infrastructure (`IntegrationTestFactory`) where two real EF-Core-8+ behaviours diverged from the plan: the dual-provider trap and the `InternalsVisibleTo` -vs- `public` encapsulation choice. Both are documented above; both should now be plan-template patterns for any future EF-Core-host-replacement task.

Reflexion — what worked / what didn't

What worked

The 21-task plan executed across one session with one rate-limit interruption (mid-Task-12) and zero `BLOCKED` / `NEEDS_CONTEXT` escalations. The Sonnet/Haiku dispatch defaults from the user's standing memory feedback continued to pay off — Haiku handled the eight csproj/justfile/markdown tasks for low cost, and Sonnet handled all TDD work without needing Opus escalation. Skipping per-task code-quality reviews (also from the standing feedback) saved roughly 60% of the token burn that Slice B's full two-stage-review pattern would have consumed; the single final branch-wide code review caught the one semantically wrong UX message that no per-task review could have surfaced. TDD discipline against EF Core's `InMemory` provider produced 13 unit tests that caught the cursor-pagination edge cases (limit+1 probe, ordering by `CreatedAt DESC`, `Id DESC`, filter composition) before any of the integration tests touched real persistence.

The slice's strategic choice — landing Block-4 (persistence) and Block-5 (real Skills + Integrations) work *early* to validate the architecture against a real homelab Wallabag/Vikunja before doing the rubric-driven Block 4 and Block 5 work — proved its value in two ways. First, it forced an honest trade-off conversation about what the Beta scope is *not*: SQLite not Postgres, no Repository layer, in-process MigrationRunner instead of the 12-Factor-XII separate init container, only `ICaptureService` swapped to EF (Skill/Integration health stay stubs). Each of those will be re-addressed in Block 4 with the Beta as the working baseline rather than designing from a blank page. Second, it revealed an interface-shape decision (`ISkillIntegration.HandleAsync` returning a `SkillResult` with `ExternalRef`) that the brainstorming spec hadn't fully landed — a real downstream system has an id, the consumer needs to persist it, the lifecycle needs a `Completed` terminal state. None of that surfaced in Slice C because Slice C's classifier doesn't have an external persistence target. The Beta MVP forced these contracts into existence.

The `CHANGELOG`, vault Block-3 cross-reference, and Block-4 Nachbereitung partial-tick documentation landed alongside code rather than as one-shot end-of-block writeups. The Block 4 Nachbereitung file went from 0% → 16% via the Beta MVP cross-reference alone, with each ticked item carrying an inline scope note (full vs partial vs deferred). When Block 4 work begins formally on 2026-05-09 the writeup is already ~1/6 done.

What needed correction

The plan's encapsulation directive (`internal sealed class CaptureEntity`) was overridden by the Task 8 implementer's first instinct to widen visibility to `public` when test seeding wouldn't compile. The fix (`InternalsVisibleTo` via `AssemblyInfo.cs`) was a one-commit follow-up, but the lesson is that when the plan specifies a non-default visibility, it should also specify the hatch (which `InternalsVisibleTo` target, where the `AssemblyInfo.cs` lives) so implementers don't default to widening. The same issue resurfaced in Task 10 when the `IntegrationTestFactory` needed to seed `CaptureEntity` directly — adding `InternalsVisibleTo("FlowHub.Api.IntegrationTests")` was the right call (event-bus-free seeding) but accumulates internal-visibility surface area across two test assemblies. Worth tracking; not yet a problem.

The `IntegrationTestFactory` rewrite (Task 10) consumed disproportionate implementer-and-controller attention relative to its size (one file, ~80 lines). Three EF-Core-8+ behaviours collided in it: the dual-provider trap, the `MigrationRunner` removal, and the seeding strategy. Each was a real catch the implementer worked through, but the cumulative time spent on test infrastructure was longer than the time spent on the actual Wallabag/Vikunja integrations. For Block 4's Postgres switch this is the file to touch *first* — the seeding strategy and the provider-removal pattern will need to evolve with whatever Testcontainers approach lands.

The captive- `HttpClient` anti-pattern in `AddHttpClient<T> + AddSingleton<ISkillIntegration>(sp => sp.GetRequiredService<T>())` was flagged by the final code review as a Block-4/5 follow-up. The pattern was inherited from Slice C's `AddFlowHubAi` — copy-pasted without re-evaluating whether it makes sense for the Skills surface (where the `HttpClient` is held over real network calls, not just used briefly for a single request). The handler-rotation lifecycle that `IHttpClientFactory` provides is effectively defeated by the singleton wrapping. Acceptable for a Beta with single-user load; the right time to fix is when `IHttpClientFactory`-named clients are being introduced for any other reason (e.g. when Telegram lands).

`EventId 2012 SkillIntegrationFailed` was reserved in the spec namespace but never wired. Both `WallabagSkillIntegration` and `VikunjaSkillIntegration` accept `ILogger<T>` and store it in `_log` without any callers. The right place is around `EnsureSuccessStatusCode()` with a `try / catch /log/ throw` — but the rethrow path's catch-and-rethrow-just-to-log feels artificial when the `SkillRoutingConsumer` already logs `2010 LogIntegrationCalled` and `2011 LogIntegrationSucceeded`. Decision: leave deferred until Block 4 produces a concrete observability requirement that `2012` actually serves.

What didn't work / where humans had to intervene

The Anthropic plan rate-limit cap during Task 12 forced a manual recovery: the Wallabag implementer subagent had created the three files (production + options + tests) but hadn't reached the `git commit` step before the cap. The controller (this main session) had to verify each file matched the plan content, run the tests (7/7 Wallabag, 138/138 full suite), and commit by hand. Recovery was clean — the filesystem boundary is durable across rate-limit cutoffs, only the commit step needs to be replayed. But it does mean any long-running plan execution should expect at least one such interruption per multi-hour session, and that the recovery cost is roughly equivalent to a 2–3 minute manual review of the files the subagent claimed to write.

The spec error table contained one runtime-incorrect claim that survived through plan execution: *"Skill not configured → Orphan"*. The actual runtime path for an unconfigured skill is `Unhandled` (because the AI classifier still returns a `MatchedSkill`, `CaptureClassified` is published, `SkillRoutingConsumer` finds no integration registered for that name, and routes to `MarkUnhandledAsync`). The boot-log message in `SkillsBootLogger` (`EventId 4021`) correctly says `Unhandled`, but the spec-doc table had `Orphan`. Code reviewer caught the contradiction during the final whole-branch review; documented as a docs-only correction (no code change). Lesson: the brainstorming spec's error-table rows should be cross-checked against the consumer's actual `Mark*Async` calls during plan-write, not after.

Honest note on Task 21 (operator demo)

The 21-task plan reserved Task 21 for the operator: set user-secrets (Anthropic key, Wallabag PAT, Vikunja PAT + project id), boot the app, verify EventIds 3020/4020/5010/5011 in the boot log, browser-test URL → Wallabag / `todo:` → Vikunja / nonsense → Orphan, restart-survival check, optional `just test-beta` against the live homelab, then push. That step is not part of this writeup — it was deliberately gated for human-driven validation against the real homelab services. The 138/138 unit-and-integration-test count above is the upper bound of what the automated pipeline can guarantee; the Task-21 evidence (screenshots, log lines, live Wallabag/Vikunja entries) will land in this section when the operator runs it.

Block 4 — Persistence (May 2026)

Tools

Tool	Role
Claude Code (Sonnet 4.6, controller)	Spec/plan authoring, subagent dispatch controller, spec+quality review
Claude Sonnet 4.6 (implementer subagents)	Implementation, TDD, entity scaffolding, test writing
Claude Haiku 4.5 (mechanical subagents)	Mechanical tasks: doc writing, YAML edits, spec docs
GitHub Copilot	Inline completions during session review

Workflow

Same `superpowers:subagent-driven-development` loop as Block 3:

- `superpowers:brainstorming` — `docs/superpowers/specs/2026-05-05-block4-nachbereitung-design.md` (Block 4 scope spec)
- `superpowers:writing-plans` — 5-slice plan at `docs/superpowers/plans/2026-05-05-block4-nachbereitung.md`
- `superpowers:subagent-driven-development` — fresh implementer + spec-compliance reviewer + code-quality reviewer subagent per task. Dispatch model: Sonnet for TDD/judgment, Haiku for docs/YAML.

Per-Slice AI Usage

Slice 1 (PostgreSQL + Repository Foundation): AI generated `EfCaptureRepository` from the interface contract; initial `FindAsync` (tracking) corrected to `.AsNoTracking().FirstOrDefaultAsync()` by reviewer. `ICaptureRepository` interface design was a human decision. AI share: ~90%.

Slice 2 (Channel + Skill Entities): AI produced entities, configurations, and repositories in one pass. `varchar(128)` Name field corrected to `varchar(64)` in review. `EfSkillRegistry` two-line class was human-designed; AI implemented. AI share: ~88%.

Slice 3 (Full Domain Model): AI generated 4 entities + configs + repos. FK strategy corrected by human (AI defaulted to CASCADE everywhere; SkillRun→Skill changed to RESTRICT, Capture→Channel/Skill soft FK). N+1 on Tags nav property caught by human in code review. AI share: ~85%.

Slice 4 (Tests + Docker): `PostgresFixture` per-test isolation pattern scaffolded by AI; human validated `NpgsqlConnectionStringBuilder` admin connection. All 16 test methods AI-generated. Docker Compose `flowhub.migrations` service was human-designed (12-Factor XII pattern). AI share: ~90%.

See `docs/insights/block-4.md` for full AI metrics table and KI-Reflexion.

Notable Corrections (Block 4)

- **Npgsql version:** Plan specified 9.0.4 (EF Core 9 era); actual project uses EF Core 10. Implementer resolved to 10.0.1. Plan version pins should be verified against installed SDK before execution.
- **N+1 on Tags:** `EfCaptureRepository.ListAsync` initially generated without `.Include(c => c.Tags)`. Caught in code-quality review. Fix: add `Include` before `AsNoTracking` in query chain.
- **FK cascade defaults:** AI defaults to CASCADE DELETE on all FKs. Requires explicit human review of each FK against domain semantics (owned entity vs. referenced entity).
- **CA1707/CA1825 in generated migrations:** Added `.editorconfig` to suppress these for `Migrations/` folder — numeric migration name prefix (`0001_`) causes a leading underscore in the generated class name, triggering CA1707.

Block 5 — Deployment & Abgabe

Tools Used

- **Claude Code (claude-sonnet-4-6):** Primary tool throughout Block 5. Used via the CLI in the project repository.
- **Skill:** `superpowers:writing-plans` : Generated the full implementation plan from the design spec.
- **Skill:** `superpowers:test-driven-development` : Enforced write-tests-first for all new services and endpoints.

- **Skill:** `cas-aise-grade-self-check` : Verified rubric coverage before declaring completion.
- **Skill:** `superpowers:subagent-driven-development` : Executed the plan via fresh subagents per task with spec and quality review gates.

What AI Generated

Artifact	AI Contribution	Human Review
GitHub Actions workflows	Full YAML scaffold	Corrected permissions, trigger conditions, artifact paths
<code>AiEmbeddingService</code>	Class scaffold + <code>GenerateAsync</code> implementation	Verified MEAI API call signature
<code>SearchEndpoints</code>	Full implementation	Verified 503 path and ProblemDetails shape
<code>AdminEndpoints</code>	Full implementation	Added probe call to verify embedding service before rebuild
<code>DemoAuthHandler</code>	Refactored from <code>DevAuthHandler</code>	Verified config-presence logic
Docker Compose stack	Full YAML	Adjusted healthcheck commands, volume mounts
<code>docs/insights/block-*.md</code>	Full content	Verified factual accuracy against git log and vault notes
pgvector migration	Code scaffold	Corrected <code>Pgvector.EntityFrameworkCore</code> requirement vs built-in

Reflection

The subagent-driven development workflow (plan → subagent per task → spec review → quality review) significantly reduced implementation errors. Having exact file paths and complete code in the plan meant each subagent had no ambiguity about what to implement.

The biggest AI assist in Block 5: generating the GitHub Actions workflows. YAML verbosity makes it easy to get details wrong (action versions, permission scopes). AI-generated YAML was 90% correct and required only minor corrections.

The biggest AI limitation: pgvector EF Core integration. AI initially planned `UseVector()` as a built-in Npgsql 10 feature without an extra package, but `Pgvector.EntityFrameworkCore` was actually required. The lesson: AI training data may not reflect very recent library releases; verify package changelogs.

Ultrareview-driven correction: sync-on-submit → async consumer

The first AI-drafted design wired embedding generation directly into `EfCaptureService.SubmitAsync` — generate the embedding, then persist the Capture with the vector in one transaction. A subsequent multi-agent review run (`/ultrareview` , branch-wide) flagged the failure mode: a slow or unavailable embedding provider would either tank the NF-09 p95 < 200 ms submit budget or block capture submission entirely. The accepted design moved embedding generation into `CaptureEmbeddingConsumer` (`IConsumer<CaptureCreated>`) — Captures persist immediately, the embedding lands a moment later, and a provider outage degrades search quality without affecting the write path. The same review prompted adding `POST /api/v1/admin/embeddings/rebuild` for backfilling captures that landed before the provider was configured. This is exactly the class of architectural mistake a single-pass AI generation tends to make — the local code looks fine, but the system-wide latency / availability invariants are wrong.

Smoke-driven correction: empty-string env defaults

The first run of `just smoke-prod` against the compose stack (2026-05-12) crashed `flowhub.web` at startup with `System.ArgumentException: Value cannot be an empty string. (Parameter 'model')`. Root cause: AI-generated configuration code used `configuration["Embeddings:Model"] ?? "mistral-embed"`, which fails for compose's `${X:-}` interpolation (substitutes empty string, not null). Same trap on `Ai__<Provider>__Model`. The fix (`is { Length: > 0 } m ? m : default`) is a one-character change in semantics but invisible to an AI focused on the in-memory case. Lesson: deployment-shape failure modes are systematically under-tested in AI-generated code; the smoke target is what closed this gap.

Custom Skills (self-authored)

In addition to the upstream `superpowers:*` and `frontend-design:*` skills (bundled with Claude Code / Anthropic skill marketplaces), the following project-specific skills were authored by the student and used throughout the CAS project. They live in a dedicated plugin repo and are loaded via the Claude Code plugin marketplace.

Source / plugin repo: <https://github.com/freaxnx01/agent-skills>

Skill	Purpose	Where it was used
<code>cas-aise-todo-list:cas-aise-todo-list</code>	Computes today's Block (1–5) and Phase (Vorbereitung / PVA / Nachbereitung) from the current date, reads the matching page from the cas-aise Obsidian vault, and renders a 5×3 Block×Phase progress matrix plus open - [] items grouped by section.	Daily orientation at session start: "what's left for this block/phase". Drove the day-by-day work order across all five blocks.
<code>cas-aise-grade-self-check:cas-aise-grade-self-check</code>	Walks the canonical CAS AISE Moodle Bewertungskriterien rubric (18 scored items, 5 buckets, max 100 pts), classifies each item against evidence in the FlowHub repo + vault, and surfaces the highest-leverage gaps. Read-only — produces a gap report only.	Before declaring each Block-Nachbereitung "done" (Blocks 3, 4, 5). Highest-leverage use was on the KI / Sub-Systeme / Reflexion bucket (30/100 pts) — caught missing evidence in <code>docs/ai-usage.md</code> (this file) before submission.
<code>sync-ai-instructions:sync-ai-instructions</code>	Assembles or refreshes the AI agent instruction files (<code>CLAUDE.md</code> , <code>.github/copilot-instructions.md</code> , <code>SKILL.md</code> , <code>.ai/base-instructions.md</code> , <code>.ai/stacks/<stack>.md</code> , shared files under <code>.ai/skills/</code> and <code>.claude/commands/</code>) by pulling a base + chosen stack overlay from a	One-shot project bootstrap (Block 1) and periodic refresh whenever the base instructions evolved. Most recent sync: commit <code>5a59074</code> (chore: sync base + dotnet-blazor overlay from ai-instructions@1071724).

Skill	Purpose	Where it was used
	separate <code>ai-instructions</code> repo. Idempotent.	

Companion repo (instruction source for `sync-ai-instructions`):

<https://github.com/freaxnx01/ai-instructions>

Why custom skills (and not just upstream)

The three skills above cover concerns that the upstream `superpowers` plugin doesn't address:

- **Course-specific rubric grounding.** `cas-aise-grade-self-check` encodes the Moodle Bewertungskriterien as a programmable checklist. Without it, rubric coverage was easy to drift away from during long implementation sessions.
- **Schedule-driven prioritisation.** `cas-aise-todo-list` ties the agent to the real Block/Phase calendar in the Obsidian vault rather than letting the agent pick the next task on vibes.
- **Cross-project instruction reuse.** `sync-ai-instructions` keeps `CLAUDE.md` and friends in sync with a single upstream source — the same overlay drives the FlowHub repo, the `poc/FlowHub.AI.Classification` POC, and any future projects.

All three skills are open-source and reusable by other CAS AISE students who fork the rubric/vault structure.

Fazit — across all five blocks

The per-Block and per-Slice "Reflexion" sections above capture what happened block-by-block. This section steps back and names the patterns across the five months — what shifted, what kept failing, what kept working, and what I'd change.

Workflow evolution (Block 1 → Block 5)

Each block traded ad-hoc chat for more structured, more isolated agent dispatch. The cost of structure paid back in fewer mid-task escalations.

Block	Workflow shift	Decisive tools
1 — Einführung	Ad-hoc Claude.ai chat for architecture; Copilot inline for code. Hand-rolled <code>CLAUDE.md</code> .	Claude.ai , Copilot
2 — Frontend	First sustained Claude Code CLI use. Phase discipline via <code>/ui-brainstorm</code> → <code>/ui-flow</code> → <code>/ui-build</code> → <code>/ui-review</code> .	Claude Code (Opus 4.7)
3 — Service	First end-to-end <code>superpowers:brainstorming</code> → <code>writing-plans</code> → <code>subagent-driven-development</code> slice. Two-stage review (spec + quality) per task.	Superpowers SDD; MEAI
4 — Persistence	Sonnet defaults + Haiku for mechanical tasks; per-task quality review dropped for one branch-wide review at end — ~60% token-burn cut, no quality loss.	Testcontainers; first custom skill (<code>cas-aise-todo-list</code>)
5 — Deployment	<code>/ultrareview</code> (multi-agent branch review) added; rubric-grounded <code>cas-aise-grade-self-check</code> skill closes each block.	<code>/ultrareview</code> , custom skill repo

By Block 5 the implementer subagents own full slices. The human work concentrates at two endpoints: the spec (where contracts get nailed down) and the review (where system-wide invariants get checked). The bottleneck moved from typing to those two places.

Recurring failure modes

- 1. Single-pass AI review misses system-wide invariants.** The Block-5 embedding-on-submit rework is canonical: the AI-drafted code was locally correct (one transaction) but globally wrong — a slow provider tanks the NfA-09 p95 < 200 ms submit budget. Per-task code-quality review couldn't catch this. `/ultrareview` was introduced for exactly this class of mistake and found it on the first run.
- 2. Deployment-shape failures are systematically under-tested.** Compose `${X:-}` interpolates empty strings (not nulls), so `??` defaults silently no-op. `.editorconfig` not COPY-ed into Docker turned analyzer suppressions into

build errors. Casing mismatch `EMBEDDINGS__APIKEY` vs `Embeddings__ApiKey` shadowed a whole feature. None surface in unit or integration tests; all five were caught by `just smoke-prod` in one afternoon. AI-generated deployment glue needs its own gate.

3. Training-data lag bites on recent libraries.

`Pgvector.EntityFrameworkCore` was assumed in-the-box with Npgsql 10 (it isn't — separate package). `MassTransit.Testing` was assumed to be a separate package (it isn't — `AddMassTransitTestHarness()` is in the main `MassTransit` assembly). Verify every AI-stated API against the actually-installed package, especially for libraries that shipped after the model's cutoff.

4. Open-ended prompts inflate scope. "Add X" without a plan produces a feature suite. The antidote was `superpowers:writing-plans` — bite-sized TDD-ordered tasks with full inline code, exact file paths, exact commit messages. Subagents dispatched against well-specified tasks reported DONE; subagents dispatched against open-ended prompts went exploring.

5. Plans authored from incomplete repo knowledge. Block 3 Slice B's plan missed `CaptureServiceStubTests.cs` and the existing `ChannelKind` values; implementers adapted mid-task. The plan-writing skill now reads every file the plan references *before* freezing it.

6. AI defaults to widening visibility when tests don't compile. Block 4 Task 8: plan said `internal sealed class CaptureEntity`; first instinct on a compile error was `public`. The right answer (`InternalsVisibleTo` in `AssemblyInfo.cs`) has to be in the plan, not implied as a footnote.

Recurring wins

1. **The plan is the contract.** When the plan contains exact file paths, exact code, exact commit messages, and exact verification commands, the implementer operates with narrow judgment space and reports DONE rather than NEEDS_CONTEXT. The 95%+ "AI-drafted" share on production code is only reachable because the plan removed the ambiguity that would otherwise force the agent to guess.
2. **Review cadence has to match the failure mode it catches.** Per-task spec + per-task quality made sense for Block 3 Slice B (first SDD run, no muscle memory). By Block 4 the per-task quality review caught nothing new while the branch-wide `/ultrareview` at end-of-slice caught the architectural mistakes per-task review never could. Dropping per-task quality review cut ~60% of the token burn with no observed quality loss.

3. **Custom skills are the project's memory layer.** `cas-aise-todo-list`, `cas-aise-grade-self-check`, and `sync-ai-instructions` cover concerns upstream `superpowers` doesn't: rubric grounding, schedule-driven prioritisation, cross-project instruction reuse. Without `cas-aise-grade-self-check` the rubric drifted during long sessions — the skill caught missing evidence in this very file before submission.
4. **AI for rigid-structure artifacts, human for architectural judgment.** GitHub Actions YAML, ADRs, OpenAPI ProblemDetails shapes, EF migrations — rigid structure where AI's 90%-correct first draft saves real time. Architectural decisions (hexagonal split for AI, in-process vs. RabbitMQ transport, dispatcher-event-carries-description vs. inline-skill-call) stayed with the human; AI's role there was to surface alternatives, not to choose.
5. **Correction stories are the rubric evidence that counts.** "Defects Found by the Smoke Run" in `docs/insights/block-5.md` and the "Ultrareview-driven correction" / "Smoke-driven correction" sections of this file are auditable: each names a concrete defect, the commit that fixed it, the lesson. Hand-wavy "AI helped overall" claims close no rubric loop; "AI produced X, smoke caught Y, fix landed in commit Z" does.

What I'd change next time

- **Rubric grounding from Block 1, not Block 3.** `cas-aise-grade-self-check` was authored mid-project; Blocks 1–2 had no formal use-cases / nfa / acceptance-criteria docs. The retroactive doc-work in late Block 3 cost real time. Formalise the spec docs in the first week.
- **`/ultrareview` from Block 3, not Block 5.** The branch-wide review catches the class of architectural mistake per-task review can't. Cheaper to catch it after Slice B than after a 21-task Beta MVP.
- **A just smoke recipe per block, not just Block 5.** Five latent bugs found in one afternoon by `just smoke-prod`. A simpler smoke (boot the app, curl a feature path) would have caught at least three of them earlier.
- **Codify `InternalsVisibleTo` and EF-Core-test-host patterns in the plan template.** Both Block 3 Slice D and Block 4 had implementers default-widen visibility when tests didn't compile. The pattern is mechanical once spelled out; the recurring friction was avoidable.

Did the original hypothesis hold?

The Block 1 assumption was that Claude would be "a fast typist with good library knowledge" — useful for boilerplate, suspect on architecture, basically a

productivity multiplier rather than a workflow change. By Block 5 that's wrong by a wide margin in both directions.

Boilerplate is faster than expected. The 95%+ AI-drafted share on production code isn't "AI did 95% of the work" — the 5% human input is concentrated at the high-judgment endpoints (scope, contracts, trade-offs). But the typing genuinely is a small fraction of the elapsed time. The bottleneck moved from implementation to spec writing and review — exactly the inversion SDD optimises for.

Architecture is more dangerous than expected. The embedding-on-submit case is a one-paragraph rewrite that's obvious in retrospect: AI produced locally-correct code that broke a system-wide invariant, and a single-pass review would have shipped it. The deployment-shape failures are the same shape: locally correct, globally wrong. The human role at the architecture boundary didn't shrink; it grew, because the agent now produces enough code fast enough that the only meaningful filter is at design and review.

AI didn't replace any one role in the workflow. It moved which role is the bottleneck. For the next project: spec docs first, `/ultrareview` from day one, a smoke target before the first feature. Implementation throughput is no longer the constraint — design and review throughput are.

References

- ADR 0003: `docs/adr/0003-async-pipeline.md`
- Spec: `docs/superpowers/specs/2026-04-30-async-pipeline-design.md`
- Plan: `docs/superpowers/plans/2026-04-30-async-pipeline.md`
- Bewertungskriterien: `vault/0rganisation/Bewertungskriterien.md`

2026-05-24 — Capture file upload (PVA 4 follow-up)

Driver: PVA 4 (2026-05-23) feedback — add browse/upload to the Capture input so a future paperless-ngx skill can ingest documents directly; cap demo uploads at 2 MB.

Workflow (Claude Opus 4.7 controller + Sonnet/Haiku subagents under the `superpowers:subagent-driven-development` skill):

1. **Brainstorming skill** — controller asked four bundled multiple-choice questions to pin scope (entry points, domain shape, storage, config) and two follow-ups (Content rule, MIME allowlist). Result: design spec at `docs/superpowers/specs/2026-05-24-capture-file-upload-design.md` .
2. **Writing-plans skill** — produced `docs/superpowers/plans/2026-05-24-capture-file-upload.md` with 16 bite-sized TDD tasks (failing test → impl → green → commit per task).
3. **Subagent-driven execution** — one implementer subagent per task, followed by a combined spec-compliance + code-quality reviewer subagent. Haiku for mechanical tasks (records, gitignore, config wiring), Sonnet for orchestration / UI / EF mapping. One review caught a security smell on Task 12 (`QuickCaptureField.ValidateFile` short-circuited the allowlist check on empty list) — fixed in a separate `fix(web):` commit so the bisect log preserves the lesson.
4. **Auto-verification approximation** for Task 14 (manual smoke): controller booted the app, asserted `/health/live=Healthy` and `/captures/new=200` . Full browser interaction smoke deferred to the user on PR review.

What worked:

- **Implementer + reviewer subagents on small, well-specified tasks** — review caught one real security issue (`Count > 0` short-circuit hiding an unrestricted upload path on misconfig) that would have shipped under a single-agent flow. Worth the cost.
- **TDD per task** — every task started with a failing test; no test was modified to pass.
- **Splitting EF entity mapping (Task 7) and migration generation (Task 8)** — left the suite intermittently red on commit `85f2240` . Acceptable for bisect, but next time bundle the two when `PendingModelChangesWarning` is configured as a hard error.

Spec: `docs/superpowers/specs/2026-05-24-capture-file-upload-design.md` Plan: `docs/superpowers/plans/2026-05-24-capture-file-upload.md`

Learnings (personal)

Source: `vault/Projektarbeit/Learnings.md`

tags:

- claude-generated
 - claude-updated updated: 2026-06-05
-

Learnings CAS AISE

Persönliche Lessons Learned aus der Arbeit am FlowHub-Projekt mit Claude Code, Codex und weiteren AI-Coding-Agents. Fokus liegt auf den praktischen Erkenntnissen rund um Agent-Tooling, die sich über alle fünf Blöcke hinweg als wirklich tragfähig erwiesen haben — nicht auf einer vollständigen Tool-Übersicht.

AI-Instructions & Write Skills

Die wichtigste strukturelle Erkenntnis war: AI-Agents brauchen genauso ein gepflegtes "Onboarding-Dokument" wie ein neuer Teamkollege. In FlowHub wurde das in mehrere Ebenen aufgeteilt — `CLAUDE.md` für die unmittelbar harten Regeln, `.ai/base-instructions.md` als kanonische Konventionsreferenz und `.ai/cas-instructions.md` für den CAS-spezifischen Kontext (Block-Schedule, Bewertungskriterien, Implementierungs-Rhythmus). Ohne diese Schichten driftet jeder Agent in seine eigenen Defaults ab; mit ihnen werden Instruktionen einmal geschrieben und von Claude Code, Codex, Copilot und Gemini gleichermassen befolgt. Das `update-ai-instructions` Skill hält diese Files synchron mit dem Upstream-Template, damit Verbesserungen aus anderen Projekten zurückfließen.

Parallel dazu hat sich das Schreiben **eigener** Skills als stärkster Hebel herausgestellt, um wiederkehrende Workflows verlässlich zu machen. Skills sind kleine, scharf umrissene Markdown-Definitionen mit einer Trigger-Description und einem expliziten Vorgehen — und genau das macht den Unterschied zu blossen Prompt-Templates: der Agent erkennt selbst, *wann* ein Skill anzuwenden ist, und

folgt dann einer geprüften Checkliste statt zu improvisieren. Das `superpowers:writing-skills` Skill diene als Meta-Werkzeug für genau diesen Prozess (Verifikation vor Deployment, klare Triggerbedingungen). Konkret entstanden so für FlowHub u.a. die `/ui-brainstorm → /ui-flow → /ui-build → /ui-review` Pipeline, der `/commit` und `/push` Flow sowie das `cas-aise-grade-self-check` Skill, das jede Block-Nachbereitung gegen die Moodle-Bewertungskriterien prüft.

Skills in Plugins nach Thema aufteilen

Skills landen — zumindest mit ihrer Description — im System Prompt des Agents. Das wird unterschätzt: jeder zusätzliche Skill kostet Tokens *bevor* die erste User-Nachricht überhaupt gelesen wird, und je mehr Skills geladen sind, desto schwerer fällt es dem Modell, den jeweils richtigen zu triggern. Im FlowHub-Setup zeigte sich das, sobald die persönliche Skill-Sammlung über ein gutes Dutzend hinauswuchs: Trigger wurden unschärfer, Kontext-Budget wurde knapper, und Skills für komplett fremde Domänen (Homelab-Routing, Movie-Sync, ...) belasteten die CAS-Sessions ohne jeden Mehrwert.

Die Lösung war, Skills konsequent in **thematische Plugins** zu zerlegen und jeweils nur das Plugin zu aktivieren, das zur aktuellen Arbeit passt. Für FlowHub-CAS-Sessions sind das z.B. `cas-aise-grade-self-check`, `cas-aise-todo-list`, `flowhub` (mit den UI-/Commit-Sub-Skills) sowie das schmale `superpowers`-Set; alles andere bleibt ausserhalb der Session unsichtbar. Das hält den System-Prompt-Footprint klein, schärft die Triggerwahrscheinlichkeit der wirklich relevanten Skills und macht es zudem möglich, Plugins einzeln zu versionieren und über Marketplace-Repos mit anderen zu teilen, ohne ständig die persönliche Default-Konfiguration mitzuschleppen.

Context Hygiene: Logs via File

Der grösste single Token-Fresser in interaktiven Sessions war anfangs der Browser-Console-Output. Wenn der Agent einen Bug im Blazor-UI nachverfolgt und dafür komplette Console-Streams ins Conversation-Window kippt — Fetch-Logs, MudBlazor-Render-Warnings, Hot-Reload-Meldungen, Stack-Traces — sind nach wenigen Iterationen mehrere zehntausend Tokens verbraucht, die anschliessend bis zur Compaction im Kontext kleben bleiben. Das verlangsamt

die Antworten, treibt die Kosten hoch und drückt ältere, oft noch relevante Information aus dem Fenster.

Die deutlich tragfähigere Variante: den Console-Output (oder `dotnet watch`-Logs, Test-Output, OTEL-Traces) in eine **Datei** schreiben und den Agent diese Datei gezielt mit `Read` plus `offset / limit` lesen lassen — oder per `grep / rg` nur die wirklich relevanten Zeilen extrahieren. Damit landet im Kontext exakt der Ausschnitt, der für die nächste Entscheidung gebraucht wird, statt des gesamten Streams. Dieselbe Logik gilt für Build-Output, Migration-SQL oder Telemetrie-Dumps: alles, was potenziell gross werden kann, gehört zuerst auf die Disk und wird dann selektiv in den Kontext geholt. Diese eine Gewohnheit hat den Token-Verbrauch pro Debug-Session subjektiv um den Faktor 5–10 reduziert.

Code-Exploration: LSP / ripgrep / grepai

Es gibt nicht *das eine* richtige Werkzeug, um sich im Code zu orientieren — die Wahl hängt von Projektgrösse und Frage-Typ ab, und das macht der Agent erst gut, wenn man ihm das explizit beibringt. Bei kleinen Repos (FlowHub ist mit wenigen tausend LOC noch handlich) reicht oft ein direktes `Read` oder ein gezielter `rg`-Aufruf; das ist schnell, deterministisch und kostet praktisch keinen Overhead. Sobald aber strukturelle Fragen ins Spiel kommen ("wo wird `IClassificationPort` implementiert?", "welche Aufrufer hat diese Methode?"), wird ein **LSP**-basierter Zugriff überlegen, weil er Symbole statt Strings auflöst und Renames, Inheritance oder Partial Classes korrekt versteht — etwas, das `grep` prinzipiell nicht leisten kann.

Bei grösseren Codebases oder unscharfen Suchen lohnt sich dann der nächste Sprung: **grepai** (oder vergleichbare semantische Such-Tools), die nicht nach exakten Tokens, sondern nach Bedeutung suchen ("wo behandeln wir Validierungsfehler an der API-Boundary?"). Die praktische Heuristik, die sich bewährt hat: erst `rg` versuchen — wenn drei Suchanläufe nichts Brauchbares liefern, ist die Frage vermutlich nicht lexikalisch, sondern semantisch, und es ist Zeit für LSP oder grepai. Diese Hierarchie bewusst zu wählen — statt reflexartig das mächtigste Tool zu nehmen — spart Tokens, hält Antworten schnell und vermeidet das typische "Agent versinkt in Suchergebnissen"-Pattern.

Context Hygiene: Superpowers-Workflow mit /clear -Schnitten

Eng verwandt mit der Logs-via-File-Disziplin ist die Frage, wie man **mehrphasige Aufgaben** sauber durch Spec → Plan → Implementierung führt, ohne dass sich der Kontext bis zur Unkenntlichkeit aufbläht. Das `superpowers`-Plugin hat dafür einen Drei-Phasen-Workflow etabliert, der zwischen den Phasen jeweils ein hartes `/clear` setzt — also den Conversation-Kontext bewusst auf null zurücksetzt:

1. **Specs / Design** — Anforderungen und Entwurf erarbeiten (z. B. via `superpowers:brainstorming`), Ergebnis als `docs/superpowers/specs/xyz.md` ablegen. `/clear`
2. **Plan** — neue Session: `> Read docs/superpowers/specs/xyz.md and run superpowers:writing-plans to produce the implementation plan.` `/clear`
3. **Implement** — wieder neue Session: `> Execute the xyz plan via subagent-driven-development`

Der Trick liegt im `/clear`: jede Phase produziert ein **Artefakt auf Disk** (Spec-MD, Plan-MD, Code-Diff), das in der nächsten Phase als reiner Input zurückgelesen wird. Den vorherigen Hin-und-her-Dialog mitzuschleppen liefert keinen Mehrwert — die Entscheidungen sind im Artefakt festgehalten — kostet aber konstant Tokens und lenkt das Modell mit veralteten Zwischenstands-Diskussionen ab. Diese Disziplin ist die strukturelle Variante des Logs-via-File-Patterns: was zwischen Phasen weiterleben muss, gehört in eine Datei; was nur Gesprächs-Begleitmaterial war, darf gelöscht werden. Praktisch macht das bei nicht-trivialen Features den Unterschied zwischen einer fokussierten Implementierung und einem zähen, immer langsamer werdenden Mega-Thread.

Vom CAS in die Projektarbeit — was konkret eingeflossen ist

Rückblickend lässt sich für jeden der fünf CAS-Blöcke ein konkreter Niederschlag in FlowHub benennen — die Projektarbeit war nicht parallele Spielwiese, sondern der Ort, an dem die jeweiligen Block-Inhalte unmittelbar zur Entscheidung wurden. Der rote Faden ist dabei bewusst, dass fast jede dieser Entscheidungen in einem **ADR** festgehalten ist: Das im CAS gelehrt Abwägen von Optionen wurde so vom Vorlesungsthema zur dokumentierten Projektpraxis.

CAS-Block	Mitgenommener Inhalt	Konkrete FlowHub-Entscheidung
1 — Einführung	Architekturoptionen (Monolith / Modular Monolith / Microservices) begründet gegeneinander abwägen; KI-Tooling für ein neues Repo aufsetzen	Entscheidung für Modular Monolith mit hexagonaler Schichtung statt verteilter Microservices (ADR 0001) — und die ADR-Praxis selbst als Format. Das Tooling-Fundament (<code>CLAUDE.md</code> , <code>.ai/</code> -Instructions, eigene Skills) entstand direkt aus diesem Block.
2 — Frontend	Frontend-Architektur und Render-Modelle	Blazor Interactive Server statt WASM-SPA (ADR 0001), MudBlazor als einzige Komponenten-Bibliothek, und die vierstufige UI-Pipeline <code>/ui-brainstorm → /ui-flow → /ui-build → /ui-review</code> , die Entwurf vor Code erzwingt.
3 — Service	Service-/Microservice-Architekturen, Protokolle (REST/gRPC), KI-gestützte Services	Bewusste Wahl einer asynchronen In-Process-Pipeline (RabbitMQ + MassTransit, ADR 0002/0003) statt eines synchronen Microservice-Geflechts — konsequent zur Modular-Monolith-Entscheidung. REST-Minimal-API für die Aussengrenze; die KI-Klassifikation als erster „intelligenter Service" (ADR 0004).
4 — Persistence	Geeignete Persistenzform wählen, DB-Zugriff über ORM abstrahieren, dynamische Abfragen	EF Core + PostgreSQL mit Port-/Repository-Abstraktion (ADR 0005). Das im Block geschärfte „passende

CAS-Block	Mitgenommener Inhalt	Konkrete FlowHub-Entscheidung
		Persistenzform wählen" führte direkt zum Entscheid, die KI-Suche über pgvector auf derselben Postgres zu lösen statt eine separate Vector-DB einzuführen (ADR 0006).
5 — Deployment	Containerisierung, CI/CD-Automatisierung, Monitoring/Observability	Multi-Stage-Dockerfile + Compose-Stack aus unabhängig deploybaren Containern , drei GitHub-Actions-Workflows (CI, Release → GHCR, Migrations) und OpenTelemetry + Prometheus + Grafana als Observability-Schicht. Die geschärfte Betriebssicht schlug sich zusätzlich in den Policy-ADRs 0007-0009 (LLM-Hosting, Logging-/Telemetry-PII) nieder.

Über die einzelnen Blöcke hinaus war das eigentliche Querschnittsthema des CAS — **KI-assistierte Software-Entwicklung** — zugleich die Arbeitsweise der gesamten Projektarbeit. Mitgenommen habe ich daraus weniger einzelne Tools als die in diesem Dokument beschriebenen Disziplinen: gepflegte Agent-Instructions, eigene Skills für wiederkehrende Workflows, Context-Hygiene über Datei-Artefakte und den Spec→Plan→Implement-Rhythmus mit harten `/clear` - Schnitten. Diese Praktiken sind der direkteste Transfer aus dem CAS in die Projektarbeit — und zugleich das, was ich über FlowHub hinaus in zukünftige Projekte mitnehme. Die vollständige, blockweise Aufschlüsselung der konkreten KI-Nutzung liegt in [docs/ai-usage.md](#) .

Block 1 — Nachbereitung

Source: `vault/Blöcke/01 Einführung/01 Einführung - c) Nachbereitung.md`

tags:

- claude-updated updated: 2026-05-17
-

Block 1 — Einführung · Nachbereitung

Phase budget: 28 h **PVA war:** 2026-02-21 **Nächste PVA:** 2026-03-21

Lernziel

- Ich kann ein KI-gestütztes Software-Engineering-Projekt initial aufsetzen und die wichtigsten Architekturentscheidungen begründen.
- Ich kann Software-Architekturoptionen (Monolith, Modular Monolith, Microservices) für ein Projektszenario gegeneinander abwägen.
- Ich kann das Tooling für KI-assistierte Entwicklung (Claude Code, IDE-Integration, Skills) für ein neues Repository einrichten.

Auftrag (Moodle)

- Sichtung der Materialien und Ergebnisse der Präsenzveranstaltung 1
- Projektbeschreibung verfassen (Idee, Solution Vision, fokussierte Sub-Systeme, Beispiel-Eingabedaten)
- Architektur-Optionen evaluieren und eine begründete Entscheidung dokumentieren
- Initiales Repository-Setup mit dem gewählten Stack

FlowHub-Kontext: Stack-Entscheidung gegen Quarkus / Jakarta EE und für .NET 10 + Blazor Web App in der Projektbeschreibung dokumentiert (siehe `vault/Projektarbeit/Dev.md`). ADR 0001 (Frontend Render Mode and

Architecture, Accepted 2026-02-15) und ADR 0002 (Service Architecture and Async Communication, Accepted 2026-03-01) halten die beiden tragenden Architekturentscheidungen fest.

Bewertungskriterien (Block 1)

Snapshot am Ende der Block-1-Phase (~2026-03-21) gegen die offizielle Moodle-Rubrik aus [[Bewertungskriterien]]. Punkteangaben in Klammern zeigen Max-Score. **Block 1 ist Foundational** — viele Items sind hier erwartungsgemäss noch nicht adressiert; die Rubrik wird kumulativ erst bei Block 5 abschliessend bewertet.

Sektion nachträglich ergänzt (2026-05-17) — damit das Pattern aus Block 3 / 4 / 5 auch in Block 1 / 2 vorhanden ist.

Quarkus/Jakarta-EE-Item ist für FlowHub (.NET-Stack) **nicht relevant** — bewusst ausgeklammert.

Spezifikation

- [-] **Use Cases & fachliche Anforderungen (5)** — informelle Liste in `vault/Projektarbeit/Idee FlowHub.md` (Capture-Idee, Channel-Konzept); formales UC-Dokument folgt in Block 2/3
- ✓ **NfA SMART (5)** — nicht in Block 1; □ geliefert Block 3 (NF-09..NF-13 in `docs/spec/use-cases.md`)
- ✓ **Solution Vision (5)** — `vault/Projektarbeit/Idee FlowHub.md` + Projektbeschreibung-Draft + ADR 0001/0002 (Vision der modularen Lösung)

Entwurf

- [-] **Lösungsansatz & Architektur textuell + bildlich (7)** — ADR 0001 + ADR 0002 textuell vorhanden; Diagramme (C4, Sequence) folgen ab Block 2/3
- [-] **Struktur / Verhalten / Interaktion (7)** — Struktur über ADR 0002 Modul-Split angerissen; Verhalten/Interaktion noch nicht modelliert
- ✓ **DB-Modell (3)** — out of scope Block 1; □ geliefert Block 4 (ADR 0005, `docs/design/db/er.md`)

Programmierung

- [-] **Code lesbar / strukturiert (7)** — Solution-Skelett `FlowHub.slnx` mit `FlowHub.Core` Domain-Typen (`Capture` , `LifecycleStage` , `ChannelKind`); strukturell sauber, aber kaum Code-Volumen
- ✓ **Quarkus/Jakarta EE** — N/A (Stack: .NET 10)
- [-] **Erkenntnisse dokumentiert (3)** — Reflexion-Sektion in dieser Datei + ADR 0001/0002 §"Alternatives considered"; ausführliche Insights folgen ab `docs/insights/block-1.md`
- ✓ **Source in Git (2)** — `github.com/freaxnx01/FlowHub-CAS-AISE` aufgesetzt, Initial-Commits gepusht

Validierung

- ✓ **Abnahmekriterien (5)** — out of scope Block 1; □ geliefert Block 3 (UC-08..UC-11 mit Acceptance in `docs/spec/use-cases.md`)
- ✓ **Test-Strategie + Technologien (5)** — out of scope Block 1; □ geliefert ab Block 2 (`docs/spec/testing-strategy.md`)
- ✓ **Unit-Tests (3)** — kein Test-Projekt in Block 1; □ geliefert Block 2 (bUnit, 31/31 grün)
- ✓ **Test-Ergebnisse dokumentiert (3)** — out of scope Block 1; □ geliefert ab Block 2 (CHANGELOG `[Unreleased]` + `dotnet test` Counts)

KI, Sub-Systeme & Reflexion

- [-] **KI-Werkzeug-Nutzung beschrieben (12)** □ — Reflexion-Sektion dieser Datei listet eingesetzte Tools (Claude Code, Brainstorming-Skill) und konkrete KI-Korrekturen; konsolidierte Doku in `docs/ai-usage.md` folgt ab Block 3
 - ✓ **Intelligente Services mit KI (6)** — out of scope Block 1; □ geliefert Block 3 (`FlowHub.AI` : `KeywordClassifier` + `AiClassifier`)
 - ✓ **Sub-Systeme als Container (5)** — out of scope Block 1; □ geliefert Block 5 (docker-compose Stack)
 - [-] **KI-Reflexion / Fazit (7)** — Block-1-Reflexion-Sektion enthält Tooling-Erfahrungen; finales Fazit gehört in `vault/Projektarbeit/Learnings.md` und das Submission-PDF (Block 5)
-

TODO

□ Done

- ✓ Sichten Materialien und Ergebnisse der Präsenzveranstaltung 1
 - ✓ Projektbeschreibung ändern auf .NET, Upload
 - ✓ Architekturen erklären lassen und entscheiden, dokumentieren — ADR 0001 (Frontend / Modular Monolith) + ADR 0002 (Service-Architektur)
 - ✓ Focus on a few Services / Categories: Vikunja, Wallabag —
`vault/Projektarbeit/External Services.md` + `FlowHub.Integrations/`
Placeholder
 - ✓ Sample input data: Books, Movies, ... — Bogus-Stubs in
`source/FlowHub.Web/Stubs/` (umgesetzt in Block 2)
 - ✓ Solution-Skelett `FlowHub.slnx` mit `FlowHub.Core` Domain-Modell (`Capture` ,
`LifecycleStage` , `ChannelKind` , `ICaptureService` , `ICaptureRepository`)
 - ✓ `Directory.Build.props` (Warnings-as-Errors, embedded PDB) +
`Directory.Packages.props` (zentrale Paketverwaltung) eingerichtet
 - ✓ `global.json` SDK-Pin auf .NET 10 für reproduzierbare Builds über Maschinen hinweg
-

Reflexion

Was hat geklappt

- **Stack-Entscheid früh und begründet.** Die Entscheidung gegen Quarkus / Jakarta EE und für .NET 10 wurde vor jeglichem Code getroffen und in der Projektbeschreibung dokumentiert. Damit war für alle Folgeblöcke klar, welche Äquivalenzen (EF Core ↔ Hibernate / Panache, MassTransit ↔ in-prozess-äquivalent zu Queue-basiertem Messaging) zu wählen sind.
- **Solution-Layout zuerst.** Das Modular-Monolith-Layout (`FlowHub.Core` Domain, `FlowHub.Web` Host, geplante `FlowHub.Persistence` / `FlowHub.Api` / `FlowHub.AI` Projekte) wurde im Block 1 grob gesetzt — Folgeblöcke konnten Projekte addieren, ohne das Grundgerüst zu rebrokern.
- **Zentrale Paketverwaltung.** `Directory.Packages.props` zahlt sich ab dem zweiten Projekt aus — Versions-Bumps an genau einer Stelle.

Was war schwieriger als erwartet

- **PKM-Domain-Modellierung.** `Capture` , `LifecycleStage` , `ChannelKind` als minimale Domain wirken trivial, aber jeder spätere Block hat Felder ergänzt (`SkillId` , `FailureReason` , `Embedding` ...). Mehr Zeit in der initialen Domain-Modellierung hätte spätere Migrationen verkleinert.
- **Architektur-Trade-offs.** Modular Monolith vs. Microservices wurde länger diskutiert als geplant. ADR 0002 §"Alternatives considered" hält die verworfenen Varianten fest.

Verwendete KI-Werkzeuge

- **Claude Code** für Solution-Scaffolding, Domain-Typ-Skizzen und ADR-Drafts.
- **Brainstorming-Skill** für die Architektur-Diskussion vor ADR 0001 / ADR 0002.

Konkrete KI-Korrekturen in diesem Block:

- Vorschlag `IAsyncEnumerable<Capture>` in `ICaptureRepository` verworfen — Overkill für die initiale Skala; `Task<IReadOnlyList<Capture>>` reicht.
- `CancellationToken` -Propagation in `ICaptureService` von KI vorgeschlagen und übernommen.

Verweise

- ADR 0001 — Frontend Render Mode and Architecture (`docs/adr/0001-frontend-render-mode-and-architecture.md`)
- ADR 0002 — Service Architecture and Async Communication (`docs/adr/0002-service-architecture-and-async-communication.md`)
- Projektbeschreibung — `docs/projektbeschreibung/FlowHub_Projektbeschreibung_v4.md`
- Insights — `docs/insights/block-1.md`

Block 2 — Nachbereitung

Source: `vault/Blöcke/02 Frontend/02 Frontend - c) Nachbereitung.md`

tags:

- claude-updated updated: 2026-05-17
-

Block 2 — Frontend · Nachbereitung

Phase budget: 26 h **PVA war:** 2026-03-21 **Nächste PVA:** 2026-04-25

Auftrag (Moodle)

- Tech-Entscheid Präsentationsschicht (CSR/SSR/Mix) + Begründung
- Wireframes für FlowHub-Frontend
- Page-Flow-Diagramme
- Frontend implementieren gegen Stub-Services mit Faker-Daten
- KI-generierte Unit-Tests, alle grün
- Master-Detail-Übung in mehreren Varianten

Volltext: `_files/Moodle/Modul/2-Frontend/pdf/W4B-C-AS001.AISE.ZH-Sa-1.PVA.FS26_Projektarbeit_Frontend_Moodle.pdf`

Bewertungskriterien (Block 2)

Snapshot am Ende der Block-2-Phase (~2026-04-25) gegen die offizielle Moodle-Rubrik aus [[Bewertungskriterien]]. Punkteangaben in Klammern zeigen Max-Score. Die Rubrik wird kumulativ erst bei Block 5 abschliessend bewertet — diese Sektion dokumentiert den Stand zum Block-2-Abschluss.

Sektion nachträglich ergänzt (2026-05-17) — damit das Pattern aus Block 3 / 4 / 5 auch in Block 1 / 2 vorhanden ist.

Quarkus/Jakarta-EE-Item ist für FlowHub (.NET-Stack) **nicht relevant** — bewusst ausgeklammert.

Spezifikation

- [-] **Use Cases & fachliche Anforderungen (5)** — informelle UCs über die 6 implementierten Pages (Dashboard, New Capture, Captures, Capture-Detail, Skills, Integrations) + Glossary (`vault/Projektarbeit/Glossary.md`); formales UC-Dokument folgt in Block 3 (`docs/spec/use-cases.md`)
- [-] **NfA SMART (5)** — Frontend-NfAs implizit über ADR 0001 (Render-Mode-Latenz, Stub-Service-Performance); SMART-Formalisierung folgt in Block 3 (`docs/spec/nfa.md`)
- ✓ **Solution Vision (5)** — ADR 0001 hält Frontend-Vision (Web UI als Channel, Stub-First, OIDC-Pfad) fest; Idee FlowHub + Glossary konkretisieren die Domäne

Entwurf

- ✓ **Lösungsansatz & Architektur textuell + bildlich (7)** — ADR 0001 (Render Mode + Architektur) + Wireframes (`docs/design/dashboard/wireframe.md` u.a.) + Page-Flow-Diagramme via `/ui-flow`
- [-] **Struktur / Verhalten / Interaktion (7)** — Struktur über MudLayout-Shell + Component-Hierarchie; Verhalten/Interaktion über Mermaid-Page-Flows; voller C4-Container-View folgt in Block 3
- ✓ **DB-Modell (3)** — out of scope Block 2 (in-memory Bogus-Stubs); □ geliefert Block 4 (ADR 0005, `docs/design/db/er.md`)

Programmierung

- ✓ **Code lesbar / strukturiert (7)** — `source/FlowHub.Web/` sauber gegliedert (Components/Layout, Components/Pages, Components/DashboardCards, Components/Shared, Stubs/); MudBlazor-Konventionen, code-behind via partial class
- ✓ **Quarkus/Jakarta EE** — N/A (Stack: .NET 10)
- [-] **Erkenntnisse dokumentiert (3)** — Block-2-Insights in `docs/insights/block-2.md` (UI-Workflow-Phasen-Disziplin, ADR-Lifecycle); CHANGELOG `[Unreleased]` mit Block-2-Deliverables

- ✓ **Source in Git (2)** — `github.com/freaxnx01/FlowHub-CAS-AISE` , Block-2-Commits gepusht

Validierung

- [-] **Abnahmekriterien (5)** — die 17 + 14 bUnit-Tests dienen als executable acceptance per Page; formalisierte AC-Liste folgt Block 3 (`docs/spec/acceptance-criteria.md`)
- [-] **Test-Strategie + Technologien (5)** — bUnit + xUnit + FluentAssertions + NSubstitute etabliert; dediziertes `docs/spec/testing-strategy.md` folgt Block 3
- ✓ **Unit-Tests (3)** — 17 bUnit-Tests + 14 Smoke-Tests (`SmokeTests.cs`), 31/31 grün
- ✓ **Test-Ergebnisse dokumentiert (3)** — CHANGELOG [Unreleased] mit Test-Counts; `dotnet test` voll grün

KI, Sub-Systeme & Reflexion

- [-] **KI-Werkzeug-Nutzung beschrieben (12)** □ — `/ui-brainstorm` → `/ui-flow` → `/ui-build` → `/ui-review` Pipeline als KI-gestützter Workflow etabliert und in Block-2-Insights dokumentiert; konsolidiertes `docs/ai-usage.md` folgt Block 3
- ✓ **Intelligente Services mit KI (6)** — out of scope Block 2 (Stubs-only); □ geliefert Block 3 (`KeywordClassifier` + `AiClassifier`) mit `KeywordClassifier` + `AiClassifier`
- ✓ **Sub-Systeme als Container (5)** — out of scope Block 2; □ geliefert Block 5 (docker-compose Stack)
- [-] **KI-Reflexion / Fazit (7)** — Block-2-Insights enthalten Workflow-Reflexion (Phase-Disziplin, KI-generierte Tests, MudBlazor-Adaptionen); finales Fazit gehört ins Submission-PDF (Block 5)

TODO

□ Done

- ✓ Master-Detail-Übung (separate Moodle-Aufgabe — siehe `docs/master-detail.html` im Repo)

- ✓ **Tech-Entscheid** — ADR 0001 im Repo: `docs/adr/0001-frontend-render-mode-and-architecture.md`
 - Blazor Interactive Server als Default Render Mode
 - OIDC gegen bestehende Authentik (Homelab SSO)
 - Web UI ist selbst ein Channel (`WebChannel`) neben Telegram
 - REST API nur für Non-UI Consumer (Telegram, Integrationen, Automation)
 - Bogus für Faker-Testdaten
- ✓ **FlowHub Glossary** — `Projektarbeit/Glossary.md` : Capture, Skill, Channel, Integration, Page/Component/Card/Widget, Render Mode
- ✓ **Dashboard wireframe (Phase 1)** — Repo: `docs/design/dashboard/wireframe.md`

□ Scaffolding

- ✓ Scaffold `source/FlowHub.Web/` aus dem Blazor Web App Template (Interactive Server, kein WASM)
- ✓ MudBlazor verkabeln: `Program.cs` , `App.razor` , `_Imports.razor` , `MainLayout.razor`
- ✓ `MudLayout` Shell bauen — `MudAppBar` + Mini `MudDrawer` (click-to-expand) + `MudMainContent` + User Menu
- ✓ AppBar **Quick-Capture Field** an Stub `CaptureService.Submit(...)` verdrahtet (WebChannel-Eingang, sichtbar auf jeder Page)
- ✓ `DevAuthHandler` — fixer `ClaimsPrincipal` "Dev Operator", nur registriert wenn `IsDevelopment()`
- ✓ **Bogus** Dependency in `Directory.Packages.props`
- ✓ Stub-Service-Interfaces in `FlowHub.Core` : `ICaptureService` , `ISkillRegistry` , `IIntegrationHealthService` + Bogus-basierte Stubs in `source/FlowHub.Web/Stubs/`
- ✓ Test-Projekt `tests/FlowHub.Web.ComponentTests/` mit bUnit

□ MVP Path — Per-Page UI Workflow

Page 1 — Dashboard (/)

- ✓ Phase 1 — `/ui-brainstorm` (Wireframe)
- ✓ Phase 2 — `/ui-flow` (Mermaid Diagrams)
- ✓ Phase 3 — `/ui-build` (Shell + Shared Components + Stubs)
- ✓ Phase 4 — `/ui-review` (12 bUnit Tests)

Page 2 — New Capture (/captures/new)

- ✓ Phase 1 — /ui-brainstorm
- ✓ Phase 2 — /ui-flow
- ✓ Phase 3 — /ui-build
- ✓ Phase 4 — /ui-review (3 bUnit Tests)

Page 3 — Captures list (/captures)

- ✓ Phase 1 — /ui-brainstorm
- ✓ Phase 2 — /ui-flow
- ✓ Phase 3 — /ui-build (with lifecycle/channel filter chips + text search + pagination)
- ✓ Phase 4 — /ui-review (2 bUnit Tests)

☐ Stretch — alle 3 geschafft

Page 4 — Capture detail (/captures/{id})

- ✓ Alle 4 Phasen — inkl. Orphan-Retry / Unhandled-Reassign Action Stubs (snackbar "Coming in Block 3")

Page 5 — Skills (/skills)

- ✓ Implementiert (MudDataGrid + HealthDot, read-only)

Page 6 — Integrations (/integrations)

- ✓ Implementiert (MudDataGrid + HealthDot, read-only)

☐ Cross-cutting Verification

- ✓ dotnet test voll grün — 17 Tests
- ✓ CHANGELOG.md [Unreleased] Section mit Block-2 Deliverables
- ✓ Diese TODO-Liste final tick-marken
- ✓ ~~Manueller Durchlauf~~ → automatisiert via 14 bUnit Smoke Tests (SmokeTests.cs), 31/31 grün

☐ Out of Scope (Block 2) — geparkt

Sind bereits in ADR 0001 als "out of scope" dokumentiert:

- Settings / Preferences Page
 - Skill Suggestion Review Queue
 - Audit Log Viewer
 - Multi-User / RBAC
 - Real Authentik Client Registration → Block 5
 - Real Persistence → Block 4 (aktuell In-Memory Bogus Stubs)
 - Live SignalR Push neuer Captures (nice-to-have, ggf. Block 3)
 - Charts / Metrics Visualisierungen
-

Verweise

- Repo: [[Repository]] — github.com/freaxnx01/FlowHub-CAS-AISE
- Glossary: [[Glossary]] — Capture, Skill, Channel, Integration, UI Vocabulary
- ADR 0001: [docs/adr/0001-frontend-render-mode-and-architecture.md](#) (im Repo)
- Dashboard Wireframe: [docs/design/dashboard/wireframe.md](#) (im Repo)
- Konzept: [[Idee FlowHub]]

Block 3 — Nachbereitung

Source: `vault/Blöcke/03 Service/03 Service - c) Nachbereitung.md`

tags:

- claude-generated
 - claude-updated updated: 2026-05-04
-

Block 3 — Service · Nachbereitung

Phase budget: 26 h **PVA war:** 2026-04-25 **Nächste PVA:** 2026-05-23

Lernziel

- Ich kann Microservices- und Service-based Architekturen entwerfen.
- Ich bin fähig, die verschiedenen Protokolle wie SOAP, REST und gRPC zu nutzen.
- Ich kann Service-Discovery und Service-Mesh für GenAI entwickeln.
- Ich kann mit KI flexible Microservice-Architekturen bauen.
- Ich kann mit Spring-AI und Koog AI Agenten bauen.

Auftrag (Moodle)

Wahrscheinlich sind Sie mit einer monolithischen oder einer einfachen Schichtenarchitektur im letzten Block gestartet, um das Frontend zu implementieren. Nun möchten wir die Services schrittweise in unabhängige Microservices auftrennen und verschiedene Technologien verwenden, diese flexibler und resilienter zu machen; resilienter — beispielsweise durch eine asynchrone Kommunikation mittels einer Queue. Wir wollen aber auch die KI selbst in den Services verwenden, um diese intelligenter zu gestalten. Nutzen Sie OpenAPI und MicroProfile REST Clients, um die Konsistenz zwischen Server und Client sicherzustellen. Am Ende dieser Nacharbeit haben Sie eine verteilte Lösung, in der Ihre Microservices fiktive oder statische Daten zurückgeben.

Termin: Bis zur nächsten Präsenzveranstaltung (2026-05-23).

Reflexion & Auswertung: Das Ergebnis dieser Projektarbeit ist Grundlage für den nächsten Block, wo die Lösung um die Persistenzschicht erweitert wird.

FlowHub-Kontext: ADR 0002 (Accepted, 2026-04-17) hält fest, dass FlowHub ein Modular Monolith bleibt. Statt physischem Service-Split wird die Auftrennung über eine **MassTransit-basierte Async-Pipeline** (Capture-Enrichment, Skill-Routing) und eine eigene **REST-API in** `source/FlowHub.Api/` für Nicht-UI-Clients realisiert. gRPC und Multi-Prozess-Deployment sind explizit out of scope. Die Lernziel-Bullets zu Spring-AI / Koog werden als Stack-neutrale Konzepte interpretiert und mit .NET-Äquivalenten ([Microsoft.Extensions.AI](#) / Semantic Kernel) abgedeckt.

Bewertungskriterien (Block 3)

Pflichtcheck am Ende jeder Nachbereitung — die offizielle Moodle-Rubrik aus [[Bewertungskriterien]] für **diesen** Block durchgehen, bevor "fertig" geclaimed wird. Punkteangaben in Klammern zeigen Max-Score (volle Erfüllung).

Quarkus/Jakarta-EE-Item ist für FlowHub (.NET-Stack) **nicht relevant** — bewusst ausgeklammert.

Spezifikation

- ✓ **Use Cases & fachliche Anforderungen benannt (5)** — UC-01..UC-07 (Web UI / Telegram / Dashboard) + UC-08..UC-11 (REST submit, async pipeline, AI fallback, retry) in `docs/spec/use-cases.md`
- ✓ **NfA SMART spezifiziert (5)** — NF-01..NF-08 (Block 2 UI) + NF-09..NF-13 (API latency, retry budget, AI fallback rate, AI cost, OpenAPI versioning SLA) in `docs/spec/use-cases.md`
- ✓ **Solution Vision beschrieben (5)** — Modular Monolith + REST + MassTransit-Async-Pipeline (ADR 0002 + ADR 0003); REST-API-Surface in `docs/design/api/api-surface.md`

Entwurf

- ✓ **Lösungsansatz & Architektur textuell + bildlich (7)** — ADR 0003 (Async Pipeline) □, ADR 0004 (KI-Integration) □ (landed at `docs/adr/0004-ai-integration-in-services.md`, Slice C)

✓ **Struktur / Verhalten / Interaktion (7):**

- Struktur: System-Context (docs/spec/system-context.md , C4 L1) + ADR 0002 module split
- Verhalten: Sequence-Diagramme in docs/design/sequences/ (Capture-Enrichment happy + AI fallback; Skill-Routing happy + retry exhaustion)
- Interaktion: API-Surface (docs/design/api/api-surface.md , OpenAPI/Scalar) + Event-Kontrakte (CaptureCreated , CaptureClassified , Fault<CaptureClassified>)

✓ **DB-Modell spezifiziert (3)** — Block-4-preview Entity-Skizze in docs/design/db/entities.md (Mermaid ER + indexes + deferred-decisions list)

Programmierung

✓ **Code lesbar, dokumentiert, nach Layer/Modul/Sub-System strukturiert (7)** — Slice B: hexagonale Ports (IClassifier ,

ISkillIntegration), Pipeline-Modul, EventId-Namespacing (1000-1999 / 2000-2999), LoggerMessage source-gen

✓ ~~Quarkus / Jakarta EE / moderne Java-Konzepte~~ — N/A (Stack: .NET 10), siehe [[Bewertungskriterien]]; consciously skipped, will be noted in submission PDF

✓ **Erkenntnisse aus der Programmierung dokumentiert (3)** — ADR 0003, docs/ai-usage.md (Reflexion-Sektion mit △/□ Tabelle); CHANGELOG noch offen

✓ **Source in Git-Repository (2)** — github.com/freaxnx01/FlowHub-CAS-AISE ; Slice-B-Branch feat/block3-async-pipeline lokal committed, Push folgt nach Final Code Review

Validierung

✓ **Abnahmekriterien definiert (5)** — UC-08..UC-11 in docs/spec/use-cases.md carry their own acceptance criteria; 17 integration tests in tests/FlowHub.Api.IntegrationTests/ + 25 unit tests in tests/FlowHub.Web.ComponentTests/Ai/ serve as executable acceptance per endpoint / classifier behaviour

✓ **Test-Strategie + Technologien spezifiziert (5)** — docs/spec/testing-strategy.md updated to Block-3 reality: 99 default-suite + 4 trait-gated live tests; tools (xunit, bUnit, NSubstitute, FluentAssertions, WebApplicationFactory<Program> , MassTransit Test Harness, Xunit.SkippableFact) and Slice-C-specific patterns (NSubstitute on IChatClient , [SkippableFact] , Category=AI trait filter) documented

- ✓ **Unit-Tests programmiert (3)** — 99 default-suite tests passing (82 component + 17 API integration) + 4 trait-gated live tests (Slice C). Build clean under warnings-as-errors
- ✓ **Test-Ergebnisse dokumentiert (3)** — `CHANGELOG.md` `[Unreleased]` populated with Slice A/B/C deliverables; ADR 0003 / 0004 reference test counts + EventId namespacing; `docs/ai-usage.md` per-slice retros include test outcomes

KI, Sub-Systeme & Reflexion

- ✓ **KI-Werkzeug-Nutzung beschrieben (12)** □ höchstgewichtetes Kriterium — `docs/ai-usage.md` (Living Doc): Tool-Stack, Workflow (brainstorming → writing-plans → SDD), generierter vs. handgeschriebener Anteil, notable Adaptationen die Subagents gefunden haben (z.B. Captive-Dependency-Trap)
- ✓ **Intelligente / flexible Services mit KI gebaut (6)** — `KeywordClassifier` (deterministische Heuristik) □ + `AiClassifier` (Anthropic Haiku 4.5 / OpenRouter Llama 3.1 70B via MEAI) □; graceful fallback bei AI-Fehlern dokumentiert (ADR 0004, EventId 3010)
- ✓ **Sub-Systeme als unabhängige Container deploybar (5)** — Sketch-Level: `docker-compose.yml` (web + rabbitmq + future api) im Repo; finale Block-5-Variante folgt
- ✓ **KI-Erfahrungen als Fazit reflektiert (7)** — Reflexion-Sektion in `docs/ai-usage.md` mit □ / △ / □ Auflistung; PVA-Folien-Material noch zu destillieren

TODO

Architektur & Entscheide

- ✓ ADR 0003 — Async Messaging Pipeline (MassTransit Topology, Retry/DLQ, Fault-Observer; Outbox auf Block 4 verschoben)
- ✓ ADR 0004 — KI-Integration in Services (Provider, Abstraction, Prompt-/Cost-Strategie, Stack-Mapping Spring-AI → .NET) — landed at `docs/adr/0004-ai-integration-in-services.md`
- ✓ OpenAPI-Versionierungs- und Konsistenzstrategie festlegen (Scalar als UI, Refit/Kiota als Client-Generator) — **partial**: Scalar wired and serving `/scalar` ; Refit/Kiota client generation deferred to Slice D

REST-API (`source/FlowHub.Api/`)

- ✓ Projekt scaffolden + in `FlowHub.slnx` registrieren
- ✓ Endpoints: `GET/POST /api/captures`, `GET /api/captures/{id}`, `POST /api/captures/{id}/retry`, `GET /api/skills`, `GET /api/integrations` — **partial**: captures endpoints (`GET /api/v1/captures`, `POST /api/v1/captures`, `GET /api/v1/captures/{id}`, `POST /api/v1/captures/{id}/retry`) all landed; `GET /api/skills` and `GET /api/integrations` **deferred to v2** per Slice A spec D7
- ✓ FluentValidation am Boundary, ProblemDetails (RFC 9457) für Fehler
- ✓ OpenAPI + Scalar UI (`/scalar`)
- ✓ Stub-Daten-Quelle (Bogus/In-Memory aus Block 2 wiederverwenden) — keine echte Persistenz (Block 4)

Async-Pipeline (MassTransit)

- ✓ MassTransit + In-Memory-Transport in `FlowHub.Web` registriert — `FlowHub.Api` - Wiring kommt mit Slice A
- ✓ Event-Kontrakte definiert: `CaptureCreated`, `CaptureClassified` — `SkillRoutingRequested` + `IntegrationCallFailed` bewusst gestrichen (ADR 0003 D2: Routing-Consumer schreibt inline; Failures sind `Fault<CaptureClassified>`)
- ✓ Consumers: `CaptureEnrichmentConsumer`, `SkillRoutingConsumer` (statt `... Handler`), `LifecycleFaultObserver`
- ✓ Retry-/DLQ-Policy + Test Harness Tests — per-Consumer-Retry (100/500ms enrichment, 500/2000/5000ms routing); Fault-Observer ohne Retry; 6 Tests
- ✓ RabbitMQ-Profil vorbereiten — Docker-Compose-Sketch im Repo; lokales Aktivieren wird in Block 5 fertiggestellt

KI in Services

- ✓ `Microsoft.Extensions.AI` (oder Semantic Kernel) als Abstraktion einbinden — landed in `source/FlowHub.AI/` via `IChatClient` ; Anthropic + OpenRouter adapters (Slice C)
- ✓ Klassifikator-Stub: `Capture` → vorgeschlagene Tags / Skill — `KeywordClassifier` (deterministische URL/todo-Heuristik, 4 Tests grün); Mock-Provider via `NSubstitute` in Pipeline-Tests
- ✓ Reflexion zu `Microsoft.Extensions.AI` / Semantic Kernel im .NET-Stack — captured in ADR 0004 §"Reflexion" + `docs/ai-usage.md` Slice-C Reflexion-Sektion

Typed Clients (Konsistenz Server↔Client) — deferred

Decision (2026-05-04): Deferred indefinitely. Not on the Beta-MVP path — the Blazor UI talks to `ICaptureService` in-process (per ADR 0001 D2), so a typed REST client is only needed by future external consumers (Telegram bot, mobile). When that lands, evaluate Refit vs Kiota then; today there is no live consumer to wire.

- ✓ ~~Refit (oder Kiota) als MicroProfile-REST-Client-Äquivalent evaluieren und entscheiden~~ — deferred (no live consumer)
- ✓ ~~Generierter/typed Client für `FlowHub.Api` aus OpenAPI~~ — deferred (no live consumer)
- ✓ ~~Blazor-Pages über typed Client gegen die REST-API verdrahten (statt direkten Stub-Service-Calls)~~ — deferred; ADR 0001 D2 specifies UI uses in-process services, not the REST API

Validierung & Tests

- ✓ Test-Strategie als Dokument — lives at `docs/spec/testing-strategy.md` ; refreshed for Block 3 reality (99 default-suite + 4 trait-gated; tools, patterns, deferrals)
- ✓ Akzeptanzkriterien je Use Case — UC-08..UC-11 in `docs/spec/use-cases.md` carry Postcondition / Error sections; 17 API integration tests + 25 Slice-C unit tests serve as executable acceptance per endpoint / classifier behaviour
- ✓ Unit-Tests für Handlers / Validators — Slice B: KeywordClassifier (4), CaptureServiceStub (6), Pipeline-Consumers (6); Validators kommen mit Slice A
- ✓ Component-Tests für API-Endpoints (`Microsoft.AspNetCore.Mvc.Testing`) — Slice A: 17 integration tests via `WebApplicationFactory<Program>` in `tests/FlowHub.Api.IntegrationTests/`
- ✓ MassTransit Test Harness für Consumers — `PipelineTestBase` + 6 Harness-basierte Tests (Enrichment, Routing, Fault-Observer)
- ✓ `dotnet test` voll grün — 47/47, Build mit warnings-as-errors clean; CHANGELOG-Eintrag noch offen

Spezifikation & Dokumentation

- ✓ Use-Case-Liste ergänzen / aktualisieren — UC-08..UC-11 in `docs/spec/use-cases.md`

- ✓ SMART NfAs für API + Async-Pipeline festhalten — NF-09..NF-13 in `docs/spec/use-cases.md`
- ✓ DB-Modell-Skizze (Vorbereitung für Block 4) — `docs/design/db/entities.md`
- ✓ CHANGELOG `[Unreleased]` mit Block-3-Deliverables füllen — Slice A/B/C entries landed

□ Out of Scope (Block 3) — geparkt

- Physischer Microservice-Split / Multi-Prozess-Deployment (out of scope per ADR 0002)
- gRPC in FlowHub (POC-only, siehe `poc/restful-api-playground/`)
- Echte Persistenz → Block 4
- Authentik / OIDC → Block 5
- Service Mesh / Service Discovery in echtem Cluster → Block 5

Verweise

- Repo: `[[Repository]]` — `github.com/freaxnx01/FlowHub-CAS-AISE`
- Block 3 Vorbereitung: `[[03 Service - a) Vorbereitung]]`
- Block 2 Nachbereitung: `[[02 Frontend - c) Nachbereitung]]`
- ADR 0001: `docs/adr/0001-frontend-render-mode-and-architecture.md`
- ADR 0002: `docs/adr/0002-service-architecture-and-async-communication.md`
- POC: `poc/restful-api-playground/` (REST + gRPC + RabbitMQ Sandbox)
- Bewertungskriterien: `.ai/cas-instructions.md` (Section "Grading")

Beta MVP follow-up (2026-05-04): Persistence + real Skill integrations landed early via the Beta-MVP slice (see `docs/superpowers/specs/2026-05-04-beta-mvp-design.md` for design and `docs/superpowers/plans/2026-05-04-beta-mvp.md` for the plan). The Block-3 rubric closure already covered the Block-3 criteria; this is just an FYI cross-reference.

Block 4 — Nachbereitung

Source: `vault/Blöcke/04 Persistence/04 Persistence - c) Nachbereitung.md`

tags:

- claude-generated
 - claude-updated updated: 2026-05-06
-

Block 4 — Persistence · Nachbereitung

Phase budget: 22 h **PVA war:** 2026-05-23 **Nächste PVA:** 2026-06-20

Lernziel

- Ich kann für die jeweilige Problemstellung die geeignete Persistenzform bestimmen.
- Ich kann mit Spezifikationen wie ORM, JPA und Alternativen den DB-Zugriff abstrahieren.
- Ich kann dynamische Abfragen effizient programmieren.
- Ich kann mit Quarkus Panache Datenzugriffe auf verschiedene Datenbankmodelle realisieren.

Auftrag (Moodle)

Wir wenden uns nun der dritten und letzten Schicht einer klassischen Enterprise Applikation zu, um Daten effizient zu speichern und für die entsprechenden Geschäftsprozesse flexibel zu nutzen. Da Daten jede Technologie überlebt und den wertvollen Teil Ihrer Applikation darstellt, ist das zugrunde liegende Datenmodell mit Bedacht zu wählen. Daten in ein neues Datenmodell zu migrieren oder fehlende Daten nachzuliefern, ist aufwändig. Entsprechend soll Ihr Datenmodell zukünftige Bedürfnisse antizipieren und daran anpassbar sein.

Entwerfen Sie hier nun das geeignete Datenmodell und implementieren Sie dessen Abstraktion über Hibernate ORM, Panache und Jakarta Data. Nutzen Sie die Möglichkeiten von Criteria API, um dynamische Abfragen zu realisieren.

Termin: Bis zur nächsten Präsenzveranstaltung (2026-06-20). Dieser Auftrag ist Grundlage für die weitere Arbeit in der Präsenzveranstaltung.

Reflexion & Auswertung: Die Applikation speichert nun die Daten persistent.

FlowHub-Stack-Mapping (.NET statt JVM):

- Hibernate ORM / Jakarta Persistence (JPA) → **EF Core 10**
(`Microsoft.EntityFrameworkCore` + `Npgsql.EntityFrameworkCore.PostgreSQL`)
- Panache (Active-Record / Repository) → EF Core `DbSet<T>` + Repository-Pattern oder `DbContext` direkt
- Jakarta Data → keine direkte Entsprechung; konzeptuell deckt EF Core's `IQueryable<T>` + LINQ den Bereich ab
- Criteria API (typsichere dynamische Abfragen) → **LINQ + Expression Trees** (`Expression<Func<T, bool>>` für dynamische Filter)
- Migrations → `dotnet ef migrations add` / `dotnet ef database update`
(Workflow steht in `CLAUDE.md`)

Datenbank: PostgreSQL (Docker), Connection-String über `ConnectionStrings__Default` ENV-Variable.

Beta-MVP follow-up (2026-05-04): Ein vertikaler Beta-Slice

(`docs/superpowers/specs/2026-05-04-beta-mvp-design.md` + `docs/superpowers/plans/2026-05-04-beta-mvp.md`) hat einen Teil des Block-4-Scopes vorgezogen, um die Architektur an einer realen Homelab-Demo zu validieren:

- **Geliefert:** `FlowHub.Persistence` aktiv, `FlowHubDbContext` (Captures-`DbSet`), `EfCaptureService`, `Initial-Migration`, `AddFlowHubPersistence` -Extension, `MigrationRunner` (`IHostedService`), Cursor-Pagination in `ListAsync` .
- **Bewusst aufgeschoben für Block 4: PostgreSQL** (Beta nutzt SQLite), separate `EntityTypeConfiguration<T>` -Klassen, Repository-Pattern-Layer, vollständiges Domänenmodell (Skill/SkillRun/Channel/Integration/Tag), `make db-up` / `db-migrate` , Testcontainers-Tests, ADR 0005,

`docs/insights/block-4.md` , Migrations als separater Init-Container (12-Factor XII).

Die unten getickten Items sind vom Beta-Slice abgedeckt; alle anderen bleiben offen und werden in der regulären Block-4-Phase (ab 2026-05-09) bearbeitet.

Bewertungskriterien (Block 4)

Pflichtcheck am Ende jeder Nachbereitung — die offizielle Moodle-Rubrik aus [[Bewertungskriterien]] für **diesen** Block durchgehen, bevor "fertig" geclaimed wird. Punkteangaben in Klammern zeigen Max-Score.

Quarkus/Jakarta-EE-Item ist für FlowHub (.NET-Stack) **nicht relevant** — bewusst ausgeklammert.

Spezifikation

- ✓ **Use Cases (5)** — datenseitige Use Cases benannt: Capture-CRUD, Such-/Filter-Abfragen über Lifecycle/Channel/Tags, Skill-Run-Historie, Integration-Health-Verlauf
- ✓ **NfA SMART (5)** — Persistenz-NfAs: Query-Latenz, Index-Anforderungen, Datenvolumen-Annahmen, Backup/Restore-Ziele, Migrations-Strategie (Zero-Downtime?)
- ✓ **Solution Vision (5)** — aktualisiert um Persistenzschicht: PostgreSQL via EF Core, Migrations-First-Ansatz, Repository-Abstraktion-Entscheid

Entwurf

- ✓ **Lösungsansatz & Architektur textuell + bildlich (7)** — ADR 0005 (`docs/adr/0005-persistence.md` , Accepted, 2026-05-04): Provider-Wahl (SQLite Beta → PostgreSQL Block 4), ORM EF Core 10, kein Repository-Layer (`EfCaptureService` = `ICaptureService` -Adapter direkt), Migrations-Workflow (`dotnet-ef` Tool-Manifest + `MigrationRunner` `IHostedService` Beta → separate init-container Block 5), `internal sealed` + `InternalsVisibleTo` , Cursor-Pagination keyset, Index-Strategie
- ✓ **Struktur / Verhalten / Interaktion (7):**
 - Struktur: `FlowHub.Persistence` -Layer, `DbContext` , Entities, Migrations, Repositories

- Verhalten: CRUD-Flows, Migrations-Run, Query-Plans für Hot-Path-Queries
- Interaktion: Aufrufkette `FlowHub.Web/.Api` → `Application Service` → `Repository` → `DbContext` → `DB`

- ✓ **DB-Modell vollständig (3)** □ in diesem Block der Schwerpunkt — ER-Diagramm mit Entitäten (Capture, Skill, SkillRun, Channel, Integration, IntegrationHealthSample, Tag, ...), Beziehungen, Indizes, Constraints. Antizipation zukünftiger Bedürfnisse begründen

Programmierung

- ✓ **Code lesbar/dokumentiert/strukturiert (7)** — `FlowHub.Persistence` als eigenes Projekt, sauber getrennt von Domain (`FlowHub.Core`)
- ✓ ~~Quarkus/Jakarta EE/moderne Java-Konzepte~~ — N/A (Stack: .NET 10)
- ✓ **Erkenntnisse dokumentiert (3)** — `docs/ai-usage.md` Block-4-prep / Beta-MVP-Sektion: dual-provider EF-Core-8+ trap, `InternalsVisibleTo` -Pattern für Test-Seeding, surgical `MigrationRunner` -Removal in `IntegrationTestFactory` , `IDesignTimeDbContextFactory` für Tooling-Discovery, captive- `HttpClient` -Anti-Pattern (gemerkt für Block 4 Cleanup); ADR 0005 §"Alternatives considered" deckt Dapper/NHibernate/Repository-Pattern Trade-offs ab
- ✓ **Source in Git (2)** — alle Block-4-Commits gepusht

Validierung

- ✓ **Abnahmekriterien (5)** — pro Use Case Datenflüsse spezifiziert, inkl. Edge Cases (leere Resultate, Konkurrenz, Migrations-Rollback)
- ✓ **Test-Strategie (5)** — Erweiterung von `docs/test-strategy.md` : Repository-Tests gegen Testcontainers PostgreSQL; Integration-Tests gegen reale DB; Migrations-Tests
- ✓ **Unit-Tests (3)** — Repository-Implementierungen, Query-Builder, Migrations-SmokeTest
- ✓ **Test-Ergebnisse dokumentiert (3)** — `dotnet test` voll grün; Counts + Coverage in CHANGELOG `[Unreleased]`

KI, Sub-Systeme & Reflexion

- ✓ **KI-Werkzeug-Nutzung beschrieben (12)** □ — wie wurde KI im Datenmodell-Entwurf, in Migrations-Generierung, in Query-Optimierung eingesetzt? Doku-Update in `docs/ai-usage.md`

- ✓ **Intelligente Services mit KI (6)** — KI-Klassifikator nutzt jetzt persistente Daten (statt Stubs); ggf. Embedding-Speicherung für Suche (Vorbereitung Block 5)
- ✓ **Sub-Systeme als Container (5)** — PostgreSQL als eigener Container; [FlowHub.Web/.Api](#) gegen DB-Container deploybar; Compose-Profil aktualisiert (auch Migrations als separater Init-Container — siehe 12-Factor XII in `CLAUDE.md`)
- ✓ **KI-Reflexion / Fazit (7)** — Block-4-Reflexion: KI bei Datenmodellierung — Stärken (Boilerplate, Migrations) und Schwächen (Schema-Antizipation, Performance-Blindheit)

TODO

Datenmodell

- ✓ ER-Diagramm der FlowHub-Domäne (`docs/design/db/er.md` oder Mermaid in ADR 0005)
- ✓ Entity-Klassen in `FlowHub.Persistence/Entities/` (Capture, Skill, SkillRun, Channel, Integration, IntegrationHealthSample, Tag) — Domain-POCOs (`Capture` , ...) bleiben in `FlowHub.Core` ; EF-spezifische `*Entity` + `*EntityTypeConfiguration` liegen in `FlowHub.Persistence` (per ADR 0005 §3)
- ✓ Indizes / Constraints / Unique-Keys spezifizieren
- ✓ Soft-Delete-Strategie entscheiden (Lifecycle vs. echte Löschung) — **Hard-Delete via FK CASCADE**; Soft-Delete-Semantik wird vom `LifecycleStage` -State-Machine getragen (`Orphan` / `Unhandled` / `Completed`). Doku: `docs/design/db/er.md` § Delete Strategy.
- ✓ Audit-Felder (`CreatedAt` , `UpdatedAt` , `CreatedBy` , ...) konsequent — **Block-4-Scope: nur `CreatedAt`** auf allen Entities. `UpdatedAt` / `CreatedBy` bewusst aus dem Scope: FlowHub ist Single-User, alle Mutationen verlaufen über klar typisierte Lifecycle-Übergänge (`stage = Wechsel`), und die Bus-Events (`CaptureCreated` / `CaptureClassified`) sind die Audit-Quelle. Wenn Multi-User in einer zukünftigen Iteration relevant wird, ist Nachrüsten über eine `0005_AddAuditFields` -Migration trivial.

Architektur & Entscheide

- ✓ ADR 0005 — Persistence (Provider, ORM, Repository-Pattern-Entscheid, Migrations-Workflow) — `docs/adr/0005-persistence.md` (Accepted, 2026-05-

04)

- ✓ Stack-Mapping-Notiz: Hibernate/Panache/Jakarta Data → EF Core (kurze Doku, warum, was äquivalent) — landed im Auftrag-Intro oben (FlowHub-Stack-Mapping)

Implementierung (`source/FlowHub.Persistence/`)

- ✓ Projekt scaffolden, in `FlowHub.slnx` registrieren — Beta MVP (`source/FlowHub.Persistence/FlowHub.Persistence.csproj`)
- ✓ `FlowHubDbContext` mit `DbSet<T>` für alle Entities — **partial**: Beta MVP nur `Captures` -`DbSet`; weitere Entities (Skill, SkillRun, Channel, Integration, Tag) in Block 4
- ✓ `EntityTypeConfiguration<T>` pro Entity (Fluent API statt Annotations) — Beta nutzt inline `OnModelCreating` ; Refactor zu separaten Config-Klassen in Block 4
- ✓ PostgreSQL-Connection via `ConnectionStrings__Default` ENV — Beta nutzt SQLite; PostgreSQL-Switch in Block 4
- ✓ Initial-Migration `0001_Initial` generieren (`dotnet ef migrations add`) — Beta MVP (`Migrations/20260504120638_Initial.cs`)
- ✓ Repository-Interfaces in `FlowHub.Core` , Implementierungen in `FlowHub.Persistence` — Beta MVP nutzt `ICaptureService` direkt gegen `DbContext` ; Repository-Layer-Entscheid in ADR 0005
- ✓ DI-Registration als `IServiceCollection` -Extension (`AddFlowHubPersistence(connectionString)`) — Beta MVP (`AddFlowHubPersistence(IConfiguration)` mit `ConnectionStrings:Default` -Lookup)

Dynamische Abfragen

- ✓ LINQ + Expression Trees als "Criteria-API"-Äquivalent — Beispiel: dynamischer Capture-Filter (Lifecycle, Channel, Tags, Search) — Beta MVP nur statische Filter; voller dynamischer Filter inkl. Tag/Search in Block 4
- ✓ Pagination-Helper (Skip/Take + Cursor-basiert für lange Listen) — Beta MVP (`EfCaptureService.ListAsync` mit `CaptureCursor` keyset-pagination, `limit+1` -Probe)
- ✓ N+1-Problem aktiv vermeiden (`Include` / Projektion / Read-Models) — `EfCaptureRepository` projiziert direkt auf das `Capture` -Record (kein Lazy-Loading der Owned-Entities Tags/SkillRuns); Tag-Filter ist als `EXISTS` -Subquery formuliert, nicht als Join + Group. Verifiziert durch `EfCaptureRepositoryTests` (29 Tests, Testcontainers).

Migrations & Deployment-Vorbereitung

- ✓ Migrations laufen separat (nicht in `app.Run()`) — 12-Factor XII
- ✓ `make db-up` / `make db-migrate` Targets oder Skript
- ✓ Docker-Compose-Snippet für PostgreSQL (lokal testen)
- ✓ EF-Migrations-Bundle / SQL-Script-Generierung getestet (für Block 5 Production-Deployment)

Stub-Replacement

- ✓ Bogus-basierte Stubs aus Block 2 durch DB-gestützte Implementierungen ersetzen — *aber* Bogus-Seed für Dev-Mode behalten (`appsettings.Development.json`-Flag) — Beta MVP hat Bogus-Stubs entfernt; deterministisches Seeding nur im `IntegrationTestFactory` für API-Tests; Dev-Mode-Seed offen
- ✓ `ICaptureService`, `ISkillRegistry`, `IIntegrationHealthService` jetzt gegen Repository statt In-Memory — **partial**: nur `ICaptureService` swap (Beta MVP); `ISkillRegistry` + `IIntegrationHealthService` bleiben Stubs (per Beta-Spec D3 — orthogonal zu Architektur-Validierung)

Tests

- ✓ Test-Strategie ergänzen (`docs/spec/testing-strategy.md`) — Pfad wurde im Repo unter `docs/spec/` einsortiert (mit dem Rest der Spezifikations-Artefakte), nicht direkt unter `docs/`. Inhaltlich vollständig (Unit / Component / Integration / E2E / MassTransit-Harness).
- ✓ Repository-Tests (Testcontainers PostgreSQL bevorzugt vor SQLite-In-Memory wegen Provider-Quirks)
- ✓ Integration-Tests für API-Endpoints gegen reale DB
- ✓ Migrations-Smoke-Test (auf leerer DB rauf, runter, rauf)

Spezifikation & Doku

- ✓ CHANGELOG [Unreleased] mit Block-4-Deliverables — **partial**: Beta-MVP-Persistence-Eintrag landed (commits `d7b7af0..c5340e9`); reine Block-4-Erweiterungen (PostgreSQL, Repository-Pattern, ADR 0005) kommen nach
- ✓ Use-Case-Liste um datenseitige Use Cases erweitern
- ✓ Performance-NfAs SMART formulieren (mit Mess-Methodik)
- ✓ `docs/insights/block-4.md` — Erkenntnisse Datenmodellierung mit KI

□ Out of Scope (Block 4)

- Read-Replicas / Sharding
- Echte Embedding-/Vector-Suche → Block 5 (KI-Suche)
- Production-Backup-Tooling → Block 5
- Authentik / OIDC → Block 5

Verweise

- Repo: [[Repository]] — github.com/freaxnx01/FlowHub-CAS-AISE
- Block 4 Vorbereitung: [[04 Persitence - a) Vorbereitung]]
- Block 3 Nachbereitung: [[03 Service - c) Nachbereitung]]
- Block 5 Nachbereitung: [[05 Deployment - c) Nachbereitung]]
- ADR 0001: [docs/adr/0001-frontend-render-mode-and-architecture.md](#)
- ADR 0002: [docs/adr/0002-service-architecture-and-async-communication.md](#)
- ADR 0005: [docs/adr/0005-persistence.md](#) — Persistence (Beta MVP)
- Bewertungskriterien: [[Bewertungskriterien]]

Block 5 — Nachbereitung

Source: `vault/Blöcke/05 Deployment/05 Deployment - c) Nachbereitung.md`

tags:

- claude-generated updated: 2026-05-17
-

Block 5 — Deployment & Abgabe Projektarbeit · Nachbereitung

Phase budget: 22 h **PVA war:** 2026-06-20 **Abgabe-Deadline:** Samstag, 2026-07-04, 24:00 (PDF-Upload mit Repo-URL)

Final Block. Hier wird die Bewertungs-Rubrik aus [[Bewertungskriterien]] **vollständig** geprüft — alle 5 Buckets müssen Abgabe-fähig sein. Was im jeweiligen Block-Nachbereitung nicht erledigt wurde, gehört hier nachgezogen.

Lernziel

- Ich bin fähig, meine Applikation zu containerisieren und in Docker und Kubernetes zu betreiben.
- Ich kann GitHub und Copilot sowie die GitLab-Agent-Plattform einsetzen, um CI/CD-Pipelines aufzusetzen und den Deployment-Prozess zu automatisieren.
- Ich bin fähig, entsprechendes Monitoring und Observation aufzusetzen, Systeme zu überwachen und zu optimieren.
- Ich bin in der Lage, mit Quarkus KI-gestützte Applikationen zu bauen.

Auftrag (Moodle)

In der letzten Nachbearbeitungsphase geht es nun darum, die Lösung zu containerisieren und für den Betrieb zu verteilen. Nutzen Sie die Möglichkeit Ihrer Git-Host-Lösung, um den Deployment-Prozess weitgehend zu automatisieren.

Erweitern Sie Ihre Applikation um KI-basierende Suche und Workflows. Schliessen Sie Ihre Arbeit ab und laden Sie diese als PDF hoch. Die Arbeit enthält die URL auf das Git-Repository Ihrer Lösung.

Termin: Bis zwei Wochen nach der letzten PVA — **konkret: 2026-07-04, 24:00.**

Reflexion & Auswertung: Reflexion über die anschliessende Bewertung der Arbeit und der Lösung.

FlowHub-Stack-Mapping (.NET + GitHub statt Quarkus + GitLab):

- Containerisierung → Multi-Stage Dockerfile (Build: `mcr.microsoft.com/dotnet/sdk:10.0-alpine`, Runtime: `mcr.microsoft.com/dotnet/aspnet:10.0-alpine`, non-root, siehe `CLAUDE.md` § Docker)
- CI/CD → **GitHub Actions** (Repo liegt auf `github.com/freaxnx01/FlowHub-CAS-AISE`); GitLab-Agent-Plattform/-Runner als Lerninhalt zur Kenntnis, Implementierung in GitHub
- Monitoring/Observability → **OpenTelemetry** (Traces, Metrics, Logs) + Prometheus + Grafana (`/metrics` Endpoint ist im Health-Plan); strukturiertes Logging mit Serilog → stdout (12-Factor XI)
- KI-gestützte Apps "mit Quarkus" → mit `Microsoft.Extensions.AI` / Semantic Kernel; KI-Suche via Vector-DB-Provider (z.B. pgvector auf bestehender PostgreSQL aus Block 4)
- Kubernetes → Manifests / Helm-Chart (oder lediglich Docker-Compose, falls K8s-Aufwand sprengt — dann begründen)
- Authentik / OIDC → finaler Schritt (siehe ADR 0001 Plan)

Bewertungskriterien (Final / Block 5)

⚠ **Hier zählt's:** Alle 18 Items aus [[Bewertungskriterien]] müssen Abgabe-fähig sein. Punkte in Klammern = Max-Score.

Quarkus/Jakarta-EE-Item ist für FlowHub (.NET) **nicht relevant** — bewusst ausgeklammert.

Spezifikation

- ✓ **Use Cases (5)** — finale, vollständige Use-Case-Liste der gesamten Applikation (alle Blöcke konsolidiert)
- ✓ **NfA SMART (5)** — vollständige NfA-Liste: Performance, Verfügbarkeit, Sicherheit, Skalierbarkeit, Betrieb (Logs, Monitoring), Deployment-NfAs (Build-Zeit, Image-Grösse)
- ✓ **Solution Vision (5)** — finale Vision: Modular Monolith + Async-Pipeline + KI-Services + PostgreSQL + Container-Deployment + Observability

Entwurf

- ✓ **Lösungsansatz & Architektur textuell + bildlich (7)** — finale Architektur-Doku mit allen ADRs (0001-0006+), C4-Diagramm (Context, Container, Component)
- ✓ **Struktur / Verhalten / Interaktion (7)** — vollständige Sicht: Modul-Struktur, Hot-Path-Sequenzen, Interaktion mit Channels/Integrationen
- ✓ **DB-Modell vollständig (3)** — finales ER-Diagramm + Indizes + ggf. Vector-Spalten für KI-Suche

Programmierung

- ✓ **Code lesbar/dokumentiert/strukturiert (7)** — alle Module sauber, README pro Hauptprojekt, Inline-Doku wo Why nicht obvious
- ✓ **Quarkus/Jakarta EE** — N/A (Stack: .NET 10)
- ✓ **Erkenntnisse dokumentiert (3)** — `docs/insights/` mit Block-1 bis -5 Erkenntnissen
- ✓ **Source in Git (2)** — alles auf `main` gepusht; Tag `v0.1.0` für die Abgabeverision (Directory.Build.props ships 0.1.0; Bewertungskriterien-Item ist 0/2 binär — erreicht)

Validierung

- ✓ **Abnahmekriterien (5)** — vollständige Liste über alle Use Cases, im Submission-PDF aufgeführt
- ✓ **Test-Strategie (5)** — finales `docs/spec/testing-strategy.md` (im Repo unter `docs/spec/` einsortiert, nicht direkt `docs/`): Unit (xUnit/FluentAssertions/NSubstitute), Component (bUnit), Integration ([ASP.NET](#) Mvc.Testing + Testcontainers), E2E (Playwright), MassTransit Test Harness

- ✓ **Unit-Tests (3)** — Coverage über alle Module
- ✓ **Test-Ergebnisse dokumentiert (3)** — CI-Run-Ergebnisse, Coverage-Reports, im Submission-PDF zitiert

KI, Sub-Systeme & Reflexion

- ✓ **KI-Werkzeug-Nutzung beschrieben (12)** □ höchstgewichtetes Kriterium — finaler `docs/ai-usage.md` : alle eingesetzten Tools (Claude Code, Copilot, ChatGPT, Cursor, ...), pro Block welche Aufgaben, Prompt-Strategien, generiert-vs-handgeschrieben-Quote, beobachtete Fehlerklassen
- ✓ **Intelligente Services mit KI (6)** — Capture-Klassifikation + KI-basierte Suche (Embeddings + Vector-Search) + ggf. KI-gestützte Workflows
- ✓ **Sub-Systeme als unabhängige Container (5)** — finale Compose- und/oder K8s-Manifests: [FlowHub.Web](#), FlowHub.Api, PostgreSQL, RabbitMQ, ggf. Authentik, Prometheus/Grafana, alle als getrennte Container; CI baut + pusht Images
- ✓ **KI-Reflexion / Fazit (7)** — finales Kapitel im Submission-PDF: was hat KI im gesamten Projekt geleistet, wo waren Grenzen, persönliche lessons learned, Empfehlungen

TODO

Containerisierung

- ✓ Multi-Stage Dockerfile für FlowHub.Web (build → publish → runtime, non-root user `appuser`)
- ✓ ~~Multi-Stage Dockerfile für FlowHub.Api~~ — **N/A:** FlowHub.Api ist ein Library-Projekt (Minimal-API-Endpoint-Definitionen), kein eigenständig gehosteter Service. Endpoints werden via `MapFlowHubApi()` in FlowHub.Web registriert; das Web-Image enthält die Api.
- ✓ `.dockerignore` sauber (bin/, obj/, node_modules, .git/, ...)
- ✓ EF-Migrations-Container (separater Init-Container — 12-Factor XII)
- ✓ Image-Grösse minimieren (`-alpine` Base, AOT/trim wenn praktikabel)
- ✓ Image-Tags: `<sha>` + `latest` auf main + `vX.Y.Z` auf Release-Tags

Compose / Kubernetes

- ✓ `docker-compose.yml` (Production-orientiert) + `docker-compose.override.yml` (Dev)
- ✓ Services: web, api, postgres, rabbitmq, prometheus, grafana, ggf. authentik
- ✓ Volumes für Postgres-Daten, RabbitMQ, Grafana-Dashboards
- ✓ Healthchecks (Compose) gegen `/health/live` (und `/health/ready` deferred) — Compose-Healthcheck nutzt `/health/live`; `/health/ready` mit DB+RabbitMQ-Probes ist bewusst zurückgestellt (Liveness deckt die Block-5-Smoke-Anforderung; Readiness-Probes brauchen Kubernetes-Kontext, der per PE-4 für später).
- ✓ ~~(Optional) Helm-Chart oder Customize-Manifests für Kubernetes~~ — **bewusst out of scope** (Vault-Wortlaut: "falls Zeitbudget reicht; sonst bewusst out of scope dokumentieren"). Docker Compose erfüllt den Block-5-Deployment-Anspruch; K8s-Migration ist explizit nach-CAS (siehe PE-4).

CI/CD (GitHub Actions)

- ✓ Workflow `ci.yml` : restore → build → test → coverage upload (Hauptbranch + PRs)
- ✓ Workflow `release.yml` : bei Tag `v*` → Docker-Images bauen + zu GHCR pushen + Release Notes via `git-cliff` generieren
- ✓ Workflow `migrations.yml` : separater Job, der Migrations-Bundle generiert
- ✓ Branch-Protection auf `main` (PRs grün)
- ✓ Doku in `docs/ci-cd.md`

KI-Suche & Workflows (Auftrag-Erweiterung)

- ✓ ADR 0006 — KI-Suche (Embeddings-Provider, pgvector vs. eigener Vector-Store, Index-Strategie)
- ✓ Embedding-Pipeline: Capture (Title + Body) → Embedding → Persistenz
- ✓ Such-Endpoint: `GET /api/v1/captures/search?q=...` (vector-only) — Hybrid-Match (full-text + PostgreSQL `tsvector` + Vector) deferred; aktueller Endpoint nutzt pgvector-Cosine, FluentAssertions-Integration-Tests + `tests/FlowHub.Api.IntegrationTests/SearchEndpointTests.cs`.
- ✓ KI-Workflow-Beispiel: **automatisches Skill-Routing via Klassifikation** (`AiClassifier` → `MatchedSkill` → `SkillRoutingConsumer` → `ISkillIntegration`) ist der eingebaute KI-Workflow. Embedding-Cluster-Routing und LLM-Tag-

Suggestions sind in [ROADMAP.md](#) ("Capture Enrichment") als post-CAS-Erweiterung skizziert.

Monitoring / Observability

- ✓ OpenTelemetry: Traces + Metrics (Logs via Serilog → stdout, 12-Factor XI) — `Microsoft.Extensions.AI.UseOpenTelemetry()` instrumentiert MEAI; [ASP.NET Core](#) + Runtime-Instrumentation in `Program.cs`; Logs-Pipeline durch Serilog statt OTel-Log-Records (gleicher Effekt im Container-Kontext: alles auf stdout, Collector sammelt strukturiert).
- ✓ ~~OTLP-Exporter konfigurierbar via~~ `OTEL_EXPORTER_OTLP_ENDPOINT` — **deferred:** aktueller Export-Pfad ist Prometheus-Scrape (`/metrics`) + Grafana-Dashboard; OTLP-Push wäre Doppel-Pfad. Add-on ist trivial (`AddOtlpExporter()` + Env-Var) sobald ein OTel-Collector gebraucht wird; aktuelle Topology hat keinen.
- ✓ Prometheus-Endpoint `/metrics` — `MapPrometheusScrapingEndpoint().AllowAnonymous()` in `Program.cs`; durch `make smoke-prod` Schritt [4/6] verifiziert (`dotnet_*` + `http_*` Series).
- ✓ Grafana-Dashboard JSON eingecheckt (`docs/monitoring/grafana/flowhub-dashboard.json`)
- ✓ Strukturiertes Logging mit Serilog → stdout (12-Factor XI); Korrelations-IDs via OpenTelemetry W3C Trace-Context (Activity.Current trägt Traceld/Spanld quer durch alle MEAI-/EF-/HTTP-Spans). Pro-Request-Serilog-Enrichment ist deferred — der OTel-Traceld reicht für Cross-Span-Korrelation im Grafana-Dashboard.
- ✓ Healthchecks `/health/live` (DB- und MQ-Probes deferred) — Live-Check liefert nur Process-Liveness, kein Readiness; per Block-5-Smoke deckt das den Deployment-Anspruch ("Container started + Prozess lebt"). DB-/MQ-Probes wären Block-6-Stretch im Kubernetes-Kontext.

Authentifizierung (Block-5-Stretch falls Zeit)

- ✓ OIDC gegen Authentik (Homelab SSO) statt Dev-Auth-Handler
- ✓ Client-Registration in Authentik dokumentieren
- ✓ Tests für Auth-Flow — `DemoAuthHandler` ist via `Web.ComponentTests` abgedeckt (Auth-bypass auto-signs für Tests); voller OIDC-End-to-End-Test gegen Authentik bewusst deferred (Authentik läuft homelab-only, kein public-CI-zugänglicher IdP).

Submission

- ✓ **Projektarbeit-PDF schreiben** — Inhalte: Vision, Use Cases, NfAs, Architektur (alle ADRs), DB-Modell, Programmierung-Highlights, Test-Strategie + Resultate, KI-Nutzung + Reflexion, Repo-URL
- ✓ Repo-URL prominent ins PDF: `github.com/freaxnx01/FlowHub-CAS-AISE`
- ✓ Tag `v0.1.0` setzen + zu Release pushen, CHANGELOG-Eintrag final — Tag `v0.1.0` (matches `Directory.Build.props` `<Version>0.1.0</Version>`) gepusht, `release.yml` grün, GitHub-Release veröffentlicht:
<https://github.com/freaxnx01/FlowHub-CAS-AISE/releases/tag/v0.1.0>
- ☐ PDF auf Moodle hochladen vor **2026-07-04 24:00**

Spezifikation & Doku konsolidieren

- ✓ Use-Case-Liste final
- ✓ NfA-Liste final (SMART)
- ✓ ADR-Index in `docs/adr/README.md`
- ✓ `docs/ai-usage.md` final (höchstgewichtetes Kriterium!)
- ✓ `docs/insights/` für alle 5 Blöcke vorhanden
- ✓ `docs/spec/testing-strategy.md` final (Pfad unter `docs/spec/` , nicht direkt `docs/`)
- ✓ CHANGELOG `[0.1.0]` Section (entspricht `<Version>0.1.0</Version>` aus `Directory.Build.props` ; nicht `v1.0.0` wie ursprünglich geplant)

□ Out of Scope (auch in Block 5)

- Multi-Tenancy / RBAC über Single-User hinaus
- Production-Backup-Automation jenseits von Dokumentation
- Mobile App / Native Clients
- Skill-Marketplace / Plugin-Loader

Verweise

- Repo: `[[Repository]]` — `github.com/freaxnx01/FlowHub-CAS-AISE`
- Block 5 Vorbereitung: `[[05 Deployment - a) Vorbereitung]]`
- Block 4 Nachbereitung: `[[04 Persistence - c) Nachbereitung]]`
- Bewertungskriterien: `[[Bewertungskriterien]]`

- ADR 0001: docs/adr/0001-frontend-render-mode-and-architecture.md
- ADR 0002: docs/adr/0002-service-architecture-and-async-communication.md
- 12-Factor: siehe CLAUDE.md § "12-Factor Compliance"

Per-block insights

Source: `docs/insights/block-1.md`

Block 1 — Insights

Block: 1 — Vorbereitung & Konzept **Date range:** 2026-02-01 – 2026-03-08

What We Built

Project concept and initial architecture. Chose .NET 10 / Blazor Web App over Quarkus / Jakarta EE (justified in the Projektbeschreibung — stack-neutral rubric). Defined FlowHub as a modular-monolith PKM automation hub. Created `FlowHub.Core` domain model: `Capture`, `LifecycleStage`, `ChannelKind`, `ICaptureService`, `ICaptureRepository`.

Set up solution structure: `FlowHub.slnx`, `Directory.Build.props` (warnings-as-errors, embedded PDB), `Directory.Packages.props` (central package management), `global.json` (SDK pin).

Key Decisions

- **Stack:** .NET 10 Blazor Server-interactive + MudBlazor (ADR 0001).
- **Architecture:** Modular Monolith with async pipeline boundary instead of physical Microservices split (ADR 0002).
- **Persistence placeholder:** `InMemoryCaptureRepository` stub — EF Core PostgreSQL wires in Block 4.
- **Test framework:** xUnit + FluentAssertions + NSubstitute + bUnit.

Process

1. **Architecture brainstorming** — Claude Code brainstorming session weighing Monolith / Modular Monolith / Microservices against the FlowHub scenario (single operator, homelab deployment).
2. **ADR drafting** — Two ADRs drafted in parallel (0001 frontend, 0002 service architecture), reviewed for "Alternatives considered" completeness, then

accepted.

3. **Scaffolding** — Solution + projects + central package management generated via AI; verified file-by-file.

Lessons Learned

- `global.json` SDK pin prevents silent .NET version drift across machines (caught an issue we'd otherwise have hit when moving between the WSL2 dev box and the homelab Linux runner).
- `Directory.Packages.props` pays off immediately — one version bump in one file, instead of N project files.
- MudBlazor's `IsDarkMode` must be wired via `MudThemeProvider.MudThemeChanged` event, not static config (cost ~30 min of debugging in Block 2 because of an assumption set here).
- Starting with the domain model (`Capture` as record, `LifecycleStage` enum) before any UI kept the architecture clean across all five blocks.

AI Usage by Activity

Solution scaffolding

AI generated the project layout (`FlowHub.slnx` , four projects, central package management, common props). Reviewed every generated file; no functional corrections. `Directory.Build.props` `TreatWarningsAsErrors=true` was a human addition after the AI defaulted to `Nullable=enable` only.

Domain modeling

AI proposed `Capture` as a class; human switched to a `record` to lean on value equality for test assertions. AI's initial `ICaptureRepository.GetAllAsync()` returning `IAsyncEnumerable<Capture>` was rejected — overkill for current scale; substituted `Task<IReadOnlyList<Capture>>`. AI's `CancellationToken` propagation in `ICaptureService` was accepted as-is.

ADR drafting

AI drafted Context / Decision / Consequences sections; "Alternatives considered" sections were extended by the human after pushback on "why not

Microservices?" The "Stack mismatch — N/A justification" pattern was introduced in ADR 0002 and reused in ADRs 0003 / 0004 / 0005.

AI Usage Metrics (Block 1, estimate)

Block 1 had relatively little production code — most output was configuration, ADR prose, and the Projektbeschreibung. Estimates are conservative.

Artifact	Approx. lines	AI-generated	Human lines	AI % (est.)
Solution files (<code>.slnx</code> , <code>.props</code> , <code>global.json</code>)	~80	~70	~10	88%
Domain types (<code>Capture</code> , <code>LifecycleStage</code> , <code>ChannelKind</code> , ports)	~120	~100	~20	83%
<code>InMemoryCaptureRepository</code> stub	~40	~36	~4	90%
ADRs 0001 + 0002 prose	~400	~280	~120	70%
Projektbeschreibung (initial draft)	~600	~360	~240	60%

Human share concentrates on the architectural decisions themselves (Modular Monolith over Microservices, .NET over Quarkus), the "why not" sections of ADRs, and the Projektbeschreibung narrative.

KI-Reflexion / Fazit

Stärken

- **Boilerplate.** Solution-Setup, zentrale Paketverwaltung, Domain-Typ-Skelette: alles generiert, alles auf Anhieb compilierbar.
- **ADR-Struktur.** Das Nygard-Format (Context / Decision / Consequences) liefert die KI zuverlässig — das spart die immer wiederkehrende Strukturierungsarbeit.
- **Alternativen-Exploration.** Im Brainstorming-Skill hat die KI Microservices-, Monolith- und Modular-Monolith-Optionen mit jeweils 2-3 konkreten Trade-offs

gegenübergestellt; das hat die ADR-0002-Argumentation deutlich beschleunigt.

Grenzen

- **Stack-Entscheid.** Die finale Entscheidung Quarkus vs. .NET ist eine Domain- und Kontext-Entscheidung (Vorwissen, Werkzeug-Reife, Deployment-Pfad), die die KI nicht treffen kann — sie kann nur Argumente sammeln und gegenüberstellen.
- **Über-Engineering-Tendenz.** Erstvorschläge waren regelmäßig "vollständiger" als nötig (z. B. `IAsyncEnumerable` für ein Repository ohne Streaming-Anforderung). Disziplinierte Ablehnung von Komfort-Features hat das Layer auf das Notwendige beschränkt.
- **Begründungen müssen menschlich werden.** KI-Erstdrafts der ADR-Begründungen wirkten oft generisch ("for scalability") — die Spezifika (homelab, single operator, Block-5-Deployment-Constraint) wurden manuell ergänzt.

Fazit

Block 1 hat die Grundlage für KI-assistierte Entwicklung über die gesamte Projektlaufzeit gelegt: Solution-Layout, ADR-Format, Domain-Vokabular. Die KI war hier ein effizienter Scaffolder; die Architektur-Entscheidungen selbst blieben menschlich und bewusst dokumentiert.

Block 2 — insights

Source: `docs/insights/block-2.md`

Block 2 — Insights

Block: 2 — UI & Prototyp **Date range:** 2026-03-09 – 2026-04-11

What We Built

Blazor UI prototype with MudBlazor. Six routable pages designed and implemented following the mandatory four-phase workflow (brainstorm → flow → build → review): Dashboard, New Capture, Captures List, Capture Detail, Skills, Integrations. Each page reached production quality against Bogus stub services.

`source/FlowHub.Web/Stubs/` folder: Bogus-backed `StubCaptureService`, `SkillRegistryStub`, `IntegrationHealthServiceStub` for realistic fake data without a real database. All UI components unit-tested with `bUnit`; a full smoke-test walkthrough lands in `tests/FlowHub.Web.SmokeTests/` and replaces the manual click-through that the *Nachbereitung* TODO originally called for.

`justfile` added with `run / watch / build / test / format` targets — the entry point used by every subsequent block. CHANGELOG bootstrapped. Five FlowHub Claude Code skills authored (`flowhub-capture`, `flowhub-triage`, `flowhub-issue`, `flowhub`, `flowhub-dispatcher`) to support the workflow that drives the captures themselves.

Key Decisions

- **Render mode:** Interactive Server (ADR 0001) — simplest path to real-time updates; no WASM bundle overhead.
- **UI workflow:** mandatory 4-phase gate (ASCII wireframe → Mermaid flow → build → review) prevents premature coding.
- **Stubs implement real interfaces:** swapping to real services in Block 4 reduces to a DI registration change.
- **Test tooling:** `bUnit` + `NSubstitute` covers component rendering and event callbacks without spinning up a full app; smoke tests cover the cross-page

walkthrough.

Process

The 4-phase UI workflow ran on every page:

1. **Wireframe (Phase 1)** — ASCII wireframe committed to `docs/design/`. AI-assisted brainstorming via the `ui-brainstorm` skill.
2. **Flow (Phase 2)** — Mermaid flow + sequence diagrams (`docs/design/sequences/`, `docs/design/journeys.md`). Rendered to PNG via the same toolchain that produces the submission PDF.
3. **Build (Phase 3)** — Page + components + DI wiring + bUnit tests in the same commit family.
4. **Review (Phase 4)** — Walkthrough check, then `ui-review` skill for cross-cutting checks.

This was enforced by skill design — the AI itself refuses to skip phases.

Lessons Learned

- `MudDataGrid<T>` pagination requires the explicit type parameter (`MudDataGrid<Capture>`) — implicit inference fails at runtime, not at compile time.
- `@inject` services must be registered before `WebApplicationFactory<Program>` boots — missing DI registration surfaces as a confusing 500, not a startup exception. Fixed by moving the stub registrations to a `Program` extension method used by both `Web` and `WebApplicationFactory`.
- Stubs implementing the real driving-port interfaces meant the Block 4 transition was nearly mechanical — a 2-line DI change, not a rewrite.
- Mermaid → PNG rendering at design time (rather than relying on runtime browser support) keeps the submission PDF self-contained and ensures the diagrams in the PDF match the diagrams the bUnit tests reference.

AI Usage by Page

Dashboard

Dashboard layout and the "Needs Attention" widget were AI-generated from the wireframe. The first version mixed concerns (data fetch + render inline); human

refactor extracted `IDashboardData` and a parallel-load pattern (4 parallel `Task<T>` for the 4 cards).

New Capture

Form layout and validation wiring fully AI-generated. AI's initial form used inline `Action<MouseEventArgs>` handlers; human refactor to `EventCallback<...>` for proper Blazor lifecycle and bUnit testability. Skill dropdown disable-on-error logic was a human-added robustness improvement.

Captures List

`MudDataGrid<Capture>` with filter chips, search field, and pagination was AI-scaffolded. Filter composition (lifecycle AND channel AND search) was a human refactor — AI's first version treated the filters as OR. Deep-link query-param parsing (`?lc=0rphan`) added in a second pass.

Capture Detail

Detail page + stub action buttons fully AI-generated. The "Coming in Block 3" snackbar pattern was a human design decision to make Block-2 stub state explicit rather than silently no-op.

Skills + Integrations

Mostly mechanical given the patterns established by the earlier pages — AI generated both in one pass, human reviewed for consistency.

Smoke tests

`tests/FlowHub.Web.SmokeTests` replaced the manual walkthrough TODO. AI-generated using the patterns already established by the bUnit suite; human-added the cross-page assertions that the original TODO described.

AI Usage Metrics (Block 2, estimate)

Artifact	Approx. lines	AI-generated	Human lines	AI % (est.)
Razor pages (6)	~900	~720	~180	80%
Stub services (3) + Bogus fakes	~250	~225	~25	90%
bUnit tests (~30 tests across pages)	~700	~595	~105	85%
Smoke tests	~150	~120	~30	80%
Wireframes + Mermaid flows	~400	~280	~120	70%
justfile + Make targets doc	~80	~60	~20	75%
FlowHub CC-skills (5 skills)	~500	~350	~150	70%

Human contributions concentrated in: filter-composition logic (AND vs OR), parallel data-fetch refactor, deep-link param parsing, Block-3-stub UX policy, and skill design (the FlowHub triage logic is human-authored after AI failed to converge on a workable classification policy).

KI-Reflexion / Fazit

Stärken

- **Razor / MudBlazor.** Komponentengeneration aus Wireframes ist die mit Abstand effizienteste Form der KI-Unterstützung in diesem Stack — Wireframe rein, Razor + bUnit-Skelett raus, in einem Schritt.
- **Test-Skelette.** bUnit-Test-Klassen folgen einem stark wiederholten Muster (`RenderComponent<T>`, `Find(...)`, `Markup.Should().Contain(...)`); KI generiert sie zuverlässig.
- **Mermaid-Diagramme.** Flow- und Sequence-Diagramme sind ein Format, in dem die KI besonders stark ist; das Phase-2-Artefakt war faktisch immer ein One-Shot.

Grenzen

- **Filter-Komposition.** Die "AND vs OR"-Frage in der Captures-List ist eine Spec-Lese-Frage — die KI hat den Wireframe-Hinweis (mehrere Chips = AND) übersehen und OR generiert. Ohne Acceptance-Kriterium kein verlässliches Verhalten.
- **Stub-Disziplin.** KI schlug mehrfach vor, in Stubs `throw new NotImplementedException()` für noch nicht implementierte Operationen zu setzen; Policy-Entscheidung war stattdessen "leere, aber valide Antwort + Snackbar-Hinweis", damit die UI im Demo-Modus durchläuft.
- **Component-Lifecycle.** Erstvorschläge nutzten `Action<...>` statt `EventCallback<...>` — funktional fast identisch, aber bUnit-Tests scheitern an fehlender Re-Render-Signal.

Fazit

Block 2 hat den UI-Workflow etabliert, der den Rest des Projekts trägt: Wireframe → Flow → Build → Review, mit der KI als Beschleuniger in Phase 1 und 3. Der menschliche Beitrag konzentrierte sich auf Policy-Entscheidungen (Stub-UX, Filter-Logik, Phase-Disziplin) — der reine Tipparbeitsanteil ist deutlich gesunken, der Anteil an Entscheidungs- und Review-Arbeit unverändert hoch.

Block 3 — insights

Source: `docs/insights/block-3.md`

Block 3 — Insights

Block: 3 — KI-Integration & Pipeline **Date range:** 2026-04-12 – 2026-05-03

What We Built

Three slices delivered in parallel, each backed by its own brainstorm → spec → plan → implementation cycle:

Slice A — REST API. Minimal API endpoints in `FlowHub.Api` (class library registered into `FlowHub.Web` via `AddFlowHubApi` / `MapFlowHubApi`). Endpoints: `POST /api/v1/captures` (FluentValidation), `GET /api/v1/captures` (cursor pagination + Stage / Source filters), `GET /api/v1/captures/{id}`, `POST /api/v1/captures/{id}/retry` (stage validation + republish). ProblemDetails (RFC 9457) with three typed problems (`validation.md`, `capture-not-found.md`, `capture-not-retryable.md`). OpenAPI via `Microsoft.AspNetCore.OpenApi`, Scalar UI at `/scalar`. `tests/FlowHub.Api.IntegrationTests/` with `WebApplicationFactory<Program>`.

Slice B — MassTransit pipeline. In-memory (dev) / RabbitMQ (prod) event bus. Two consumers wired in this block — `CaptureEnrichmentConsumer` (classification → `CaptureClassified`) and `SkillRoutingConsumer` (skill dispatch → `Routed` / `Orphan`). `LifecycleFaultObserver` maps `Fault<T>` events to `Unhandled` with `FailureReason`. Bus transport switches via `Bus:Transport` env var. ADR 0003 accepted with explicit EventId range allocation (`3000-3099` enrichment, `3100-3199` routing, `3200-3299` faults — aligned with the code in the same commit).

Slice C — AI classifier. `IClassifier` port in `FlowHub.Core`, `AiClassifier` adapter in `FlowHub.AI`. Supports Anthropic Claude Haiku and OpenRouter providers via the [Microsoft.Extensions.AI](#) (MEAI) `IChatClient` abstraction. Graceful fallback to `KeywordClassifier` when no API key is configured **or** when any AI call fails (ADR 0004 §D5 — `AiClassifier` never throws). EventId `3010` `AiClassifierFellBackToKeyword` logs the fallback path.

`docs/ai-usage.md` was bootstrapped this block as the living KI-Werkzeug-Nutzung document for the rubric.

Key Decisions

- **ADR 0003** — MassTransit pipeline topology, retry policy, EventId allocation.
- **ADR 0004** — AI provider abstraction via MEAI `IChatClient`; Anthropic Claude Haiku as default; graceful keyword fallback policy (§D5).
- **API-as-class-library** — `FlowHub.Api` registers into `FlowHub.Web`'s route builder. One process, two surfaces (UI + REST).
- **Env-var provider switch** — `Ai:Provider = Anthropic | OpenRouter`. No code change to switch.
- **URL versioning from day one** — `/api/v1/...` (NF-13). Breaking changes land only in a new major version.

Process

The "Slice A / Slice B / Slice C" labeling itself comes from the workflow: each slice ran through `superpowers:brainstorming` → spec doc → plan doc → subagent-driven implementation, with the slices able to land in any order (commits show them interleaved). The brainstorming / spec / plan artifacts live under `docs/superpowers/specs/` and `docs/superpowers/plans/`.

This block was the first sustained use of `superpowers:subagent-driven-development` — one agent per slice, parent context only seeing the spec / plan / commit summary. Worked well for Slice A and Slice B; Slice C required a second pass after the subagent under-specified the fallback contract (fixed by adding ADR 0004 §D5 before the second implementation pass).

Lessons Learned

- MassTransit's in-memory transport processes messages synchronously during `await bus.Publish()` — useful for integration tests (no harness needed for happy-path) but requires care not to block the request thread in production code paths.
- `IChatClient` from MEAI abstracts provider differences cleanly; the `ConfigureOptions(o => o.ModelId = model)` pattern avoids repeating the model name on every call.

- `bUnit's IRenderedComponent<T>.WaitForState()` is essential for testing components that await async service calls — without it, assertions race the renderer.
- Centralizing `ProblemDetails` type URIs (`FlowHubProblemTypes` constants + three Markdown files under `docs/api/`) keeps the wire contract and the human-readable explanation in sync — easy to drift otherwise.
- Narrowing the `LifecycleFaultObserver` catch to exclude `OperationCanceledException` (commit `faee77d`) prevented graceful shutdown from being logged as a fault.

AI Usage by Slice

Slice A — REST API

AI scaffolded the project (`FlowHub.Api` + `AddFlowHubApi` / `MapFlowHubApi`), the four endpoints, `FluentValidation` rules, `ProblemDetails` type registration, and the integration test class with `WebApplicationFactory<Program>` . Cursor pagination — `CaptureCursor` with Base64Url JSON encode / decode — was AI-generated but the human caught an off-by-one when the cursor was at exactly the page boundary; fixed before merge. The three `ProblemDetails` Markdown files (`validation.md` , `capture-not-found.md` , `capture-not-retryable.md`) were human-authored to ensure the wire `type` URIs resolve to actual documentation.

Slice B — MassTransit pipeline

AI scaffolded `CaptureEnrichmentConsumer` , `SkillRoutingConsumer` , the DI registration, and the bus transport switch. Two corrections from review: retry policy changed from fixed to exponential intervals (matches NF-10), and the `EventId` range allocation in code did not match the original ADR 0003 draft — code was authoritative, ADR was updated (commit `d2eaaba`). `LifecycleFaultObserver` was AI-generated; the `OperationCanceledException` exclusion (commit `faee77d`) was a human follow-up after seeing graceful shutdowns logged as faults.

Slice C — AI classifier

`AiClassifier` and `KeywordClassifier` were AI-generated. The fallback contract (`ClassifyAsync` must never throw) was strengthened after the first subagent pass — ADR 0004 §D5 was added explicitly, then the implementation was re-generated to match. The `AiClassifierTests` class covers the three fault modes

(network exception, JSON parse error, schema violation) — all AI-generated from a one-paragraph human spec.

AI Usage Metrics (Block 3, estimate)

Artifact	Approx. lines	AI-generated	Human lines	AI % (est.)
<code>FlowHub.Api</code> endpoints + DI	~350	~315	~35	90%
FluentValidation validators (4)	~80	~72	~8	90%
ProblemDetails type registration + types (3)	~120	~80	~40	67%
<code>CaptureCursor</code> (Base64Url JSON)	~60	~45	~15	75%
MassTransit consumers (2) + observer	~200	~170	~30	85%
<code>IClassifier</code> / <code>AiClassifier</code> / <code>KeywordClassifier</code>	~180	~160	~20	89%
<code>AiClassifierTests</code> (3 fault modes)	~150	~135	~15	90%
API integration tests	~280	~245	~35	88%
ADR 0003 + 0004 prose	~600	~390	~210	65%
ProblemDetails Markdown files (3)	~90	~30	~60	33%

Human contributions concentrated in: pagination edge case, EventId range reconciliation, fallback contract strengthening (ADR 0004 §D5), ProblemDetails human-facing prose, and the integration patterns (API-as-class-library, env-var provider switch).

KI-Reflexion / Fazit

Stärken

- **Subagent-driven development.** Pro Slice ein Subagent mit Spec + Plan, der Parent-Kontext sah nur Commit-Summaries. Hat den Hauptkontext schlank gehalten und parallele Slices ermöglicht.
- **MEAI-Provider-Abstraktion.** Die KI hat `IChatClient` korrekt verwendet und Provider-spezifische Details (Anthropic vs. OpenRouter Optionen) sauber im DI-Setup gekapselt — keine `if (provider == "Anthropic")`-Verzweigungen im Code.
- **Test-Generation.** FluentValidation-Tests, ProblemDetails-Wire-Format-Tests, MassTransit-Harness-Tests: alles in ein-zwei Sätzen spezifiziert, KI liefert lauffähigen Code.

Grenzen

- **Vertragsstärke unterspezifiziert.** Der erste `AiClassifier`-Wurf hat im Fallback geschwiegen, statt zu loggen — das Acceptance-Kriterium "EventId 3010 mit Exception-Typ" war nicht im Plan. Fix: ADR 0004 §D5 explizit, dann zweiter Implementierungspass.
- **EventId-Drift.** Code und ADR 0003 wurden parallel generiert und sind in den allokierten Bereichen auseinandergedriftet (Code-Bereiche waren feiner). Lesson: Code-First für Konstanten, ADR dokumentiert post-hoc.
- **Cursor-Pagination-Edge-Cases.** Cursor genau auf Page-Boundary hat zu Duplikat-Eintrag in der Folgeseite geführt. Wäre in einem manuellen Code Review aufgefallen; in KI-Generation muss man die Edge-Cases als Tests explizit verlangen.

Fazit

Block 3 war der größte Sprung in der KI-Produktivität — drei Slices in einem Block sind ohne den Subagent-Driven-Workflow nicht machbar gewesen. Gleichzeitig hat der Block gezeigt, dass KI-Generierung an Vertragsstärke und Edge-Cases scheitert, wenn diese nicht im Spec / Plan dokumentiert sind. Die Brainstorm → Spec → Plan → Implementation-Pipeline ist deshalb non-negotiable geworden für alle Folgeblöcke.

Block 4 — insights

Source: `docs/insights/block-4.md`

Block 4 Erkenntnisse — Persistence Layer

What Was Built

Block 4 completed the persistence layer for FlowHub:

- **PostgreSQL switch:** Replaced SQLite with PostgreSQL 17 + Npgsql EF Core provider. No data to migrate; existing SQLite migration dropped and regenerated.
- **Repository pattern:** `ICaptureRepository`, `IChannelRepository`, `ISkillRepository`, `IIntegrationRepository`, `ITagRepository`, `ISkillRunRepository` — all in `FlowHub.Core` returning domain types. EF Core implementations in `FlowHub.Persistence`.
- **Full domain model:** 6 new entities (Channel, Skill, SkillRun, Integration, IntegrationHealthSample, Tag) with real DB FKs where appropriate (hard FKs on audit trail entities, soft FKs for deregisterable references).
- **Stub retirement:** `EfSkillRegistry` replaced `SkillRegistryStub`; `EfIntegrationHealthService` replaced `IntegrationHealthServiceStub` in DI.
- **Dynamic filter:** `CaptureQueryBuilder` combines Stage, Source, Tag, and SearchTerm (ILike) predicates via expression composition.
- **Testcontainers:** 16 integration tests against real PostgreSQL (provider-parity with production).
- **Docker Compose:** Full stack with postgres service + migrations init container (12-Factor XII).

AI Usage by Slice

Slice 1: PostgreSQL + Repository Foundation

AI generated the full `EfCaptureRepository` from the interface contract. The initial output used `context.Captures.FindAsync(id)` (tracking) instead of `.AsNoTracking().FirstOrDefaultAsync()` — human corrected for read operations. AI drafted `CaptureEntityTypeConfiguration`; human added `IX_Captures_MatchedSkill` index (missed from initial AI output).

The repository interface design (`ICaptureRepository` returning domain types, not entities) was a human architectural decision. AI implemented the resulting design quickly once the contract was defined.

Estimated AI share: ~90% implementation code, ~10% interface and architectural decisions.

Slice 2: Channel + Skill Entities

AI produced `ChannelEntity`, `SkillEntity`, their type configurations, and the repositories in one pass. Field lengths were wrong on first attempt (AI used `varchar(128)` for Name; spec says `varchar(64)`) — fixed in review. `EfSkillRegistry` wrapping `ISkillRepository` (two-line class) was a human-driven design decision to keep the driving port decoupled from the driven port.

Estimated AI share: ~88% code, ~12% review corrections.

Slice 3: Full Domain Model + Dynamic Filter

AI generated 4 entity classes + configurations + repositories. Main corrections:

- FK cascade rules: AI defaulted all FKs to CASCADE DELETE. Human changed `SkillRun→Skill` to `Restrict` (preserve audit trail if skill name is updated/removed); `Capture→Channel` and `Capture→Skill` are soft FKs (no DB FK at all).
- `CaptureQueryBuilder` was AI-generated. Human noticed the N+1 on Tags (EF does not auto-load nav properties) and added `.Include(c => c.Tags)`.
- AI placed the `.Include()` call after `.AsNoTracking()` — correct order in EF Core 10; the cursor ordering precondition comment was a human addition after code review.

Estimated AI share: ~85% code, ~15% architectural decisions and corrections.

Slice 4: Tests + Docker

`PostgresFixture` with per-test isolated databases was AI-scaffolded; human reviewed `NpgsqlConnectionStringBuilder` usage and the `CREATE DATABASE` admin connection pattern. All 16 test methods were AI-generated from the method names and expected behaviors.

Docker Compose `flowhub.migrations` service (12-Factor XII) was a human design decision; AI filled in the YAML after the pattern was described.

Estimated AI share: ~92% test code, ~85% Docker YAML.

AI Usage Metrics (Block 4)

Artifact	Total Lines	AI-Generated	Human Lines	AI %
Entity classes (7)	~120	~108	~12	90%
EntityTypeConfiguration (7)	~130	~117	~13	90%
Repository impls (6)	~350	~315	~35	90%
Service impls (EfSkillRegistry, EfIntegrationHealthService)	~30	~27	~3	90%
CaptureQueryBuilder	~30	~25	~5	83%
EfCaptureService refactor	~70	~60	~10	86%
Test files (4 classes, 16 tests)	~230	~210	~20	91%
Docker Compose	~50	~40	~10	80%
Total	~1010	~902	~108	~89%

Human contributions were concentrated in: FK strategy decisions, N+1 detection, field length corrections, index additions, and Docker Compose service dependency pattern.

KI-Reflexion / Fazit

Stärken der KI-Unterstützung (Strengths)

Boilerplate-Generierung: Die sieben `IEntityTypeConfiguration<T>`-Klassen (insgesamt ~130 Zeilen) wurden vollständig von der KI generiert. Ohne KI wäre dieser Schritt zeitintensiv und fehleranfällig gewesen — die Klassen sind strukturell identisch, unterscheiden sich nur in Tabellennamen und Feldlängen.

Migrationsgenerierung: Das Scaffolding der drei EF Core Migrations (0001–0003) war vollständig automatisiert. Die KI hat den Design-Time-Factory-Ansatz korrekt angewendet.

Expression-Tree-Filter: `CaptureQueryBuilder` mit kombinierter Prädikatenkomposition (Stage, Source, Tag, ILike) war ohne KI-Unterstützung deutlich aufwendiger. Die KI hat das korrekte Muster für EF Core LINQ-Komposition auf Anhieb angewendet.

Test-Scaffolding: 16 Integrationstests wurden von der KI generiert und bestehen alle beim ersten Durchlauf. Die Teststruktur (Arrange/Act/Assert, FluentAssertions-Syntax) ist konsistent und entspricht den Projektvorgaben.

Schwächen und Grenzen (Weaknesses)

N+1-Blindheit: Die KI hat `EfCaptureRepository.ListAsync` initial ohne `.Include(c => c.Tags)` generiert. Das N+1-Problem wäre in einer Codeüberprüfung aufgefallen — KI hat keine implizite Performance-Awareness für Navigation Properties.

FK-Strategie: Die KI hat durchgängig Hard-FKs mit CASCADE DELETE vorgeschlagen, ohne die Unterscheidung zwischen "owned" Entitäten (Tags, IntegrationHealthSamples) und "referenced" Entitäten (Channel, Skill) zu erkennen. Die Soft-FK-Entscheidung für Capture→Channel und Capture→Skill war eine Domain-Entscheidung, die domänenspezifisches Verständnis erforderte.

Feldlängen: Erste Generierung verwendete `varchar(128)` für Name-Felder. Die spec schreibt `varchar(64)` vor. KI liest Spezifikationen korrekt, wenn sie explizit zitiert werden — aber ohne direkten Verweis wird auf "sichere" Standardwerte ausgewichen.

Paket-Versionen: Der Plan hatte `Npgsql 9.0.4` spezifiziert (EF Core 9 Ära). Der implementierende Subagent musste die Version auf `10.0.1` anpassen. Plan-

Templates mit versionsgepinnten Paketen degradieren mit der Zeit.

Fazit

KI-Unterstützung hat in Block 4 die Implementierungszeit für Persistence-Infrastruktur auf ~89% des Codes reduziert. Der menschliche Beitrag war konzentriert auf: Architekturentscheidungen (FK-Strategie, Repository-Interface-Design), Performance-Korrekturen (N+1), Spezifikationsabgleich (Feldlängen, Indexe) und Code Review.

Die Kombination aus Brainstorming-Skill → Spec-Dokument → Plan-Dokument → Subagent-getriebene Implementierung hat sich bewährt: Jede Phase hat die nächste informiert, ohne dass die KI unkontrolliert in die falsche Richtung implementiert hat.

Einschätzung: KI ist ein starker Accelerator für Infrastruktur-Code (Boilerplate, Migrations, Tests), erfordert aber menschliche Führung bei Architektur- und Domain-Entscheidungen.

Vom CAS in die Projektarbeit

Was wurde aus dem CAS in Block 4 konkret in die Projektarbeit übernommen — und wie:

CAS-Input (Vorbereitung / PVA #4 / Skript)	Übertragung in FlowHub (.NET-Stack)	Wo
ORM- Abstraktion (Hibernate ORM, Jakarta Persistence) als Pflicht zwischen Domain und DB	EF Core 10 als ORM, <code>DbSet<T></code> + <code>IEntityTypeConfiguration<T></code> pro Entity statt Annotations	<code>source/FlowHub.Persisten</code> ADR 0005 §1 + §8
Panache / Repository- Pattern- Diskussion (Active-Record vs. Repository)	Bewusste Entscheidung <i>gegen</i> Repository-per-Entity zu Beta- Zeit (ADR 0005 §3), <i>für</i> Repository-Layer nach realer Schmerz-Erfahrung in Block 4 (<code>ICaptureRepository</code> etc. in <code>FlowHub.Core</code> , EF-Adapter in <code>FlowHub.Persistence</code>)	<code>source/FlowHub.Core/Repo</code> Sektion
Criteria API / typsichere dynamische Abfragen	LINQ + Expression Trees als .NET-Pendant — <code>CaptureQueryBuilder</code> komponiert <code>Expression<Func<Capture,bool>></code> für Stage/Source/Tag/SearchTerm	<code>source/FlowHub.Persisten</code> <code>EfCaptureRepositoryTests</code>
Datenmodell- Antizipation ("Daten überleben jede Technologie") aus dem Moodle-Auftrag- Wording	Hard- vs. Soft-FK-Entscheidung pro Beziehung (Audit-Trail hart, deregistrierbare Referenzen soft); <code>LifecycleStage</code> als Soft- Delete-Ersatz statt <code>IsDeleted</code> - Flag	<code>docs/design/db/er.md</code> §FK §6
12-Factor- Prinzipien (Backing	PostgreSQL als ge-attachte Resource über <code>ConnectionStrings__Default</code> ;	<code>docker-compose.yml</code> (<code>flow</code> Factor, ADR 0005 §4

CAS-Input (Vorbereitung / PVA #4 / Skript)	Übertragung in FlowHub (.NET-Stack)	Wo
Services IV, Admin Processes XII) aus Block 1/Vorbereitung	Migrations als separater <code>flowhub.migrations -Init-</code> Container (nicht in <code>app.Run()</code>)	
Testcontainers / provider- parity (Diskussion in PVA #4 zu Test- Strategien gegen reale DB)	Testcontainers PostgreSQL 17 statt EF-Core <code>InMemory</code> für alle Repository- und Integration- Tests; Provider-Drift bewusst ausgeschlossen	<code>tests/FlowHub.Persistenc</code> <code>tests/FlowHub.Api.Integr</code> ADR 0005 §10
ADR-Workflow (Architecture Decision Records als CAS-vermittelte Disziplin)	ADR 0005 als Persistence- Master-Entscheid (355 Zeilen, 10 nummerierte Entscheidungen + Alternatives- Sektion), kreuzreferenziert aus Code, CHANGELOG und Block- Nachbereitung	<code>docs/adr/0005-persistenc</code>

Die Übersetzung *vom Java-/Jakarta-EE-Vokabular der CAS-Skripten in das .NET-Ökosystem* war in Block 4 ein durchgehender Aufwandsposten — aber sie hat den fachlichen Kern der Vorgaben (ORM-Abstraktion, dynamische Abfragen, Datenmodell-Antizipation, 12-Factor-Disziplin) nicht verändert, sondern nur die Werkzeuge.

Test Results (as of 2026-05-18)

`dotnet test FlowHub.slnx --filter "FullyQualifiedName!~E2ETests"` — **223 tests pass, 0 fail, 6 skipped** (live-service integration tests requiring external API keys):

Project	Passed	Skipped	Duration	
FlowHub.Persistence.Tests	29	0	~24 s	EF Core Postgre EfCapt EfCapt EfSkill EfInteg migrati
FlowHub.Web.ComponentTests	144	0	~11 s	bUnit – MassTra tests + classifi unit tes
FlowHub.Api.IntegrationTests	17	0	~11 s	WebApp Testcon pipeline Postgre
FlowHub.Skills.Tests	20	0	~1 s	Unit-lev Vikunja
FlowHub.Skills.ContractTests	13	0	~1 s	WireMo tests or socket
FlowHub.Skills.IntegrationTests	0	2	—	just to live Wa
FlowHub.AI.IntegrationTests	0	4	—	just to Ai__*_
Total	223	6		

E2E tests (`FlowHub.Web.E2ETests` , Playwright) are run on demand via `just test-e2e` and gated by the Web server + Chromium dependency. Latest green CI run: <https://github.com/freaxnx01/FlowHub-CAS-AISE/actions/workflows/ci.yml>.

Persistence-Layer Coverage

The 29 `FlowHub.Persistence.Tests` cover the full persistence surface introduced in this block:

Test class	What it verifies	
<code>EfCaptureRepositoryTests</code>	Create/Read/Update/Delete; cursor pagination ordering (<code>CreatedAt DESC</code> , <code>Id DESC</code>); filter composition (<code>Stage / Source / Tag / SearchTerm</code>); <code>.Include(Tags)</code> round-trip	Testcont 17
<code>EfCaptureServiceTests</code>	Stage transitions: <code>Raw → Classified → Routed → Completed</code> ; <code>MarkClassifiedAsync</code> now also persists <code>VikunjaProject</code> ; <code>MarkOrphan</code> , <code>MarkUnhandled</code> paths	Testcont <code>IPublis</code> substitu
<code>EfSkillRegistryTests</code>	Skill registry CRUD; deregistration; lookup-by-name	Testcont
<code>EfIntegrationHealthServiceTests</code>	Health-sample insertion; latest- per-integration query	Testcont
<code>MigrationSmokeTest</code>	All migrations from <code>0001_Initial</code> to <code>0008_AddVikunjaProjectToCapture</code> apply cleanly on an empty database	Testcont <code>db.Data</code>

The cursor-pagination edge cases (limit+1 probe, ordering precondition) and the FK-strategy choices documented in the AI-Usage section above are all covered by these 29 tests against real Postgres — no in-memory provider drift.

Block 5 — insights

Source: `docs/insights/block-5.md`

Block 5 — Insights

Block: 5 — Deployment & Abgabe

Date range: 2026-05-07 - 2026-07-06

What We Built

Containerisation: Multi-stage Docker images for `flowhub.web` (SDK → Alpine runtime, non-root user) and `flowhub.migrations` (efbundle binary). `.dockerignore` keeps images lean.

Docker Compose Stack: 6-service production topology — `flowhub.web`, `flowhub.migrations`, `postgres`, `rabbitmq`, `prometheus`, `grafana`. Migrations run as a `service_completed_successfully` dependency (12-Factor XII). Override file activates demo mode without code changes.

GitHub Actions: `ci.yml` (build+test on every push), `release.yml` (Docker image to GHCR on tags), `migrations.yml` (efbundle artifact on migration changes). Branch protection on `main` gates on CI.

OIDC Auth: `DemoAuthHandler` replaces `DevAuthHandler` — auth mode now driven by `Auth:OIDC:Authority` presence, not `IsDevelopment()`. Production uses Authentik OIDC.

Semantic Search: `pgvector` `vector(1024)` column on `Captures` with HNSW index. `AiEmbeddingService` wraps Mistral `mistral-embed` via MEAI. `IEmbeddingService` port in `FlowHub.Core`; graceful no-op when key absent. Search endpoint: `GET /api/v1/captures/search?q=...`. Admin rebuild: `POST /api/v1/admin/embeddings/rebuild`.

Key Decisions

- `pgvector` built into `Npgsql 10` via `Pgvector.EntityFrameworkCore` — no extra vector DB service

- Env-var auth switching (not environment name) — same binary works anywhere
- efbundle migrations binary — no SDK at runtime (12-Factor XII)
- DemoAuthHandler name clarifies intent (works in any environment, not just dev)

Lessons Learned

- `UseVector()` in Npgsql 10.0.1 requires `Pgvector.EntityFrameworkCore` — the built-in support requires a newer minor release than what was available; always verify package changelog before assuming a feature is in-the-box
- EF Core's `FromSqlRaw` with a float-array vector literal is injection-safe since all values are IEEE 754 floats
- Authentik's OIDC discovery endpoint makes `.AddOpenIdConnect()` configuration minimal — just `Authority`, `ClientId`, `ClientSecret`
- Building efbundle in CI and shipping it as an artifact decouples migration tooling from the runtime image
- 12-Factor XII (migrations as separate step) forces explicit migration management; `just migrate` fills the local dev gap

Test Results (as of 2026-05-12)

```
dotnet test FlowHub.slnx --filter
```

"Category!=AI&Category!=BetaSmoke&Category!=E2E" — **171 tests pass, 0 fail, 0 skip:**

Project	Passed	Duration	Coverage
FlowHub.Persistence.Tests	29	~34 s	EF Core + Testcontainers PostgreSQL — EfCaptureService , repositories, migrations
FlowHub.Web.ComponentTests	92	~11 s	bUnit — every Razor page + MassTransit harness pipeline tests
FlowHub.Skills.Tests	20	~1 s	Unit-level handler tests for Vikunja + Wallabag adapters
FlowHub.Skills.ContractTests	13	~1 s	New (Block 5): WireMock.Net wire-contract tests on a real loopback socket — path, bearer, JSON shape, 401/500 mapping for both skills
FlowHub.Api.IntegrationTests	17	~17 s	WebApplicationFactory + Testcontainers — full HTTP pipeline against real Postgres

Excluded by category in default just test (run on demand):

Trait	Project	When
Category=AI	FlowHub.AI.IntegrationTests	just test-ai — needs Ai__*__ApiKey
Category=BetaSmoke	FlowHub.Skills.IntegrationTests	just test-beta — needs live Wallabag + Vikunja
Category=E2E	FlowHub.Web.E2ETests	just test-e2e — needs running web server + Playwright Chromium

CI run reference: latest green `ci.yml` workflow at <https://github.com/freaxnx01/FlowHub-CAS-AISE/actions/workflows/ci.yml>.

Production-Stack Smoke (just smoke-prod , 2026-05-12)

End-to-end probe of the deployment claim from `vault/Blöcke/05 Deployment/05 Deployment - c) Nachbereitung.md` , wired as a make target on 2026-05-12 (commit `f0424ec`). Six steps, last run pass:

```
[1/6] docker compose up --build (detached, --wait until healthy) ✓
[2/6] verifying flowhub.migrations exited 0 ✓ (exit
[3/6] GET /health/live ✓ 200
[4/6] GET /metrics — expect dotnet_* and http_* series ✓ dotne
[5/6] POST /api/v1/captures (URL capture for embedding round-trip) ✓ id 73
[6/6] polling Captures.Embedding (up to 30s) ✓ popul
==> smoke OK — stack left running. Tear down with: just smoke-down
```

The smoke uses a `curlimages/curl:8.10.1` sidecar joined to the `flowhub.web` network namespace, so no host port needs to be published — production network topology is preserved.

Defects Found by the Smoke Run

The first attempts caught real, latent bugs that would have blocked the submission stack — exactly the kind of finding the smoke target is meant to surface:

1. **Dockerfile missed `.editorconfig`** (commit `4da7ea8`). Both `source/FlowHub.Web/Dockerfile` and `docker/migrations/Dockerfile` copied solution + props + sources but not `.editorconfig`, which silences `CA1707` / `CA1825` on EF Core migration files (auto-generated names starting with `_0001_`). `dotnet publish` inside the build container failed with `TreatWarningsAsErrors`. Fix: add `.editorconfig` to the `COPY` line.
2. **`docker-compose.yml` env-interpolation casing mismatch** (commit `8062bef`). Compose interpolated UPPERCASE `${EMBEDDINGS__APIKEY:-}` while `.env.example` defines mixed-case `Embeddings__ApiKey` (matching the .NET configuration key used everywhere else). Compose substituted an empty string regardless of `.env` contents; the embedding consumer silently no-op'd. Same trap on `Auth__OIDC__*`, plus `Ai__Provider` / `Ai__<Provider>__*` were missing from the compose env block entirely. Fix: switch all interpolations to mixed-case + add the missing `Ai__*` mappings.
3. **Empty-string model strings tripped OpenAI's `AssertNotNullOrEmpty`** (commit `867ada7`). Compose maps `Embeddings__Model:` `${Embeddings__Model:-}` which substitutes an empty string when unset. `configuration["Embeddings:Model"] ?? "mistral-embed"` never fired the fallback because the env value was `""` (not null). Same trap on `Ai__<Provider>__Model`. `flowhub.web` crashed with `System.ArgumentException: Value cannot be an empty string. (Parameter 'model')` at startup. Fix: switch to `is { Length: > 0 } m ? m : default` in both spots.
4. **Mistral rejects `dimensions` request field**. The default `.env.example` included `Embeddings__Dimensions=1024`; Mistral's `mistral-embed` returned 422 Unprocessable Entity. (`dimensions` is OpenAI- `text-embedding-3-*` only.) Fix: comment the default in `.env.example` with an explanatory note.
5. **Passbolt `passbolt: //` refs were shadowed by justfile-imported `.env`**. The justfile's `-include .env / export` pulled raw `passbolt: //<uuid>` strings into recipe env, overriding whatever a parent `passbolt exec` had already resolved. `just ai-classify` therefore *appeared* to succeed (output matched the keyword-classifier fallback exactly — `[task]`, `Wallabag`) but had never

actually reached the AI provider. Caught only by inspecting latencies. Fix: a `SECRET_EXEC` wrapper that re-sources `.env` *inside* the recipe shell and reroutes through `passbolt exec --` when present.

Each was committed independently with a one-line "caught by smoke" reference in the commit body; the smoke target paid for itself within a single afternoon.

AI Tooling

Claude Code was used throughout Block 5 for:

- Writing the implementation plan from the design spec
- Generating GitHub Actions workflow YAML (corrected permissions and artifact paths)
- Scaffolding the `AiEmbeddingService` (corrected `GenerateAsync` call signature for MEAI)
- Writing the `SearchEndpoints` and `AdminEndpoints` (endpoint registration pattern taken from existing `CaptureEndpoints`)

All generated code was reviewed, tested, and adjusted before commit. AI-suggested `Pgvector.EntityFrameworkCore` (a separate package) was the correct approach — Npgsql 10.0.1 does not include built-in vector support without it. The lesson: AI training data may not reflect very recent library releases; verify against installed package versions.

Operations — CI/CD

Source: `docs/ci-cd.md`

CI/CD

FlowHub uses three GitHub Actions workflows.

Workflows

`ci.yml` — Continuous Integration

Triggers: Every push to any branch; every PR targeting `main`.

Steps: restore → build (warnings-as-errors) → test (all projects, XPlat coverage) → upload artifacts.

Gate: PR merge to `main` is blocked if CI fails. Configure branch protection in GitHub: Settings → Branches → `main` → Require status checks → select `build-and-test`.

`release.yml` — Docker Image Release

Triggers: Tags matching `v*` (e.g. `v1.0.0`).

Steps:

1. Builds the `flowhub.web` Docker image from `source/FlowHub.Web/Dockerfile`.
2. Pushes to GHCR: `ghcr.io/freaxnx01/flowhub-web:vX.Y.Z` and `:latest`.
3. Generates release notes via `git-cliff` from Conventional Commits.
4. Creates a GitHub Release with the changelog section.

How to trigger a release:

```
git tag v1.0.0 -m "release: v1.0.0 – CAS AISE Projektarbeit Abgabe"
git push origin v1.0.0
```

migrations.yml — EF Core Migrations Bundle

Triggers: Push to `main` when files under `source/FlowHub.Persistence/Migrations/` change.

Steps: Builds a self-contained `efbundle` native binary (no .NET SDK needed at runtime). Uploads as a workflow artifact retained for 30 days.

How to apply in production: Download the artifact, copy to the target host, set `ConnectionStrings__Default` env var, run `./efbundle`.

Local Development

```
# Apply migrations to local PostgreSQL (no auto-migration at app startup)
just migrate

# Run app without migration step
just run
```

Runbook — Authentik OIDC setup

Source: `docs/runbooks/authentik-oidc-setup.md`

Authentik OIDC Setup

This runbook documents how to register FlowHub as an OAuth2/OIDC client in Authentik and configure the environment variables.

Authentik Configuration

1. Log into Authentik admin (`https://authentik.home.freaxnx01.ch/if/admin/`).
2. Navigate to **Applications → Providers → Create → OAuth2/OpenID Provider**.
3. Configure:
 - **Name:** `flowhub`
 - **Client type:** Confidential
 - **Client ID:** `flowhub` (copy this value)
 - **Client Secret:** (generate and copy)
 - **Redirect URIs:** `https://<flowhub-host>/signin-oidc`
 - **Post-logout redirect URIs:** `https://<flowhub-host>/signout-callback-oidc`
 - **Scopes:** `openid profile email`
4. Navigate to **Applications → Applications → Create**.
 - **Name:** `FlowHub`
 - **Slug:** `flowhub`
 - **Provider:** select the provider created above
5. Note the **OpenID Configuration URL**:
`https://authentik.home.freaxnx01.ch/application/o/flowhub/.well-known/openid-configuration`

FlowHub Environment Variables

Set these in your `.env` file or secrets manager (never commit):

```
Auth__OIDC__Authority=https://authentik.home.freaxn01.ch/application/o/flo
Auth__OIDC__ClientId=flowhub
Auth__OIDC__ClientSecret=<secret from Authentik>
```

Demo Mode (no OIDC)

Omit all `Auth__OIDC__*` variables. FlowHub will use `DemoAuthHandler` and auto-sign in as "Demo Operator". This is the default for `docker-compose.override.yml`.

Verification

After setting env vars and restarting:

1. Open FlowHub in browser.
2. Should redirect to Authentik login page.
3. After login, redirected back to FlowHub with the authenticated user's name in the top bar.

Runbook — Beta-MVP acceptance (Block 4 milestone, historical)

Source: `docs/runbooks/beta-mvp-acceptance.md`

Beta MVP — Operator Acceptance Runbook (Task 21)

- **Source:** `docs/superpowers/plans/2026-05-04-beta-mvp.md` §"Task 21"
 - **Goal:** Validate the Beta MVP slice end-to-end against the real homelab Wallabag + Vikunja, then push the branch.
 - **Branch:** `feat/beta-mvp` (28 commits ahead of `main`)
 - **Pre-state:** `just build` clean, `just test` 138/138 green. All operator-facing changes already committed.
-

What this runbook does

Walks through the seven-step demo path from the spec's "*Demo path (acceptance)*" section, plus the live Beta-smoke test run and the final push. Without this validation the Beta MVP slice is technically complete but architecturally unverified.

Time budget: 30–60 min depending on how many test captures you make.

Step 21.1 — Set user secrets

```
cd source/FlowHub.Web

dotnet user-secrets set "Ai:Provider" "Anthropic"
dotnet user-secrets set "Ai:Anthropic:ApiKey" "<sk-ant-...>"
dotnet user-secrets set "Skills:Wallabag:BaseUrl" "https://wallabag.home.fr
dotnet user-secrets set "Skills:Wallabag:ApiToken" "<wallabag-pat>"
dotnet user-secrets set "Skills:Vikunja:BaseUrl" "https://vikunja.home.frea
dotnet user-secrets set "Skills:Vikunja:ApiToken" "<vikunja-pat>"
dotnet user-secrets set "Skills:Vikunja:DefaultProjectId" "42"

cd ../../..
```

Where to get the tokens:

- Anthropic API key: console.anthropic.com → API keys
- Wallabag PAT: Wallabag UI → `/developer/new-token` (or "Site → Developer → My applications → Create a new client" depending on Wallabag version)
- Vikunja API token: Vikunja UI → Settings → API Tokens → New token (give it `tasks:write` scope minimum)
- Vikunja project id: open the target Vikunja project in the UI; the project id is in the URL (`/projects/<id>`). The plan example uses `42` — substitute your real one.

Sanity check after setting:

```
cd source/FlowHub.Web
dotnet user-secrets list
cd ../../..
```

You should see all 7 keys listed. Tokens are visible in plaintext (that's how user-secrets works); they live under

```
~/ .microsoft/usersecrets/<UserSecretsId>/secrets.json .
```


Step 21.2 — Boot and verify EventIds

```
just run > /tmp/flowhub-boot.log 2>&1 &  
sleep 8  
grep -E "EventId.*(3020|4020|5010|5011)" /tmp/flowhub-boot.log
```

Expected stdout (each on its own line):

- EventId 3020 – AI provider registered (provider=Anthropic, model=claude-haiku-4-5-20251001)
- EventId 4020 – Skill registered (skill=Wallabag)
- EventId 4020 – Skill registered (skill=Vikunja)
- EventId 5010 – Applying EF Core migrations...
- EventId 5011 – EF Core migrations up-to-date.

If any of those is **missing**, stop the app (`pkill -f FlowHub.Web`) and check:

Missing EventId	Likely cause	Fix
3020	<code>Ai:Provider</code> or <code>Ai:Anthropic:ApiKey</code> not set	re-run Step 21.1; <code>dotnet user-secrets list</code> to verify
4020 (Wallabag)	<code>Skills:Wallabag:BaseUrl</code> or <code>:ApiKey</code> not set	re-run Step 21.1; expect 4021 <code>SkillNotConfigured</code> instead in this case
4020 (Vikunja)	<code>Skills:Vikunja:BaseUrl</code> , <code>:ApiKey</code> , or <code>:DefaultProjectId</code> missing/zero	re-run Step 21.1; <code>:DefaultProjectId</code> must be <code>> 0</code>
5010 / 5011	<code>MigrationRunner</code> <code>IHostedService</code> not started — usually means <code>AddFlowHubPersistence</code> isn't wired in <code>Program.cs</code>	git log shows it landed in commit 27c9946 ; if branch is broken, <code>git diff main..HEAD -- source/FlowHub.Web/Program.cs</code>

Leave the app running for the next step (don't kill the background process yet).

Step 21.3 — Browser: empty Recent Captures

Open `http://localhost:5070/` in your browser. **Recent Captures grid should be empty** (the SQLite DB at `source/FlowHub.Web/flowhub.db` is fresh on first boot, *unless* it survived from earlier smoke tests).

For a strictly-clean start (optional):

```
pkill -f FlowHub.Web
rm -f source/FlowHub.Web/flowhub.db
just run > /tmp/flowhub-boot.log 2>&1 &
sleep 8
```

You can leave older captures in place — the demo only needs to add three new ones, and the matrix is identifiable by content + time.

Step 21.4 — URL capture → Wallabag

In the AppBar **Quick-Capture** field (top of every page), paste:

```
https://en.wikipedia.org/wiki/Hexagonal_architecture
```

Submit (Enter or the submit button).

Expected within ~3 seconds in the Recent Captures grid:

Field	Expected value
Stage	Completed
Skill	Wallabag
Title	Hexagonal architecture (or close — the AI's call; could be "Hexagonal architecture (software)" or similar)
Content	the URL

Cross-check: open Wallabag's homelab UI → "Unread" view. The new entry should be there with the article title and content scraped.

If Stage stuck at Raw for > 10s: AI classifier is not firing. Check `/tmp/flowhub-boot.log` for `EventId 3010 AiClassifierFellBackToKeyword` (means AI errored, fell back to deterministic) — the deterministic `KeywordClassifier` should still match URLs to Wallabag, so Stage should still progress.

If Stage Unhandled: integration call failed. Click into the capture (`/captures/{id}`) — the Capture Detail page now surfaces the `FailureReason` (per the final-review fixup in commit `c5340e9`). Common causes:

- 401 Unauthorized → token wrong or expired
- 503/timeout → Wallabag homelab not reachable from the dev box
- *Schema mismatch* (Wallabag response did not include an 'id' field) → unusual Wallabag version

Step 21.5 — Todo capture → Vikunja

Quick-capture:

```
todo: review Block 4 prep tomorrow
```

Expected within ~3 s:

Field	Expected value
Stage	Completed
Skill	Vikunja
Title	Review Block 4 prep tomorrow (or similar — AI may capitalize/normalize)
ExternalRef	a numeric id (Vikunja task id)

Cross-check: open Vikunja's project (the one matching `Skills:Vikunja:DefaultProjectId` from Step 21.1) → the new task should be visible.

The `todo:` prefix is what the `KeywordClassifier` and `AiClassifier` use to route to Vikunja. The `Title` is what the classifier produces; if it's null (e.g. AI fallback), the integration falls back to the truncated content (first 120 chars).

Step 21.6 — Nonsense capture → Orphan

Quick-capture:

```
asdfqwerty
```

Expected within ~3 s:

Field	Expected value
Stage	<code>Orphan</code>
Skill	<code>– (none matched)</code>
FailureReason	something like <code>"no skill matched"</code> or AI's explanation

The Dashboard's **"Needs Attention"** card should increment by 1.

This validates the Orphan path: AI classifier returned empty `MatchedSkill`, `CaptureEnrichmentConsumer` saw the empty match, called `MarkOrphanAsync`. The capture never reached `SkillRoutingConsumer`.

Step 21.7 — Restart, state survives

```
pkill -f FlowHub.Web
just run > /tmp/flowhub-boot.log 2>&1 &
sleep 8
```

Browser → `http://localhost:5070/`. **All three captures from steps 21.4-21.6 should still be visible** in the Recent Captures grid with their stages intact. This

is the core persistence validation — without EF Core / SQLite, the in-memory stub would lose them all.

The `flowhub.db` SQLite file at `source/FlowHub.Web/flowhub.db` is the durable state. To inspect:

```
sqlite3 source/FlowHub.Web/flowhub.db "SELECT Id, Stage, Source, Title, Ext
```

Stop the app for the next step:

```
pkill -f FlowHub.Web
```

Step 21.8 — Run live Beta tests (just test-beta)

Export the same secrets as env vars (the trait-gated tests resolve them via `Environment.GetEnvironmentVariable`):

```
export Skills__Wallabag__BaseUrl=https://wallabag.home.freaxnx01.ch
export Skills__Wallabag__ApiToken=<wallabag-pat>
export Skills__Vikunja__BaseUrl=https://vikunja.home.freaxnx01.ch
export Skills__Vikunja__ApiToken=<vikunja-pat>
export Skills__Vikunja__DefaultProjectId=42

just test-beta
```

Expected: 2 tests pass

(`WallabagLiveTests.HandleAsync_LiveWallabag_PostsUrlAndReturnsExternalRef` ,
`VikunjaLiveTests.HandleAsync_LiveVikunja_PutsTaskAndReturnsExternalRef`).

Side effect: 2 new entries appear in Wallabag and Vikunja (a Wikipedia article and a `todo: FlowHub Beta smoke test <ISO timestamp>` task). These are not cleaned up; manually delete them if you want a clean slate.

If the tests fail with `Skip:` messages, the env vars aren't being read — check the underscore convention (`Skills__Wallabag__BaseUrl` with **double underscores**, not single).

Step 21.9 — Final default suite

```
just build && just test
```

Expected: clean build (warnings-as-errors), 138/138 tests pass under `Category!=AI&Category!=BetaSmoke` .

This is your gate for Step 21.10. If anything is red here, **don't push** — fix or roll back first.

Step 21.10 — Push (final step)

Only after Steps 21.1–21.9 confirm green:

```
git push origin feat/beta-mvp
```

After push, decide:

- **Open a PR to** `main` for visibility (recommended — gives a place to attach the demo evidence).
- **Or merge directly to** `main` if no PR review is needed (single-user homelab project; PR may be ceremony).

Either way, document the homelab demo evidence in the PR description or in a follow-up commit on `main` :

- Screenshots of the three captures in Recent Captures (URL → Completed/Wallabag, todo → Completed/Vikunja, nonsense → Orphan)
- Screenshot or `sqlite3` dump showing rows survived restart
- The Wallabag entry id + Vikunja task id (matches `ExternalRef` columns)

This evidence is what feeds the *"Test-Ergebnisse dokumentiert (3)"* and *"Sub-Systeme als unabhängige Container deploybar (5)"* rubric items in Block 5's submission PDF.

Cleanup (optional)

After the demo, you can:

```
rm source/FlowHub.Web/flowhub.db # discard the demo SQLite DB
unset Skills__Wallabag__BaseUrl Skills__Wallabag__ApiToken \
Skills__Vikunja__BaseUrl Skills__Vikunja__ApiToken Skills__Vikunja__D
```

User-secrets persist on the dev machine (`~/ .microsoft/usersecrets/`) — leave them in place; they're per-project and harmless.

What this validates (rubric mapping)

- **Block 3 — KI-Werkzeug-Nutzung (12 pts):** the URL/todo flow exercises the production-runtime AI classifier; `just test-beta` evidence + screenshots strengthen the Block 5 submission section.
 - **Block 4 — Sub-Systeme als unabhängige Container deploybar (5 pts):** Wallabag + Vikunja are external homelab services; FlowHub talks to them via HTTP. Validates the integration boundary.
 - **Block 5 — Intelligente Services mit KI (6 pts):** AI classifier in the request path with graceful fallback (per ADR 0004) — exercised live.
 - **Cross-block — Demo path (acceptance):** the seven steps above are the exact gating criteria the brainstorm spec called out (`docs/superpowers/specs/2026-05-04-beta-mvp-design.md` §"Demo path (acceptance)").
-

References

- Brainstorm spec: `docs/superpowers/specs/2026-05-04-beta-mvp-design.md`
- Plan: `docs/superpowers/plans/2026-05-04-beta-mvp.md` (Task 21 is the last task)

- ADR 0005 — Persistence: docs/adr/0005-persistence.md
- ai-usage retrospective: docs/ai-usage.md §"Block 4 prep — Beta MVP"
- Bewertungskriterien: vault/0rganisation/Bewertungskriterien.md

Runbook — Public demo

Source: `docs/runbooks/public-demo.md`


Public Demo Environment

A fully open, rate-limited, self-resetting FlowHub instance at <https://demo.flowhub.freaxnx01.ch>. Designed so the operator pays €4.50/mo for the VPS and **€0** for AI inference, while still giving a visitor a working "click → see Capture classified by AI" round-trip.

Posture

Concern	Stance
Hosting	VPS-DE (existing Hetzner-class VM, Traefik-fronted)
Public URL	<code>https://demo.flowhub.freaxnx01.ch</code> (Cloudflare DNS → VPS-DE Traefik)
Auth	None — fully open; <code>Auth__0IDC__*</code> unset → <code>DemoAuthHandler</code> auto-signs every request
Rate limit	Traefik <code>RateLimit</code> middleware on the demo route: <code>average=10/min, burst=20, period=1m</code> per source IP
Data lifetime	All Captures + Tags + SkillRuns truncated every 15 minutes by <code>flowhub.demo-reset</code> sidecar; small fixture set reseeded
AI provider	OpenRouter <code>google/gemma-4-31b-it:free</code> only. KeywordClassifier auto-fallback on 429/error. <code>KeywordClassifier</code> is also the demo's safety net if the daily quota runs out.
Embeddings	Disabled — <code>Embeddings__ApiKey</code> unset → <code>AiEmbeddingService</code> is null → consumer no-ops → <code>Captures.Embedding</code> stays NULL. <code>GET /api/v1/captures/search</code> returns 503 ProblemDetails with an explanatory body (transparent about what's wired vs not).
Skill integrations	Disabled — <code>Skills__Vikunja__*</code> and <code>Skills__Wallabag__*</code> unset → captures stop at <code>Classified</code> stage and surface the <code>MatchedSkill</code> on the dashboard but never write to anyone's real Vikunja/Wallabag.
Observability	Prometheus + Grafana not exposed publicly — only <code>flowhub.web</code> + <code>postgres</code> + <code>rabbitmq</code> + <code>flowhub.demo-reset</code> go through the demo compose overlay. Operator metrics still scrapable via the VPS internal network.

Topology



Files

- `demo/docker-compose.yml` — overlay layered on top of the root `docker-compose.yml` via `docker compose -f docker-compose.yml -f demo/docker-compose.yml up`. Adds Traefik labels to `flowhub.web`, adds the `flowhub.demo-reset` sidecar, suppresses the public ports on `postgres` / `rabbitmq` / `prometheus` / `grafana`.

- `demo/.env.example` — demo-only env vars; **no real Skills_***, **no Embeddings**.
- `demo/reset/Dockerfile` + `demo/reset/reset.sh` — Alpine image with `postgresql-client` + `bash`. Sleep-loop runs `reset.sh` every 900 s. Script TRUNCATES Captures (CASCADE removes Tags + SkillRuns) and reseeds a fixture set of ~6 example captures spanning Wallabag / Vikunja / Orphan / Unhandled stages so the dashboard is never empty.

Deploy

On the VPS-DE host:

```
git clone https://github.com/freaxnx01/FlowHub-CAS-AISE.git
cd FlowHub-CAS-AISE
cp demo/.env.example .env
# Edit .env → set Ai__OpenRouter__ApiKey to the demo-only OpenRouter key
# (the one with a $1/mo hard cap configured in the OpenRouter dashboard)

docker compose -f docker-compose.yml -f demo/docker-compose.yml up --build
```

Cloudflare DNS entry (via `homelab-service-routing` skill):

- Type: `A`
- Name: `demo.flowhub`
- Target: VPS-DE public IPv4
- Proxy: ☐ Proxied (HTTPS termination by Cloudflare → origin TLS by Traefik)

OpenRouter key hygiene

The demo runs on a **dedicated OpenRouter key**, never the developer key:

1. <https://openrouter.ai/keys> → "Create Key"
2. Name: `flowhub-demo`
3. **Credit limit: \$1.00** (hard cap — exhausted key returns 402 → AiClassifier catches → KeywordClassifier fallback kicks in)
4. Rotate quarterly or after suspicious activity.

Reset cadence

Default: 900 s (15 min). To change:

```
# demo/docker-compose.yml
services:
  flowhub.demo-reset:
    environment:
      RESET_INTERVAL_SECONDS: "900" # adjust here
```

Operator runbook

Symptom	Action
<code>/api/v1/captures</code> returns 429 from Traefik	Rate-limit middleware fired — a single IP exceeded 20-req burst. No action needed.
All classifications show as <code>Keyword</code> fallback	OpenRouter daily quota exhausted OR Gemma model unavailable. Check OpenRouter dashboard; demo continues serving via KeywordClassifier.
Dashboard is empty mid-cycle	Reset just ran. Wait < 15 min; fixture seed will appear on the next reset, or trigger an immediate reset via <code>docker compose exec flowhub.demo-reset /reset.sh</code> .
Inappropriate user content visible	At most 15 min lifetime. To purge immediately: <code>docker compose exec flowhub.demo-reset /reset.sh</code> .
Need to take the demo offline	<code>docker compose -f docker-compose.yml -f demo/docker-compose.yml down</code> . Cloudflare DNS can stay; visitors get 522 Origin Unreachable.

Out of scope

- Multi-region failover
- WAF beyond Cloudflare's free tier + Traefik rate-limit
- Persistent demo-user accounts

- Captures persisting across resets

References

- `docker-compose.yml` (production stack reused by overlay)
- `homelab-service-routing` skill (Cloudflare DNS workflow)
- `docs/runbooks/test-services.md` (sibling environment — internal test, not public)
- [ROADMAP.md](#) "Additional AI Providers" (future swap to self-hosted Gemma if quota becomes painful)

Runbook — Test services

Source: `docs/runbooks/test-services.md`

Test Services Environment

A dedicated LXC container hosts disposable instances of the productivity services FlowHub integrates with (Vikunja, Wallabag, Paperless-ngx, Forgejo) — fully isolated from the production instances.

Host

- **Proxmox host:** `cirrus-pve` (192.168.1.111)
- **Container:** CT 128 `flowhub-test-services`
- **IP:** 192.168.1.105
- **Stack:** Debian + Docker + Traefik v3.6 (TLS via internal CA)
- **Storage:** `/home/admin/mydocker` (compose), `/home/admin/container-data/flowhub-test-services` (volumes, on NFS)

Start/stop from any host with SSH access:

```
ssh cirrus-pve "pct start 128"    # or: pct stop 128
```

Services & Endpoints

Service	Hostname	API base	Auth header
Vikunja	todo.flowhub-test-services.home.freaxnx01.ch	/api/v1	Authorization: Bearer tk_...
Wallabag	read-later.flowhub-test-services.home.freaxnx01.ch	/api	OAuth2 (see below)
Paperless-ngx	dms.flowhub-test-services.home.freaxnx01.ch	/api	Authorization: Token ...
Forgejo	git.flowhub-test-services.home.freaxnx01.ch (SSH :222)	/api/v1	Authorization: token ...

DNS caveat: the *.flowhub-test-services.home.freaxnx01.ch names currently resolve to the homelab Traefik dispatcher (192.168.1.101), but no dispatcher route exists yet — TLS returns unrecognized_name. Until that's fixed, clients must target the CT directly:

```
curl --resolve todo.flowhub-test-services.home.freaxnx01.ch:443:192.168.1.101 https://todo.flowhub-test-services.home.freaxnx01.ch/api/v1/info
```

Accounts

Service	Admin user	Notes
Vikunja	admin	email admin@freaxnx01.ch
Wallabag	wallabag	registration disabled
Paperless	admin	Django superuser
Forgejo	gitadmin	site admin

Wallabag OAuth2 flow

```
curl -X POST https://read-later.flowhub-test-services.home.freaxnx01.ch/oau
-d "grant_type=password" \
-d "client_id=<id>" -d "client_secret=<secret>" \
-d "username=wallabag" -d "password=<pw>"
# → { "access_token": "...", "refresh_token": "...", "expires_in": 3600 }
```

Use the returned `access_token` as `Authorization: Bearer <access_token>` against `/api`.

Secrets — Passbolt

All credentials live in Passbolt under the `flowhub-test-services/` prefix. **Never commit these.**

Resource name	Holds
<code>flowhub-test-services/vikunja API token</code>	Bearer token (full perms, expires 2027-05-12)
<code>flowhub-test-services/vikunja admin</code>	Web login for <code>admin</code>
<code>flowhub-test-services/wallabag OAuth2 client</code>	<code>client_id</code> (username) / <code>client_secret</code> (pw)
<code>flowhub-test-services/wallabag user</code>	<code>wallabag</code> web + password-grant credentials
<code>flowhub-test-services/forgejo API token</code>	API token, scope=all
<code>flowhub-test-services/paperless API token</code>	DRF Token for <code>admin</code>
<code>flowhub-test-services/paperless admin</code>	Web login for <code>admin</code>

Fetch via Passbolt CLI:

```
passbolt list resource --filter 'matches(Name, "flowhub-test-services")'  
passbolt get resource <id> -j # JSON incl. decrypted password
```

For local development, prefer `passbolt exec` to inject secrets as env vars instead of writing them to disk.

Grading rubric

Source: `vault/Organisation/Bewertungskriterien.md`

tags:

- claude-generated updated: 2026-04-29
-

CAS AISE — Bewertungskriterien (Projektarbeit)

Offizielle Moodle-Rubrik für die Projektarbeit. **Source of truth** für alle Block-Nachbereitungen — die Skala (0 / 1 / 3 / 5 , 0 / 1 / 4 / 7 , 0 / 3 / 7 / 10 , 0 / 2 / 6 , 0 / 1 / 7 / 12 , 0 / 2) zeigt die Gewichtung jedes Kriteriums.

Anwendung: Am Ende jeder Block-Nachbereitung diese Liste durchgehen. Pro Kriterium prüfen, welche Stufe heute realistisch ist, und fehlende Deliverables proaktiv erzeugen — nicht erst zum Schlussabgabe-Zeitpunkt.

Spezifikation

Spezifikation: Sind die wichtigsten Use-Cases und fachlichen Anforderungen benannt

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Spezifikation: Sind die Qualitätsanforderungen (NfA) nach SMART spezifiziert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise

- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Spezifikation: Ist die grundsätzliche Vision der Lösung beschrieben

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Entwurf

Entwurf: Ist der Lösungsansatz und die Architektur beschrieben (bildlich wie textuell)

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 4 — überwiegend bzw. mehrheitlich
- 7 — vollständig bzw. korrekt

Entwurf: Ist der Entwurf aus den verschiedenen Perspektiven (Struktur, Verhalten, Interaktion) beschrieben

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 4 — überwiegend bzw. mehrheitlich
- 7 — vollständig bzw. korrekt

Entwurf: Ist das DB-Modell spezifiziert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 2 — überwiegend bzw. mehrheitlich
- 3 — vollständig bzw. korrekt

Programmierung

Programmierung: Ist der Code lesbar, dokumentiert und nach Layer, Modulen und Sub-Systemen strukturiert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 4 — überwiegend bzw. mehrheitlich
- 7 — vollständig bzw. korrekt

Programmierung: Wurden die Konzepte von Quarkus, Jakarta EE und modernen Java Applikationen berücksichtigt

- 0 — nicht bzw. kaum
- 3 — teilweise bzw. ansatzweise
- 7 — überwiegend bzw. mehrheitlich
- 10 — vollständig bzw. korrekt

FlowHub-Kontext: Stack-Mismatch — FlowHub ist .NET 10 / [ASP.NET](#) Core / Blazor / EF Core. Dieses Kriterium wird **nicht aktiv verfolgt**. Die generischen "modernen Konzepte" (DI, Configuration, OpenAPI, async/await, Testbarkeit) sind über das Stack-äquivalente .NET-Ökosystem ohnehin abgedeckt.

Programmierung: Sind die Erkenntnisse aus der Programmierung dokumentiert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 2 — überwiegend bzw. mehrheitlich
- 3 — vollständig bzw. korrekt

Programmierung: Ist der Source-Code in einem Git-Repository verfügbar

- 0 — nicht bzw. kaum
- 2 — vollständig bzw. korrekt

Validierung

Validierung: Ist definiert, welches die Abnahmekriterien sind

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Validierung: Ist spezifiziert, wie die Applikation getestet wird und welche Technologien dazu verwendet werden

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Validierung: Sind Unit-Tests programmiert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — vollständig bzw. korrekt

Validierung: Sind die Test-Ergebnisse dokumentiert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 3 — vollständig bzw. korrekt

KI, Sub-Systeme & Reflexion

Wurden KI-unterstützende Werkzeuge verwendet und deren Nutzung beschrieben

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 7 — überwiegend bzw. mehrheitlich
- 12 — vollständig bzw. korrekt

Wurden mit Hilfe der KI intelligente und flexible Services gebaut

- 0 — nicht bzw. kaum
- 2 — teilweise bzw. ansatzweise
- 6 — vollständig bzw. korrekt

Wurde die Lösung in verschiedene Sub-Systeme aufgeteilt, die unabhängig voneinander als Container verteilt und betrieben werden können

- 0 — nicht bzw. kaum

- 1 — teilweise bzw. ansatzweise
- 3 — überwiegend bzw. mehrheitlich
- 5 — vollständig bzw. korrekt

Sind die Erfahrungen während der Projektarbeit mit KI-unterstützenden Werkzeugen als Fazit reflektiert

- 0 — nicht bzw. kaum
- 1 — teilweise bzw. ansatzweise
- 4 — überwiegend bzw. mehrheitlich
- 7 — vollständig bzw. korrekt

Punkte-Total

Bereich	Max-Punkte
Spezifikation (3 Items × 5)	15
Entwurf (7 + 7 + 3)	17
Programmierung (7 + 10 Quarkus + 3 + 2)	22
Validierung (5 + 5 + 3 + 3)	16
KI / Sub-Systeme / Reflexion (12 + 6 + 5 + 7)	30
Total	100

Mit Quarkus-Item ausgeklammert (Stack-Mismatch .NET): erreichbar 90 / 90.

Verweise

- Repo-Mirror dieser Dimensionen (Kurz-Summary): `.ai/cas-instructions.md` → Section "Grading"
- Block-Nachbereitungen sollen die Kriterien aktiv abhaken (siehe `Blöcke/<NN ...>/<NN ...>` - c) Nachbereitung.md)

Hilfsmittelverzeichnis & Eigenständigkeitserklärung

Source: `docs/submission/eigenstaendigkeitserklaerung.md`

Hilfsmittelverzeichnis & Eigenständigkeitserklärung

Modul: CAS AI-Assisted Software Engineering (AISE) · W4B-C-AS001 · ZH-Sa-1 · FS26 **Projektarbeit:** FlowHub — AI-assisted personal inbox **Studierender:** Andreas Imboden (`freaxnx01`) **Repository:** <https://github.com/freaxnx01/FlowHub-CAS-AISE> **Abgabe-Termin:** vor 2026-07-04 24:00

Dieses Dokument ist gemäss FFHS-Richtlinie *"Hinweise und Anforderungen zum Umgang mit generativer künstlicher Intelligenz"* (Stand 01.08.2025) der Projektarbeit beizulegen. Es gehört zwingend zur Abgabe jedes seit dem 1.8.2023 begonnenen unbeaufsichtigten schriftlichen Leistungsnachweises.

1 — Hilfsmittelverzeichnis

Die folgende Tabelle deklariert sämtliche Hilfsmittel, die bei der Erarbeitung dieser Projektarbeit eingesetzt wurden — für Text, Code und Dokumentation. Eine ausführliche, narrativ-strukturierte Beschreibung des KI-Workflows liegt unter `docs/ai-usage.md` im Repository; die per-Block-Reflexion findet sich in den jeweiligen Block-Nachbereitungen unter `vault/Blöcke/`.

Welches Hilfsmittel wurde eingesetzt?	Wozu wurde das Hilfsmittel eingesetzt?	Betroffene Stellen
Claude Code (Opus 4.7, 1M context)	Brainstorming, ADR-Entwürfe, Plan-Schreiben, Controller für Subagent-Dispatches, Code-Reviews, Submission Document.	docs/superpowers/specs/ , docs/superpowers/plans/ , docs/adr/ , docs/spec/ , docs/insights/ , docs/ai-usage.md , SUBMISSION.md , grosse Teile von source/ und tests/ über Subagent-Dispatches.
Claude Sonnet 4.6 (Subagents)	Implementer-, Spec-Reviewer- und Code-Quality-Reviewer-Subagents im superpowers subagent-driven-development Workflow. TDD-Ausführung, urteilsintensive Refactorings.	Code-Änderungen in source/ und tests/ ab Block 3 — konkret: Async-Pipeline (Slice B), REST-API (Slice A), AI-Classifer (Slice C), EF-Persistenz (Block 4).
Claude Haiku 4.5	Mechanische Tasks im SDD-Workflow — Projekt-Scaffolding, Datei-Moves, repetitive Registrierungen, Dokumentations-Stubs.	Diverse Scaffolding- und Registrierungs-Tasks in Block 3-5; pro-Block-Details in docs/ai-usage.md .
GitHub Copilot	Inline-Code-Vorschläge während des Editierens in VS Code.	Sparsam — Claude Code steuert ganze Sessions; Copilot nur für kurze Completions ausserhalb dedizierter Sessions.

Welches Hilfsmittel wurde eingesetzt?	Wozu wurde das Hilfsmittel eingesetzt?	Betroffene Stellen
ChatGPT (GPT-4-Familie)	Ad-hoc-Klärung von Konzepten, Quervergleiche, während Claude Code an einem anderen Thema arbeitete.	Punktuelle Cross-Checks; nicht zur Generierung abgegebenen Codes oder Texts eingesetzt.
DeepL / Google Translate	Übersetzung einzelner Passagen zwischen Deutsch und Englisch (z. B. Moodle-Originalwortlaut in §2 des Submission Document, deutsche Vault-Notizen, die in englischen Dokumenten referenziert werden).	Ausgewählte DE↔EN-Passagen in <code>SUBMISSION.md</code> , <code>docs/ai-usage.md</code> , <code>vault/Projektarbeit/Learnings.md</code> .
Mistral <code>mistral-embed</code> (via API)	Laufzeit-Bestandteil des Produkts — Embeddings für die FlowHub-Semantic-Search-Funktion (<code>Captures.Embedding</code> , ADR 0006). Nicht zur Erzeugung von Abgabe-Texten verwendet.	Nur Produktions-Code: <code>FlowHub.AI</code> Embedding-Adapter, <code>Captures</code> - Tabellen-Spalte.
OpenRouter (Gemma free tier)	Laufzeit-Bestandteil der öffentlichen Demo — Klassifikations-Fallback, solange das	Nur Demo-Umgebung (<code>demo.flowhub.freaxnx01.ch</code>).

Welches Hilfsmittel wurde eingesetzt?	Wozu wurde das Hilfsmittel eingesetzt?	Betroffene Stellen
	\$1-Budget der Demo nicht erschöpft ist. Nicht zur Erzeugung von Abgabe-Texten verwendet.	

Anteil generiert vs. handgeschrieben (Schätzung): Für Implementierungs-Code (Block 3–5) ca. 70–80 % KI-erzeugt via Subagent-Workflow, danach manuell reviewed, angepasst oder verworfen. Für Dokumentation (ADRs, Specs, Runbooks, Submission Document) liegt der Anteil ähnlich, jedes Artefakt durchlief jedoch manuelle Nachbearbeitung in Bezug auf inhaltliche Korrektheit und Stil. Pro-Slice-Aufschlüsselung mit konkreten Zahlen in `docs/ai-usage.md`. Alle Architekturentscheidungen, Scope-Entscheidungen und die finale Annahme jedes Artefakts liegen beim Autor.

2 — Eigenständigkeitserklärung

Hiermit erkläre ich,

- dass ich die vorliegende Arbeit selbstständig verfasst habe,
- dass alle sinngemäss und wörtlich übernommenen Textstellen aus fremden Quellen kenntlich gemacht wurden,
- dass alle mit Hilfsmitteln erbrachten Teile der Arbeit präzise deklariert wurden (siehe §1 Hilfsmittelverzeichnis oben),
- dass keine anderen als die im Hilfsmittelverzeichnis aufgeführten Hilfsmittel verwendet wurden,
- dass das Thema, die Arbeit oder Teile davon nicht bereits Gegenstand eines Leistungsnachweises eines anderen Moduls waren, sofern dies nicht ausdrücklich mit der Referentin oder dem Referenten im Voraus vereinbart wurde,
- dass ich mir bewusst bin, dass meine Arbeit elektronisch auf Plagiate und auf Dittautorschaft menschlichen oder technischen Ursprungs überprüft werden

kann und ich hiermit der FFHS das Nutzungsrecht so weit einräume, wie es für diese Verwaltungshandlungen notwendig ist.

Ort, Datum: Sisseln, _____

Name: Andreas Imboden

Unterschrift:
